

1. Models, Abstractions and Basic Concepts

DEFINITION (System)

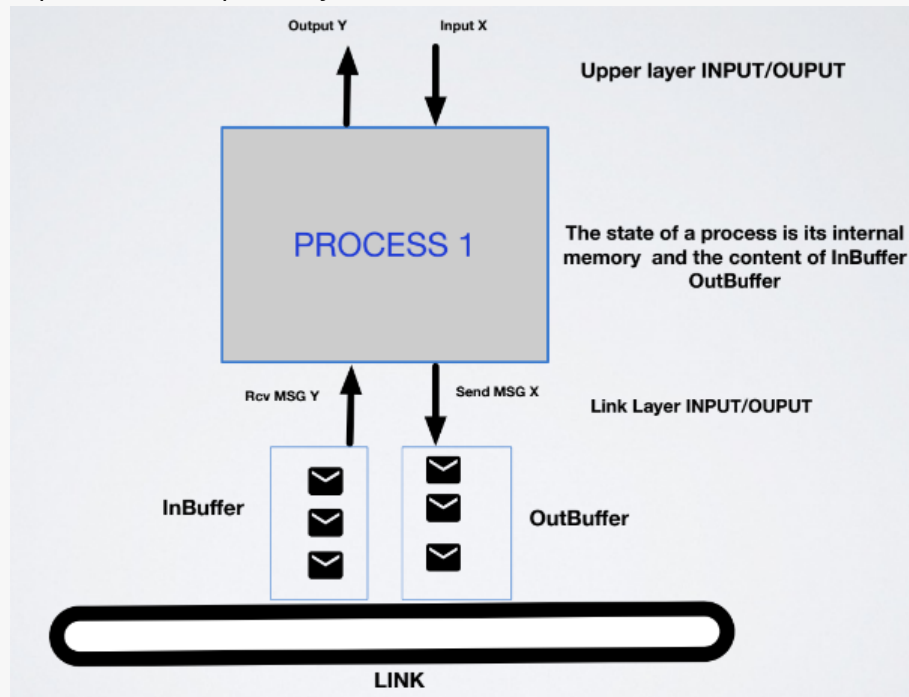
We have n processes in $\Pi : \{p_0, \dots, p_{n-1}\}$ with distinct identities.

Processes communicate with a **communication graph**: $G : (\Pi, E)$ (usually G is complete).

The communication happens by exchanging messages on communication link.

DEFINITION (Process)

A process is a (possibly infinite) State Machine (I/O Automaton).



Processes Communication is based on a link defined by an Input Buffer and an Output Buffer.

- internal states: set Q
- initial states: set $Q_i \subset Q$
- Messages: set all possible messages M in the form $\langle sender, receiver, payload \rangle$
- $InBuff_j$: multiset of delivered messages
- $OutBuff_j$: multiset of inflight messages (messages sent but not delivered)

$$P_j(q \in Q \cup Q_{in}, InBuff_j) = (q' \in Q, Send_{msg} \subset M)$$

where:

- $OutBuff_j = OutBuff_j \cup Send_{msg}$
- $InBuf = \emptyset$

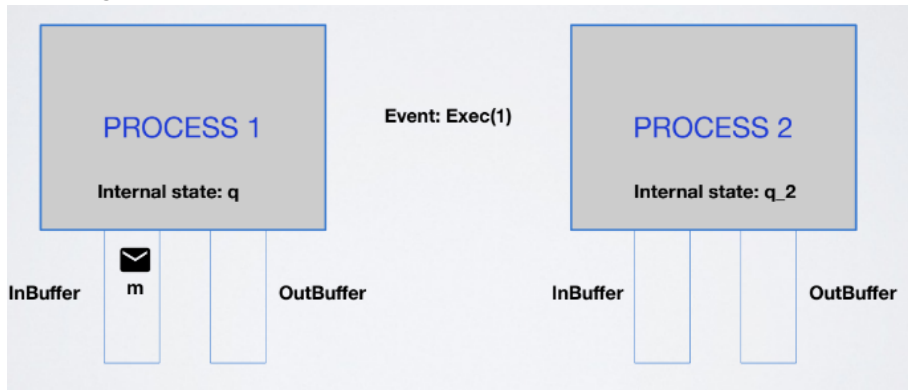
MODEL: Asynchronous Executions

DEFINITION (Execution)

Scheduling of a set of events (scheduler):

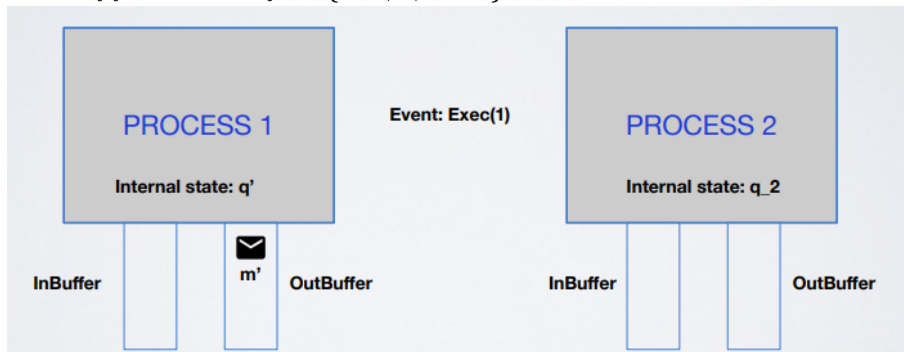
- Delivery of a message $Del(m, i, j)$: move message m from $OutBuff_i$ to $InBuff_j$
- Execution of a local step $Exec(i)$: process i executes one step of its state machine

Message is on InBuffer of the first process



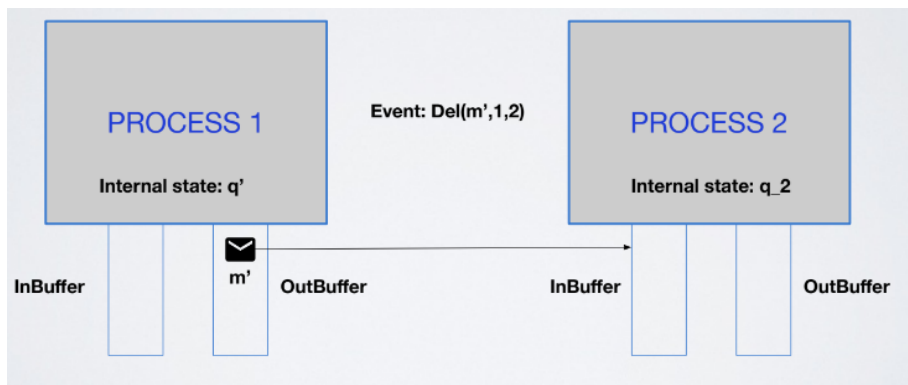
Message passes at the OutBuffer of P_1 doing:

- $P_1(q, \{m\}) = (q', \{< 1, 2, m'\})$
- $OutBuff_1 = OutBuff_1 \cup \{< 1, 2, m' >\}$

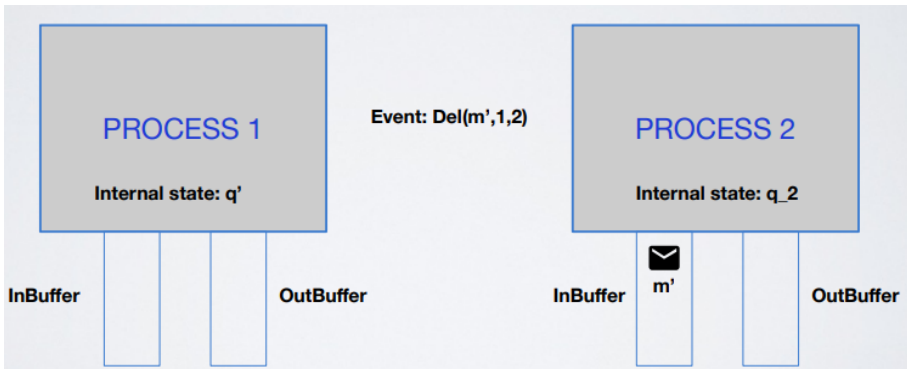


OutBuffer of P_1 send the message m' to InBuffer of P_2 doing:

- $InBuff_2 = InBuff_2 \cup \{< 1, 2, m' >\}$

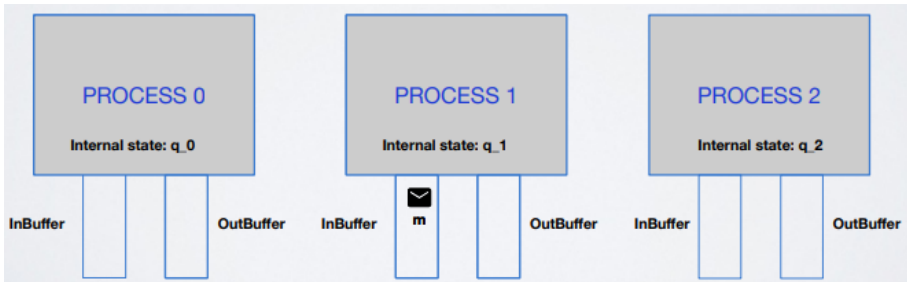


result:



Configuration C_t : is a vector of n components.

Component j is the state of process j : $C_t[j](q_i, InBuff_j, outBuff_j)$



$$C_0 = \langle (q_0, \{\}, \{\}), (q_1, \{m\}, \{\}), (q_2, \{\}, \{\}) \rangle$$



An event e is enabled in a configuration C if:

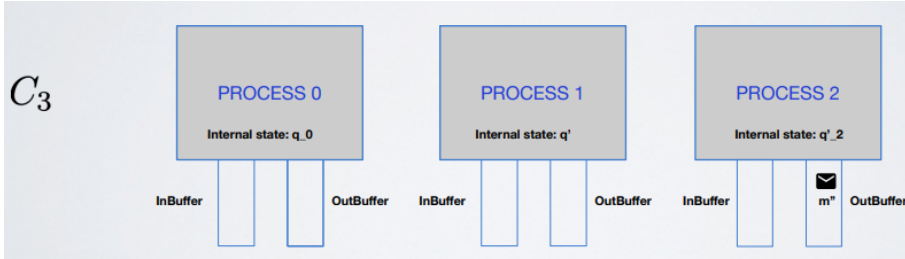
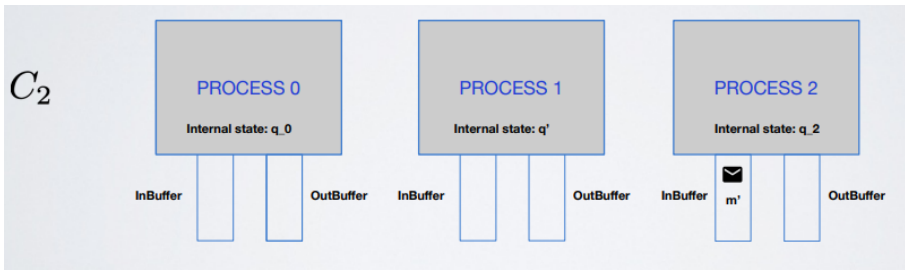
- $\text{Del}(m', 0, 2)$ is not enabled in C_0 because OutBuff_0 does not contain a message m' .
- $\text{Exec}(1)$ is enabled in $C_0 \text{ as } \text{Exec}(0)$ and $\text{Exec}(2)$.

Execution: infinite sequence that alternate configurations and events


example: $(C_0, e_0, C_1, e_1, \dots)$ such that each event e_t is enabled in configuration C_t that is obtained by applying $e_{t-1} \rightarrow C_{t-1}$

$$\varepsilon : (C_0, e_0 = \text{Exec}(1), C_1, e_1 = \text{Del}(1, 2, m'), C_2, e_2 = \text{Exec}(2), C_3, e_3 = \text{Del}(2, 3, m'), \dots)$$





Plotting this in a graph ($time \times processes$) we obtain something like:



DEFINITION (Fair)

E is fair if each process p_i executes an infinite number of local computation ($Exec(i)$ events are not finite) and each message m is eventually delivered (not possible to stall forever a message: there must exists a $Del(m, x, y)$).

Unfair executions brake any possible non-trivial algorithm so we will always consider fair executions.

When you create a process, the only thing that you know is the id of the other processes.

DEFINITION (Local Execution (local view))


Given an execution E and a process p_j , we define the local execution ($E|_{p_j}$) of p_j the subset as the subset of events in E that impact p_j .

Considered:

$$\varepsilon : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3 = Del(2, 3, m'), \dots)$$

we have:

- $\varepsilon|_{p_1} = (Del(0, 1, m), \dots)$
- $\varepsilon|_{p_2} = (Del(3, 2, m'), \dots)$



The only way to reach a status where all processes know informations about the others is communicate them with a row from p_2 to p_1 for example but in that case the event must end before the next one.


Different execution could give the same local execution:

$$\varepsilon = (C_0, e_0 = Exec(0), C_1, e_1 = Del(0, 1, m), C_2, e_2 = Exec(3), C_3, e_3 = Del(3, 2, m'), C_4, \dots)$$

$$\varepsilon' = (C_0, e_2 = Exec(3), C'_1, e_3 = Del(3, 2, m'), C'_2, e_0 = Exec(0), C'_3, e_1 = Del(0, 1, m), C_4, \dots)$$

$$\varepsilon|_{p_1} = \varepsilon'|_{p_1}$$

$$\varepsilon|_{p_2} = \varepsilon'|_{p_2}$$



To identify an execution all the process have to communicate them information to each others. In this case we say that p_1 and p_2 cannot distinguish E to E' .

DEFINITION (Indistinguishability)

$$\varepsilon = (C_0, e_0 = Exec(0), C_1, e_1 = Del(0, 1, m), C_2, e_2 = Exec(3), C_3, e_3 = Del(3, 2, m'), C_4, \dots)$$

$$\varepsilon' = (C_0, e_2 = Exec(3), C'_1, e_3 = Del(3, 2, m'), C'_2, e_0 = Exec(0), C'_3, e_1 = Del(0, 1, m), C_4, \dots)$$

if we have:

$$\forall p_j \in \Pi, \varepsilon|_{p_j} = \varepsilon'|_{p_j}$$

then we can say that E and E' are indistinguishable.

THEOREM:

In the asynch. model there is no distributed algorithm capable of reconstructing the system execution.