

MODELS, ABSTRACTIONS AND BASIC CONCEPTS

DISTRIBUTED SYSTEMS
Master of Science in Cyber Security
A.Y. 2023/2024



SAPIENZA
UNIVERSITÀ DI ROMA

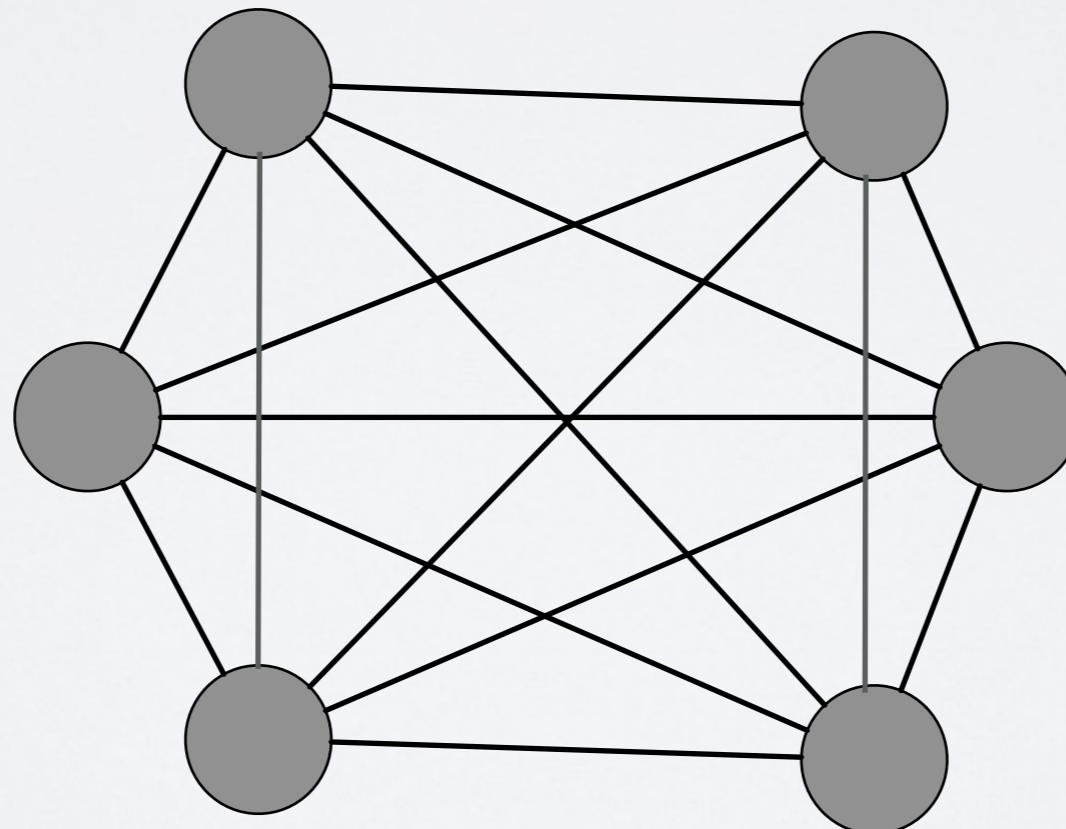


CIS SAPIENZA
CYBER INTELLIGENCE AND INFORMATION SECURITY

MODEL: SYSTEM AND PROCESSES

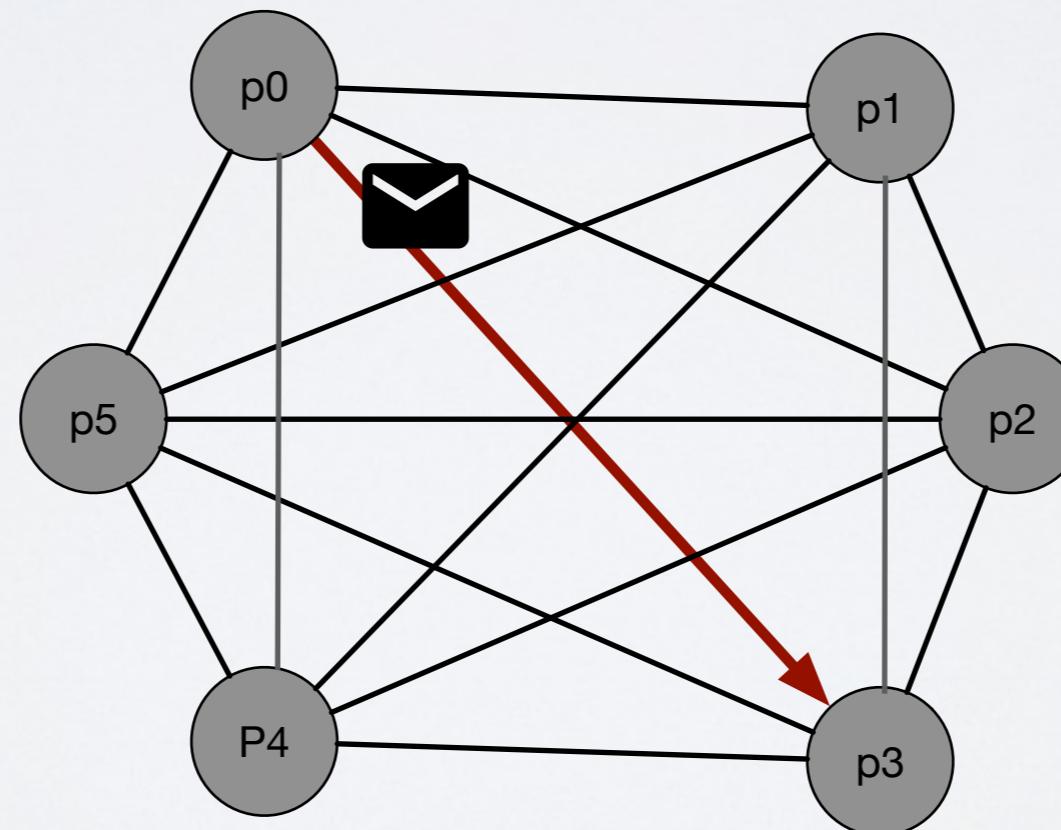
SYSTEM

- We have n processes in $\Pi : \{p_0, \dots, p_{n-1}\}$ with distinct identities.
- Processes communicate using a communication graph $G : (\Pi, E)$ (usually, G is complete).
- The communication happens by exchanging messages on communication links (edges of G).



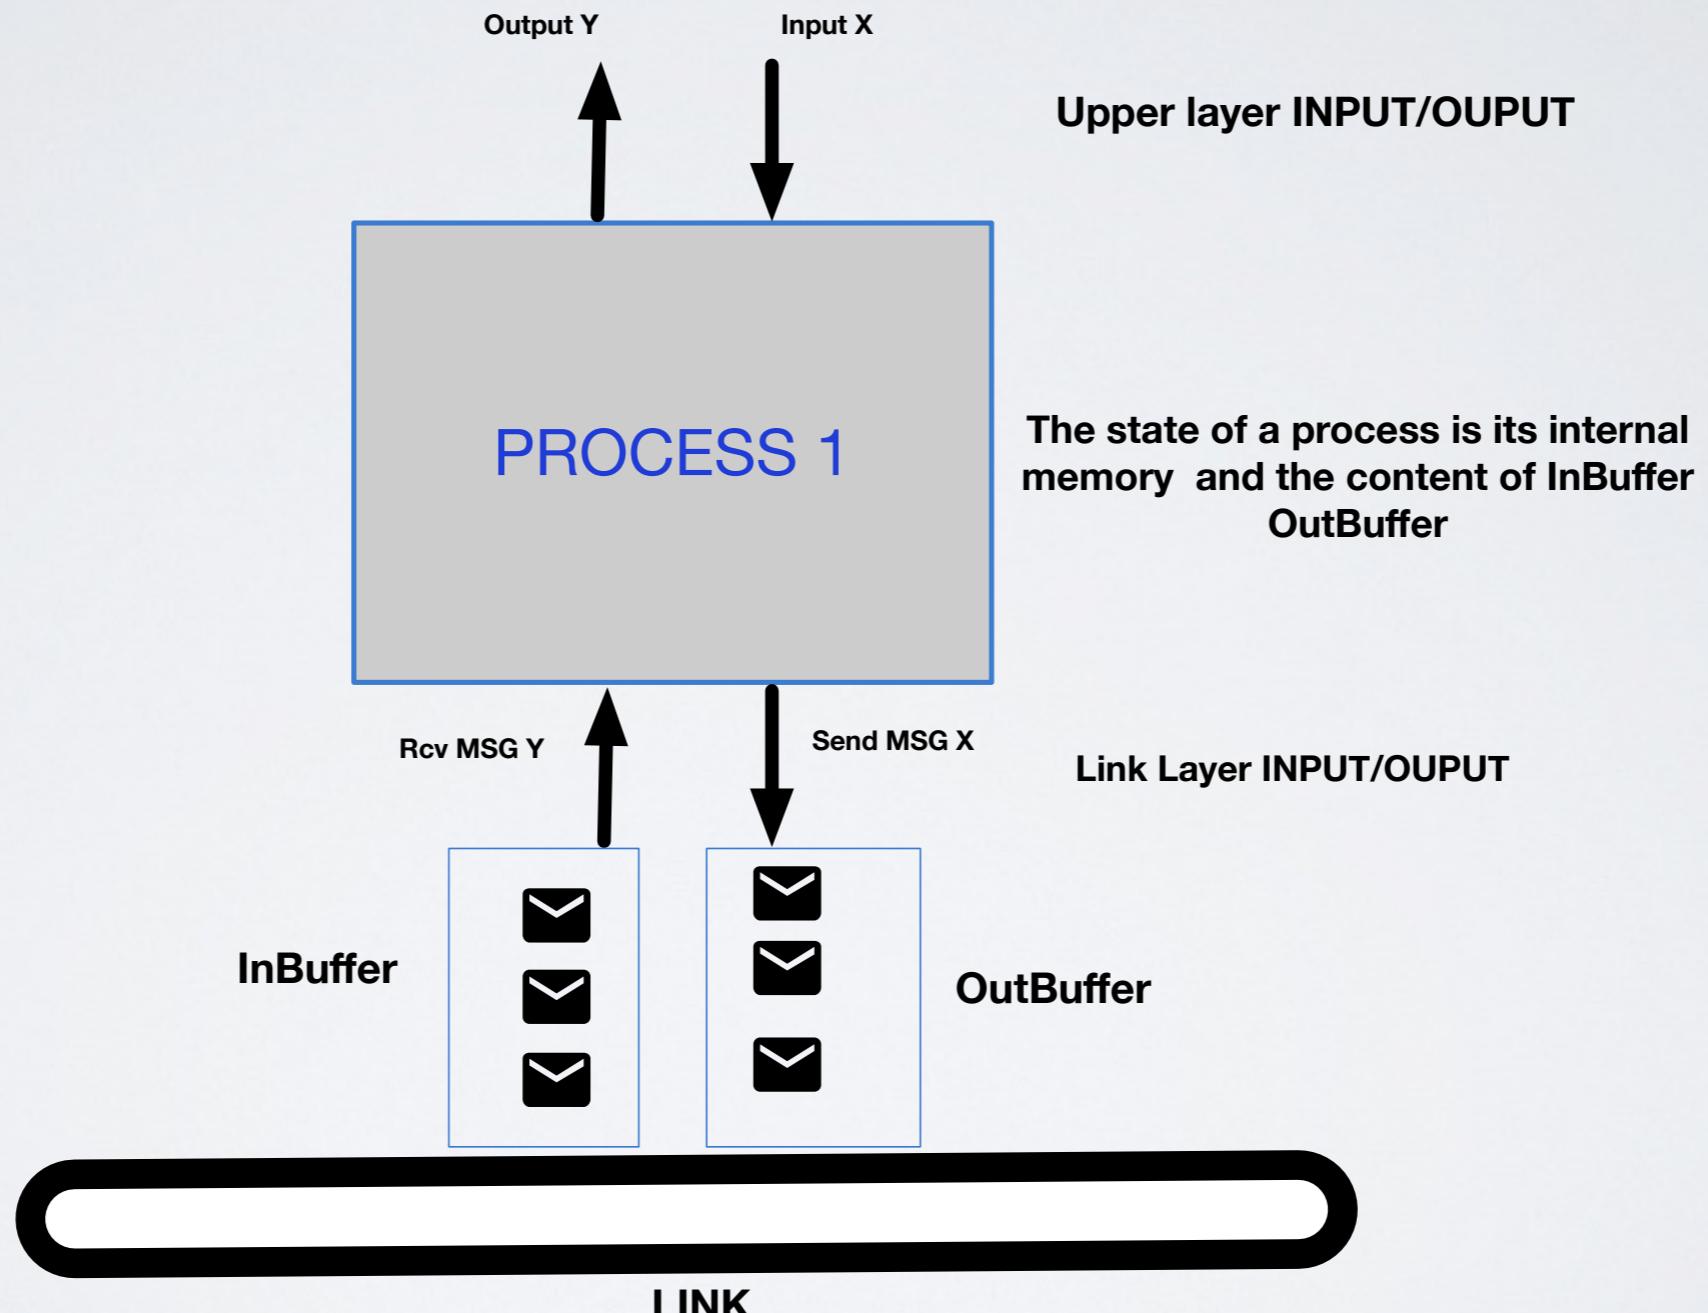
SYSTEM

- We have n processes in $\Pi : \{p_0, \dots, p_{n-1}\}$ with distinct identities.
- Processes communicate using a communication graph $G : (\Pi, E)$ (usually, G is complete).
- The communication happens by exchanging messages on communication links (edges of G).



PROCESS

- A process is a (possibly infinite) State Machine (I/O Automaton).



PROCESS

- A process is an infinite state I/O Automaton:
 - Internal states - set Q
 - Initial states - set $Q_i \subset Q$
 - Messages - set of all possible messages M in the form
 $\langle \text{sender}, \text{receiver}, \text{payload} \rangle$
 - $InBuf_j$ - multiset of delivered messages
 - $OutBuf_j$ - multiset of “inflight” messages (messages sent but not delivered).

$$P_j(q \in Q \cup Q_{in}, InBuf_j) = (q' \in Q, Send_msg \subset M)$$

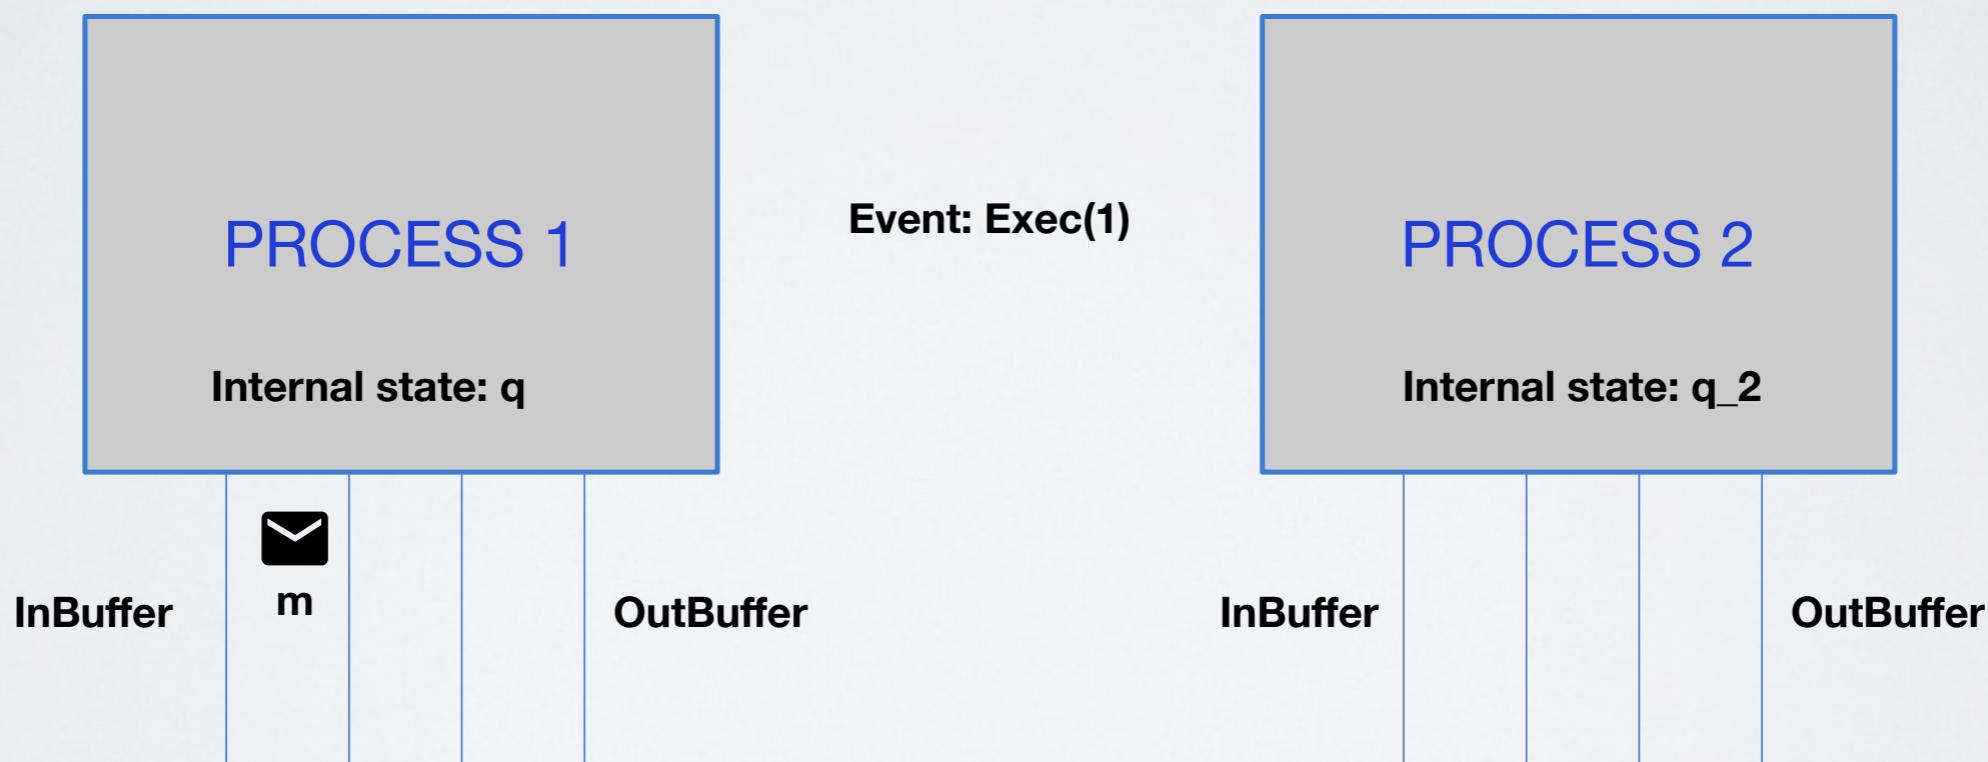
$$OutBuf_j = OutBuf_j \cup Send_msg$$

$$InBuf_j = \emptyset$$

MODEL: ASYNCHRONOUS EXECUTIONS

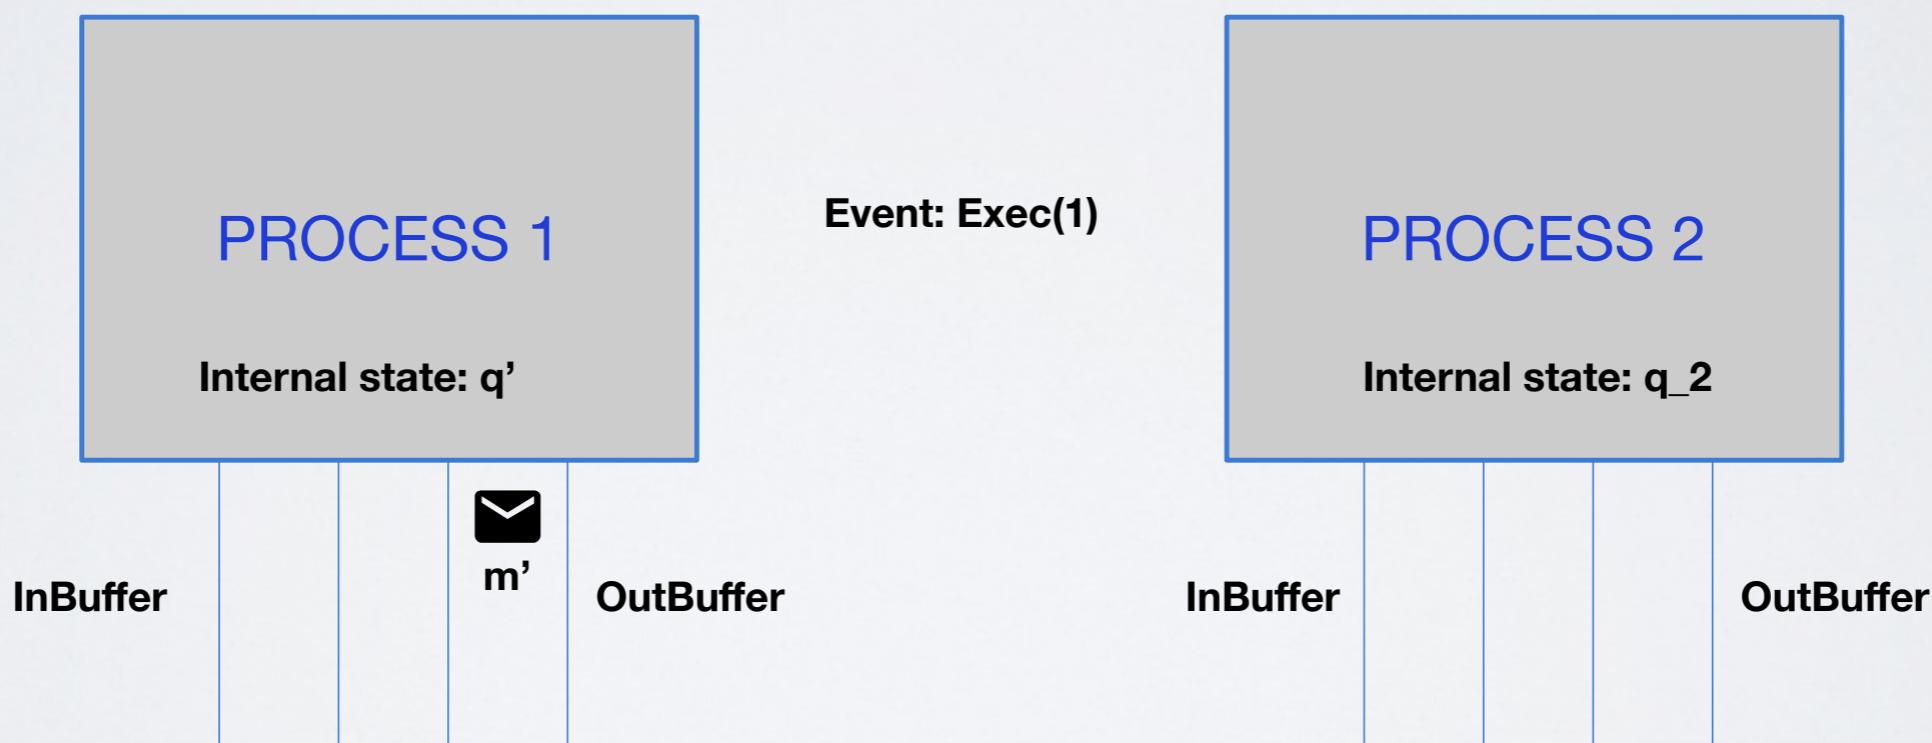
EXECUTION

- There is an adversary that schedule a set of events (scheduler):
 - Delivery of a msg - $\text{Del}(m,i,j)$: move message m from OutBuff_i to InBuff_j
 - Execution of a local step - $\text{Exec}(i)$: process i executes one step of its state machine.



EXECUTION

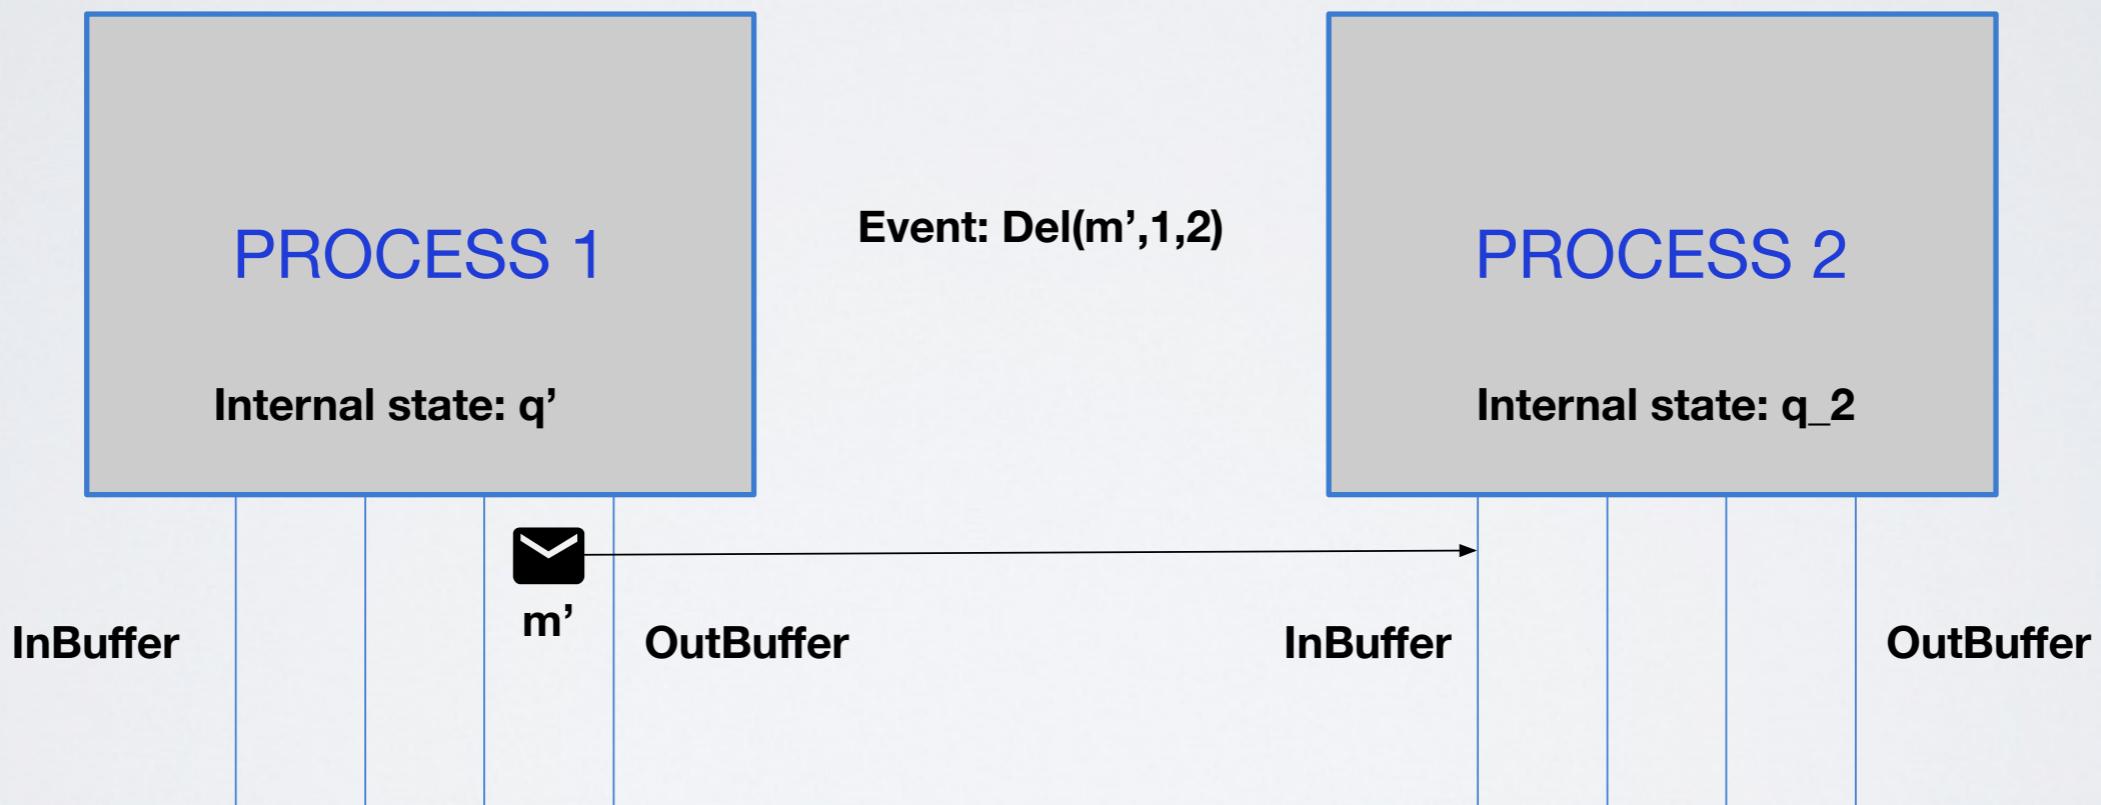
- There is an adversary that schedule a set of events (scheduler):
 - Delivery of a msg - $\text{Del}(m, i, j)$: move message m from OutBuff_i to InBuff_j
 - Execution of a local step - $\text{Exec}(i)$: process i executes one step of its state machine.



$$P_1(q, \{m\}) = (q', \{\langle 1, 2, m' \rangle\})$$
$$\text{OutBuf}_1 = \text{OutBuf}_1 \cup \{\langle 1, 2, m' \rangle\}$$

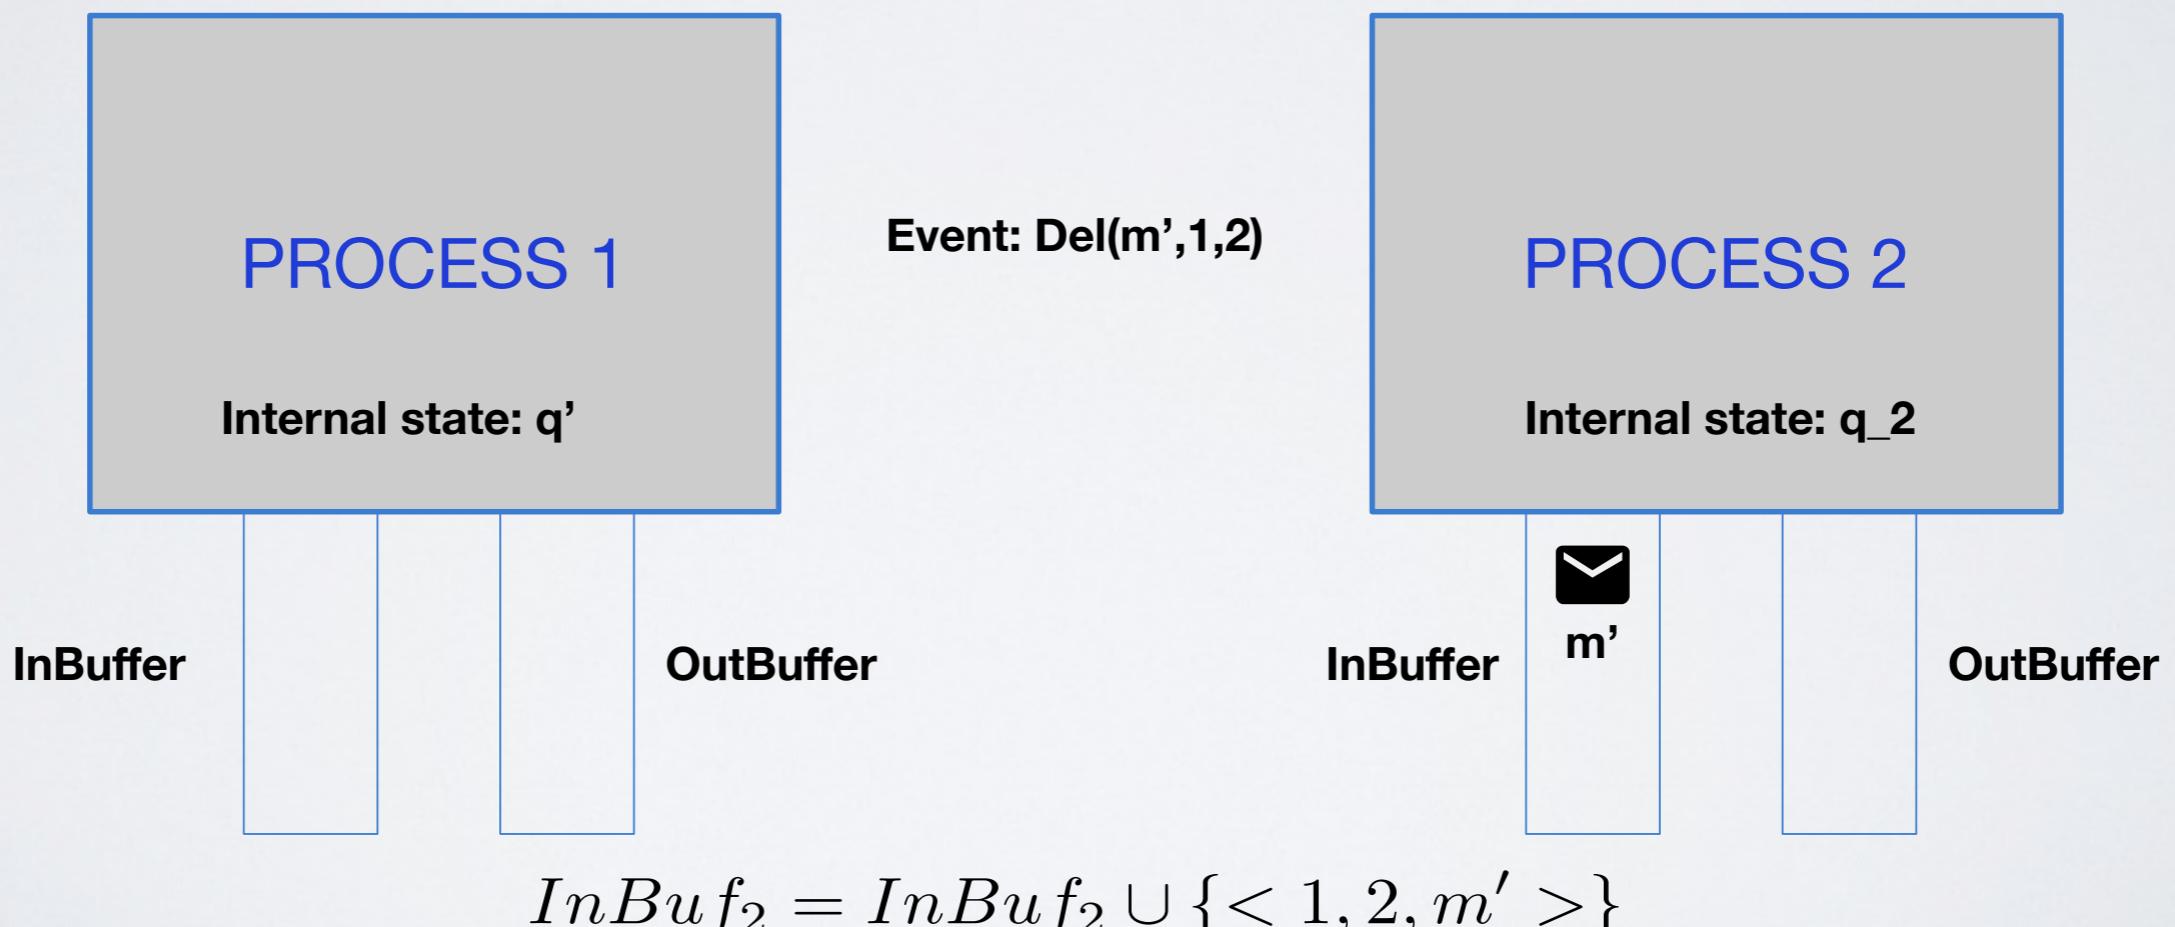
EXECUTION

- There is an adversary that schedule a set of events (scheduler):
 - Delivery of a msg - $\text{Del}(m,i,j)$: move message m from OutBuff_i to InBuff_j
 - Execution of a local step - $\text{Exec}(i)$: process i executes one step of its state machine.



EXECUTION

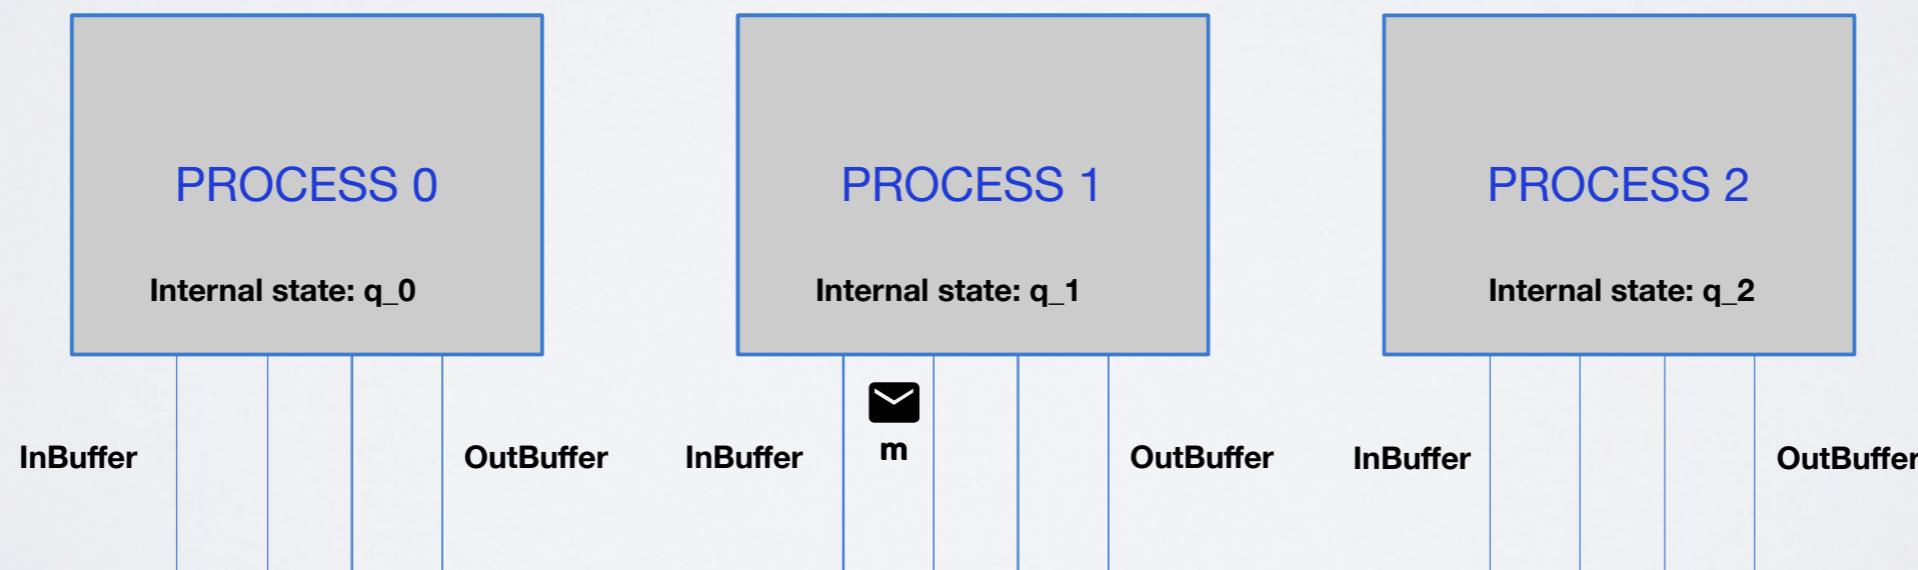
- There is an adversary that schedule a set of events (scheduler):
 - Delivery of a msg - $\text{Del}(m,i,j)$: move message m from OutBuff_i to InBuff_j
 - Execution of a local step - $\text{Exec}(i)$: process i executes one step of its state machine.



EXECUTION

- A configuration C_t is a vector of n components. Component j indicates the state of process j (that is $C_t[j]$ ($q_i, InBuff_j, OutBuff_j$)).

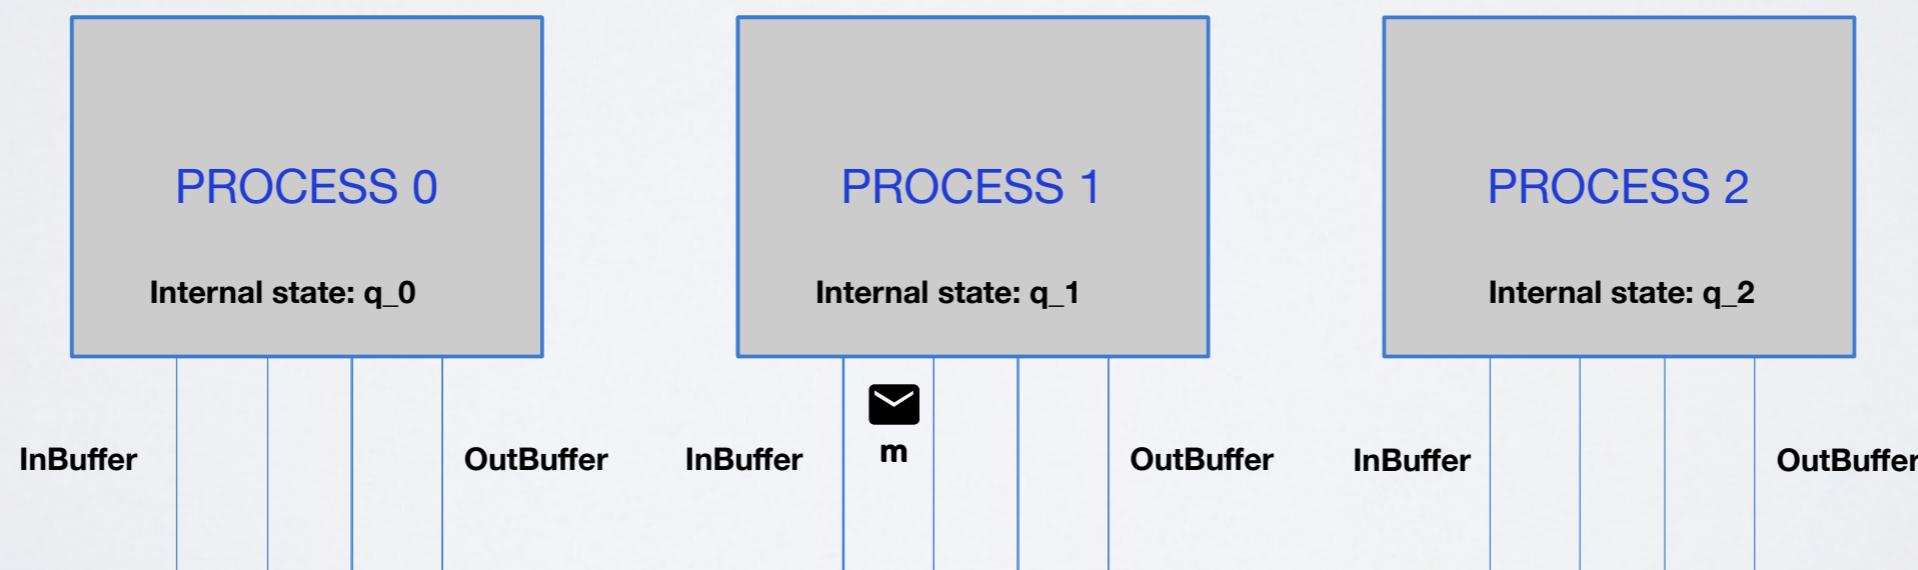
$$C_0 = < (q_0, \{\}, \{}), (q_1, \{m\}, \{}), (q_2, \{\}, \{}) >$$



EXECUTION

- An event e is enabled in a configuration C if it can happen:
 - Del(m',0,2) is not enabled in C_0 because OutBuf_0 does not contain a message m'.
 - Exec(1) is enabled in C_0, as Exec(0) and Exec(2).

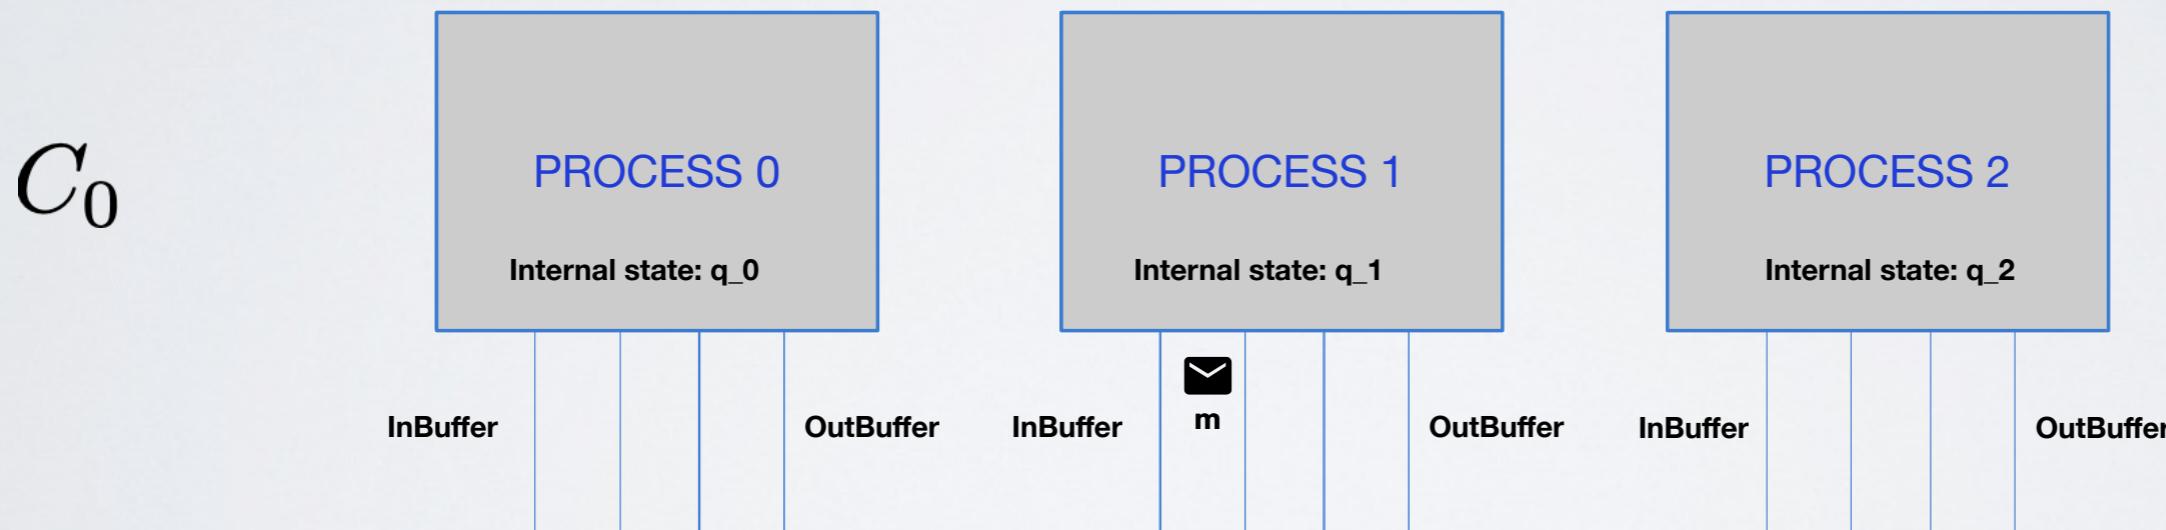
$$C_0 = < (q_0, \{\}, \{\}), (q_1, \{m\}, \{\}), (q_2, \{\}, \{\}) >$$



EXECUTION

- An execution is infinite sequence that alternates configurations and events: $(C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$ such that each event $e_{\{t\}}$ is enabled in configuration C_t and C_t is obtained by applying $e_{\{t-1\}}$ to $C_{\{t-1\}}$.

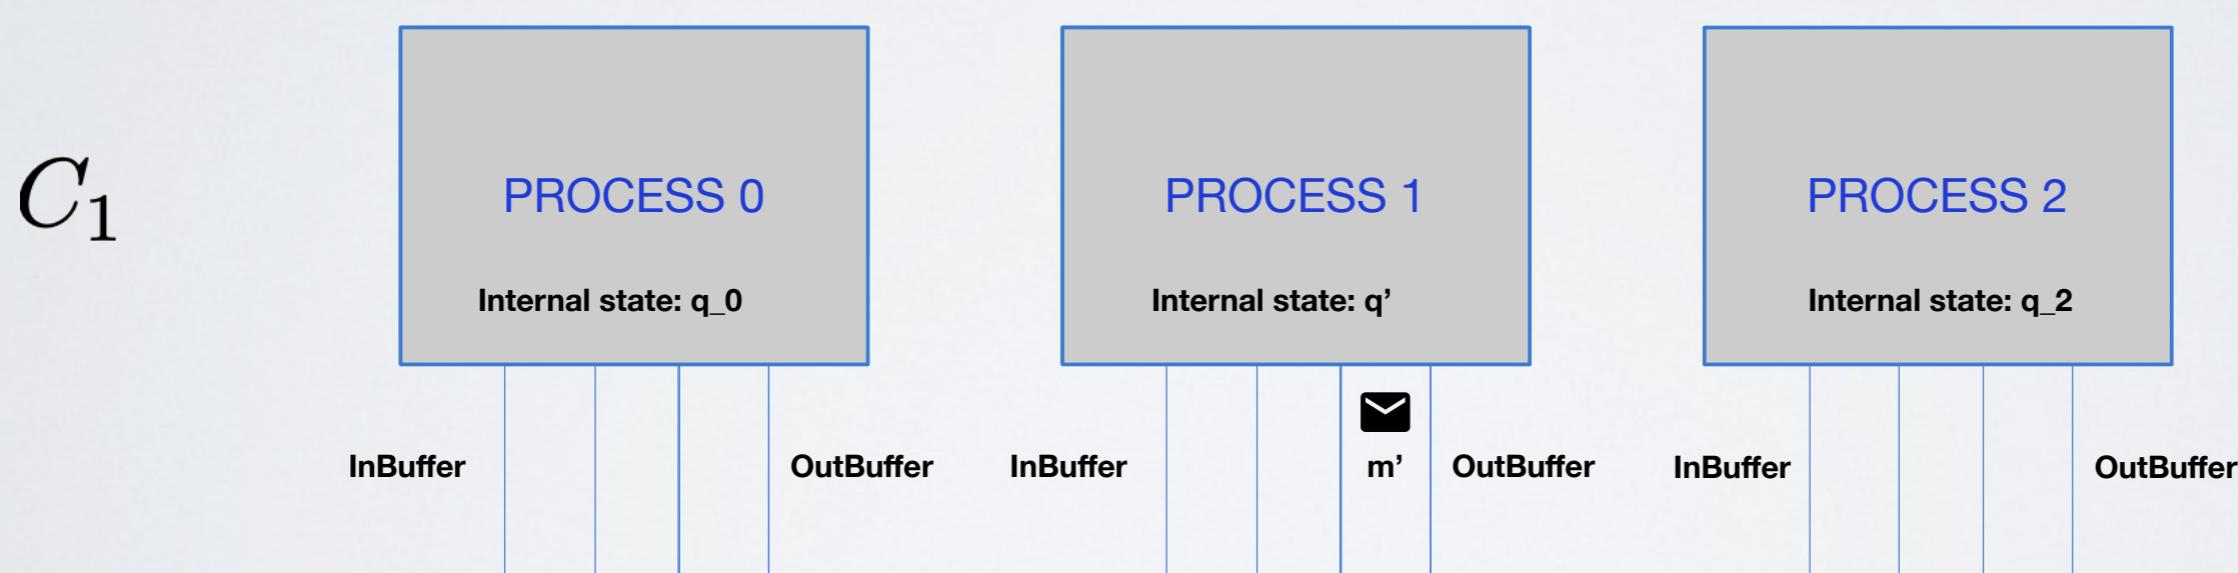
$$\mathcal{E} : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$$



EXECUTION

- An execution is infinite sequence that alternates configurations and events: $(C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$ such that each event $e_{\{t\}}$ is enabled in configuration C_t and C_t is obtained by applying $e_{\{t-1\}}$ to $C_{\{t-1\}}$.

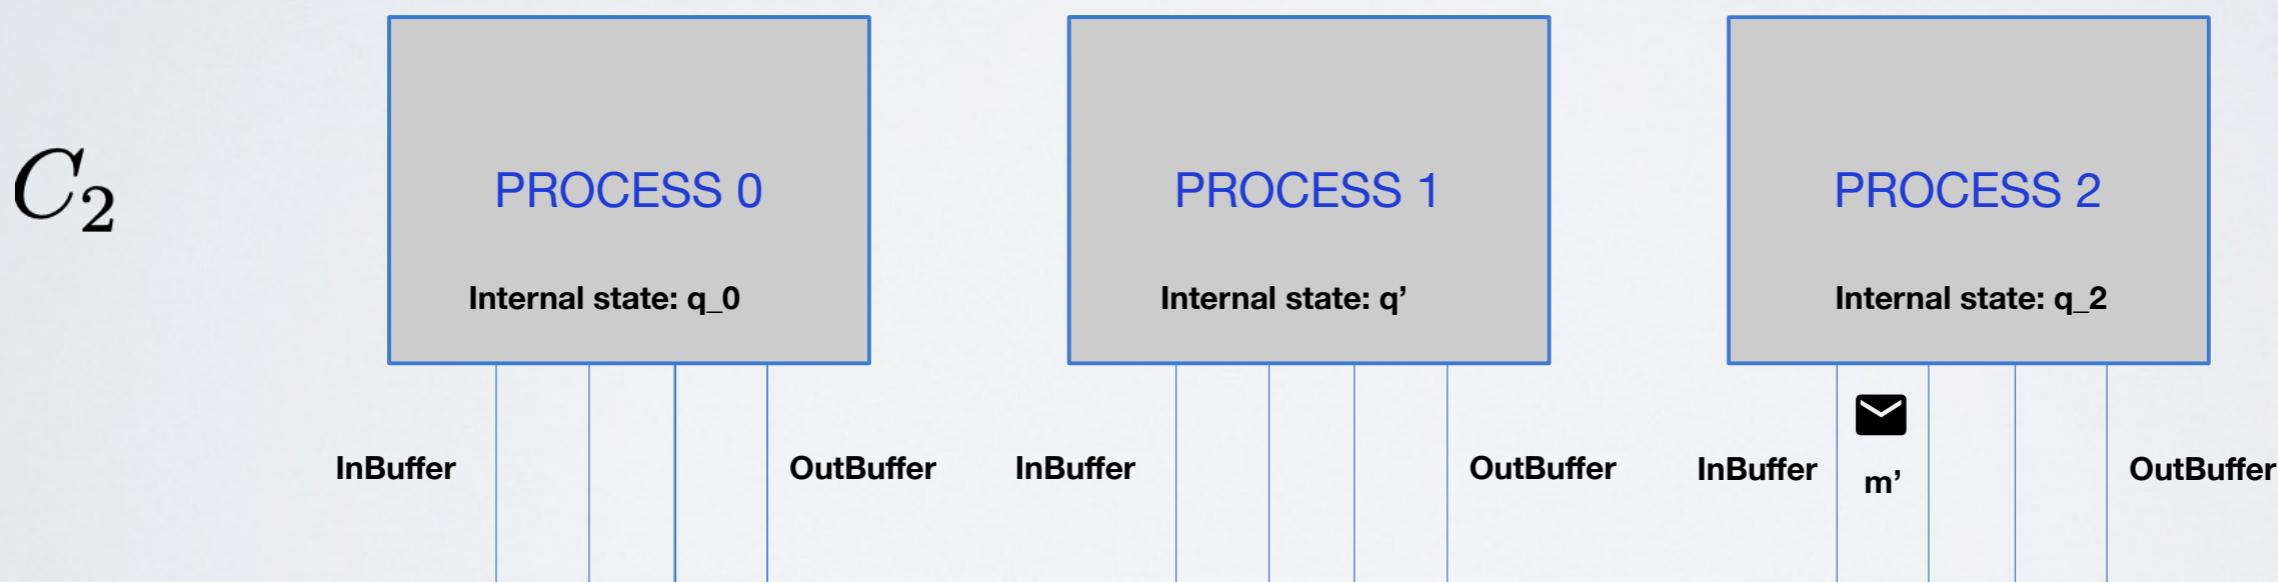
$$\mathcal{E} : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$$



EXECUTION

- An execution is infinite sequence that alternates configurations and events: $(C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$ such that each event $e_{\{t\}}$ is enabled in configuration C_t and C_t is obtained by applying $e_{\{t-1\}}$ to $C_{\{t-1\}}$.

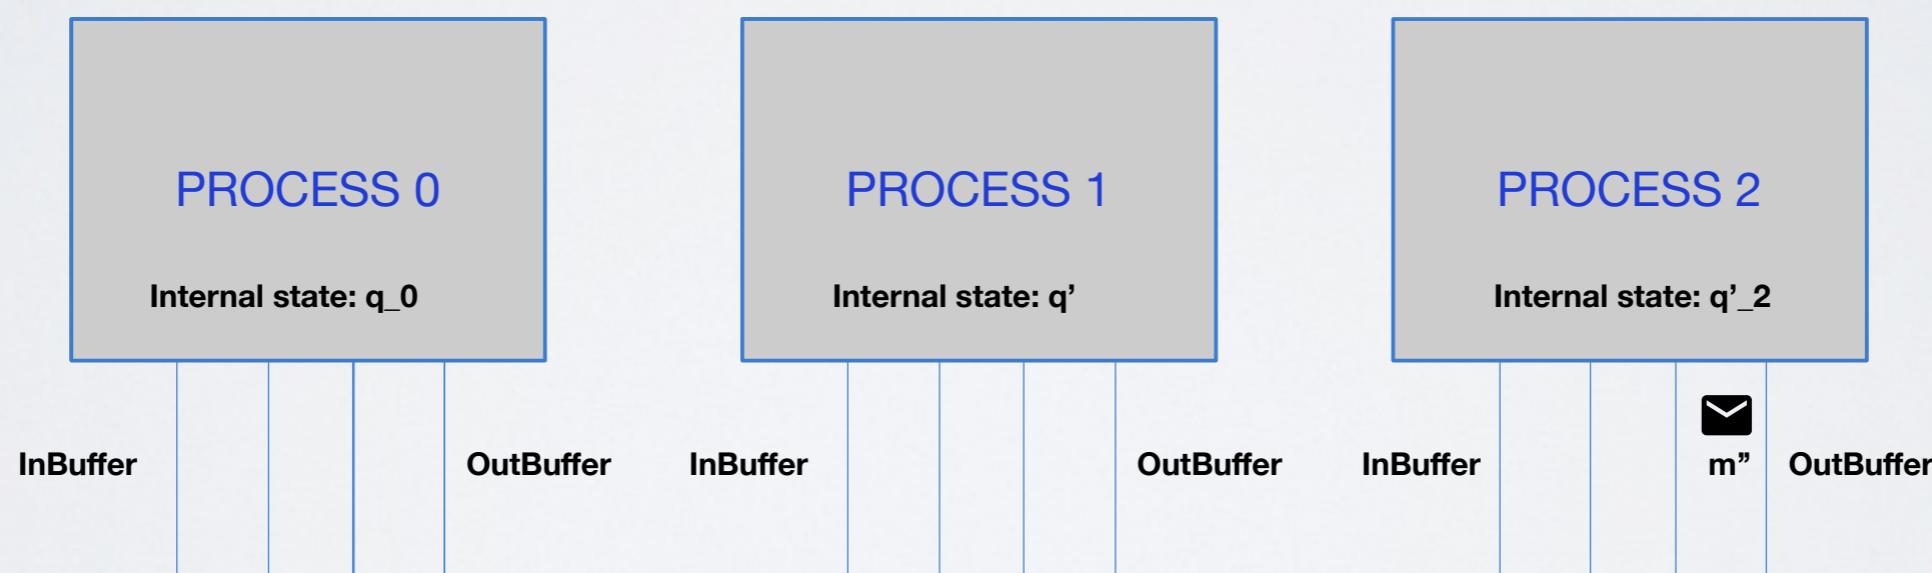
$$\mathcal{E} : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$$



EXECUTION

- An execution is infinite sequence that alternates configurations and events: $(C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$ such that each event $e_{\{t\}}$ is enabled in configuration C_t and C_t is obtained by applying $e_{\{t-1\}}$ to $C_{\{t-1\}}$.

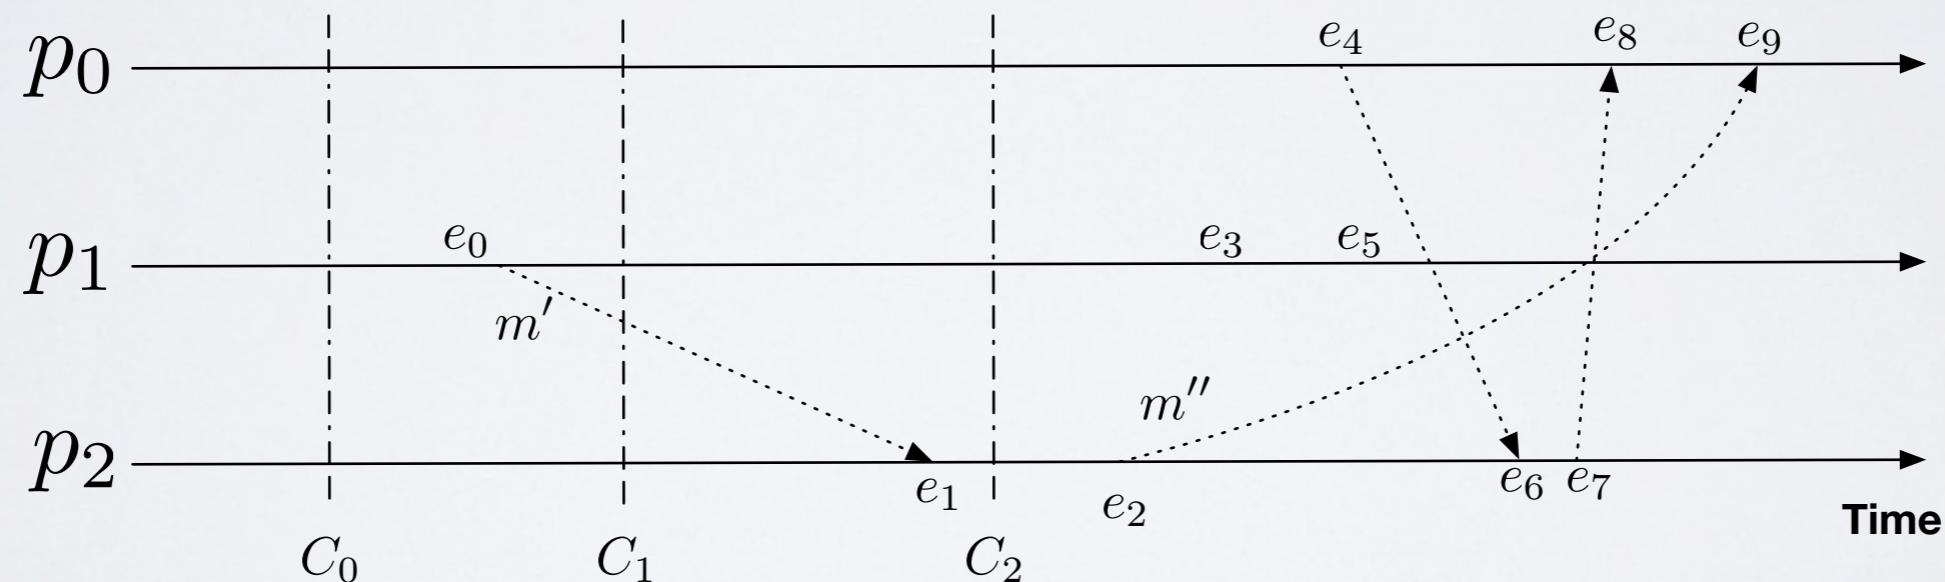
$$\mathcal{E} : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$$



EXECUTION

- An execution is infinite sequence that alternates configurations and events: $(C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$ such that each event $e_{\{t\}}$ is enabled in configuration C_t and C_t is obtained by applying $e_{\{t-1\}}$ to $C_{\{t-1\}}$.

$$\mathcal{E} : (C_0, e_0 = Exec(1), C_1, e_1 = Del(1, 2, m'), C_2, e_2 = Exec(2), C_3, e_3, C_4, \dots)$$



FAIR EXECUTION

- An execution E is fair if each process p_i executes an infinite number of local computations ($\text{Exec}(i)$ events are not finite) and each message m is eventually delivered (is not possible to stall forever a message: there must exist a $\text{Del}(m,x,y)$).
- We will always consider fair executions.
- Unfair executions break any possible non-trivial algorithm (**why?**).

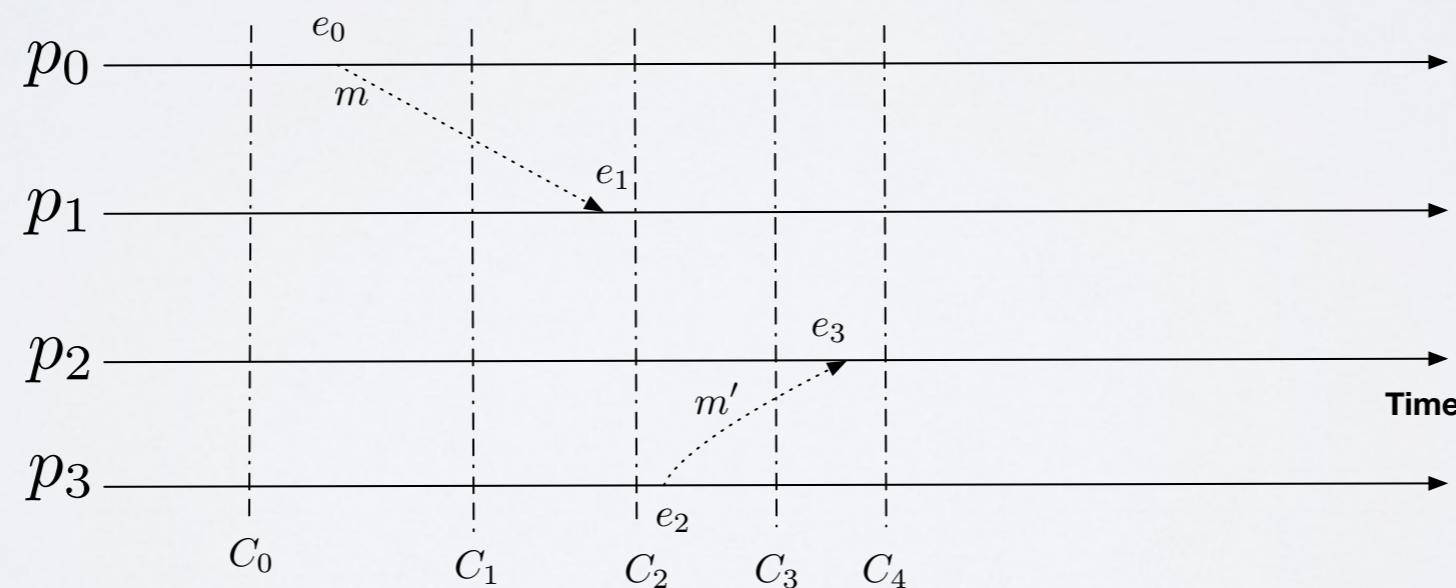
LOCAL EXECUTION (OR LOCAL VIEW)

- Given an execution E and a process p_j , we define the local execution $(E|p_j)$ of p_j as the subset of events in E that "impact" p_j .

$$\mathcal{E} = (C_0, e_0 = \text{Exec}(0), C_1, e_1 = \text{Del}(0, 1, m), C_2, e_2 = \text{Exec}(3), C_3, e_3 = \text{Del}(3, 2, m'), \dots)$$

$$\mathcal{E}|p_1 = (\text{Del}(0, 1, m), \dots)$$

$$\mathcal{E}|p_2 = (\text{Del}(3, 2, m'), \dots)$$



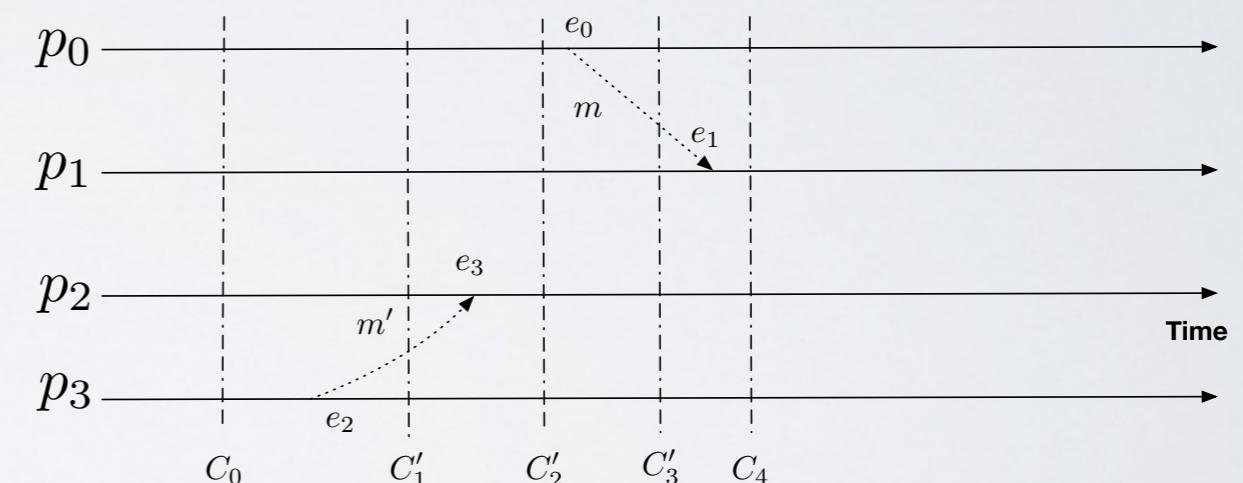
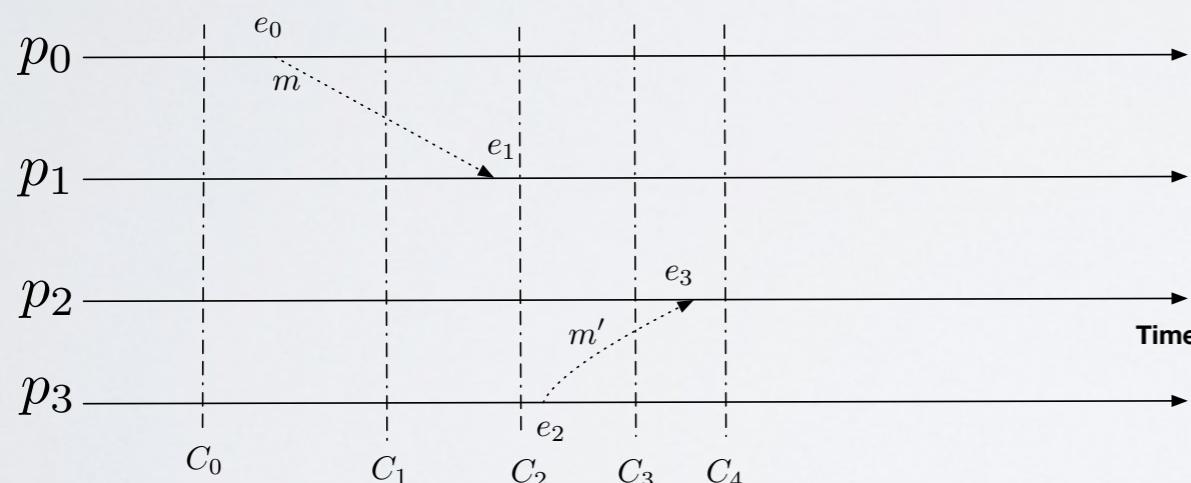
LOCAL EXECUTION

- Different executions could give the same local exec.:

$$\mathcal{E} = (C_0, e_0 = \text{Exec}(0), C_1, e_1 = \text{Del}(0, 1, m), C_2, e_2 = \text{Exec}(3), C_3, e_3 = \text{Del}(3, 2, m'), C_4, \dots)$$

$$\mathcal{E}' = (C_0, e_2 = \text{Exec}(3), C'_1, e_3 = \text{Del}(3, 2, m'), C'_2, e_0 = \text{Exec}(0), C'_3, e_1 = \text{Del}(0, 1, m), C_4, \dots)$$

$$\mathcal{E}|_{p_1} = \mathcal{E}'|_{p_1} \text{ and } \mathcal{E}|_{p_2} = \mathcal{E}'|_{p_2}$$



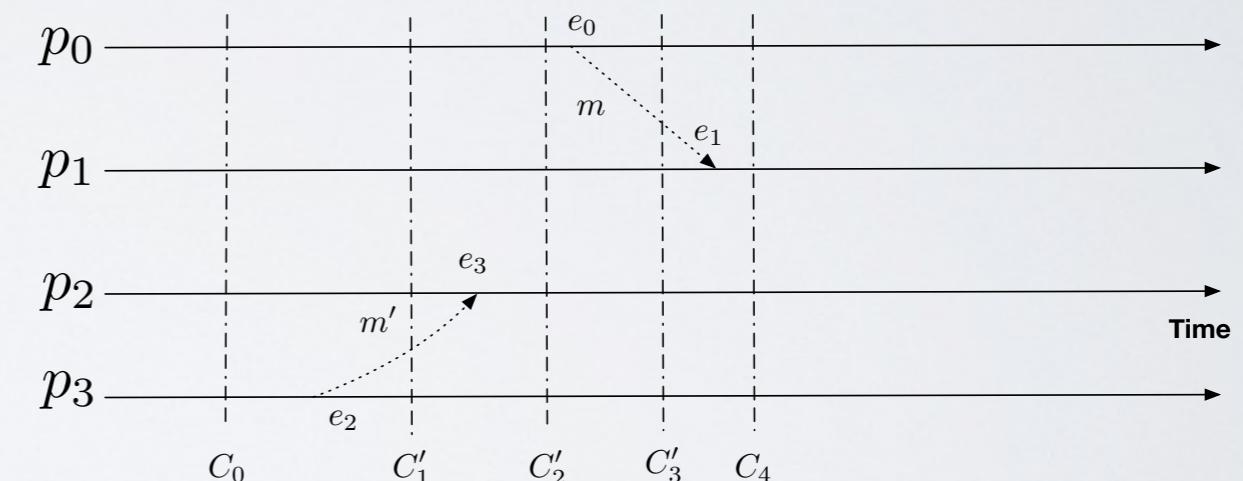
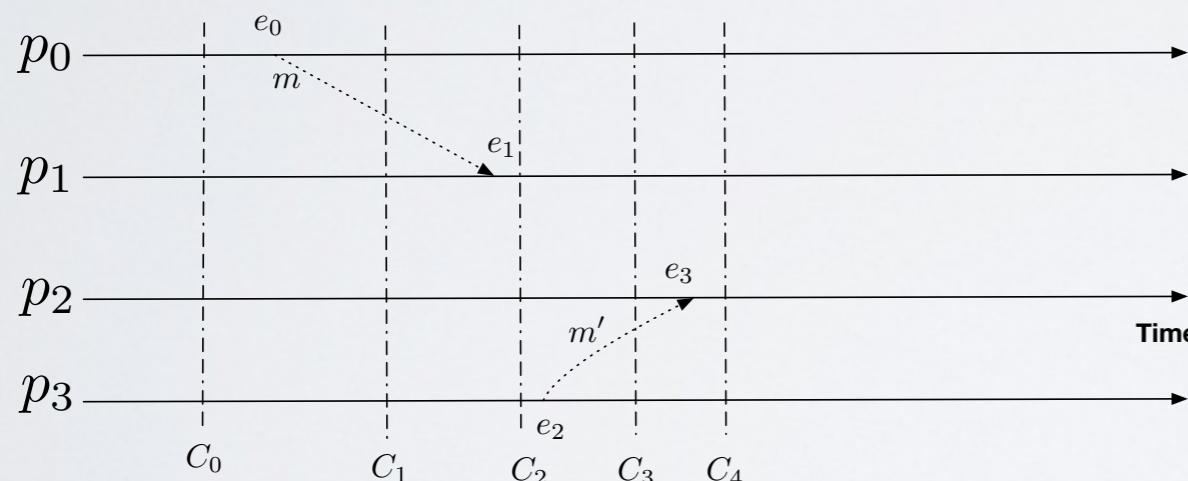
LOCAL EXECUTION

- Different executions could give the same local exec.:

$$\mathcal{E} = (C_0, e_0 = \text{Exec}(0), C_1, e_1 = \text{Del}(0, 1, m), C_2, e_2 = \text{Exec}(3), C_3, e_3 = \text{Del}(3, 2, m'), C_4, \dots)$$

$$\mathcal{E}' = (C_0, e_2 = \text{Exec}(3), C'_1, e_3 = \text{Del}(3, 2, m'), C'_2, e_0 = \text{Exec}(0), C'_3, e_1 = \text{Del}(0, 1, m), C_4, \dots)$$

$$\mathcal{E}|_{p_1} = \mathcal{E}'|_{p_1} \text{ and } \mathcal{E}|_{p_2} = \mathcal{E}'|_{p_2}$$



- We say that p_1 (and p_2) cannot distinguish \mathcal{E} from \mathcal{E}' .

INDISTINGUISHABILITY

$\mathcal{E} = (C_0, e_0 = Exec(0), C_1, e_1 = Del(0, 1, m), C_2, e_2 = Exec(3), C_3, e_3 = Del(3, 2, m'), C_4, \dots)$

$\mathcal{E}' = (C_0, e_2 = Exec(3), C'_1, e_3 = Del(3, 2, m'), C''_2, e_0 = Exec(0), C'_3, e_1 = Del(0, 1, m), C_4, \dots)$

- It is not hard to see that: $\forall p_j \in \Pi, \mathcal{E}|p_j = \mathcal{E}'|p_j$
- In this case we say that \mathcal{E} and \mathcal{E}' are indistinguishable.

INDISTINGUISHABILITY

$\mathcal{E} = (C_0, e_0 = Exec(0), C_1, e_1 = Del(0, 1, m), C_2, e_2 = Exec(3), C_3, e_3 = Del(3, 2, m'), C_4, \dots)$

$\mathcal{E}' = (C_0, e_2 = Exec(3), C'_1, e_3 = Del(3, 2, m'), C''_2, e_0 = Exec(0), C'_3, e_1 = Del(0, 1, m), C_4, \dots)$

- It is not hard to see that: $\forall p_j \in \Pi, \mathcal{E}|_{p_j} = \mathcal{E}'|_{p_j}$
- In this case we say that \mathcal{E} and \mathcal{E}' are indistinguishable.
- **Theorem:** In the asynch. model there is no distributed algorithm capable of reconstructing the system execution.

SYNCHRONOUS VS ASYNCHRONOUS

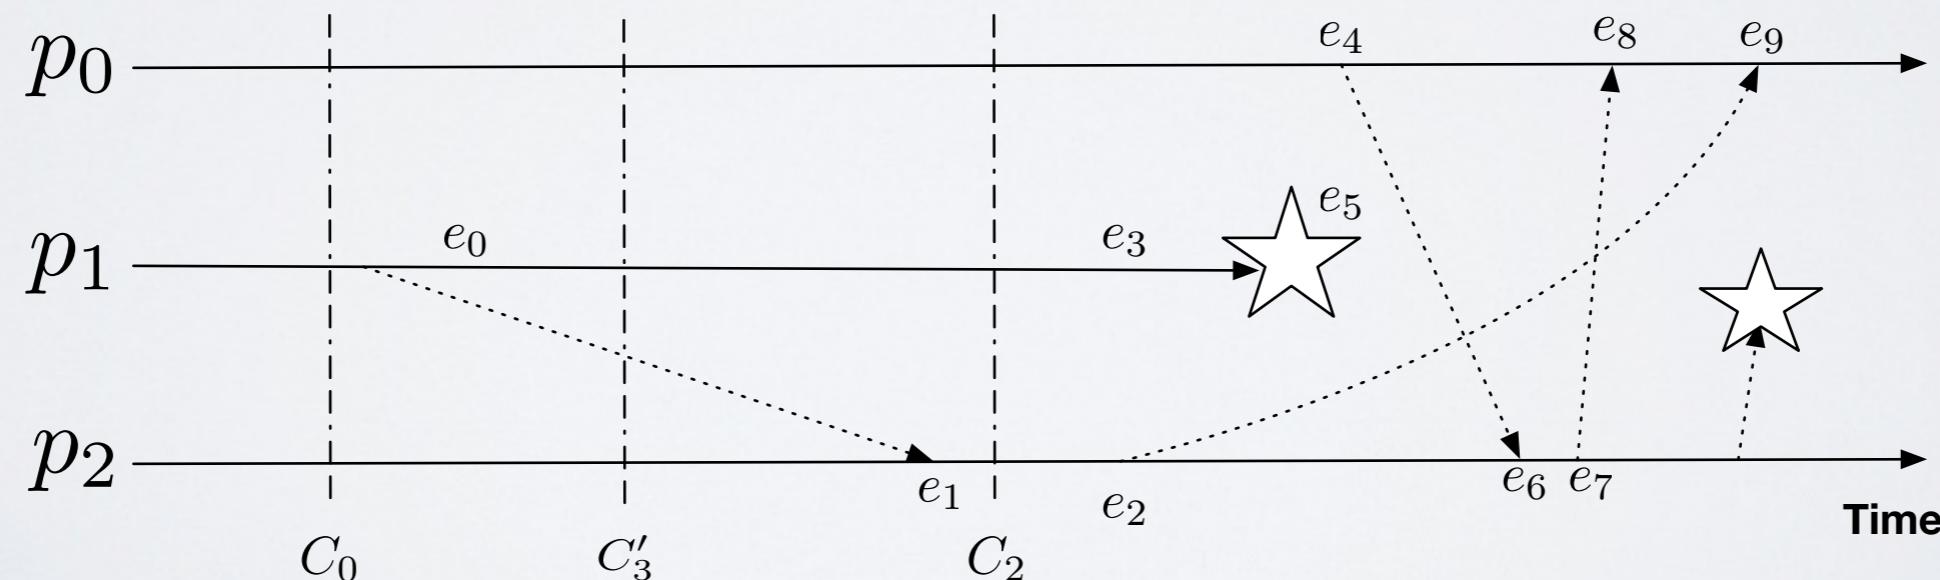
SYNCHRONOUS

- In this course we will see three definitions of synchrony:
 - Asynchronous system
 - Eventually-synchronous (in module M2)
 - Synchronous system
-
- A System is synchronous if there is fixed bound on the delay of messages.

MODEL: FAILURES

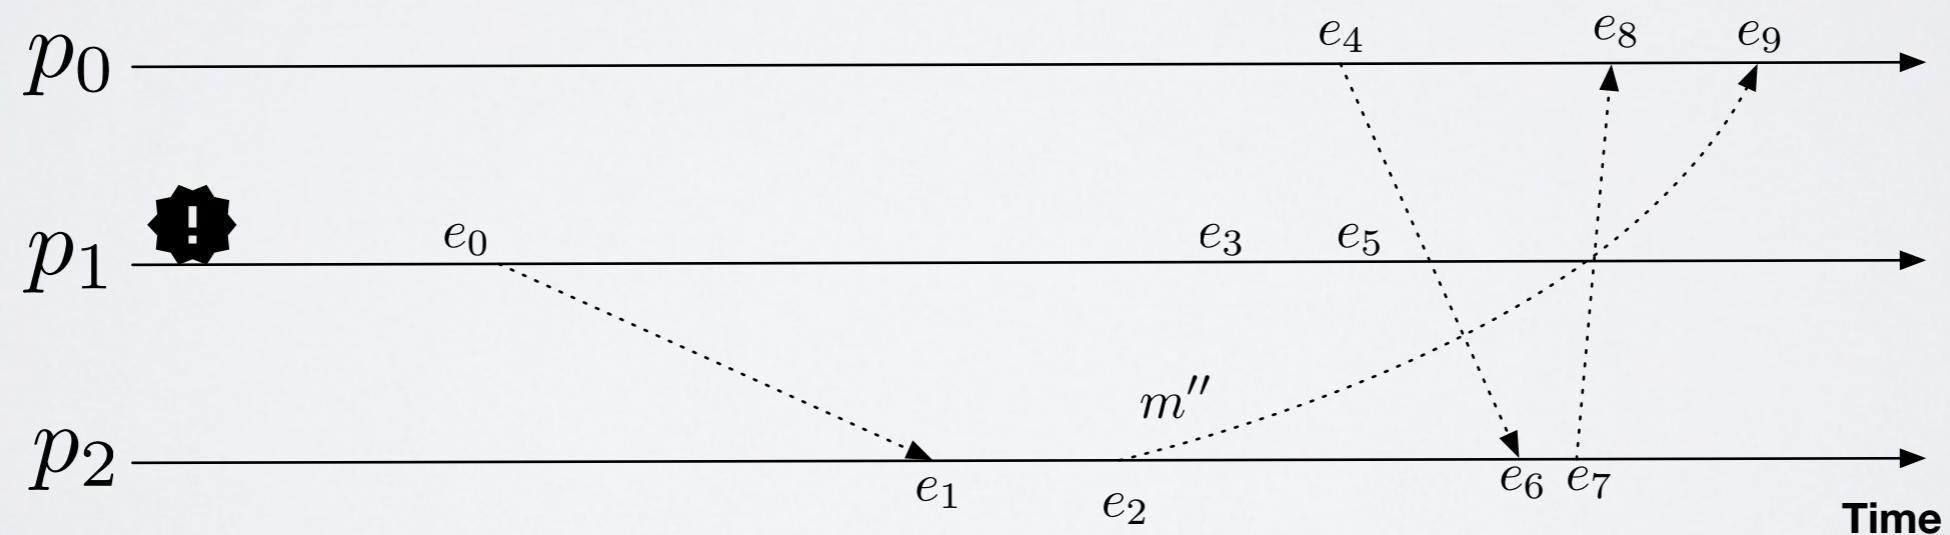
MODELS FOR PROCESS FAILURES

- Two main models:
 - Crash-Stop Failure
 - Byzantine Failure
- Crash(p_j): after this event process p_j does not execute any local computation step (no $\text{Exec}(j)$).



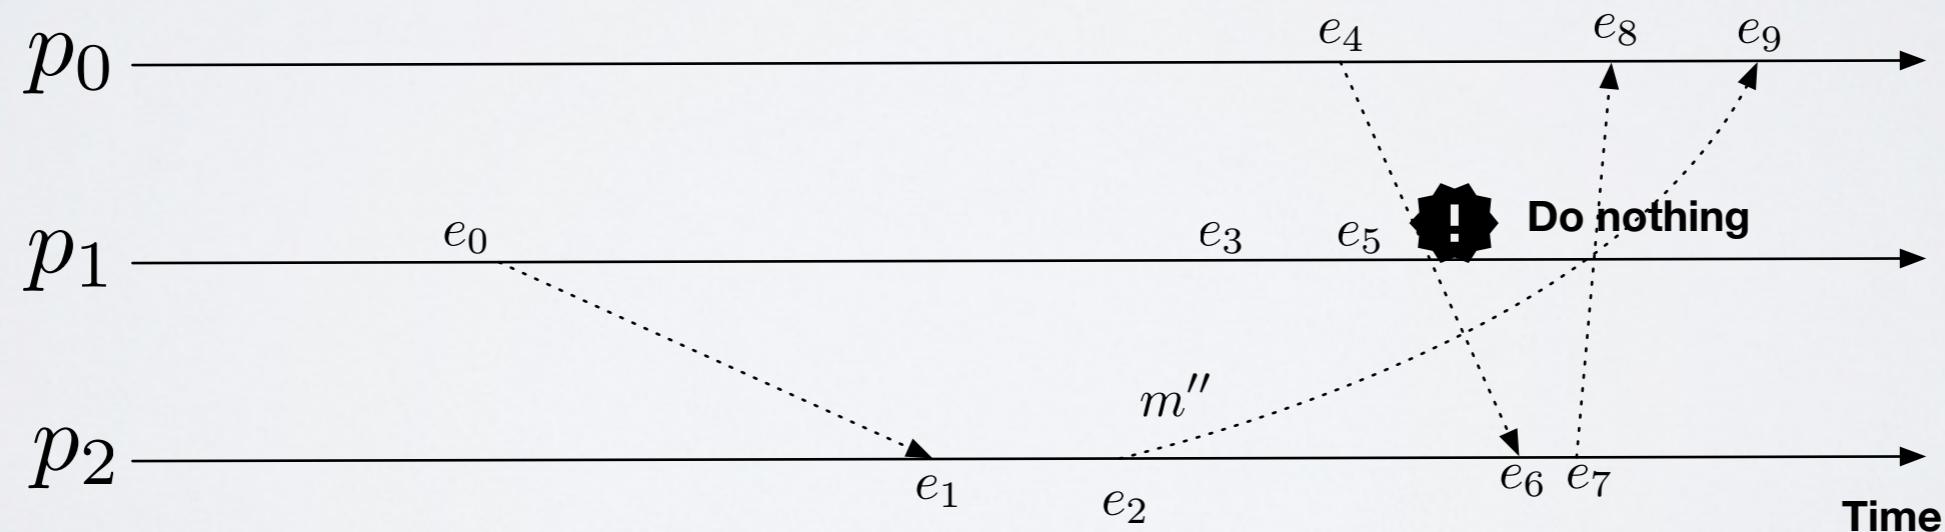
MODELS FOR PROCESS FAILURES

- Two main models:
 - Crash-Stop Failure
 - Byzantine Failure
- $\text{Byz}(p_j)$: after this event process p_j behaves in an arbitrary way (when we have $\text{Exec}(j)$ we do not follow anymore the automaton of p_j).



MODELS FOR PROCESS FAILURES

- Two main models:
 - Crash-Stop Failure
 - Byzantine Failure
- $\text{Byz}(p_j)$: after this event process p_j behaves in an arbitrary way (when we have $\text{Exec}(j)$ we do not follow anymore the automaton of p_j).



MODELS FOR PROCESS FAILURES

- Two main models:

- Crash-Stop Failure
- Byzantine Failure



Crash-stop Failure

Byzantine Failure

MODELS FOR PROCESS FAILURES

- Two main models:
 - Crash-Stop Failure
 - Byzantine Failure
- We say that a process is **correct** if it **does not experience a failure** (but we do not know who is correct).
- We indicate with **f** the **maximum number** of processes that can experience a failure event in an execution.
 - Usually we bound f. Example: f cannot be more than $n-1$.
 - We never do assumptions on when a failure event happens.

HOMEWORKS

- Model the loss of a message as an event.

ABSTRACTIONS, PSEUDO-CODE CONVENTIONS AND OUR FIRST ALGORITHM

ABSTRACTION AND PROPERTIES

THE ROAD TO BUILD A DISTRIBUTED ABSTRACTION

Step 1: define a system model

Step 2: formalize a problem/(object) with an abstraction

understand how to design an abstract formal object that captures the problem that you want to solve:

- Specify unambiguously the properties of your abstraction

Step 3: implements the abstraction with a distributed protocol

Step 4: prove that the protocol implements the abstraction

ABSTRACTING A REAL COMMUNICATION LINK

- Informal description: We have **a link** that loses messages with a certain probability pr , the channel can duplicate a message a finite number of times, and it does not create messages from thin air.
- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q .
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p .

ABSTRACTING A REAL COMMUNICATION LINK

- Informal description: We have a link that **loses messages with a certain probability pr**, the channel can duplicate a message a finite number of times, and it does not create messages from thin air.
- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q.
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p.
- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.

ABSTRACTING A REAL COMMUNICATION LINK

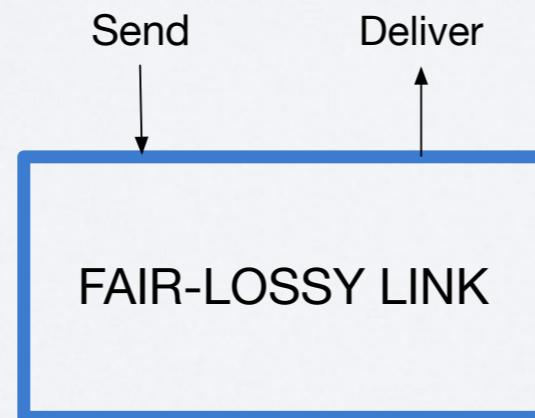
- Informal description: We have a link that loses messages with a certain probability p_r ,
the channel can duplicate a message a finite number of times, and it does not create messages from thin air.
- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q .
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p .
- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q , then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q , then q cannot deliver m an infinite number of times.

ABSTRACTING A REAL COMMUNICATION LINK

- Informal description: We have a link that loses messages with a certain probability p_r , the channel can duplicate a message a finite number of times, and **it does not create messages from thin air.**
- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q .
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p .
- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q , then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q , then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p , then m was sent by p to q .

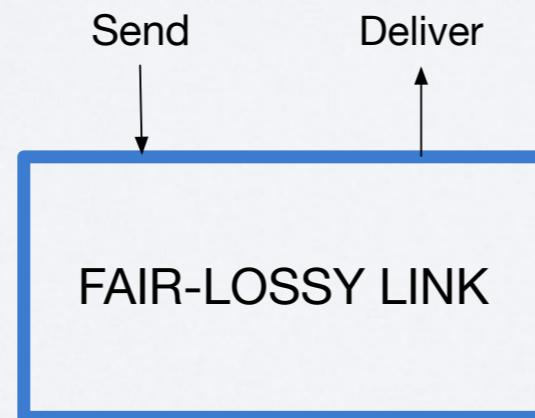
FAIR-LOSSY LINK

- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q.
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p.
- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.



FAIR-LOSSY LINK

- A link exposes two events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q.
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p.
- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.



LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- All properties fall in two classes: **Safety** or **Liveness**.
 - A **Safety property** is property that **if violated at a time t, it can never be satisfied after t**.
 - Formally, if a safety property is violated in execution E, there is a prefix E' of E such that any extension of E' also violates the property.

LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- All properties fall in two classes: **Safety** or **Liveness**.
 - A Safety property is property that if violated at a time t, it can never be satisfied after t.
 - Formally, if a safety property is violated in execution E, there is a prefix E' of E such that any extension of E' also violates the property.
 - Informally, **a safety property states that something bad cannot happen**.
 - Among FL1, FL2 and FL3 there is at least one safety property. Which one?

LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - **FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.**
- **Safety:**
 - Once a process q delivers a message m that was not sent by anyone FL3 is violated forever.
 - $E = (e_1, e_2, e_3, \textcolor{red}{e}_4, e_5, \dots)$
 - $E' = (e_1, e_2, e_3, \textcolor{red}{e}_4)$

LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- All properties fall in two classes: **Safety** or **Liveness**.
 - A Liveness property is property that cannot be violated in finite executions:
 - Formally, given any finite execution E that does not satisfy a liveness property there is an extension of E that satisfy it.
 - Informally, a liveness property specify that something good will happen.
 - Among FL1, FL2 and FL3 there is at least one liveness property. Which one?

LET US FOCUS ON PROPERTIES

- Properties:

- **FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.**
- FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.
- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

- **Liveness.**

- If p sent a message m, and q has not delivered it at time t, there is always hope that it can deliver it at time $t+\Delta$.
 - $E = (e_1, e_2, e_3, e_4, e_5)$
 - $E' = (e_1, e_2, e_3, e_4, e_5, \text{del}(p, q, m))$

LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - **FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.**
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
 - ?

LET US FOCUS ON PROPERTIES

- Properties:
 - FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q, then q delivers m an infinite number of times.
 - **FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, then q cannot deliver m an infinite number of times.**
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- **Liveness.**
 - It cannot be violated in finite execution.
 - $E = (\text{del}(p, q, m), \text{del}(p, q, m), \text{del}(p, q, m), \dots, \text{del}(p, q, m))$
 - $E' = (\text{del}(p, q, m), \text{del}(p, q, m), \text{del}(p, q, m), \dots, \text{del}(p, q, m), e_1, e_2, e_3, \dots)$
 - Such that no e_j is $\text{del}(p, q, m)$.

LET US FOCUS ON PROPERTIES

- Property:
 - **FL2: (At-most7 duplication)** If a correct process p sends m a finite number of times to q, then q cannot deliver m **more than 7 times**.
- **Liveness?**

LET US FOCUS ON PROPERTIES

- Property:
 - **FL2: (At-most7 duplication)** If a correct process p sends m a finite number of times to q, then q cannot deliver m **more than 7 times**.
- **Liveness?**
 - $E=(\text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m), \text{del}(p,q,m))$
 - E cannot be extended, it always violates At-most7 duplication.
 - At-most7 duplication is a **safety** property.

OTHER PROPERTIES

- Mutual Exclusion: if a process p is granted a resource r at time t, then no other process q is granted r at t.
- No-deadlock: if r is not already granted, eventually someone get a grant on resource r.
- No-Starvation: if a process p request a grant on resource r, it will eventually get it.

BADLY WRITTEN PROPERTIES

- If process p sends a message m to q, then q will eventually deliver it and this deliver is unique.
 - It is mixing two aspects:
 - A liveness: q will eventually deliver it
 - A safety: the deliver is unique

BADLY WRITTEN PROPERTIES

- If process p sends a message m to q, then q will eventually deliver it and this deliver is unique.
 - It is mixing two aspects:
 - A liveness: q will eventually deliver it
 - A safety: the deliver is unique
- It should be decomposed in two properties:
 - If p sends a message m to q, then q eventually delivers it.
 - If p sends a message m to q, then m is delivered at most once.

HOMEWORKS

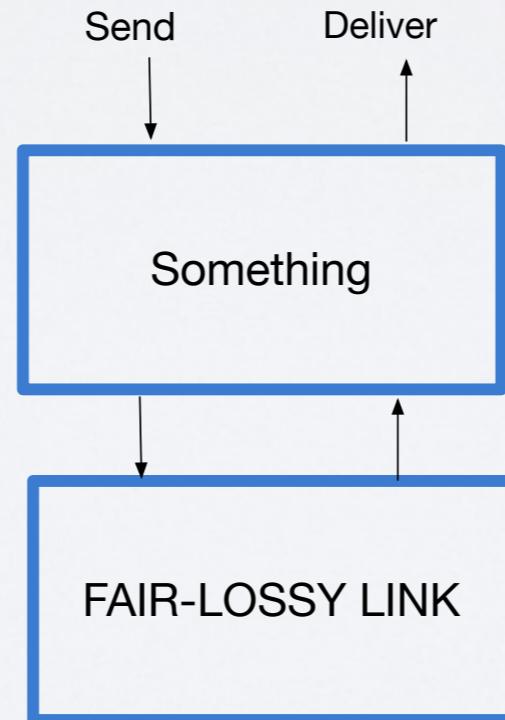
- Safety or liveness?
 - Fairness of an execution.
 - (At most- f failures) An execution E can contain at most f failures.
 - (Eventual recovery) If a process p experience a $\text{crash}(p)$, then it eventually experiences also a $\text{recovery}(p)$ (this event starts again process p).
 - (Time-bounded recovery) If a process p experiences a $\text{recovery}(p)$, then this recovery happens at most k events after its failure.
- Write three new safety properties and three new liveness properties.
- Does Time-bounded recovery implies Eventual recovery?
- How would you improve the writing of this property:
 - If a process p outputs a set of integer O_p , then any other process q outputs also outputs a set O_q . For any two processes z and t we have either O_z is a subset of O_t or that O_t is a subset of O_z .

ALGORITHMS: COMMUNICATION LINKS

LET US GO BACK ON OUR FL-LINK

- Properties:

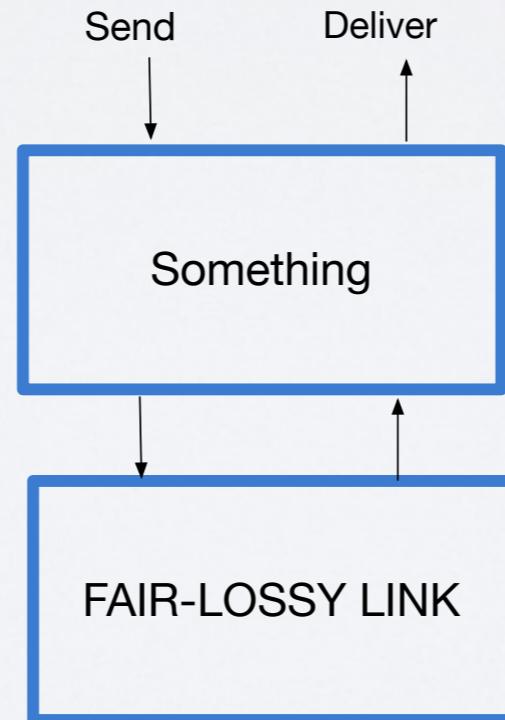
- FL1: (Fair-loss) If a correct process p sends infinitely often m to a process q , then q delivers m an infinite number of times.
- FL2: (Finite duplication) If a correct process p sends m a finite number of times to q , then q cannot deliver m an infinite number of times.
- FL3: (No creation) If some process q delivers a message m with sender p , then m was sent by p to q .



LET US GO BACK ON OUR FL-LINK

- Properties:

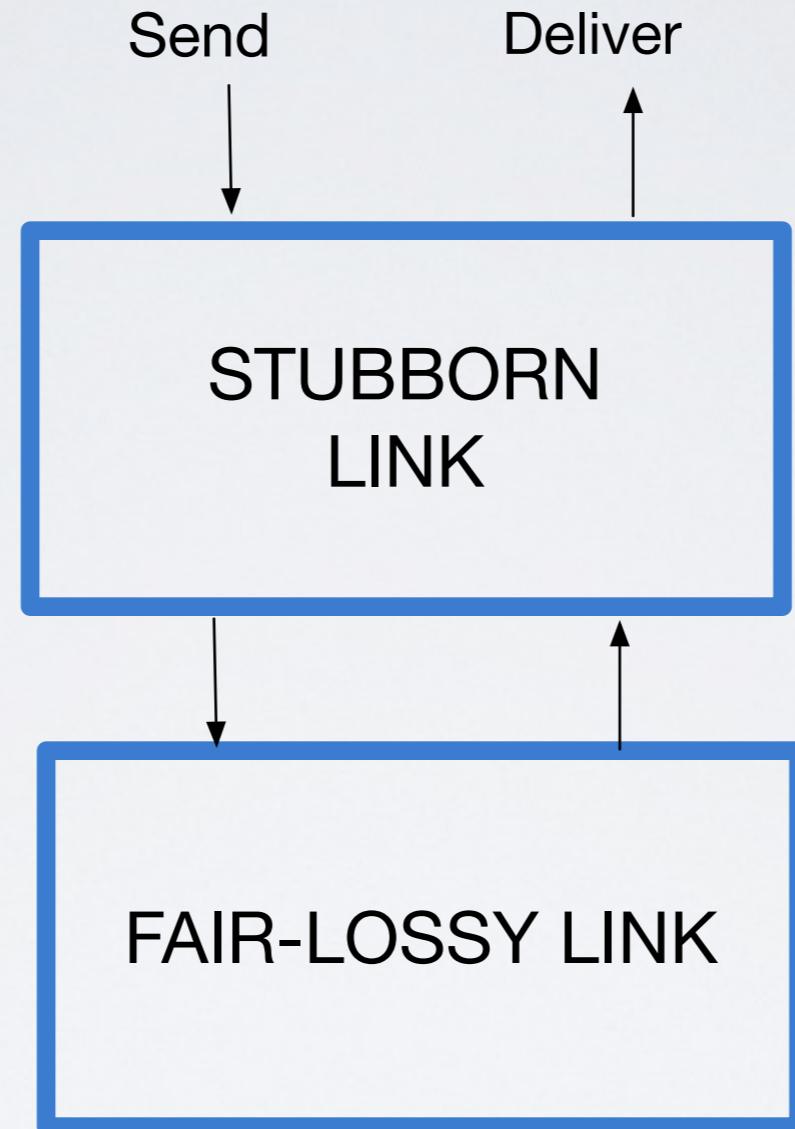
- FL1: (Fair-loss) If a correct process **p sends infinitely often m** to a process q, then q delivers m an infinite number of times.
- FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, **then q cannot deliver m an infinite number of times.**
- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.



LET US GO BACK ON OUR FL-LINK

- Let us fix:
 - FL1: (Fair-loss) If a correct process **p sends infinitely often m** to a process q, then q delivers m an infinite number of times.
 - FL2: (Finite duplication) If a correct process p sends m a finite number of times to q, **then q cannot deliver m an infinite number of times.**
- In:
 - SL1: (Stubborn-delivery) If a correct process p sends m to q, then q delivers m an infinite number of times.
- While we keep:
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

OUR FIRST ALGORITHM



COMPOSITION MODEL

- The protocols we will consider in this course are presented in pseudo-code
- The pseudo code reflects a reactive computing model where
 - components of the same process communicate by exchanging events
 - the algorithm is described as a set of event handlers
 - handlers react to incoming events and possibly trigger new events.
 - Handlers are atomic. They can only be interrupted in case of a crash.

STUBBORN

- Events:
 - Requests: $\langle \text{Send} \mid q, m \rangle$ send a message m to process q.
 - Indication: $\langle \text{Deliver}, p, m \rangle$ delivers message m from process p.
- Properties:
 - SL1: (Stubborn-delivery) If a correct process p sends m to q, then q delivers m an infinite number of times.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

STUBBORN

- Properties:

- SL1: (Stubborn-delivery) If a correct process p sends m to q, then q delivers m an infinite number of times.
- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

upon event $\langle sl, Init \rangle$ **do**

$sent := \emptyset;$

$starttimer(\Delta);$

upon event $\langle Timeout \rangle$ **do**

forall $(q, m) \in sent$ **do**

trigger $\langle fll, Send \mid q, m \rangle;$

$starttimer(\Delta);$

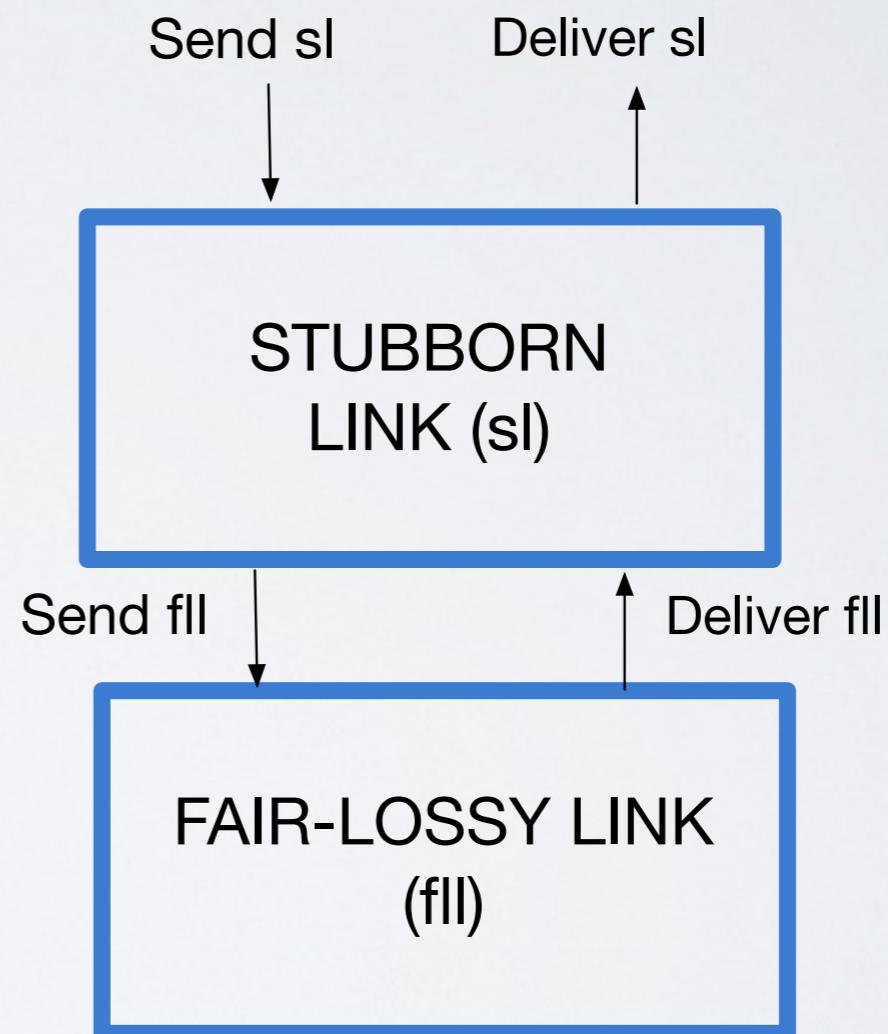
upon event $\langle sl, Send \mid q, m \rangle$ **do**

trigger $\langle fll, Send \mid q, m \rangle;$

$sent := sent \cup \{(q, m)\};$

upon event $\langle fll, Deliver \mid p, m \rangle$ **do**

trigger $\langle sl, Deliver \mid p, m \rangle;$



STUBBORN FORMAL PROOF

- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- Proof. By contradiction.
- Suppose process q executing our algorithm receives message m that was not sent by p.
- Fact 1: If q delivers a message, then it delivers here:

```
upon event ⟨ fl1, Deliver | p, m ⟩ do  
trigger ⟨ sl, Deliver | p, m ⟩;
```

STUBBORN FORMAL PROOF

- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- Proof. By contradiction.
- Suppose process q executing our algorithm receives message m that was not sent by p.
- Fact 1: If q delivers a message, then it delivers here:

```
upon event ⟨ fll, Deliver | p, m ⟩ do  
    trigger ⟨ sl, Deliver | p, m ⟩;
```

- Fact 1 implies that *fll*, is delivering a message that was not sent by p. This implies that *fll* is not fair-lossy. This contradicts our hypothesis: *fll* is fair-lossy.

STUBBORN FORMAL PROOF

- SL1: (Stubborn-delivery) If a correct process p sends m to q, then q delivers m an infinite number of times.
- Proof. By contradiction. Suppose q delivers m a finite number of times.
- Fact 1: p sends m on fll an infinite number of times:

```
upon event < Timeout > do
  forall (q, m) ∈ sent do
    trigger < fll, Send | q, m >;
    starttimer(Δ);
```

```
upon event < sl, Send | q, m > do
  trigger < fll, Send | q, m >;
  sent := sent ∪ {(q, m)};
```

- Fact 2: if q “stubborn delivers” m a finite number of times, then fll delivered m a finite number of times.

```
upon event < fll, Deliver | p, m > do
  trigger < sl, Deliver | p, m >;
```

IMPROVING STUBBORN

- Properties:
 - SL1: (Stubborn-delivery) If a correct process p sends m to q, then q **delivers m an infinite number of times.**
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- In:
 - PL1: (Reliable delivery) If a correct process p sends m to q, then q eventually delivers m.
 - PL2: (No duplication) A message is delivered at most once.
 - FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

PERFECT P2P LINK

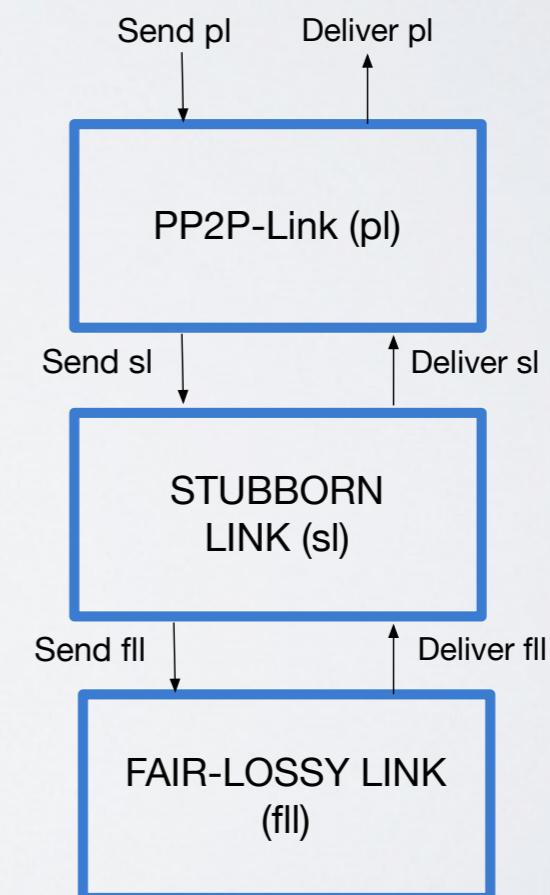
- Properties:

- PL1: (Reliable delivery) If a correct process p sends m to q, then q eventually delivers m.
- PL2: (No duplication) A message is delivered at most once.
- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.

upon event $\langle pl, Init \rangle$ **do**
delivered := \emptyset ;

upon event $\langle pl, Send \mid q, m \rangle$ **do**
trigger $\langle sl, Send \mid q, m \rangle$;

upon event $\langle sl, Deliver \mid p, m \rangle$ **do**
if $m \notin delived$ **then**
delivered := *delivered* $\cup \{m\}$;
trigger $\langle pl, Deliver \mid p, m \rangle$;



PERFECT P2P LINK - FORMAL PROOF

- FL3: (No creation) If some process q delivers a message m with sender p, then m was sent by p to q.
- Proof. By contradiction.
- Suppose process q executing our algorithm receives message m that was not sent by p.
- Fact 1: If q delivers a message, then it delivers here:

```
upon event < sl, Deliver | p, m > do
  if m  $\notin$  delivered then
    delivered := delivered  $\cup$  {m};
    trigger < pl, Deliver | p, m >;
```

- This implies that sl delivered a message that was not created. Violates the hypothesis that sl is a stubborn.

PERFECT P2P LINK - FORMAL PROOF

- PL2: (No duplication) A message is delivered at most once
- Fact 1: The pp2p-delivery of a message is “guarded” by an if $m \in \text{delivered}$:

```
upon event < sl, Deliver | p, m > do
    if m  $\notin$  delivered then
        delivered := delivered  $\cup$  {m};
    trigger < pl, Deliver | p, m >;
```

- Suppose m is delivered twice, at time t and t' (with $t < t'$).
- At time t the delivery handler is executed. Since the handler is atomic we have that $\text{delivered} := \text{delivered} \cup \{m\}$ is executed before t' . Therefore, at t' m is in delivered , this contradict the fact that $\text{trigger } < \text{pl, deliver} | P, m >$ is executed at (or after) time t' .

PERFECT P2P LINK - FORMAL PROOF

- PL1: (Reliable delivery) If a correct process p sends m to q, then q eventually delivers m.
- Suppose, p sends m and q does not deliver it. There could be two reasons for q to not deliver m:
 - Reason 1: m is in delivered when the delivery handler is executed:

```
upon event < sl, Deliver | p, m > do
    if m  $\notin$  delivered then
        delivered := delivered  $\cup$  {m};
    trigger < pl, Deliver | p, m >;
```

- If m is in delivered then, q eventually will execute **trigger** <*pl, Deliver | p, m*>. This contradicts the fact that q does not deliver m.

PERFECT P2P LINK - FORMAL PROOF

- PL1: (Reliable delivery) If a correct process p sends m to q, then q eventually delivers m.
- Suppose, p sends m and q does not deliver it. There could be two reasons for q to not deliver m:
 - Reason 2: the delivery handler is never triggered with $\langle p, m \rangle$

```
upon event < sl, Deliver | p, m > do
    if m  $\notin$  delivered then
        delivered := delivered  $\cup$  {m};
    trigger < pl, Deliver | p, m >;
```

- This means that sl is not stubborn. Violating our hypothesis.

HOMEWORKS 1

- Show that our stubborn algorithm does not work if we change the first property in:
 - SL1: (Stubborn-delivery) If a process p sends m to q, then q delivers m an infinite number of times. {Pay attention something is missing}.
- Show that there is no algorithm that satisfies the property above using a fair-lossy link.
- * Suppose you can use only a fixed amount k of memory (see it as having the possibility to store at most k different messages), could you implement a Perfect p2p link on a fair-lossy?

HOMEWORKS 1

- Show that our stubborn algorithm does not work if we change the first property in:
 - SL1: (Stubborn-delivery) If a process p sends m to q, then q delivers m an infinite number of times. {Pay attention something is missing}.
- The property does not specify the correctness of p. This means that even if p crashes, then q has to deliver m (at least once).
- If p crashes, then it cannot send an infinite number of messages on fil.
- If p sends a finite number of messages on fil, then it could be that no message is delivered to q.

HOMEWORKS 1

- Show that our stubborn algorithm does not work if we change the first property in:
 - SL1: (Stubborn-delivery) If a process p sends m to q, then q delivers m an infinite number of times. {Pay attention something is missing}.
- Show that there is no algorithm that satisfies the property above using a fair-lossy link.
- The same argument shows that no algorithm could satisfy the property above on a fair-lossy link.

REASONING TIME....

- Recall the following property of a perfect p2p link:
 - PL1: (Reliable delivery) If a correct process p sends m to q , then q eventually delivers m .
- Assume a system with crash failure, is there something wrong?

REASONING TIME....

- PL1: (Reliable delivery) If a correct process p sends m to q, then q eventually delivers m.
- Suppose, p sends m and q does not deliver it. There could be two reasons for q to not deliver m:
- Reason 1: m is in delivered when the delivery handler is executed:

```
upon event < sl, Deliver | p, m > do
    if m  $\notin$  delivered then
        delivered := delivered  $\cup$  {m};
    trigger < pl, Deliver | p, m >;
```

- If m is in delivered then, q eventually will execute trigger <pl, Deliver| p, m>. This contradicts the fact that q does not deliver m.

REASONING TIME....

- Recall the following property of a perfect p2p link:
 - PL1: (Reliable delivery) If a correct process p sends m **to q and q is correct**, then q eventually delivers m .

HOMEWORKS 2

- One of your friends, Caio, is really excited about its new discovery. Caio claims to have implemented a stubborn over a fair-lossy without the need of any internal timer or infinite loop. It shows you this algorithm:

```
upon event INIT  
    set =  $\emptyset$   
  
upon event SL, SEND( $< m, q >$ )  
    set = set  $\cup \{< m, q >\}$   
    for all  $< x, y > \in set$  do  
        invoke fl Send( $< x, y >$ )  
  
upon event FL, DELIVER( $< m, q >$ )  
    invoke sl Deliver( $< m, q >$ )
```

What is wrong with his algorithm? In which executions does it work? Could you find an argument that convinces Caio on the impossibility of implementing a stubborn (and thus a perfect link) over a fair-lossy without using a timer or an infinite loop?
(note that an infinite loop is actually a timer)

HOMEWORKS 2

What is wrong with his algorithm? In which executions does it work? Could you find an argument that convinces Caio on the impossibility of implementing a stubborn (and thus a perfect link) over a fair-lossy without using a timer or an infinite loop?

(note that an infinite loop is actually a timer)

```
upon event INIT
    set = ∅

upon event SL, SEND(< m, q >)
    set = set ∪ {< m, q >}
    for all < x, y > ∈ set do
        invoke fl Send(< x, y >)

upon event FL, DELIVER(< m, q >)
    invoke sl Deliver(< m, q >)
```

HOMEWORKS 2

What is wrong with his algorithm? In which executions does it work? Could you find an argument that convinces Caio on the impossibility of implementing a stubborn (and thus a perfect link) over a fair-lossy without using a timer or an infinite loop?

(note that an infinite loop is actually a timer)

```
upon event INIT
    set = ∅

upon event SL, SEND(< m, q >)
    set = set ∪ {< m, q >}
    for all < x, y > ∈ set do
        invoke fl Send(< x, y >)

upon event FL, DELIVER(< m, q >)
    invoke sl Deliver(< m, q >)
```

The algorithm does not implement PL1: (Reliable delivery). Suppose, p stubborn sends message m, and then it stops sending messages. Then the fair-lossy send is triggered only once. This implies that q could not deliver m.

HOMEWORKS 2

What is wrong with his algorithm? **In which executions does it work?** Could you find an argument that convinces Caio on the impossibility of implementing a stubborn (and thus a perfect link) over a fair-lossy without using a timer or an infinite loop?

(note that an infinite loop is actually a timer)

```
upon event INIT
    set = ∅

upon event SL, SEND(< m, q >)
    set = set ∪ {< m, q >}
    for all < x, y > ∈ set do
        invoke fl Send(< x, y >)

upon event FL, DELIVER(< m, q >)
    invoke sl Deliver(< m, q >)
```

HOMEWORKS 2

What is wrong with his algorithm? **In which executions does it work?** Could you find an argument that convinces Caio on the impossibility of implementing a stubborn (and thus a perfect link) over a fair-lossy without using a timer or an infinite loop?

(note that an infinite loop is actually a timer)

```
upon event INIT
    set = ∅

upon event SL, SEND(< m, q >)
    set = set ∪ {< m, q >}
    for all < x, y > ∈ set do
        invoke fl Send(< x, y >)

upon event FL, DELIVER(< m, q >)
    invoke sl Deliver(< m, q >)
```

The algorithm works only if all messages in set are fully sent an infinite number of times. Therefore, it only works in executions in which p stubborn sends an infinite number of messages.

HOMEWORKS 2

What is wrong with his algorithm? In which executions does it work? **Could you find an argument** that convinces Caio **on the impossibility of implementing a stubborn** (and thus a perfect link) **over a fair-lossy without using a timer or an infinite loop?**

(note that an infinite loop is actually a timer)

HOMEWORKS 2

What is wrong with his algorithm? In which executions does it work? **Could you find an argument** that convinces Caio **on the impossibility of implementing a stubborn** (and thus a perfect link) **over a fair-lossy without using a timer or an infinite loop?**

(note that an infinite loop is actually a timer)

Proof: Suppose you can, then in your algorithm you must have an infinite number of internal events in which you send message m over the fair-lossy link.

If you do not have an infinite loop (or a timer) and p does not receive an infinite number of input events, then it means that after a certain number of local execution steps, p will never change its internal state neither execute any actions.

Therefore, there exists at least an execution in which p executes a finite number of actions, in particular p sends m on the fair-lossy a finite number of times. This implies that is impossible for the algorithm to ensure the delivery of m by q

HOMEWORKS BIS

From the question of one of you colleagues:

- We have seen that the fairness property requires a process to be scheduled forever, so if a process crashes the execution is not fair. The reason is that a failed process will never be scheduled again for execution.

How this is **consistent** with the definition of fairness?

- The reason is that the definition of Fairness has been given for a model without failures! In case of processes failures, we have to modify the definition of fairness. That becomes:

HOMEWORKS BIS

From the question of one of your colleagues:

- We have seen that the fairness property requires a process to be scheduled forever, so if a process crashes the execution is not fair. The reason is that a failed process will never be scheduled again for execution.

How this is **consistent** with the definition of fairness?

- An execution E is fair if each **correct** process p_i executes an infinite number of local computations ($\text{Exec}(i)$ events are not finite) and each message m sent **by a correct process or to a correct process** is eventually delivered.

HOMEWORKS 3

- Main book: page 67 exercises 2.2, 2.3
- You built a COUNTING SERVER that communicates with you using our Perfect-p2p algorithm over a stubborn link. The server stores a variable v (initially equal to 0) and it exposes three operations:
 - INC: that increases of 1 the value of v .
 - DEC: that decreases of 1 the value of v .
 - GET: that invokes $\langle pl, Send \mid v, client \rangle$ (client is your name as a process).

The server executes an operation when it delivers the corresponding message. As example, the increase happens when the server sees the event $\langle pl, Deliver \mid INC, client \rangle$.

You act as a client and you invoke exactly three operations in the following order: $\langle pl, Send \mid INC, server \rangle, \langle pl, Send \mid INC, server \rangle, \langle pl, Send \mid GET, server \rangle$.

You are really puzzled when, after a while, on your client you see $\langle pl, Deliver, 1, server \rangle$. Startled, you asked again $\langle pl, Send \mid GET, server \rangle$. Unfortunately, you are still waiting for an answer.

Questions: What is happening? Is there any way to fix it?

Note that, there are two possible explanations, try to find both. (Hint: you can see one of the explanation by answering to the question: How much will you wait for an answer?)