

4. LOGICAL CLOCK


In a Distributed System, each system has its own **logical clock**.

If clocks are not aligned it is not possible to order events generated by different processes


GOAL: find a way to timestamp event that follows out intuitive notion of causality

CAUSAL RELATIONSHIP

1. two events occurred at same process p_i happened in the same order as p_i observes them



2. when p_i sends a message to p_j , the send event happens before the receive event:




Lamport introduced the happened-before relation that capture the causal dependencies between events (**causal order relation**):

- we denote with \rightarrow_i the ordering relation between events in a process p_i .
- we denote with \rightarrow the happened-before relation between any pair of events.

Happened-Before RELATION:

Two event e and e' related by happened-before relation ($e \rightarrow e'$) if:

- **Local ordering:** $\exists p_i | e \rightarrow_i e'$
- **snd-rcv ordering:** $\forall m, send(m) \rightarrow receive(m)$
 - $e = send(m)$ is the event of sending a message m .
 - $e' = receive(m)$ is the event of receipt of the same message m
- **Transitivity:** $\exists e'' : (e \rightarrow e'') \wedge (e'' \rightarrow e') \Rightarrow e \rightarrow e'$
 - the happened-before relation is transitive




Applying these three rules is possible to define a causal ordered sequence of events

e_1, e_2, \dots, e_n .


Notes:

- the sequence e_1, e_2, \dots, e_n **may not be unique**.
- it **may exists a couple of events such that e_1 and e_2 are not in happenedbefore relation**.
- **if e_4 and e_3 are not in happened-before relation then they are concurrent** ($e_4 || e_3$).

- for any two events e_x and e_y in the execution history of a distributed system, either $e_x \rightarrow e_y$, $e_y \rightarrow e_x$ or $e_y || e_x$.



EXAMPLE:



$S_1 = \langle e^1_1, e^1_2, e^2_2, e^2_3, e^3_3, e^3_1, e^4_1, e^5_1, e^4_2 \rangle$

$S_2 = \langle e^1_3, e^2_1, e^3_1, e^4_1, e^5_1 \rangle$

Concurrent Elements: (e^1_3, e^1_2)

Logical/Lamport/Scalar Clock: monotonically increasing software counting register (not related to physical clock).

Each process p_i employs its logical clock L_i to apply a timestamp to events.


$L_i(e)$ is the **logical timestamp** assigned, using the logical clock, by a process p_i to event e .

Property:

if $e \rightarrow e'$ then $L(e) < L(e')$


Each process p_i initializes its logical clock $L_i = 0$

- when p_i sends a message m :
 - creates an event $send(m)$
 - increases L_i
 - timestamps m with $t = L_i$
- when p_i receives a message m with timestamp t
 - updates its logical clock $L_i = \max(t, L_i)$
 - produces an event $receive(m)$
 - increases L_i



because of the property (if $e \rightarrow e'$ then $L(e) < L(e')$)

Example:



Limits of Scalar Logical Clock:

Scalar Logical clocks can guarantee the property

- if $e \rightarrow e'$ then $L(e) < L(e')$

But it is **not possible to guarantee**:

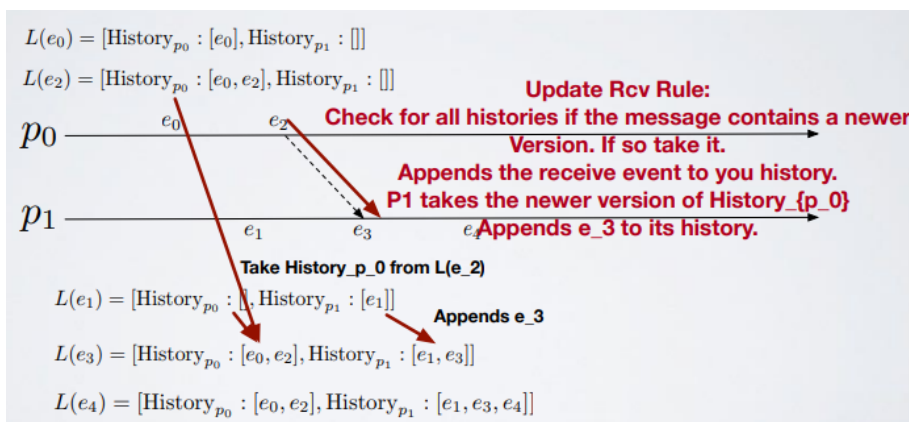
- if $L(e) < L(e')$ then $e \rightarrow e'$ IS not true everytime

So it is **not possible to determine**, analyzing only scalar clocks, **if two events are concurrent or correlated by the happened-before relation**.

VECTOR CLOCK


GOAL: capture causality (if $L(e) < L(e')$ then $e \rightarrow e'$)

$L(e)$ has not to be a single number. what if $L(e)$ is a history of events that happened before e (including e)?



$$L(e_i) > L(e_j) \Leftrightarrow \forall k : L(e_j)_{\text{History}_k} \subseteq L(e_i)_{\text{History}_k} \wedge \exists x : L(e_j)_{\text{History}_x} \subset L(e_i)_{\text{History}_x}$$


(there is a causal path form $e \rightarrow e'$ and the structure that contains e is less than the structure that contains e')



$History_x \subset History'_x \rightarrow History_x$ is a proper prefix of $History'_x$

We can say that:

$$History_x \subset History'_x \rightarrow \text{len}(History_x) < \text{len}(History'_x)$$



An event e is in happened-before relation with an event e' if in his *History* there is a tuple of elements that \subseteq and a tuple that is strictly \subset .

A vector clock for a set of N processes is an array of N integer counters:

- Each process p_i maintains a vector clock V_i and timestamps events by mean of it.
- Similarly to scalar clock, a vector clock is attached to message m (in this case we attach an array of integer).

Implementation:

- each process p_i **initializes its clock** $V_i = 0$
- p_i **increases** $V_i[i] + 1$ **when it generates a new event** e .
- when p_i **sends a message** m then:
 - creates an event $send(m)$.
 - $V_i[i] + 1$.
 - timestamps m with $t = V_i$.
- when p_i **recives a message** m containing timestamp $V - t$ then:
 - updates its logical clock: $V_i[j] = \max(V_t[j], V_i[j] \ \forall j \in \{1, \dots, N\})$.
 - generates an event $recv(m)$.
 - increases V_i .

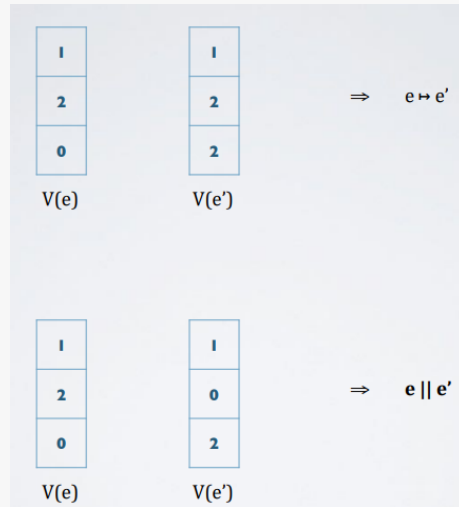
where:

- $V_i[i]$ represents the number of events produced by p_i .
- $V_i[j]$ with $i \neq j$ represents the number of events generated by p_j that p_i knows.


Properties:

- $V = V' \Leftrightarrow V[j] = V'[j] \quad \forall j \in \{1, \dots, N\}$
- $V \leq V' \Leftrightarrow V[j] \leq V'[j] \quad \forall j \in \{1, \dots, N\}$
- $V < V' \Leftrightarrow V \leq V' \wedge \exists j \in \{1, \dots, N\} | V[j] < V'[j]$

Example 1:



Example 2:



$$[-, -, 1] < [4, 3, 2] \Rightarrow e_3^1 \rightarrow e_1^4$$

$$[4, 3, 2] \nless [3, -, 3] \Rightarrow e_1^4 \nless e_3^3$$

Each mechanism can be used to solve different problems:

- Scalar Timestamp \rightarrow Lamport's Mutual Exclusion
- Vector Timestamp \rightarrow Causal Broadcast

MUTUAL EXCLUSION ABSTRACTION

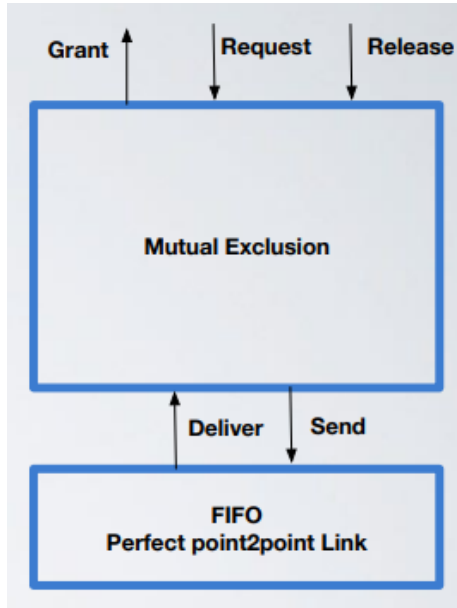
Events:

- **Request:** from upper layer - **requests access to Critical Section (CS).**
- **Grant:** to upper layer - **grant the access to CS.**
- **Release:** from upper layer - **release the CS.**

Properties:

- **(Mutual Exclusions)** at any time t , only one process is inside the CS.

- (**Liveness**) if a process p requests access, then it eventually enters the CS.
- (**Fairness**) if the request of process p happens before the request of process q , then q cannot access the CS before p .



The algorithm assumes **no crashes** ($F = 0$):

when a process wants to enter the CS (critical section) it sends a request message to all the other (using **scalar clocks**). The algorithm assume a FIFO link.

Lamport's Algorithm:

Algorithm 1 Lamport's ME Algorithm on process p_i - MSGS are REQ, ACK, RLS

```

1: upon event INIT
2:   Requests = Acks =  $\emptyset$ 
3:   scalar_clock = 0
4:   my_req =  $\perp$ 
5:    $\Pi = \{p_0, p_1, \dots, p_{n-1}\}$  ▷ Set of all processes

6: ▷ Request access to CS from upper layer
7: upon event REQUEST
8:   scalar_clock = scalar_clock + 1
9:   my_req = (REQ, ts = < i, scalar_clock >)
10:  for all  $p_j \in \Pi$  do
11:    SEND FIFOPERFECTLINK( $p_j$ , req_msg)  ▷ Send a REQ containing my ID (i) and ts (scalar_clock) to all  $p \in \Pi$ 

12: ▷ Release CS from upper layer
13: upon event RELEASE
14:   my_req =  $\perp$ 
15:   scalar_clock = scalar_clock + 1
16:   for all  $p_j \in \Pi$  do
17:     SEND FIFOPERFECTLINK( $p_j$ , (RLS, ts = < i, scalar_clock >))

18: ▷ ts(x) < ts(y) when scalar_clock of x is less than the one of y, or they are equal and the id that sent x is less than the id that sent y
19: upon event  $\nexists req \in Requests : ts(req) < ts(my\_req) \wedge \forall p \in \Pi : \exists m \in Acks | ts(m) > ts(my\_req) \wedge sender(m) = p$ 
20:   trigger event GRANTED

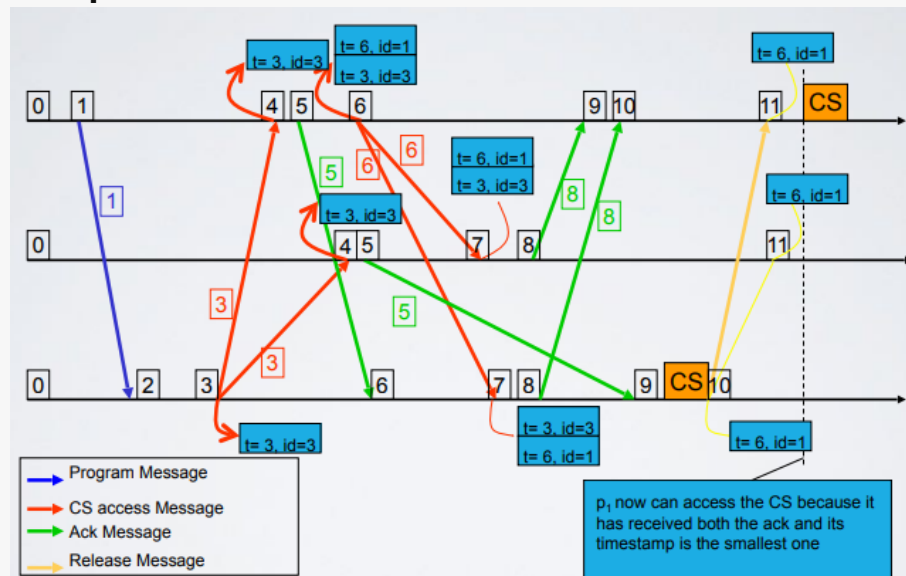
21: upon event DELIVER MESSAGE(m)
22:   scalar_clock = max(clock(m), scalar_clock) + 1
23:   if m is a REQ then
24:     Request_set = Request_set  $\cup \{m\}$ 
25:     scalar_clock = scalar_clock + 1
26:     SEND FIFOPERFECTLINK(sender(m), (ACK, ts = < i, scalar_clock >))
27:   else if m is a ACK then
28:     Acks = Acks  $\cup \{m\}$ 
29:   else if m is a RLS  $\wedge \exists req \in Request\_set : sender(req) = sender(m)$  then
30:     Requests = Requests  $\setminus \{req\}$ 

```

****Local data structures to each process p_i :**

- ck is the counter for process p_i .
- *Request*: a set maintained by p_i where CS access requests are stored.
- **Algorithm rules for a process p_i :
- **access** the CS:
 - p_i sends a **request message** (attaching ck) to all the other processes.
 - p_i adds its request to *Requests* structure.
- **request reception** from process p_j :
 - p_i puts p_j request (including the timestamp) in its *Requests*.
 - p_i sends back an *ACK* message to p_j including its local timestamp ck .
- p_i **enters the CSS iff**:
 - the request of p_i is the one with smallest timestamp in its *Requests*.
 - p_i has already received an *ACK* with timestamp t' from any other processes and $t' > t$.
- **release** of the CS:
 - p_i send a *Release* message to all the other processes.
 - p_i deletes its request from *Requests*.
- **reception of release message** from a process p_i :
 - p_i deletes p_j request from *Requests*.

Example:



DEMONSTRATION (by **contradiction**): why only one can enter the CS:

1. **Mutual Exclusion**: assume that p_1 and p_2 enter CS at the same time.
 - you cannot enter CS if you have not received acks from everyone, and such acks happened after your request (when you create a new request its ts is greater than the tss of all old acks).
 - \Rightarrow both the process have received an ACK from any other process and each my_{req} is the smallest in the respective queue:
 - p_i received the ack from p_j . When p_j sends the ack to p_i it inserts in its set the req of p_i .
 - p_j received the ack from p_i . When p_i sends the ack to p_j it inserts in its set the req of p_j .

- if p_j Requests after acking p_i , then $ts(req, p_j) > ts(req, p_i)$ and that's a **contradiction**. (The same for p_i)
- So p_j Requests before acking p_i , then by FIFO the (req, p_j) reaches p_i before the ack of p_j . Thus p_i has (req, p_j) in its set and p_j has (req, p_i) in its set. **Since request are total ordered then we have a contradiction.**

2. **Fairness:** different requests are satisfied in the same order as they are generated (such order comes from the happened-before relation).

- **Proof:** suppose p_i enters before p_j , even if (req, p_i) happened after (req, p_j) . Since p_i enters only after the ack of p_j , by FIFO it sees (req, p_j) before receiving the ack that allows him to enter. Since (req, p_j) happens before (req, p_i) we have $ts((req, p_j)) < ts((req, p_i))$, thus (req, p_i) is not the request with minimal timestamp in the set of p_i .

Algorithm Cost:

Number of operation for a CS execution: $3(N - 1)$:

- $N - 1$ requests.
- $N - 1$ acks.
- $N - 1$ releases

Deelay to enter the CS: $2 \leq deelay \leq N + 2$:

- $\Omega(2)$
- $O(N + 2)$