

DISTRIBUTED REGISTERS II

ATOMIC REGISTERS

CONSISTENCY PROPERTIES

- REGULAR CONSISTENCY: every reads on register x outputs the last value written on x, or the value that someone is concurrently writing on x.
 - We have seen algorithms for (1,N) Regular register with P (fail-stop model) and without P (fail-silent) using quorums and majority of corrects.
- SEQUENTIAL CONSISTENCY: there exists a global ordering of all writes and reads that is consistent with the local order of operations that each process sees.
- ATOMIC CONSISTENCY/(LINEARIZABILITY): Each operation appears to take effect instantaneously at some moment between its start and end.

(1,N) ATOMIC REGISTER: SPECIFICATION

Properties:

Termination. If a correct process invokes an operation, then the operation eventually receives the corresponding confirmation.

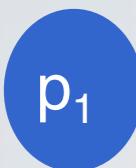
Validity. A read operation returns the last value written or the value concurrently being written.

Ordering. If a read returns v₂ after a read that precedes it has returned v₁, then v₁ has not been written after v₂

(Ordering is stated for (1,N) and it is equivalent to the more general definition with Linearization Points).

(1,N) ATOMIC REGISTER: SCENARIO

es.1



write(5)



write(6)



1. **Regular but not atomic register:** Write(5) precedes write(6). But process p₂ read first the value 6 and then the value 5

es.2



write(5)



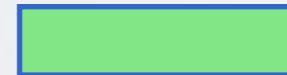
write(6)



2. **The register is atomic**



read()→5



read()→6



read()→6



(1,N) ATOMIC REGISTER: SCENARIO

Non atomic register: the precedence relation also refers to read operations issued by different processes



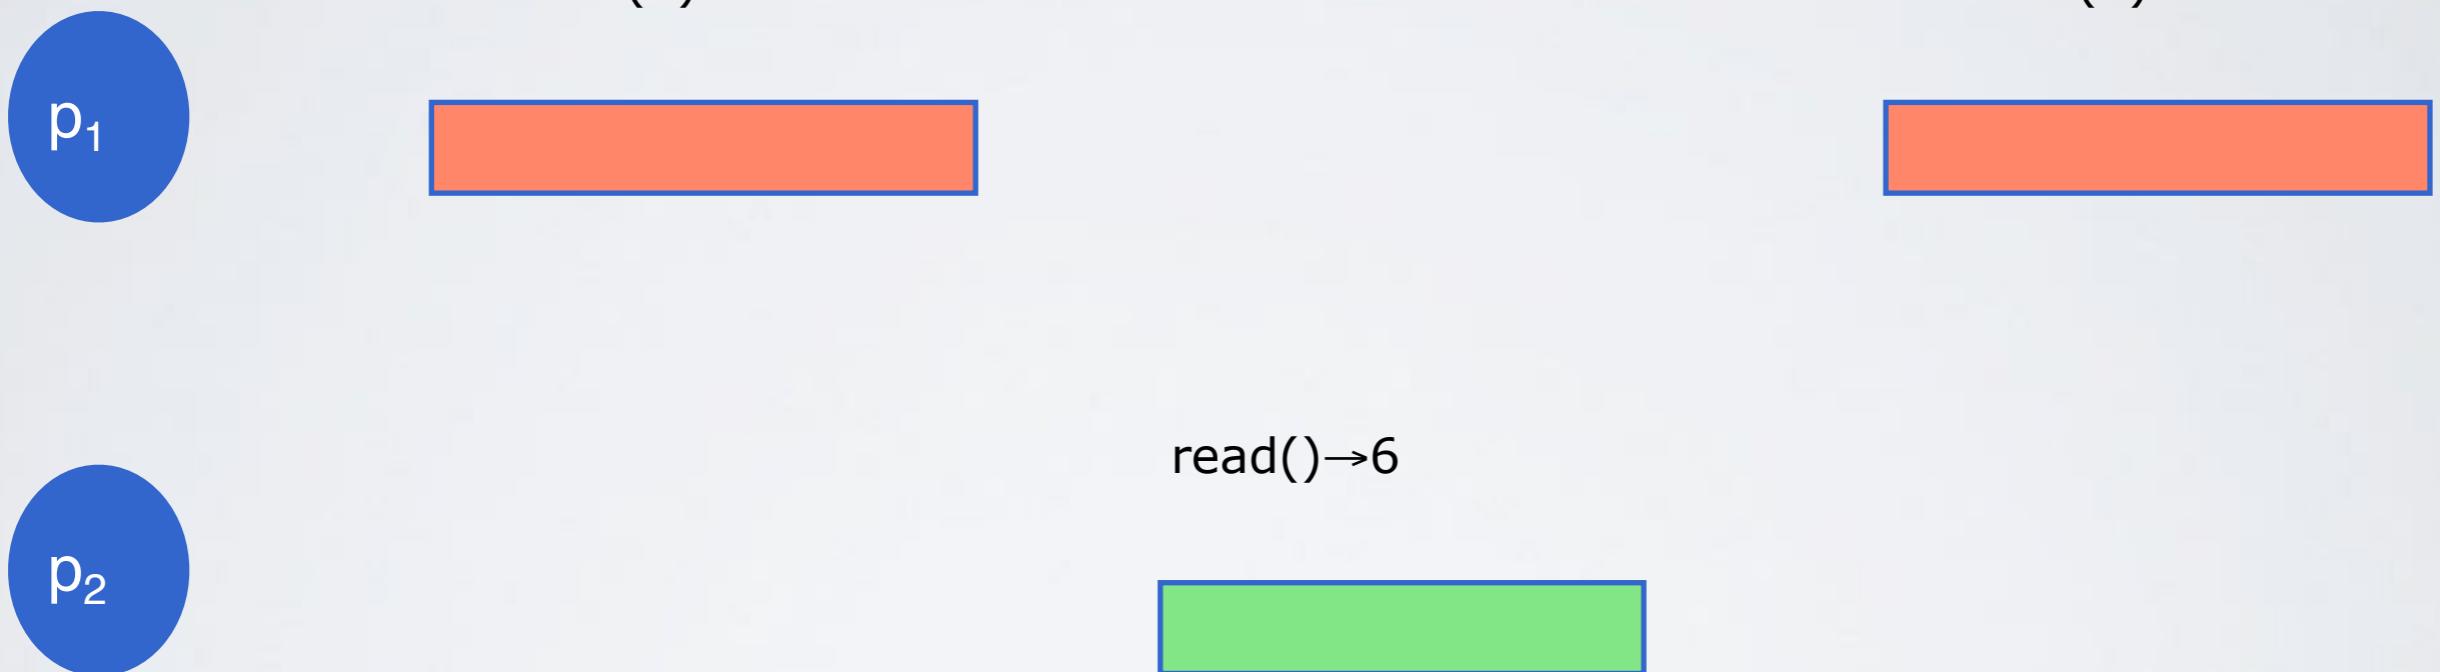
SCENARIO 1



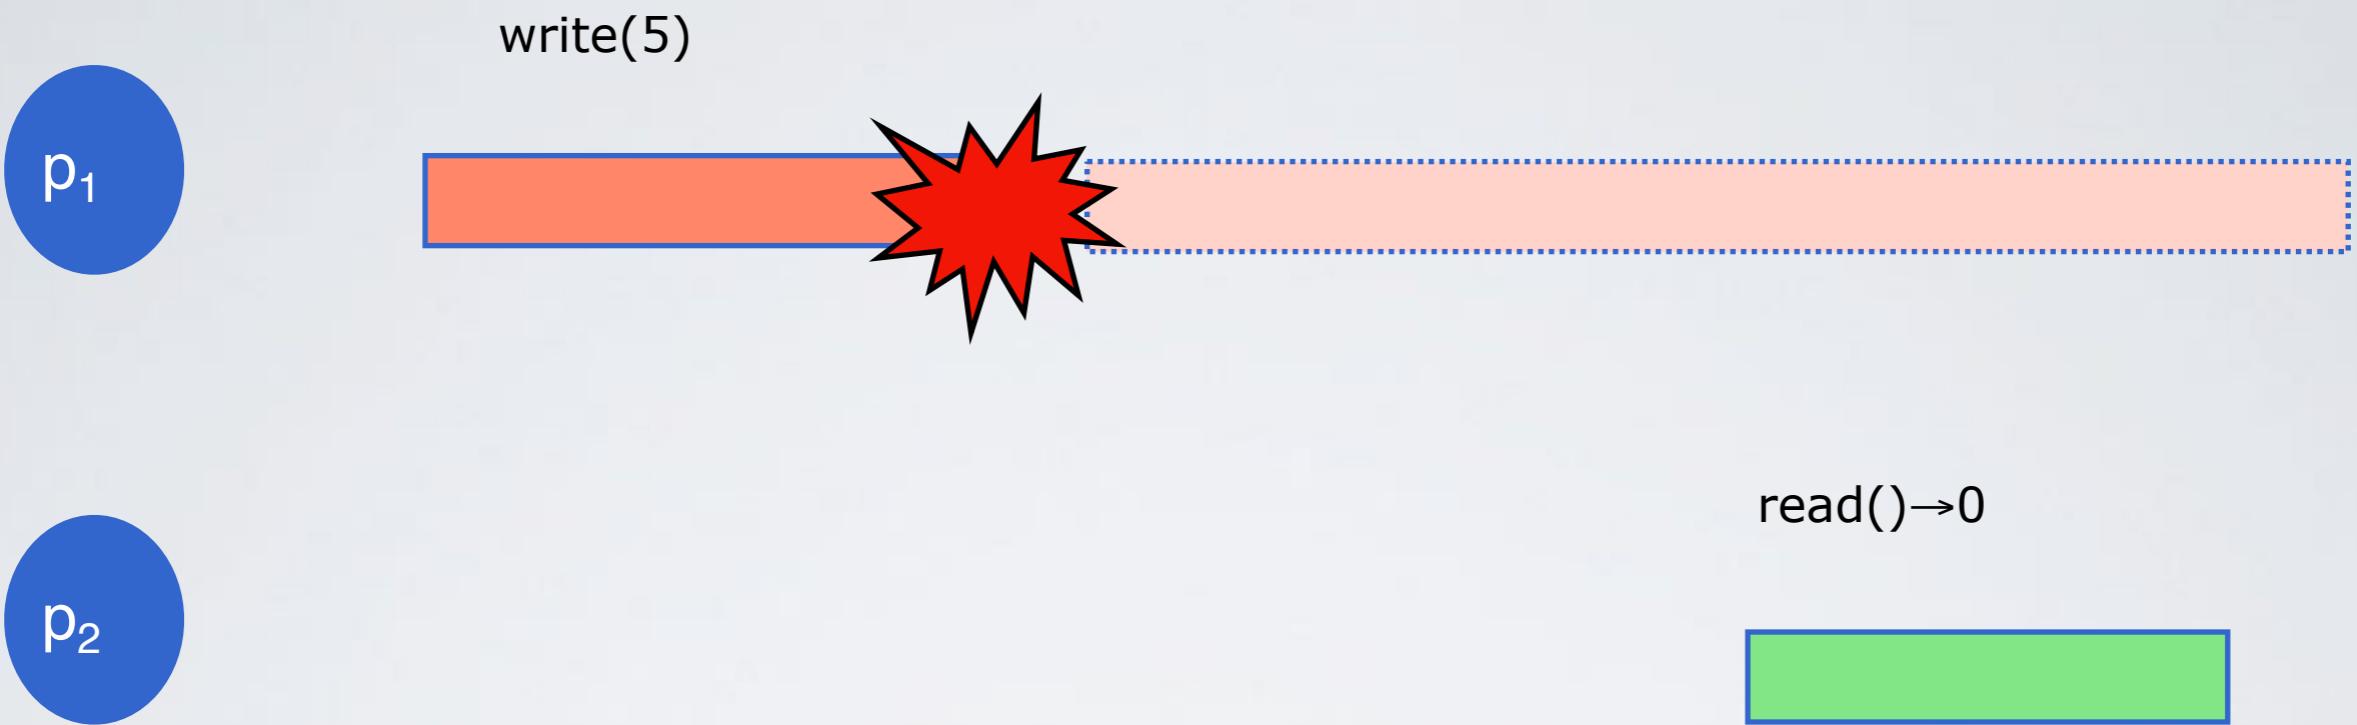
ATOMIC and **REGULAR**.

SCENARIO 2

NOT ATOMIC and NOT REGULAR, but SEQUENTIAL CONSISTENT

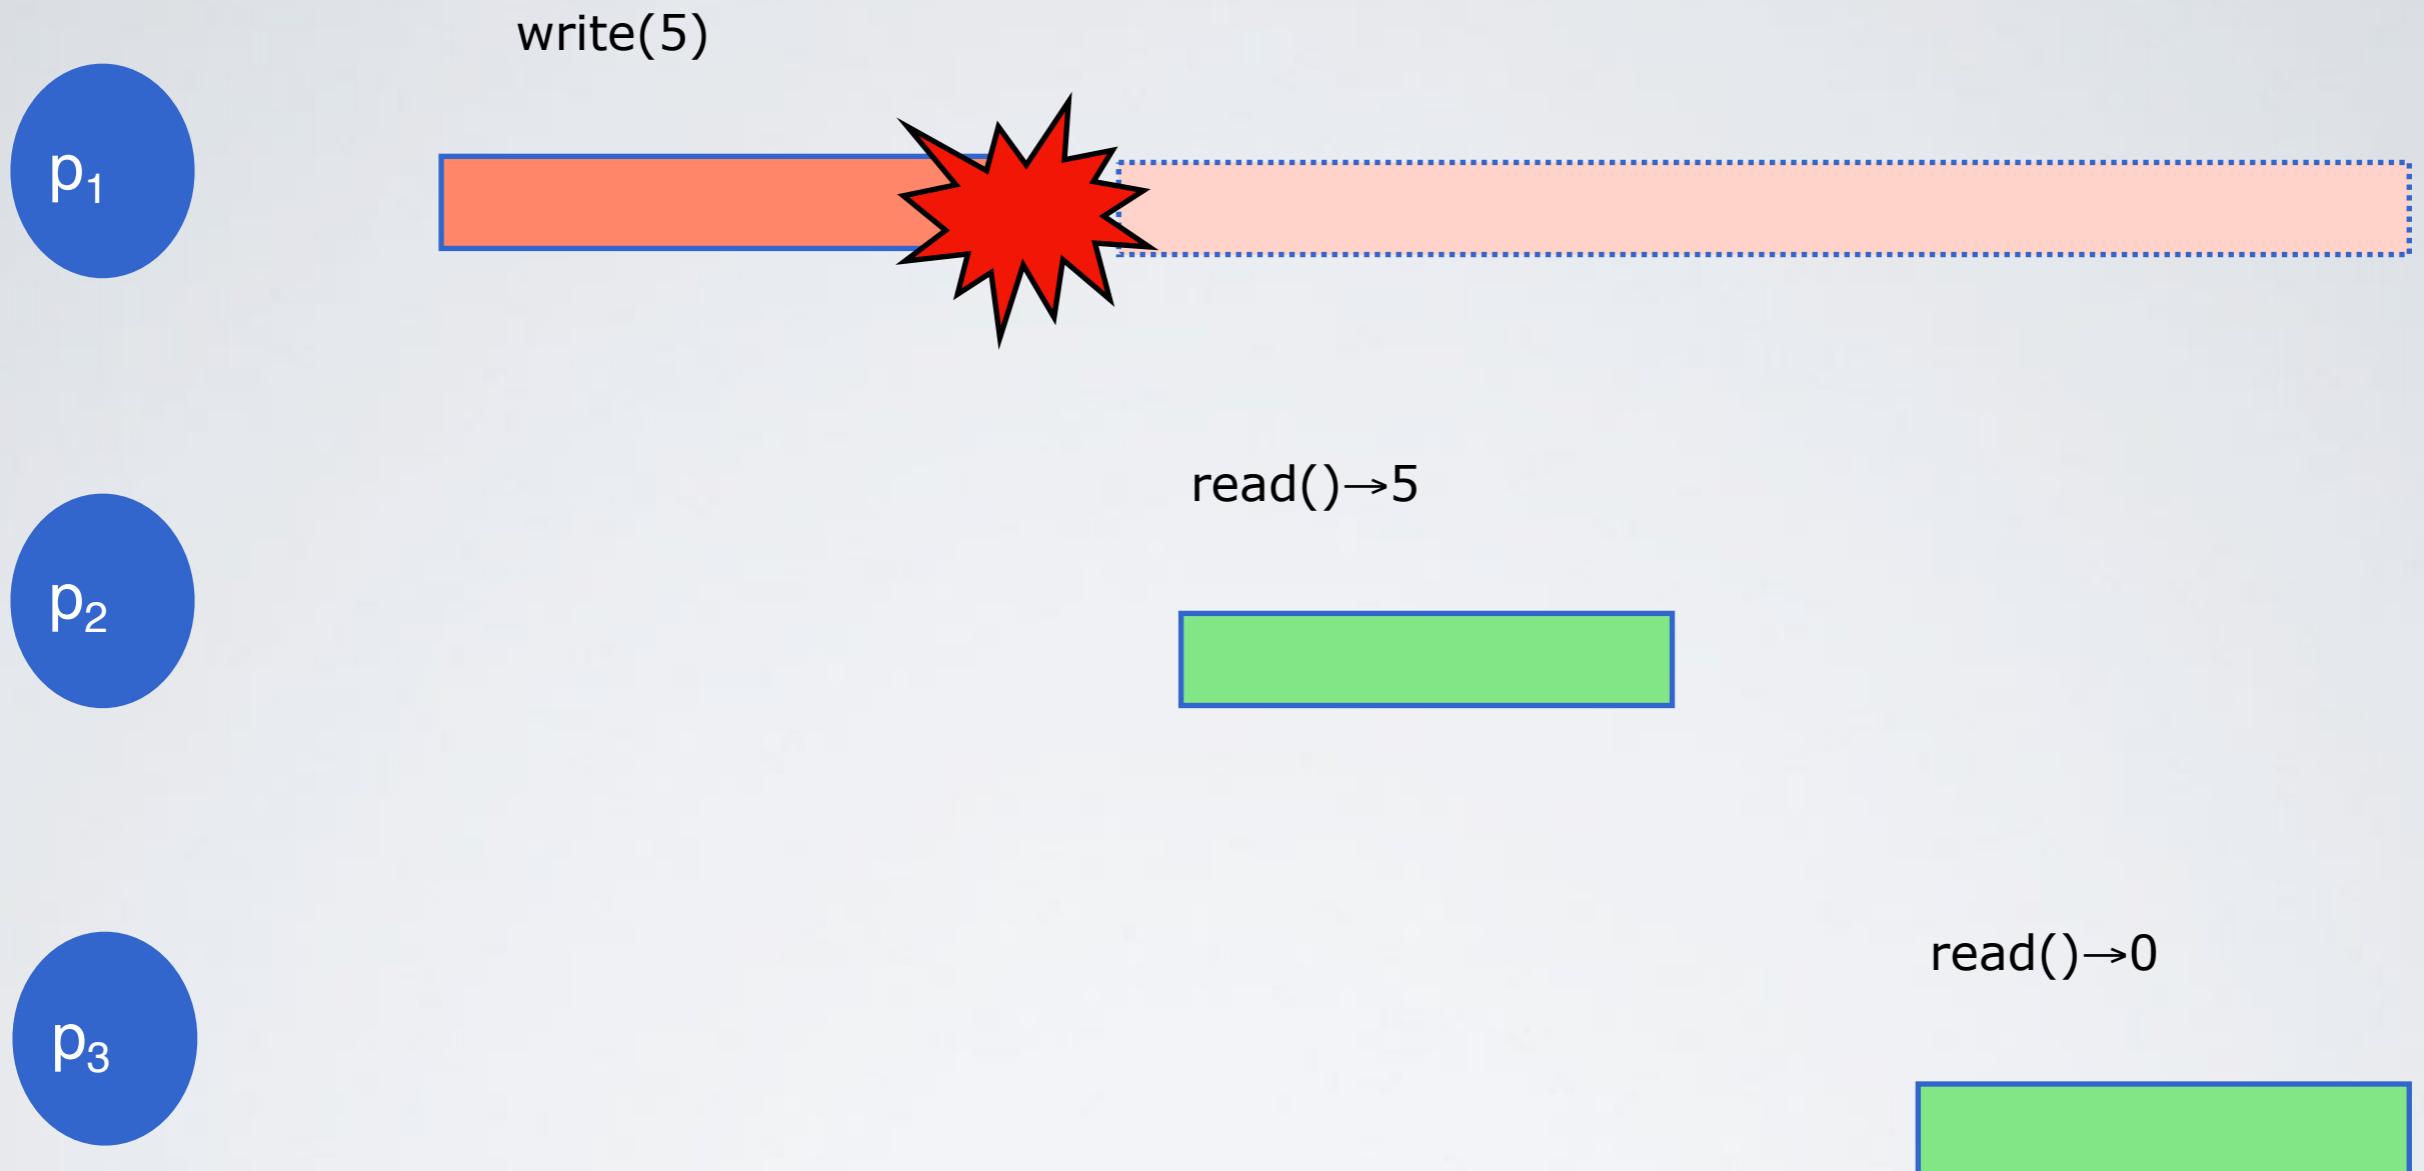


SCENARIO 3



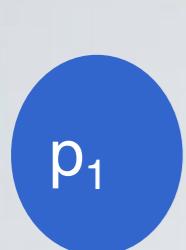
ATOMIC and **REGULAR**. $\text{write}(5)$ executed by p_1 fails.
So, it does not complete and it is concurrent with the
read by p_2 . Validity is respected.

SCENARIO 4



REGULAR but non **ATOMIC**. The ordering property is violated.

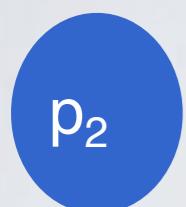
SCENARIO 5



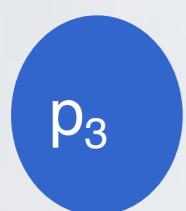
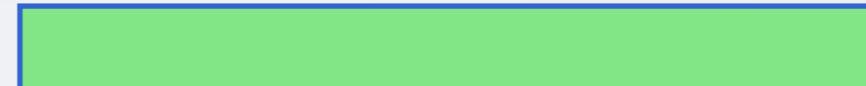
write(5)



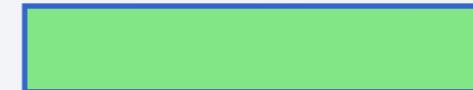
write(6)



read()→6



read()→5



ATOMIC and **REGULAR**

(1, N) ATOMIC REGISTER: INTERFACE

Module 4.2: Interface and properties of a $(1, N)$ atomic register

Module:

Name: $(1, N)$ -AtomicRegister, **instance** *onar*.

Events:

Request: $\langle \text{onar}, \text{Read} \rangle$: Invokes a read operation on the register.

Request: $\langle \text{onar}, \text{Write} \mid v \rangle$: Invokes a write operation with value v on the register.

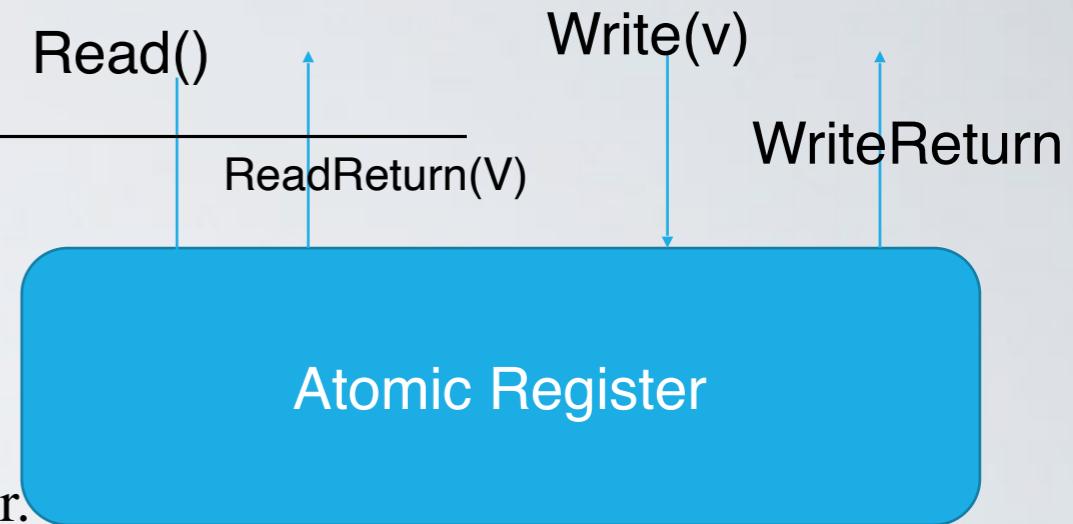
Indication: $\langle \text{onar}, \text{ReadReturn} \mid v \rangle$: Completes a read operation on the register with return value v .

Indication: $\langle \text{onar}, \text{WriteReturn} \rangle$: Completes a write operation on the register.

Properties:

ONAR1–ONAR2: Same as properties ONRR1–ONRR2 of a $(1, N)$ regular register (Module 4.1).

ONAR3: *Ordering*: If a read returns a value v and a subsequent read returns a value w , then the write of w does not precede the write of v .



ALGORITHMS

We will show algorithms to implement (1,N) Atomic Register using (1,N) Regular.

- WE WILL NOT USE MESSAGES TO COMMUNICATE, JUST Regular Registers

We will show algorithms to implement (1,N) Atomic Register using directly message passing.

- We will only use messages to communicate

(1,N) ATOMIC REGISTER USING REGULAR REGISTER

(1,N) ATOMIC REGISTER -

The algorithm consists of two phases

PHASE 1. We use a (1,N) regular register to build a (1,1) atomic register

PHASE 2. We use a set of (1,1) atomic registers to build a (1,N) atomic register

NOTATION:

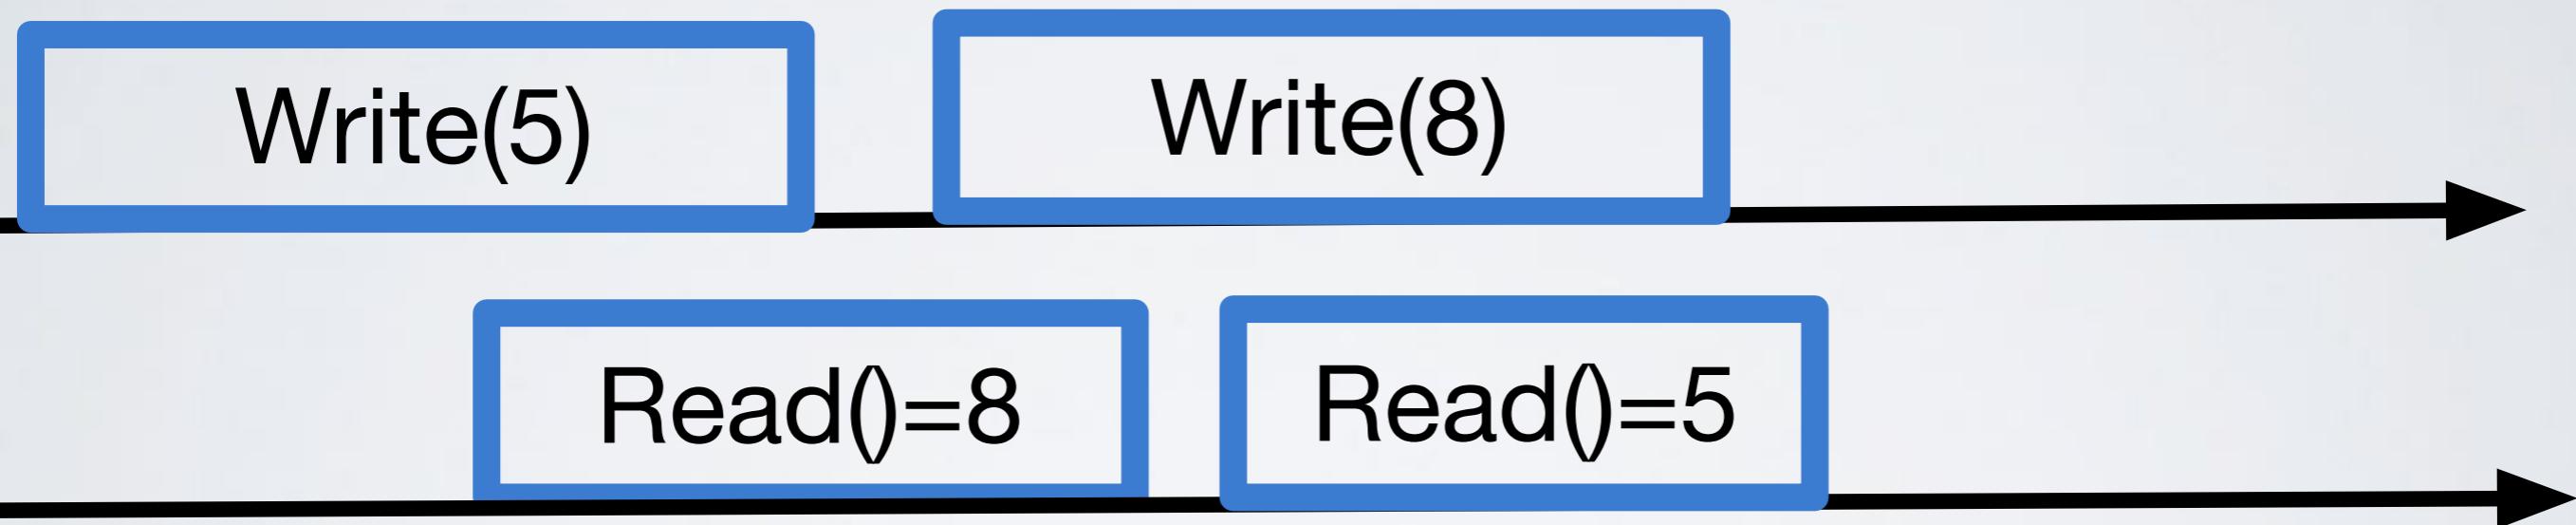
Hereafter, **RR** and **AR**, will be sometimes used to respectively denote Regular Register and Atomic Register

PHASE 1 - FROM (1,N)-RR TO (1,1)-AR



PHASE 1 - FROM (1,N)-RR TO (1,1)-AR

IDEA 1: We have just one writer and one reader, what is the pathological case with a RR that is not acceptable in AR?



PHASE 1 - FROM (1,N)-RR TO (1,1)-AR

IDEA:

- p_1 is the writer and p_2 is the reader of the (1,1) atomic register, we aim to implement
- We use a (1,N) regular register where p_1 is the writer and p_2 is the reader
- Each write operation on the atomic register writes the pair *(value, timestamp)* into the underlying regular register
- The reader tracks the timestamp of previously read values to avoid to read something old

PHASE 1 - FROM $(1, N)$ -RR TO $(1, 1)$ -AR

Algorithm 4.3: From $(1, N)$ Regular to $(1, 1)$ Atomic Registers

Implements:

$(1, 1)$ -AtomicRegister, **instance** $ooar$.

upon event $\langle onrr, \text{WriteReturn} \rangle$ **do**
trigger $\langle ooar, \text{WriteReturn} \rangle$;

Uses:

$(1, N)$ -RegularRegister, **instance** $onrr$.

upon event $\langle ooar, \text{Read} \rangle$ **do**
trigger $\langle onrr, \text{Read} \rangle$;

upon event $\langle ooar, \text{Init} \rangle$ **do**

$(ts, val) := (0, \perp)$;

$wts := 0$;

upon event $\langle onrr, \text{ReadReturn} \mid (ts', v') \rangle$ **do**
if $ts' > ts$ **then**
 $(ts, val) := (ts', v')$;

trigger $\langle ooar, \text{ReadReturn} \mid val \rangle$;

upon event $\langle ooar, \text{Write} \mid v \rangle$ **do**

$wts := wts + 1$;

trigger $\langle onrr, \text{Write} \mid (wts, v) \rangle$;

PHASE 1 - FROM $(1,N)$ -RR TO $(1,1)$ -AR

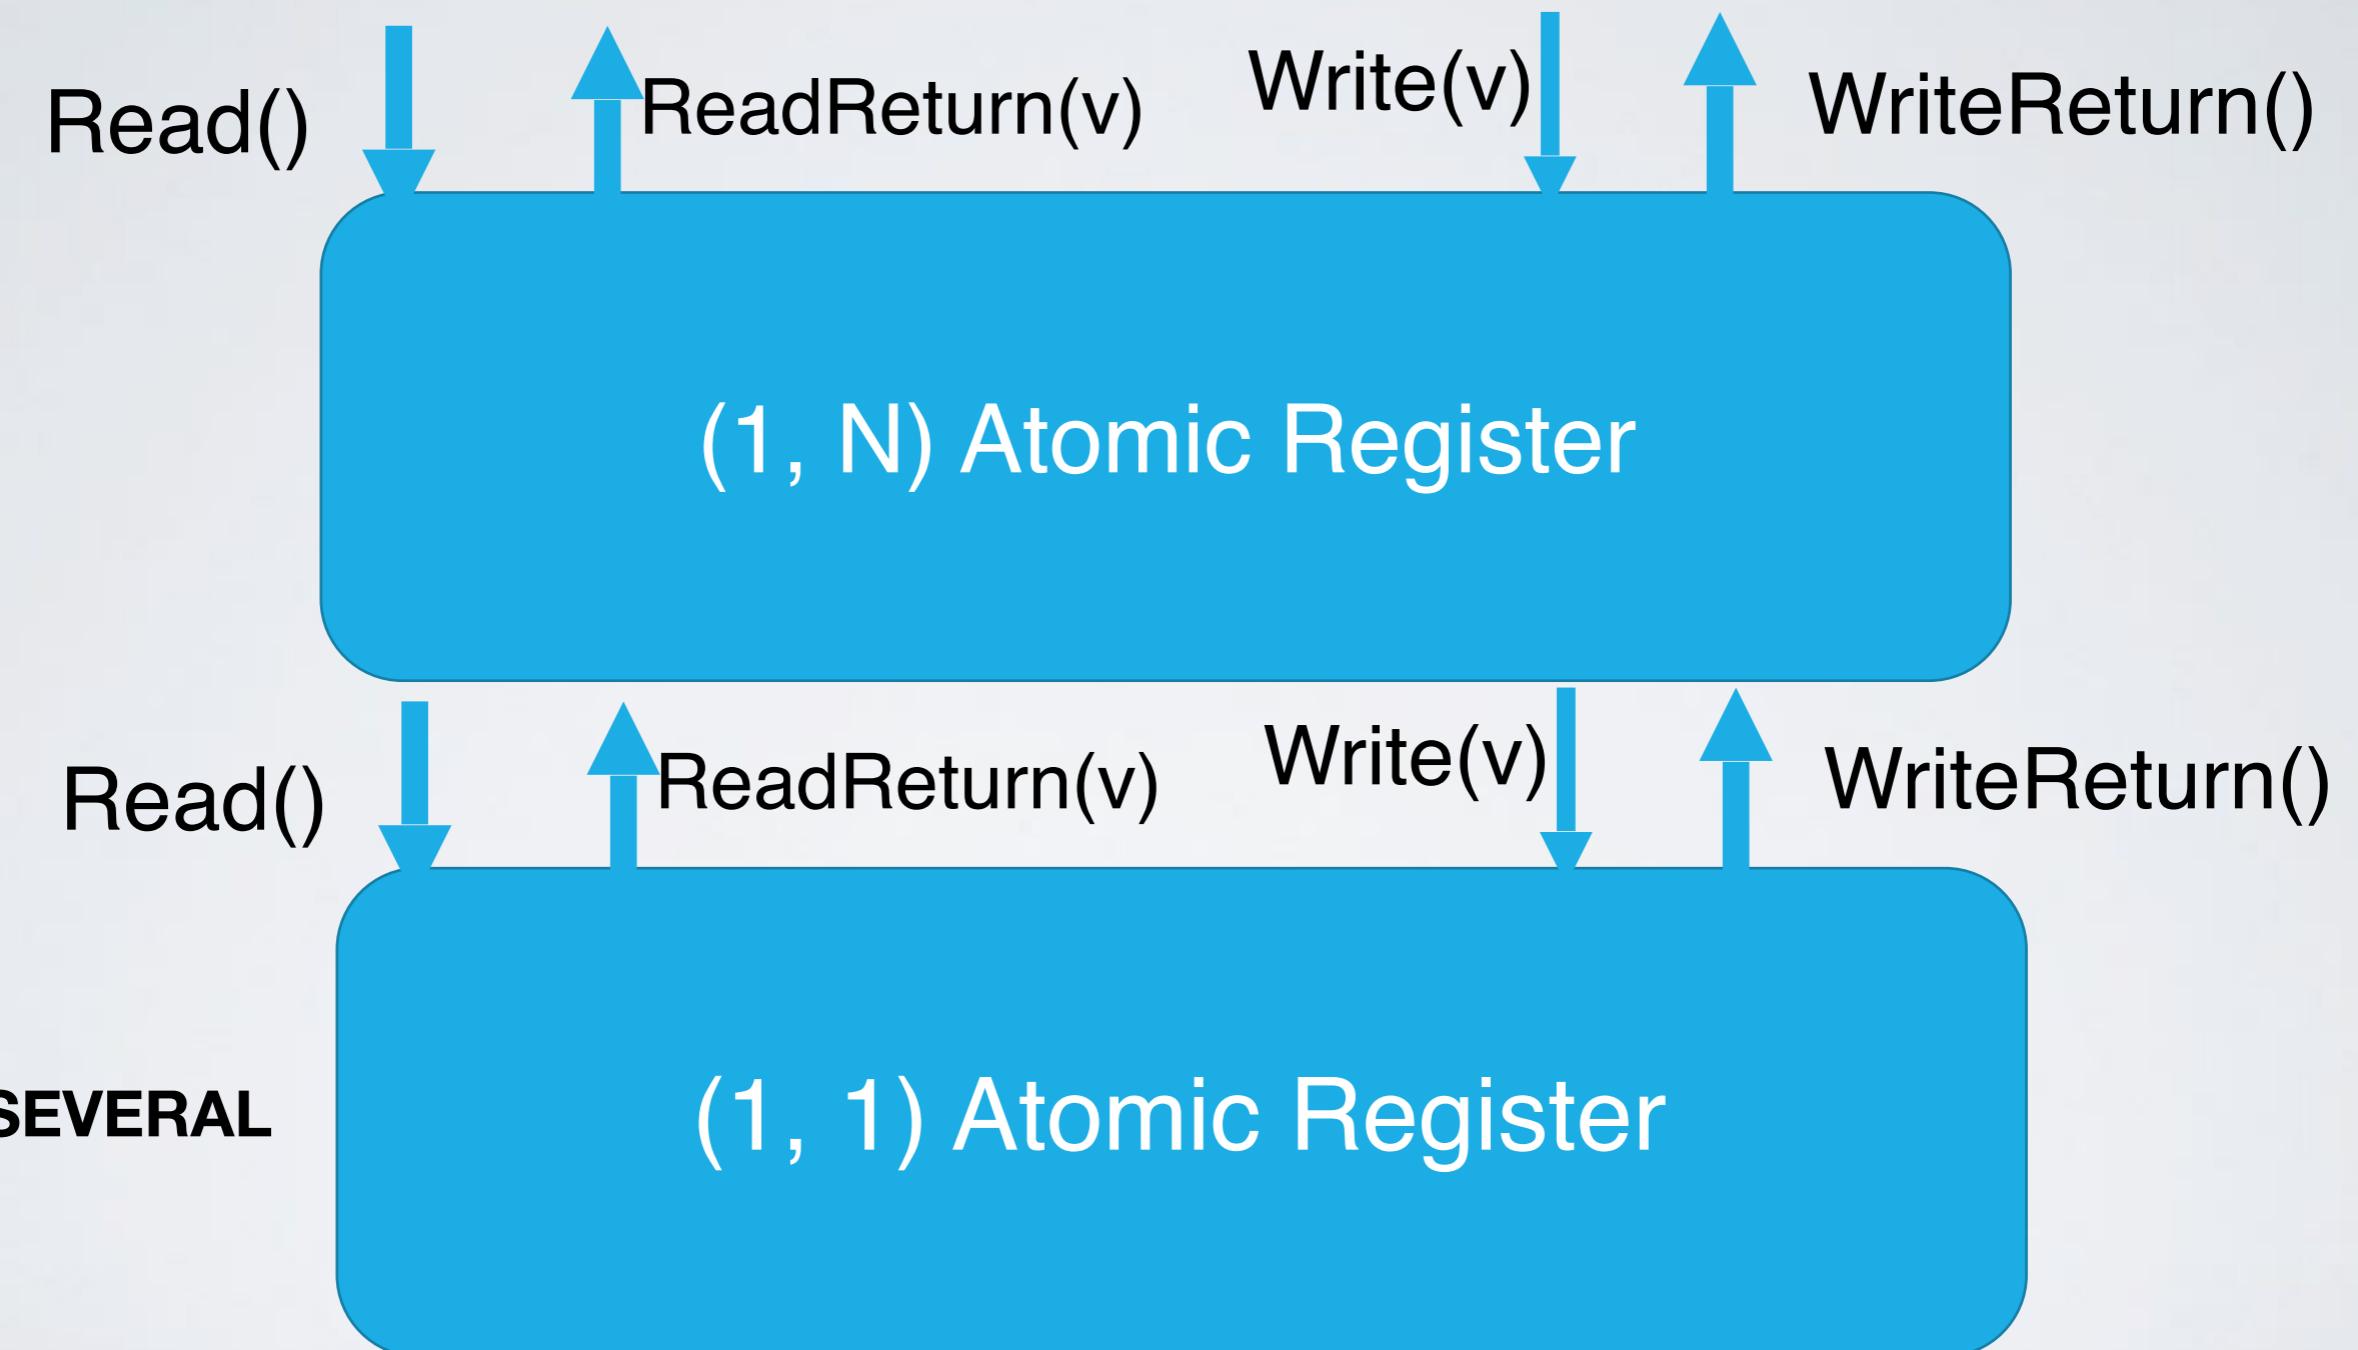
Correctness:

- **Termination** – from the termination property of the regular register
- **Validity** – from the validity property of the regular register
- **Ordering** – from the validity property and from the fact that the read tracks the last value read and its timestamp. A read operation always returns a value with a timestamp greater or equal to the one of the previously read value

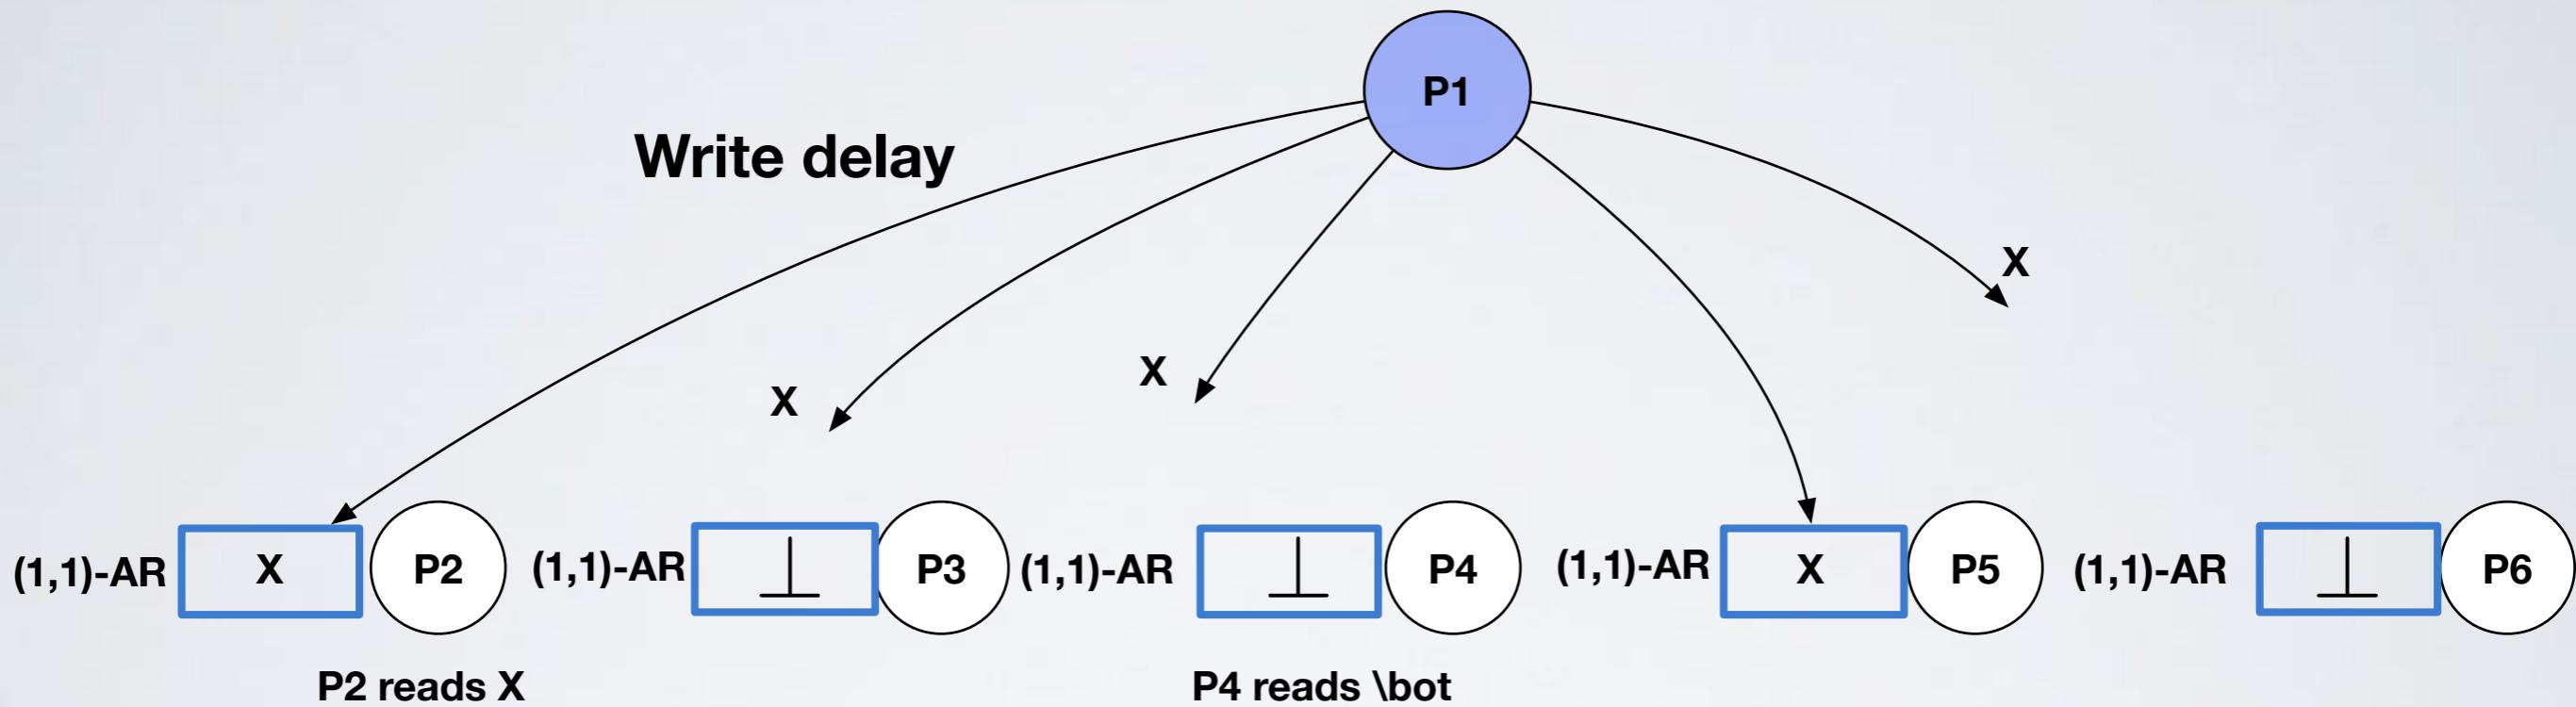
Performance:

- **Write** – Each write operation performs a write on a $(1,N)$ regular register
- **Read** - Each read operation performs a read on a $(1,N)$ regular register

PHASE 2 - FROM $(1, 1)$ -AR TO $(1, N)$ -AR



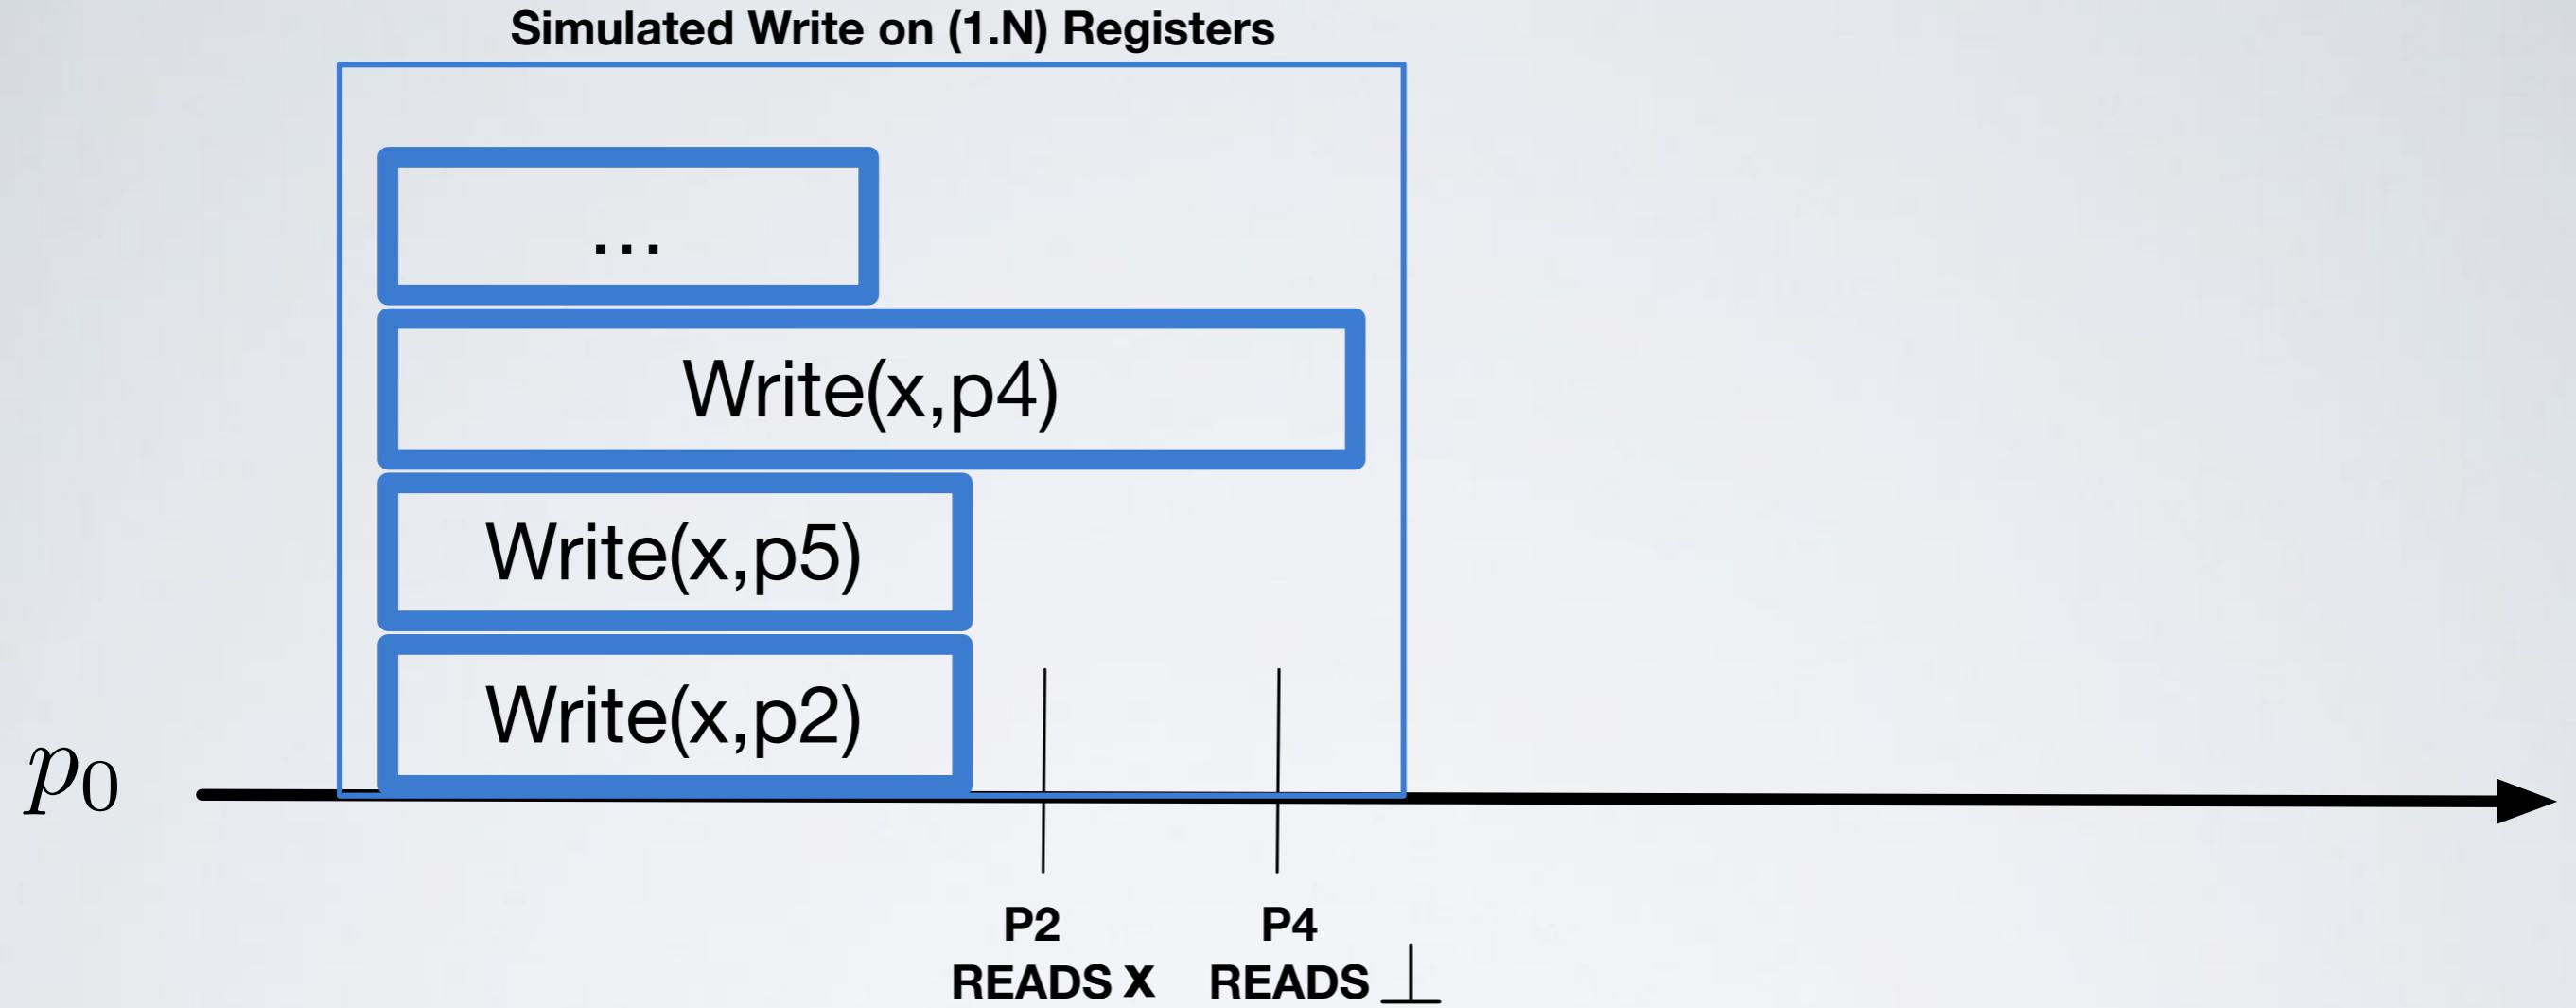
PHASE 2 - FROM (1,1)-AR TO (1,N)-AR



Each one has a (1,1)-register on which P1 writes.

Problem: multiple writes on different registers are not instantaneous, so a write can appear on P2 but not on P4.

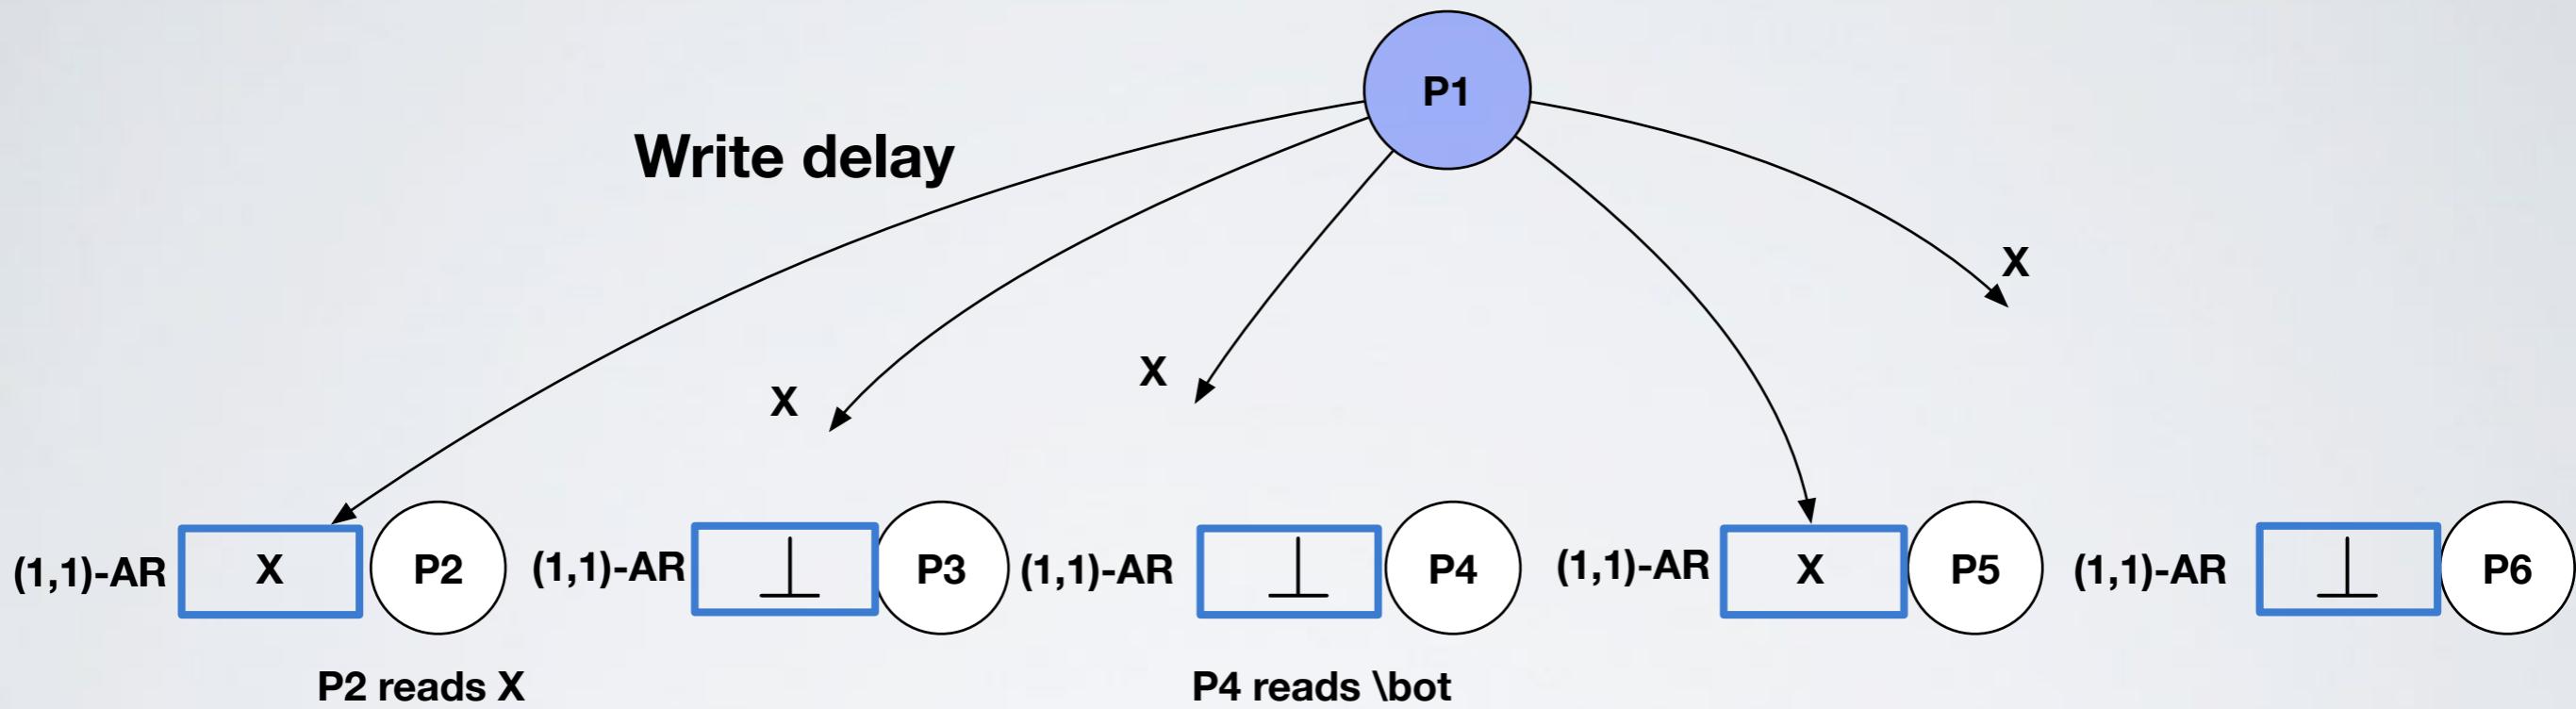
PHASE 2 - FROM (1,1)-AR TO (1,N)-AR



Each one has a (1,1)-register on which P1 writes.

Problem: multiple writes on different registers are not instantaneous, so a write can appear on P2 but not on P4. This does not satisfy an atomic registers

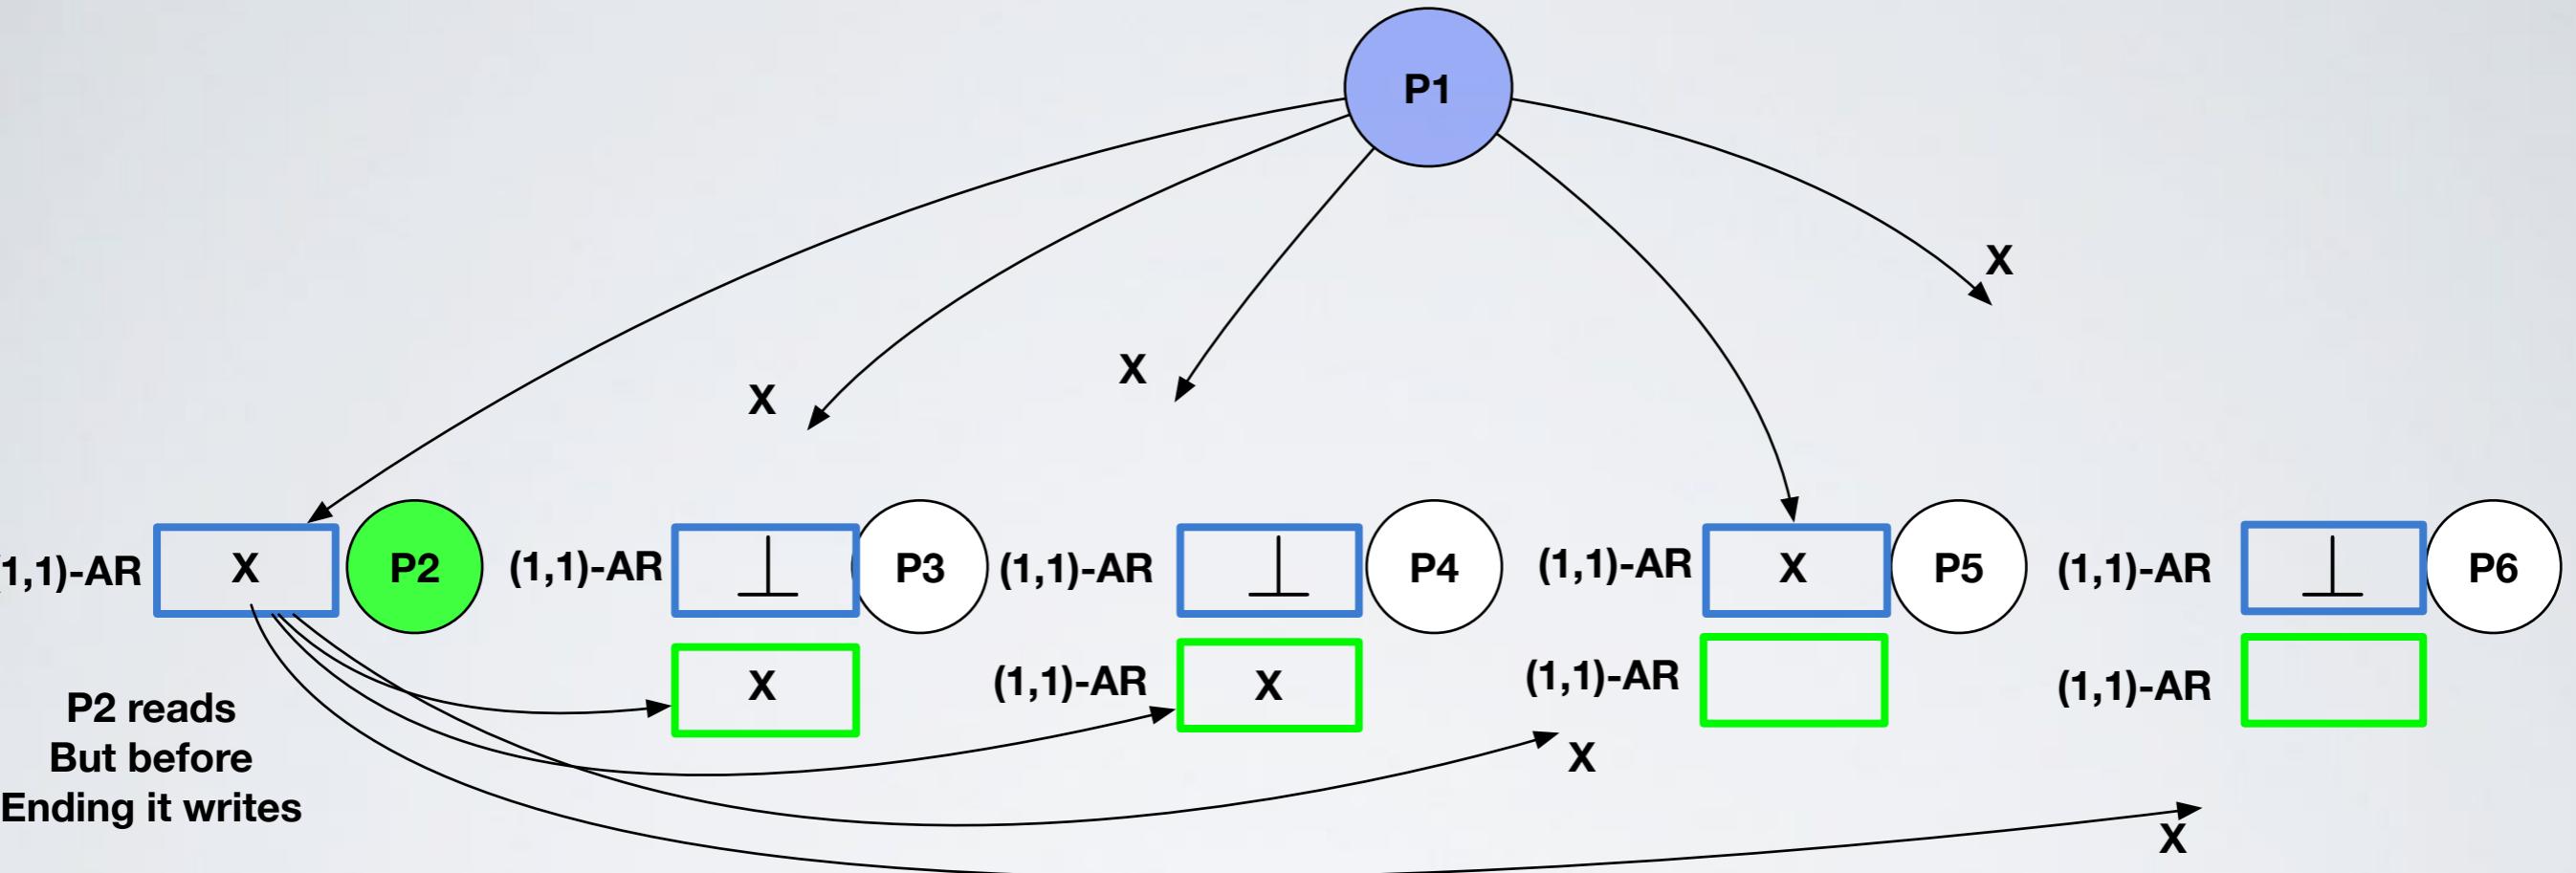
PHASE 2 - FROM (1,1)-AR TO (1,N)-AR



Problem: multiple writes on different registers are not instantaneous, so a write can appear on P2 but not on P4.

How to fix it?

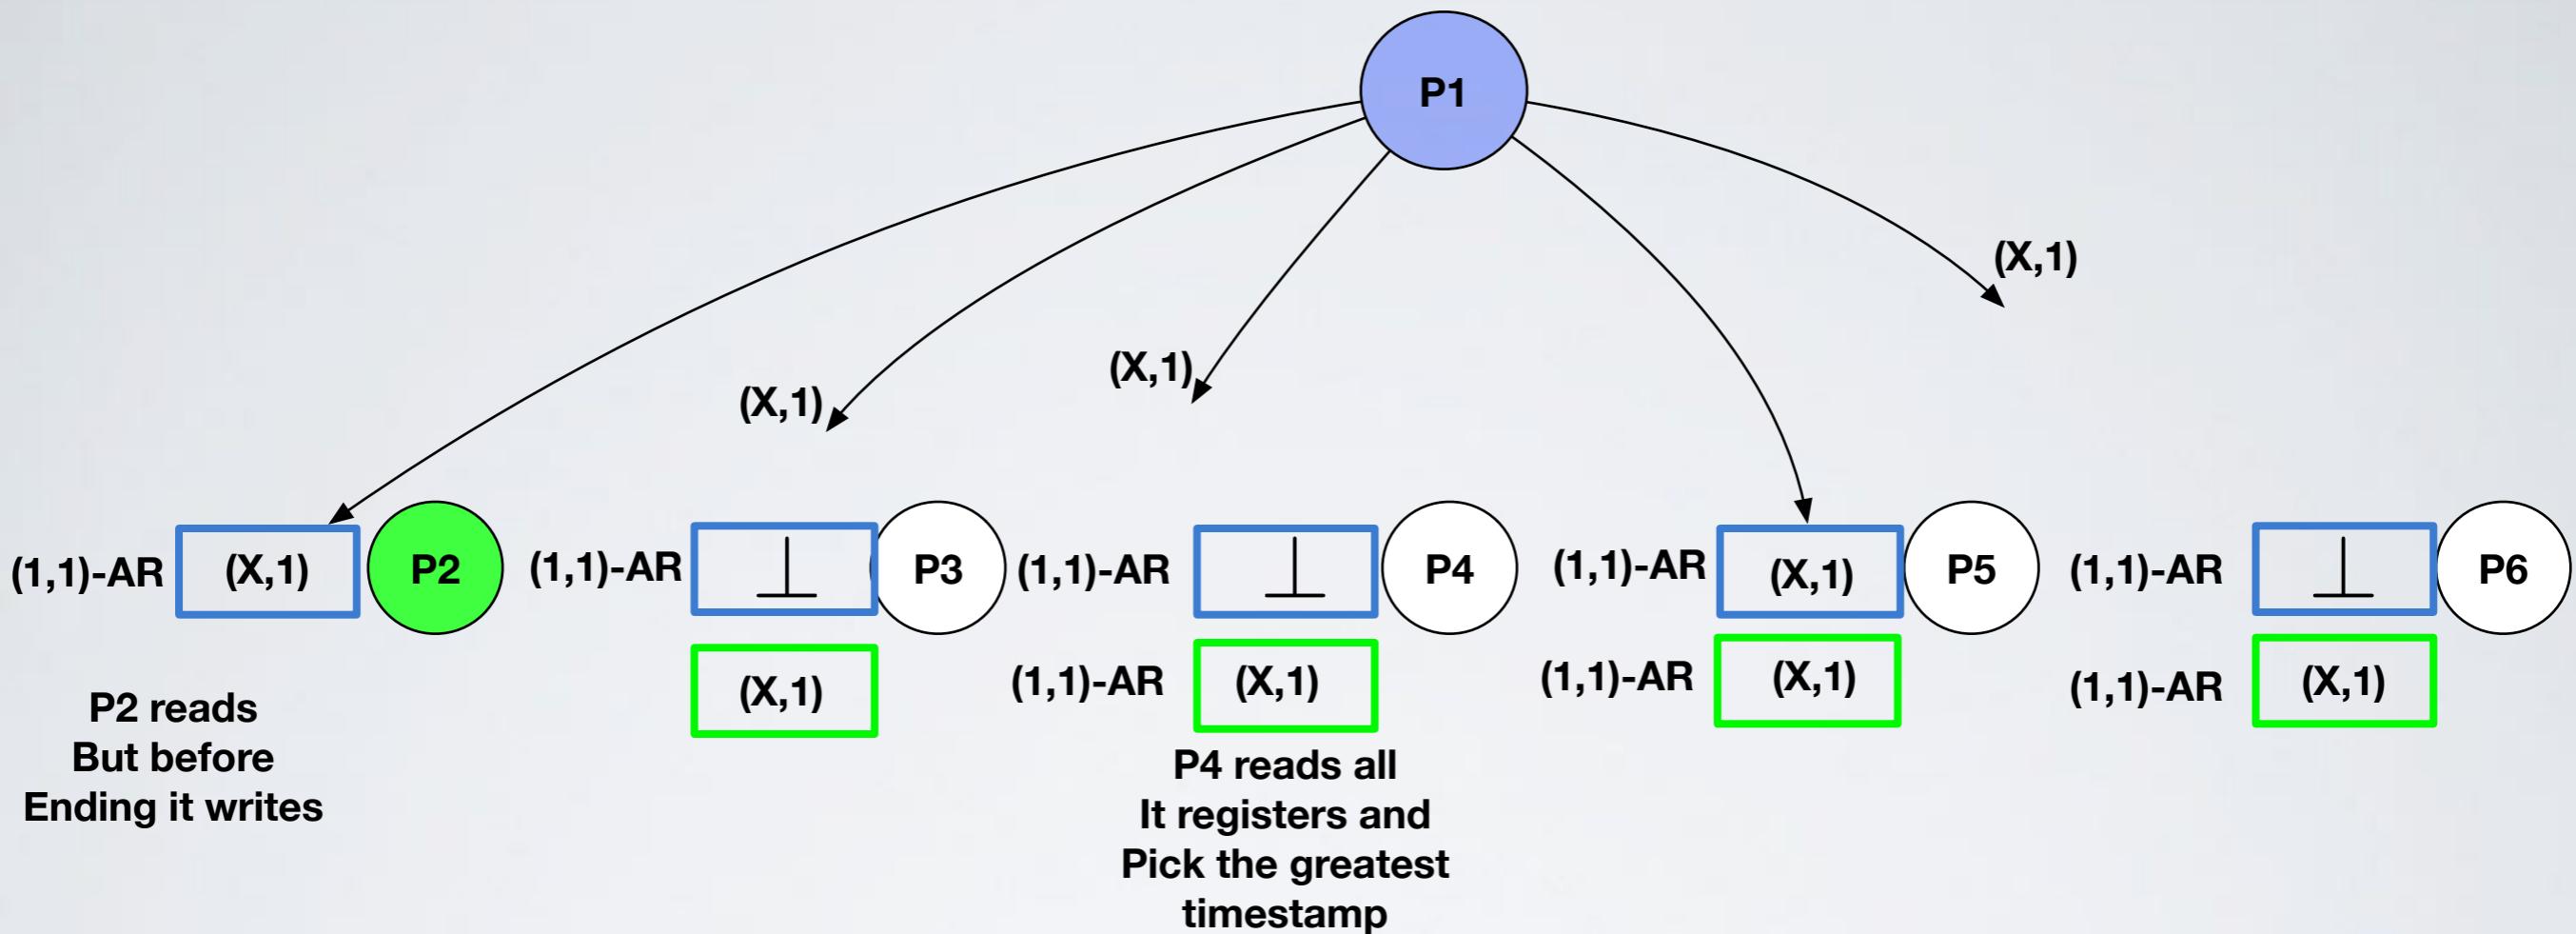
PHASE 2 - FROM (1,1)-AR TO (1,N)-AR



Problem: multiple writes on different registers are not instantaneous, so a write can appear on P2 but not on P4.

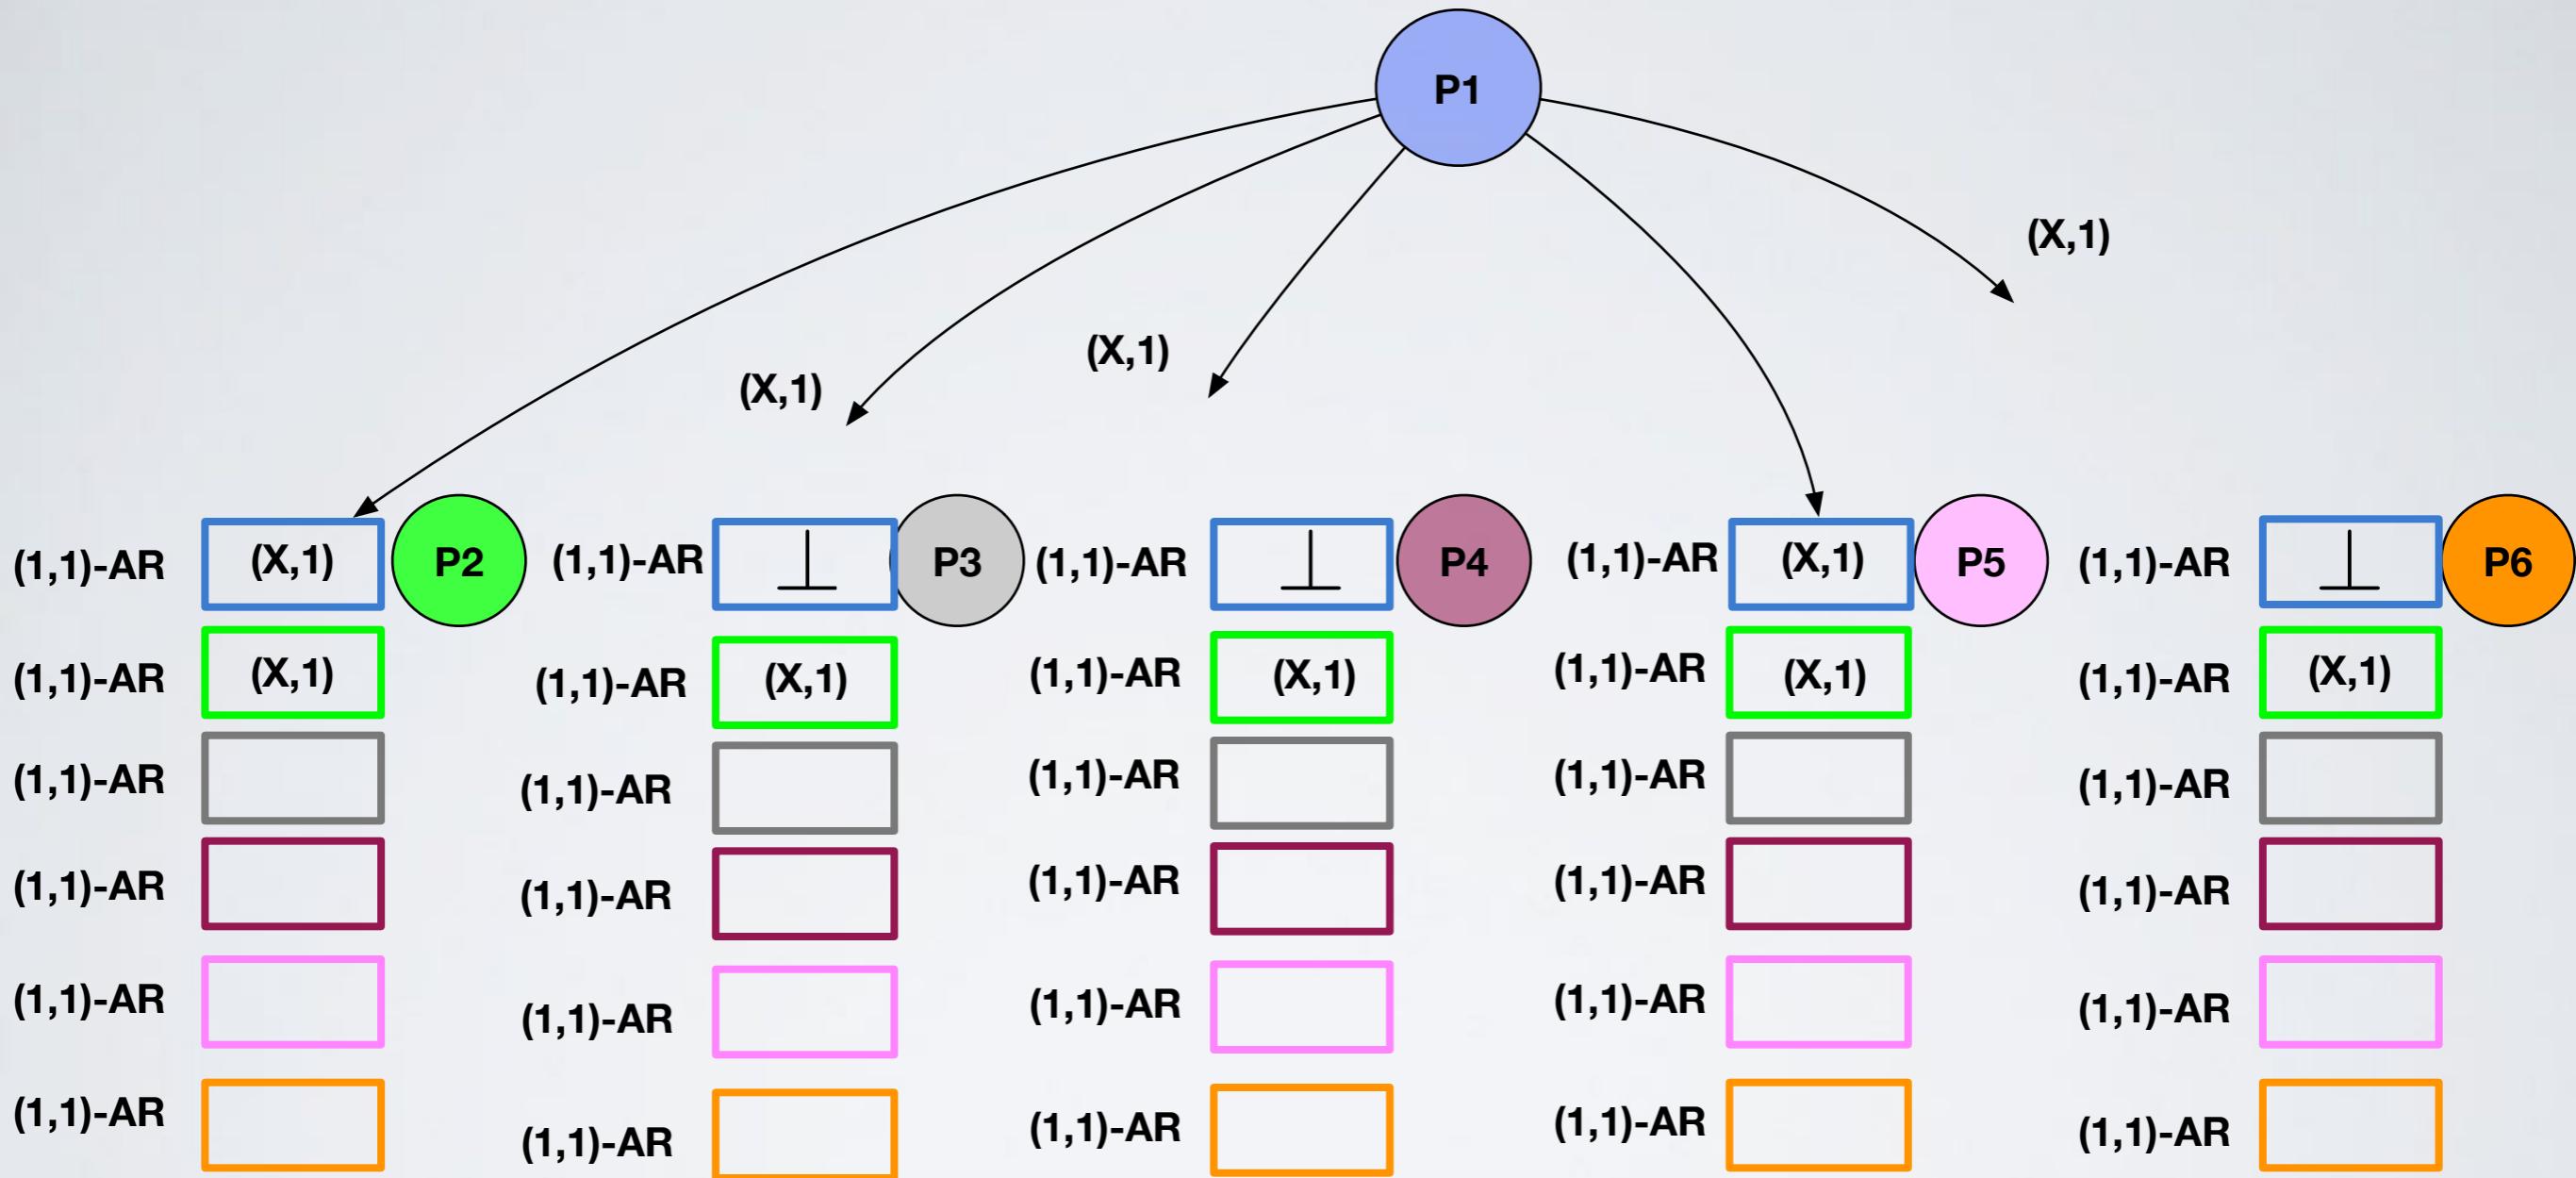
How to fix it? When P2 reads, before completing its operation writes the Value X on the other processes using another register

PHASE 2 - FROM (1,1)-AR TO (1,N)-AR



We need timestamp to decide the last value. A process reads the incoming registers In which P_1 and P_2 can write.

PHASE 2 - FROM $(1, 1)$ -AR TO $(1, N)$ -AR



The same case could happen for other processes, so we need a matrix of registers.

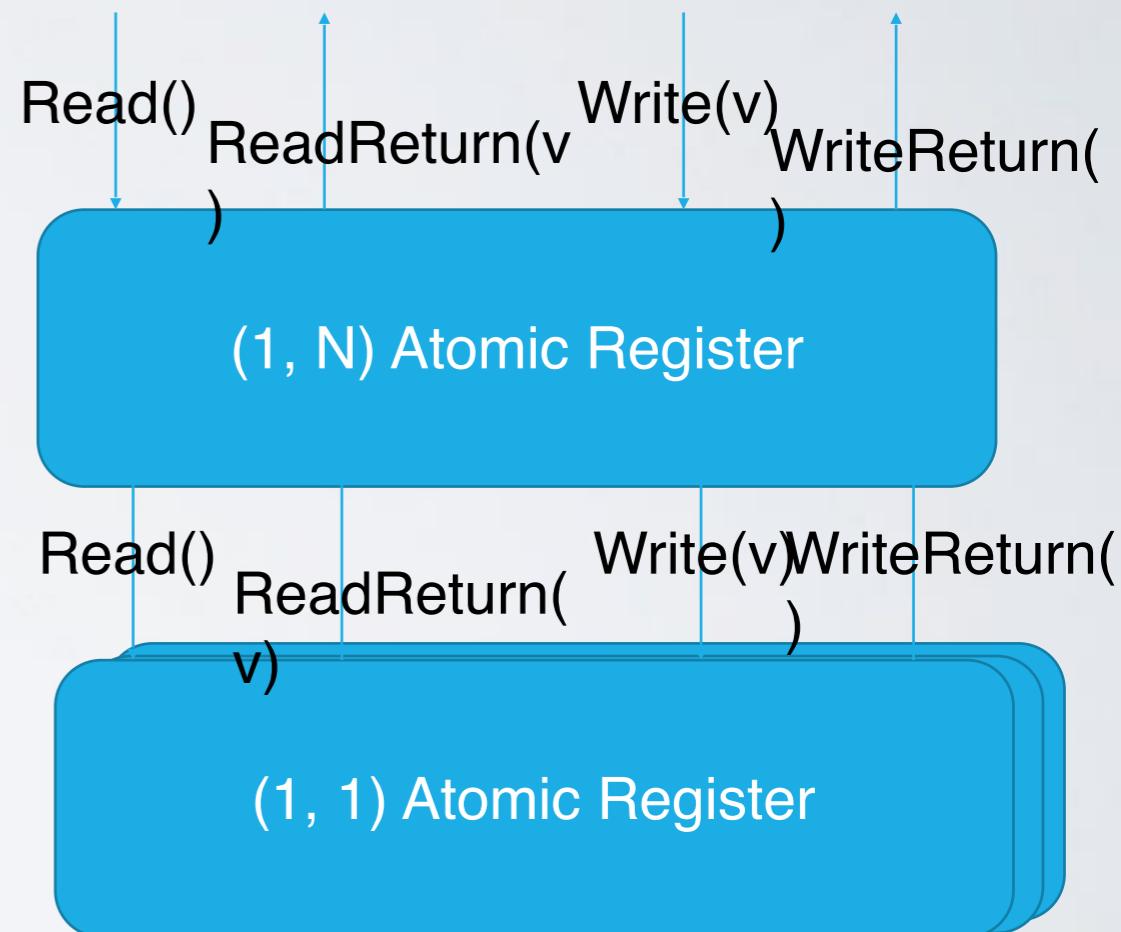
PHASE 2 - FROM (1,1)-AR TO (1,N)-AR

(1,1) atomic \rightarrow (1,N) atomic

- R_A : N (1,1) registers that connect the writer to the readers
- R_B : N^2 (1,1) registers that connect readers

Idea:

- Write writes val in registers R_A
- Reader
 - reads from 1 R_A register and from N R_B registers
 - chooses the value val with the largest timestamp
 - writes val in N R_B registers
 - returns val



PHASE 2 - FROM (1,1)-AR TO (1,N)-AR

upon event $\langle onar, Init \rangle$ do

$ts := 0;$

$acks := 0;$

$writing := \text{FALSE};$

$readval := \perp;$

$readlist := [\perp]^N;$

forall $q \in \Pi, r \in \Pi$ do

 Initialize a new instance $ooar.q.r$ of (1,1)-AtomicRegister
 with writer r and reader q ;

OOAR.X.Y is a single register in a matrix
X is the reader
And Y is the writer

PHASE 2

upon event $\langle onar, \text{Write} \mid v \rangle$ **do**
 $ts := ts + 1;$
 $writing := \text{TRUE};$
forall $q \in \Pi$ **do**
 trigger $\langle ooar.q.self, \text{Write} \mid (ts, v) \rangle$;

upon event $\langle ooar.q.self, \text{WriteReturn} \rangle$ **do**
 $acks := acks + 1;$
if $acks = N$ **then**
 $acks := 0;$
 if $writing = \text{TRUE}$ **then**
 trigger $\langle onar, \text{WriteReturn} \rangle$;
 $writing := \text{FALSE};$
else
 trigger $\langle onar, \text{ReadReturn} \mid readval \rangle$;

upon event $\langle onar, \text{Read} \rangle$ **do**
forall $r \in \Pi$ **do**
 trigger $\langle ooar.self.r, \text{Read} \rangle$;

upon event $\langle ooar.self.r, \text{ReadReturn} \mid (ts', v') \rangle$ **do**
 $readlist[r] := (ts', v');$
if $\#(readlist) = N$ **then**
 $(maxts, readval) := \text{highest}(readlist);$
 $readlist := [\perp]^N;$
forall $q \in \Pi$ **do**
 trigger $\langle ooar.q.self, \text{Write} \mid (maxts, readval) \rangle$;

PROOF

Correctness:

Termination – from the termination of the (1,1) atomic register

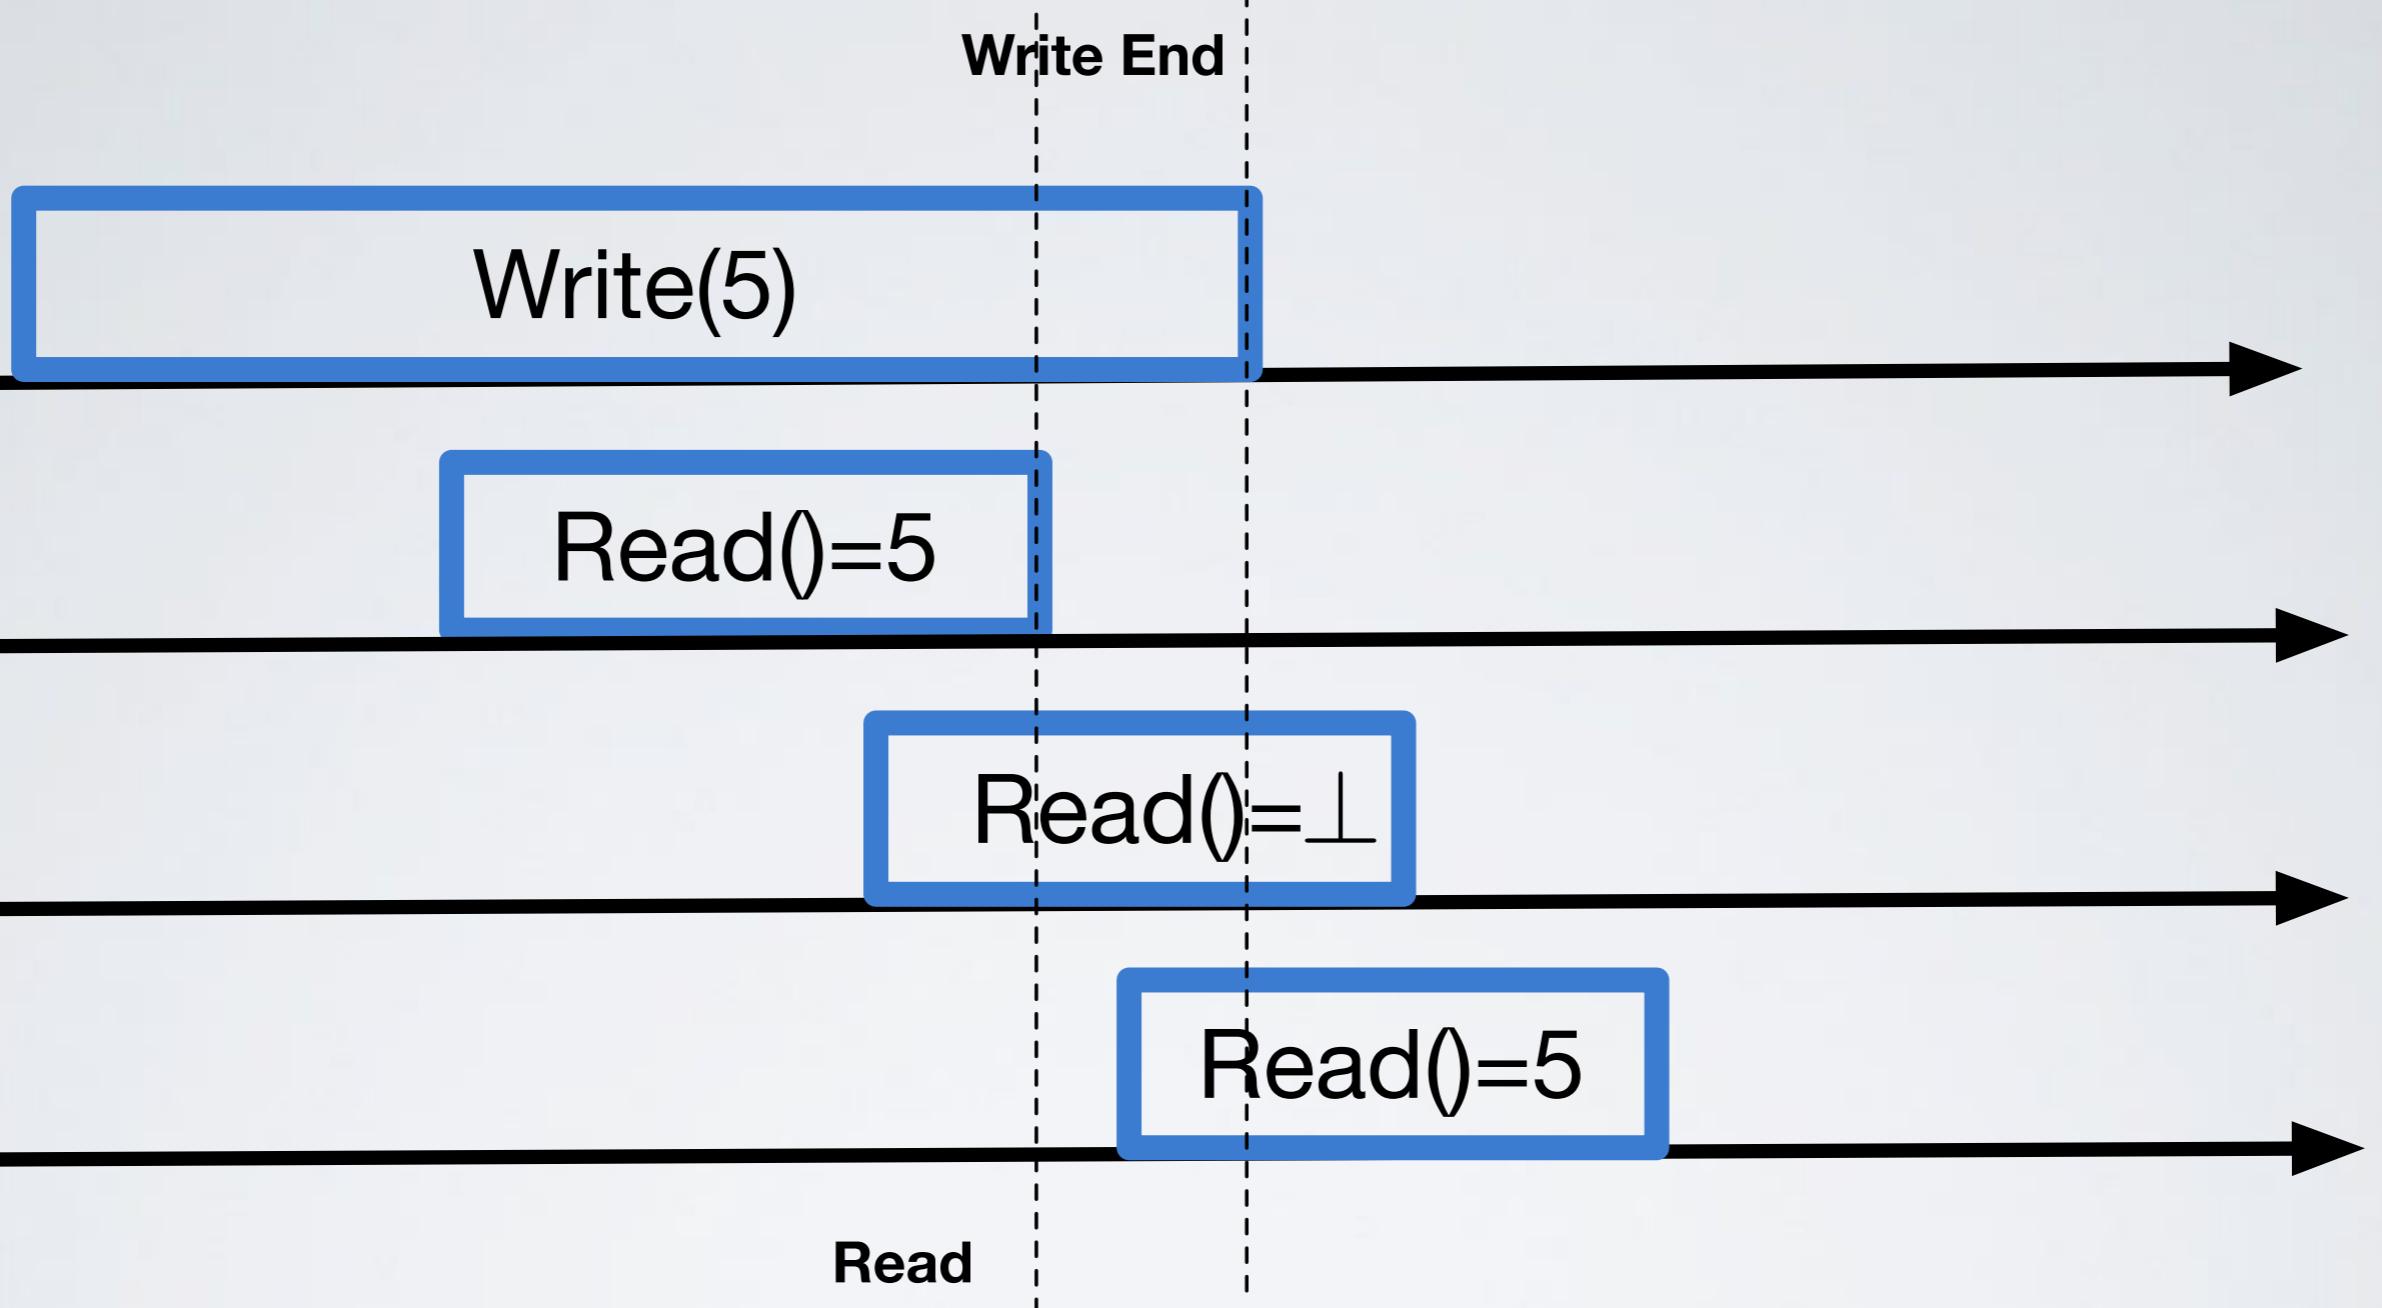
Validity – from the validity of the (1,1) atomic register.

Ordering - Consider a write operation w_1 which writes value v_1 with timestamp s_1 . Let w_2 be a write that starts after w_1 . Let v_2 and s_2 ($s_1 < s_2$) be the value and the timestamp corresponding to w_2 .

Let assume that a read returns v_2 : by the algorithm, for each j in $[1;N]$, p_i has written (s_2, v_2) in $\text{readers}[r; i; j]$.

For the ordering property of the underlying (1,1) atomic registers, each successive read will return a value with timestamp greater or equal to s_2 . Then s_1 cannot be returned.

LIN. POINTS OF THE ALGORITHM

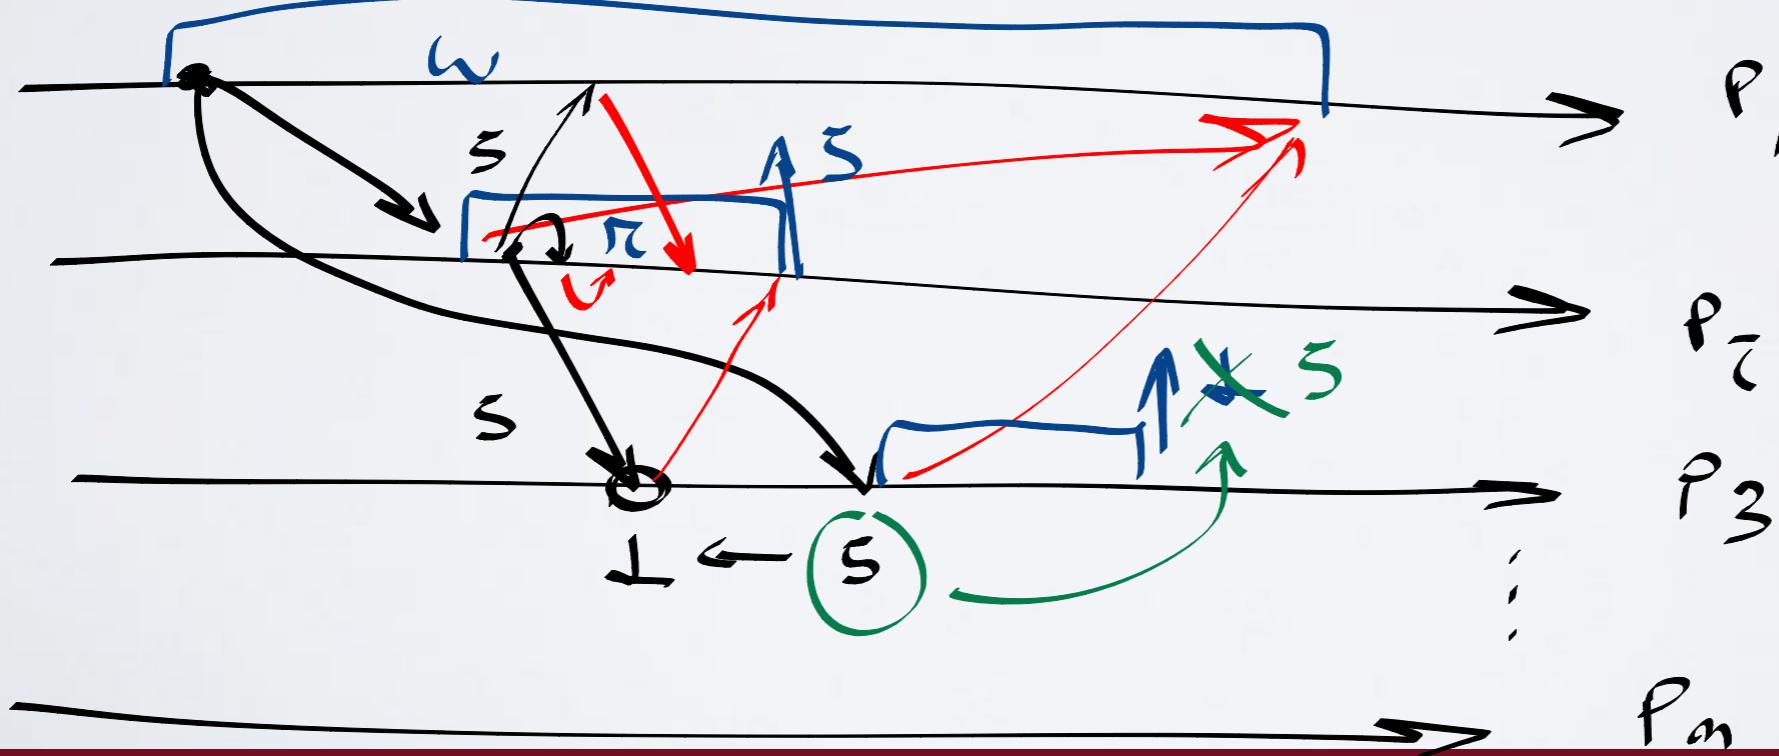


The LP of the write is the minimum between the end of the write and the end of the first read that sees the value of the write.

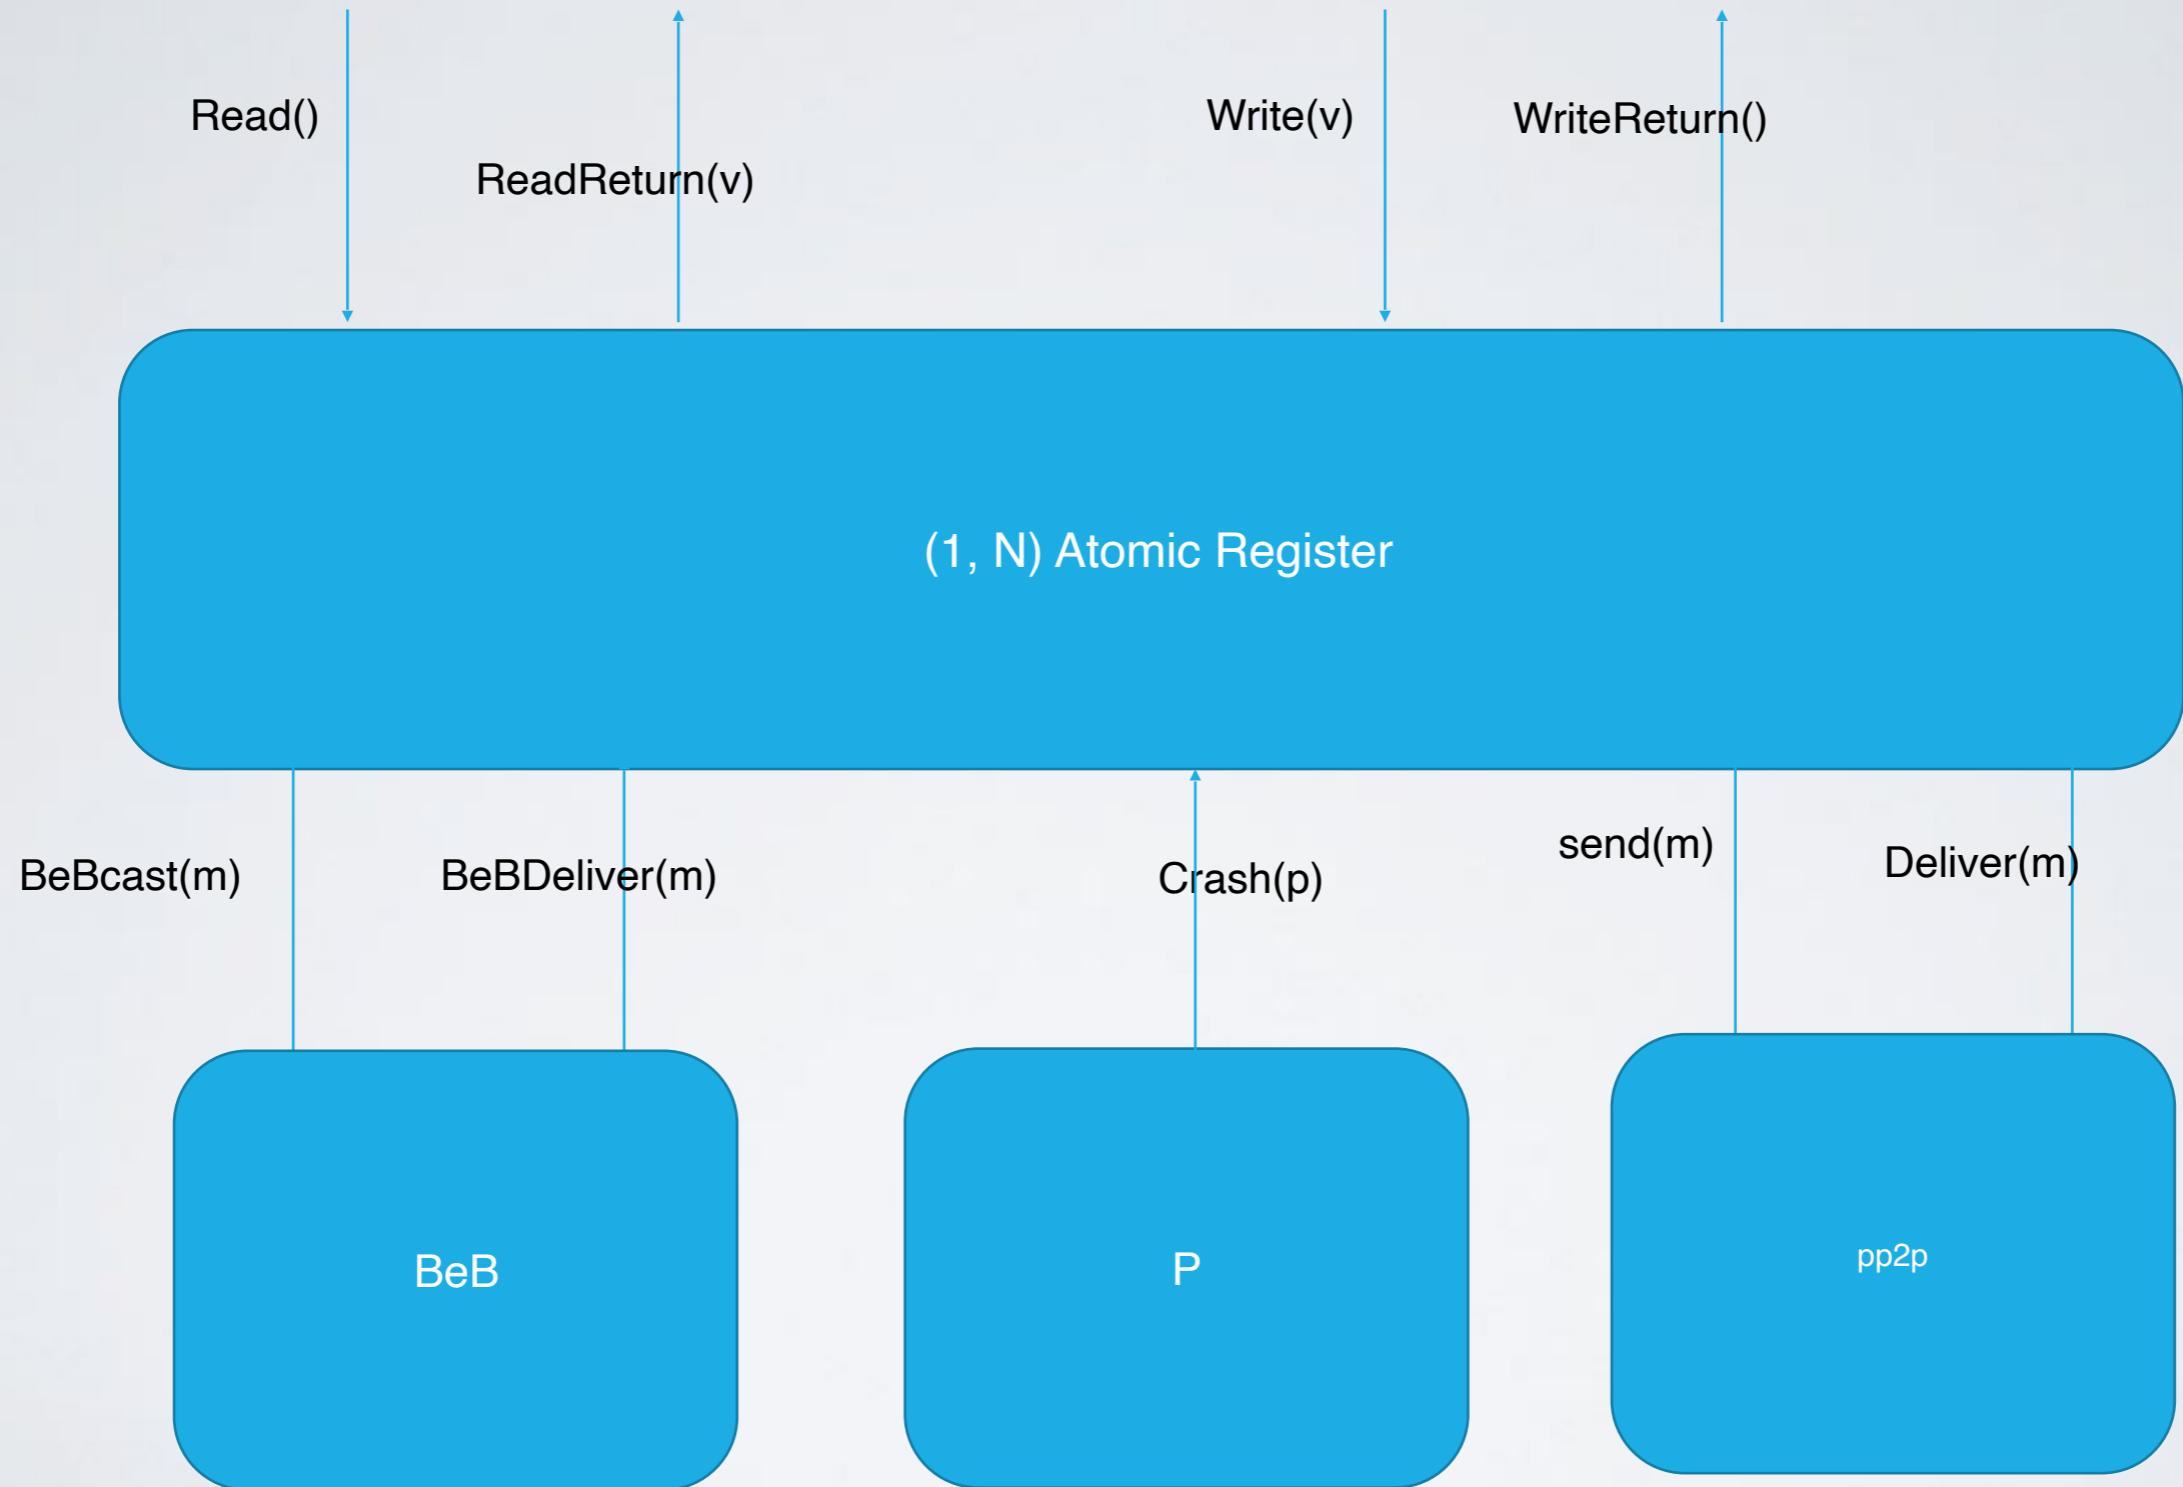
can happen that the broadcasting message ω arrives after $p_3(\text{read})$ but in that case $p_2(\text{read})$ can't terminate because doesn't receive all ACKs $\Rightarrow p_2(\text{read})$ and $p_3(\text{read})$

ATOMIC REGISTERS IMPLEMENTATIONS USING MESSAGE PASSING.

to ensure that a read with an older value disseminate faulty messages, we use timestamps



READ-IMPOSE WRITE-ALL



READ-IMPOSE WRITE-ALL

The algorithm is a modified version of the Read-One Write-All (1,N) Regular Register

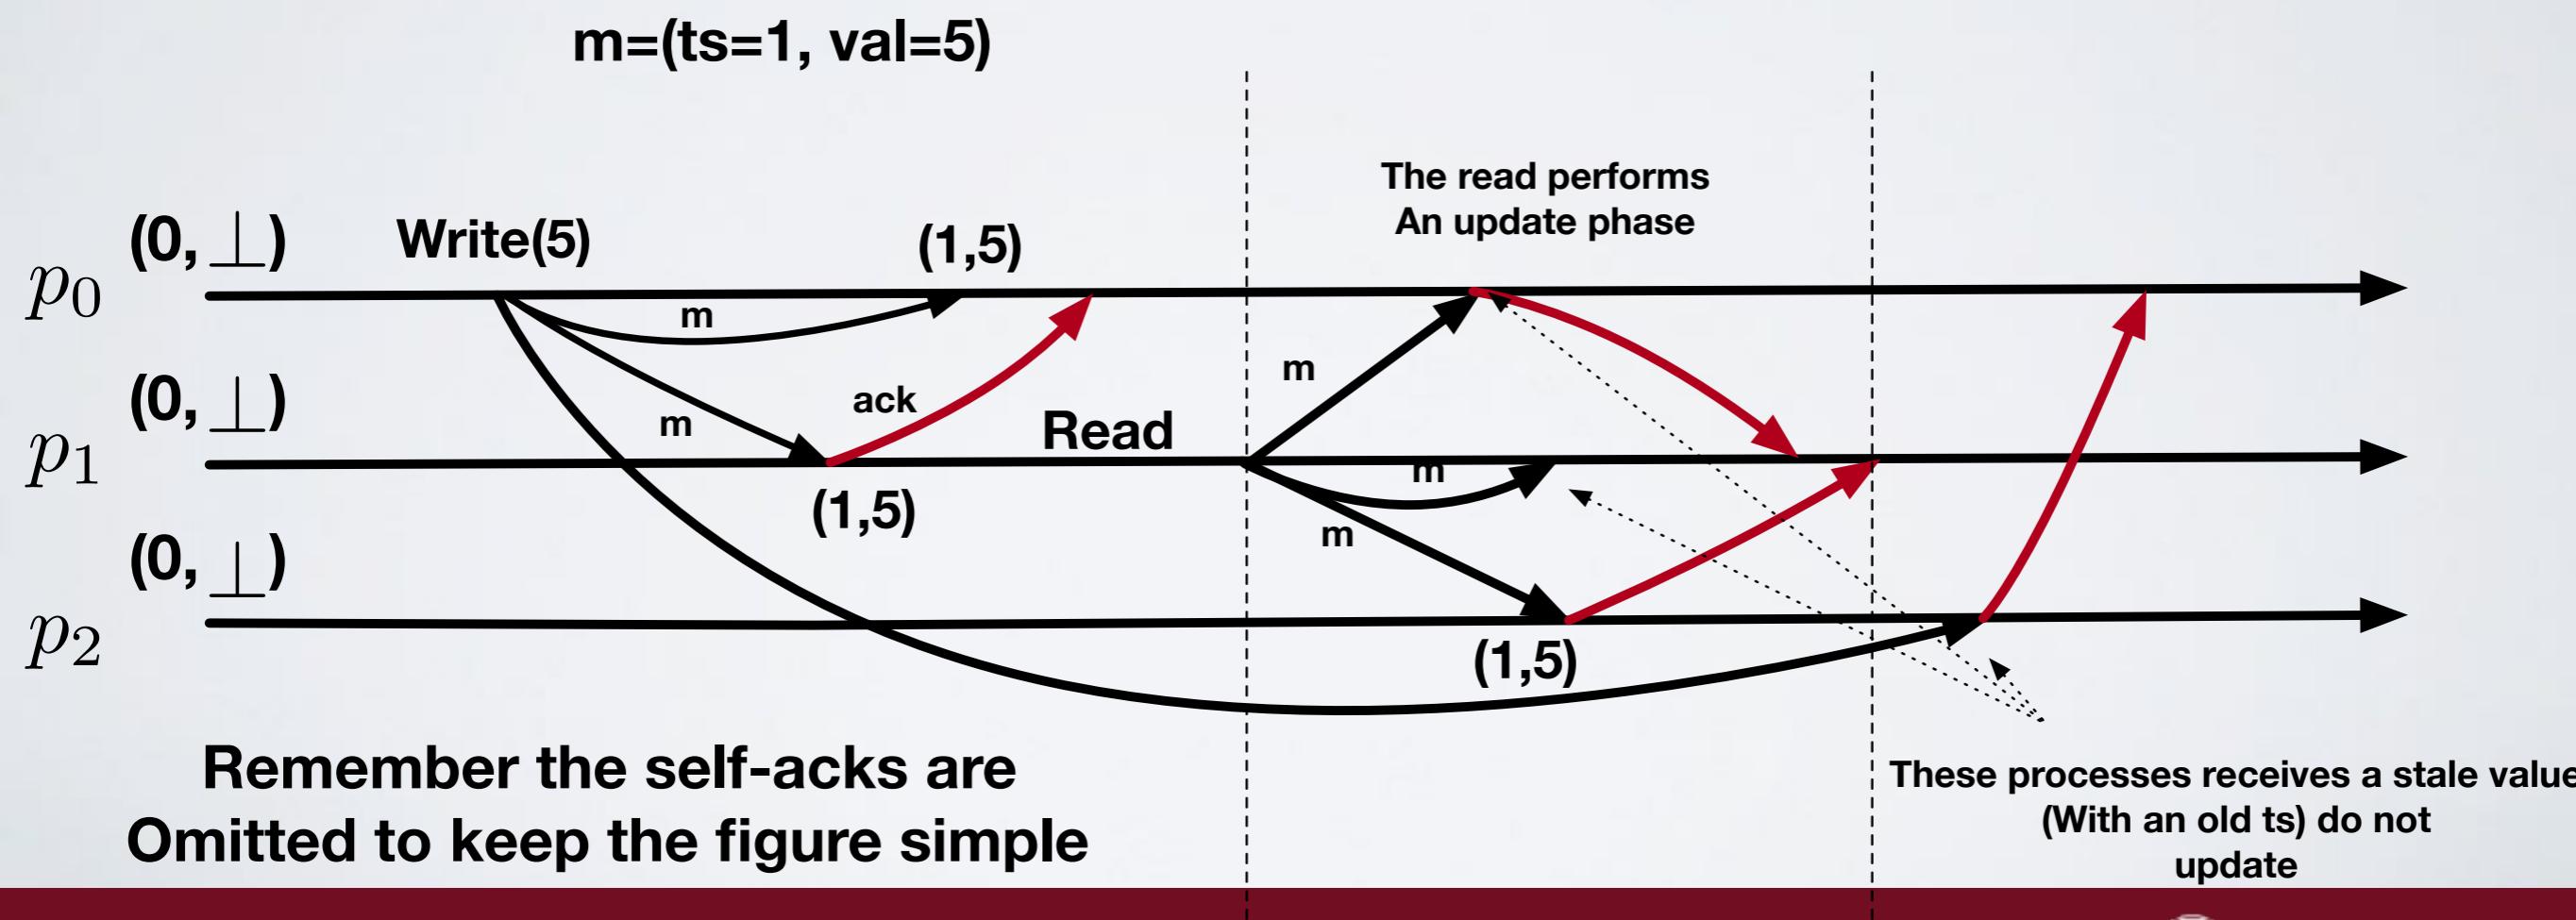
IDEA: “the read operation writes”

The algorithm is called “Read-Impose Write-All” because a read operation imposes to all correct processes to update their local copy of the register with the value read, unless they store a more recent value

READ-IMPOSE WRITE-ALL

Algorithm structure:

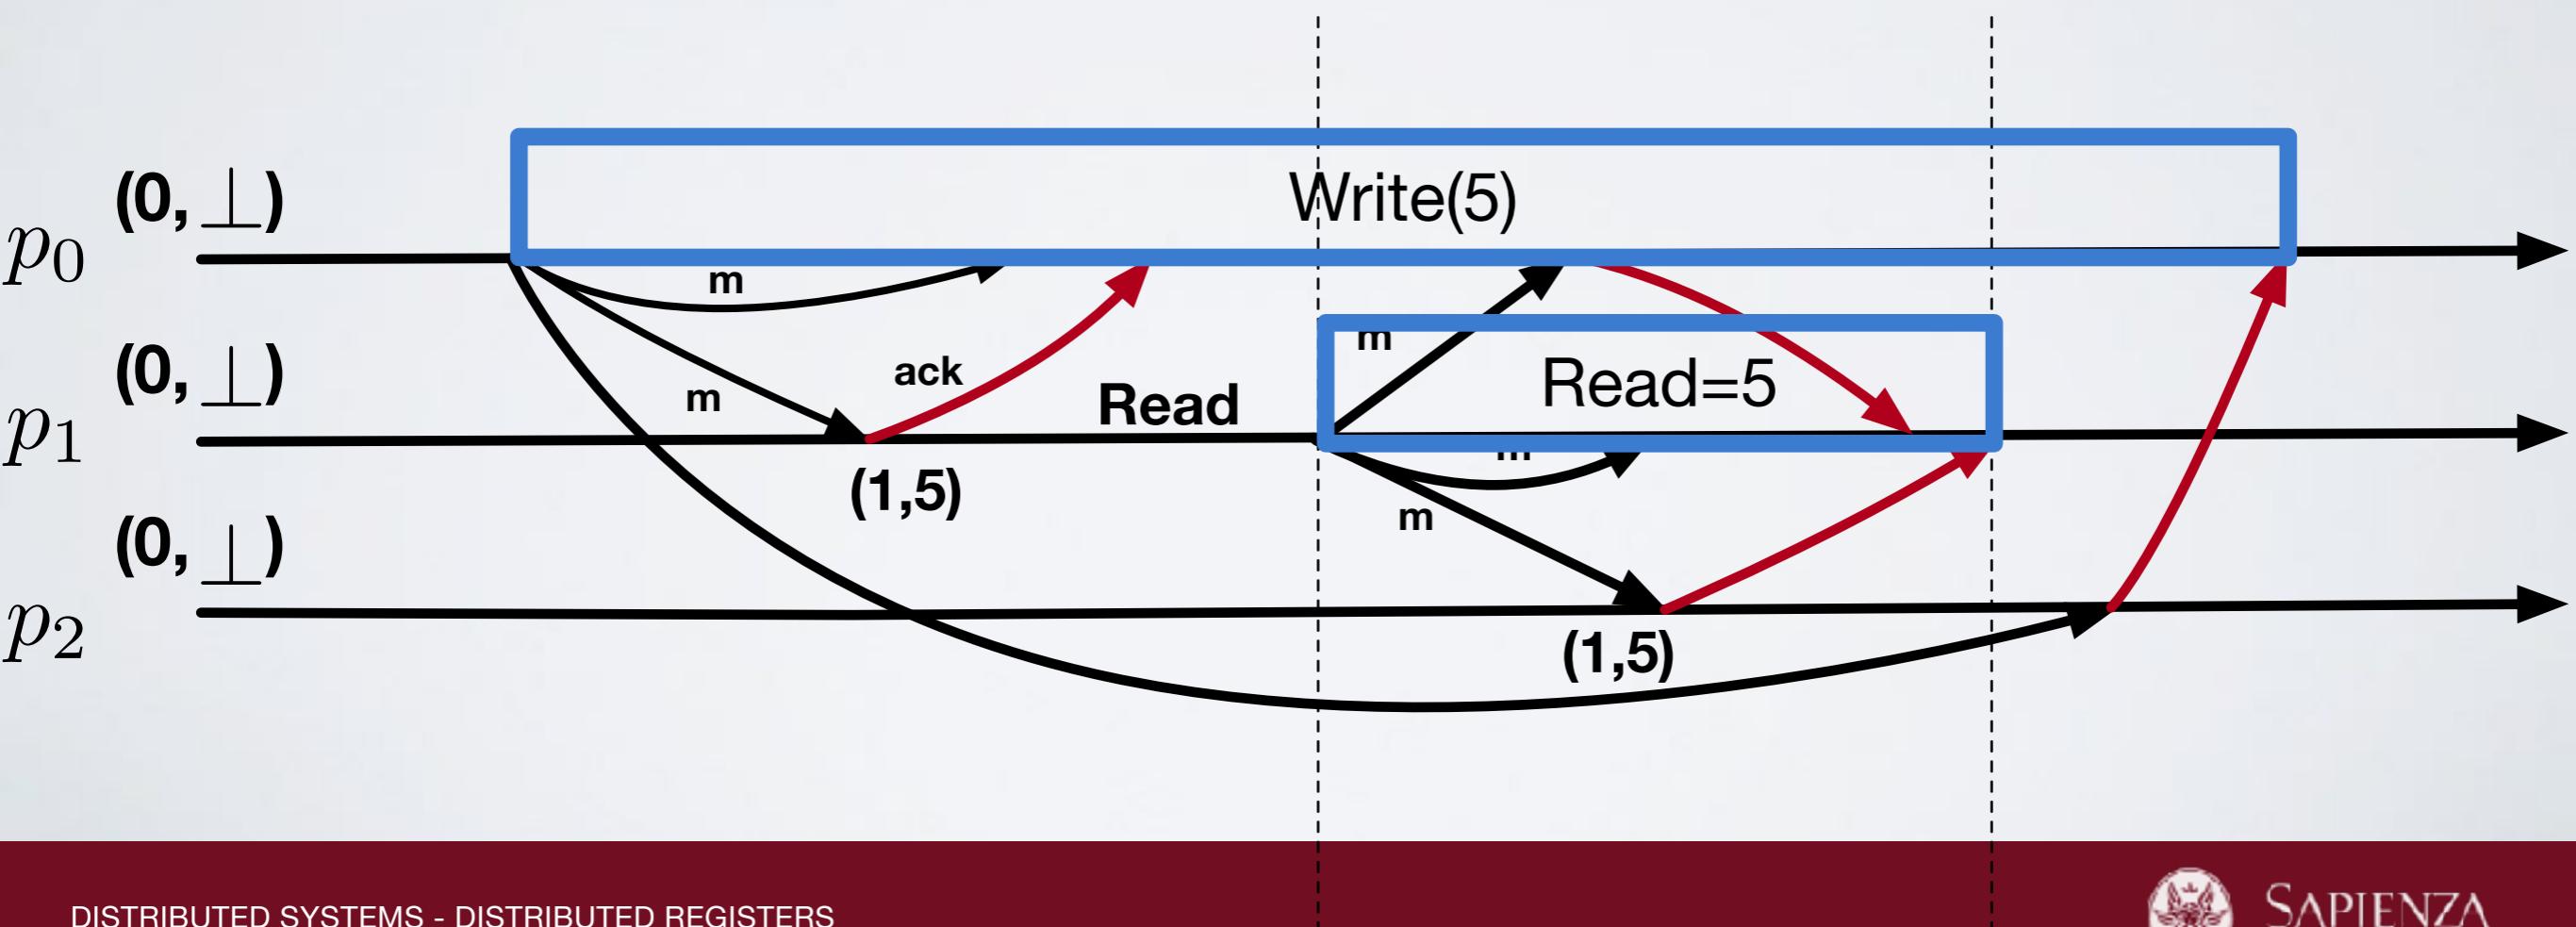
- Write: increase wts, and disseminate a write order to any one
- Read: read locally a (wts, val) , and disseminate a write order to any one
- Process: accept a write if your ts is less than the one in the order. Send ack in any case



READ-IMPOSE WRITE-ALL

Algorithm structure:

- Write: increase wts, and disseminate a write order to any one
- Read: read locally a (wts, val) , and disseminate a write order to any one
- Process: accept a write if your ts is less than the one in the order. Send ack in any case



READ-IMPOSE WRITE-ALL

2n ops WRITE

2n ops READ

upon event $\langle onar, Init \rangle$ **do**

$(ts, val) := (0, \perp);$

$correct := \Pi;$

$writeset := \emptyset;$

$readval := \perp;$

$reading := \text{FALSE};$

upon event $\langle \mathcal{P}, Crash \mid p \rangle$ **do**

$correct := correct \setminus \{p\};$

upon event $\langle onar, Read \rangle$ **do**

$reading := \text{TRUE};$

$readval := val;$

trigger $\langle beb, Broadcast \mid [\text{WRITE}, ts, val] \rangle;$

$ts + 1$

upon event $\langle onar, Write \mid v \rangle$ **do**

trigger $\langle beb, Broadcast \mid [\text{WRITE}, ts + 1, v] \rangle;$

if is sended an older value the message is discarded

upon event $\langle beb, Deliver \mid p, [\text{WRITE}, ts', v'] \rangle$ **do**

|| if $ts' > ts$ **then**

$(ts, val) := (ts', v');$

trigger $\langle pl, Send \mid p, [\text{ACK}] \rangle;$

upon event $\langle pl, Deliver \mid p, [\text{ACK}] \rangle$ **then**

$writeset := writeset \cup \{p\};$

upon $correct \subseteq writeset$ **do**

$writeset := \emptyset;$

if $reading = \text{TRUE}$ **then**

$reading := \text{FALSE};$

trigger $\langle onar, ReadReturn \mid readval \rangle;$

else

trigger $\langle onar, WriteReturn \rangle;$

*



violates
visibility
(not a
Regular
Register)
and Ordering

Correctness:

- Termination – as for the Read-One Write-All (1,N) Regular Register.
- Validity - as for Read-One Write-All (1,N) Regular Register.
- Ordering – to complete a read operation, the reader process has to be sure that every other process has in its local copy of the register a value with timestamp bigger or equal of the timestamp of the value read. In this way, any successive read could not return an older value.

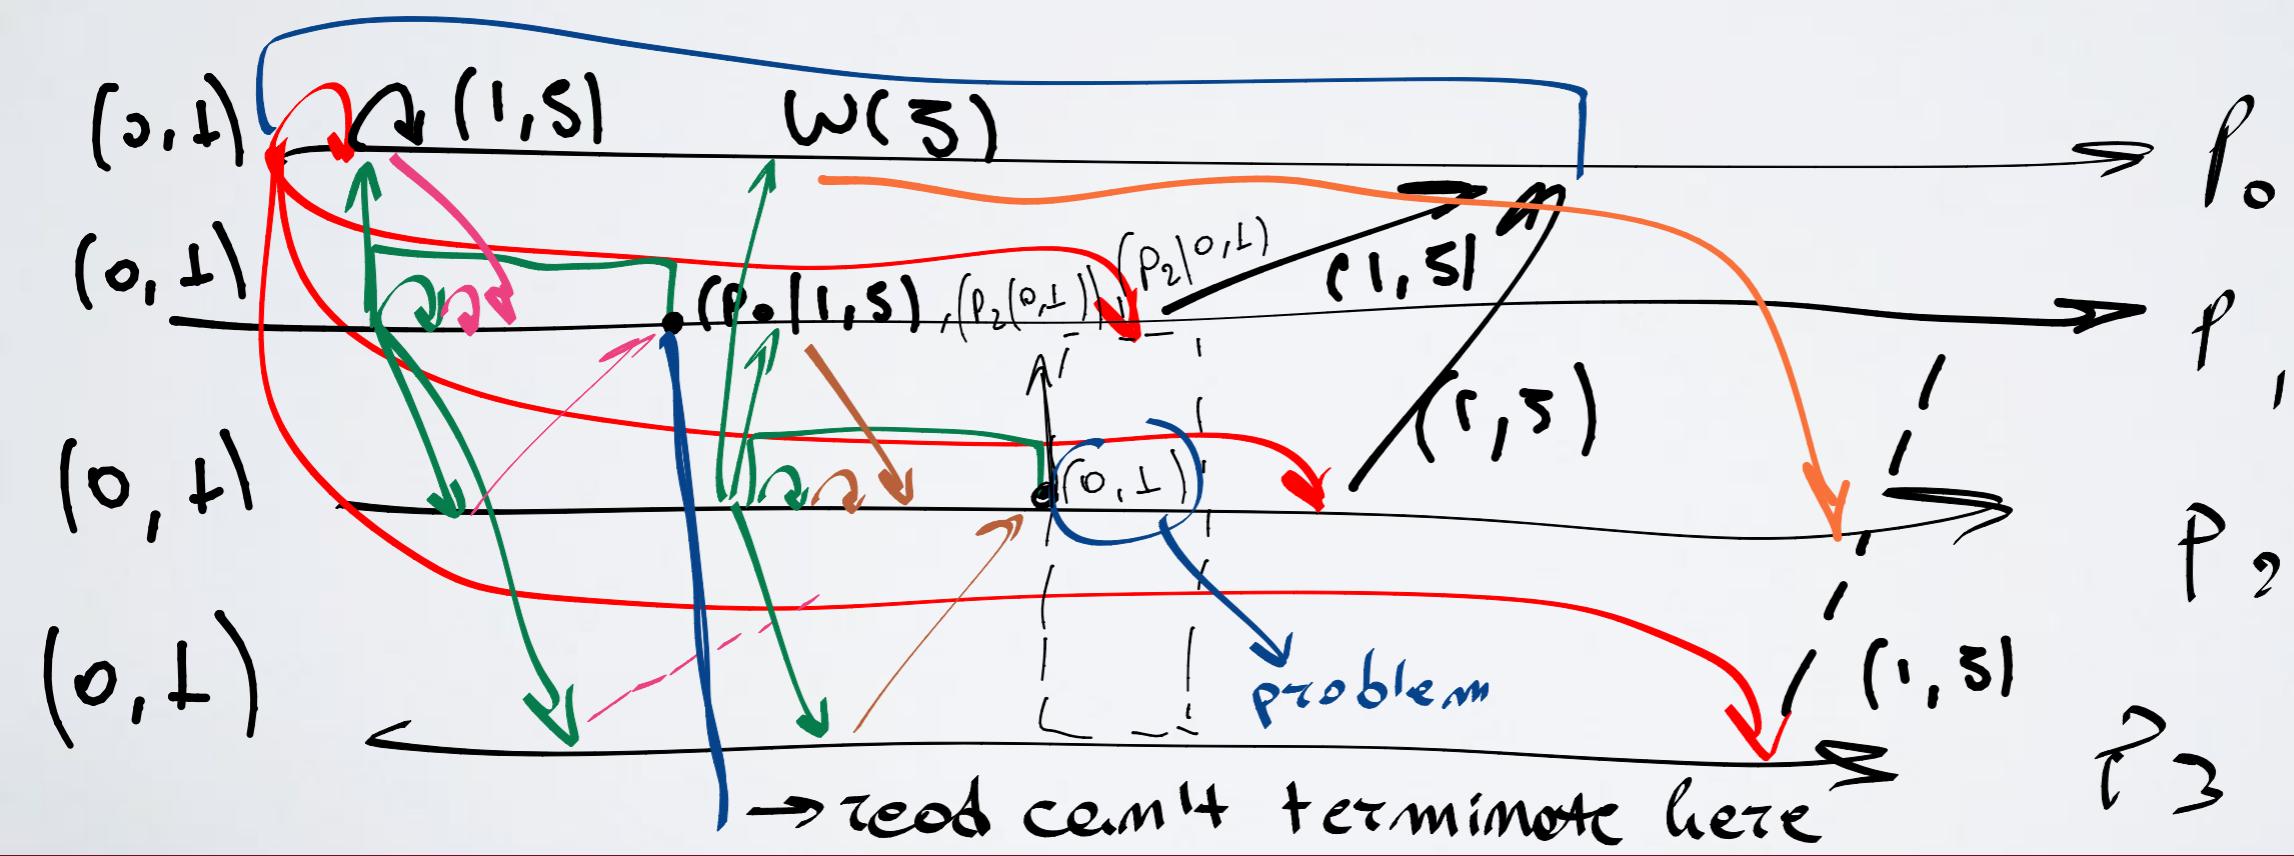
Performance:

- Write - a write requests at most $2N$ messages
- Read - a read requests at most $2N$ messages
- Number of steps? 2 for write or read

$1, N$ - Atomic Register



IMPLEMENTATION FOR FAIL-SILENT

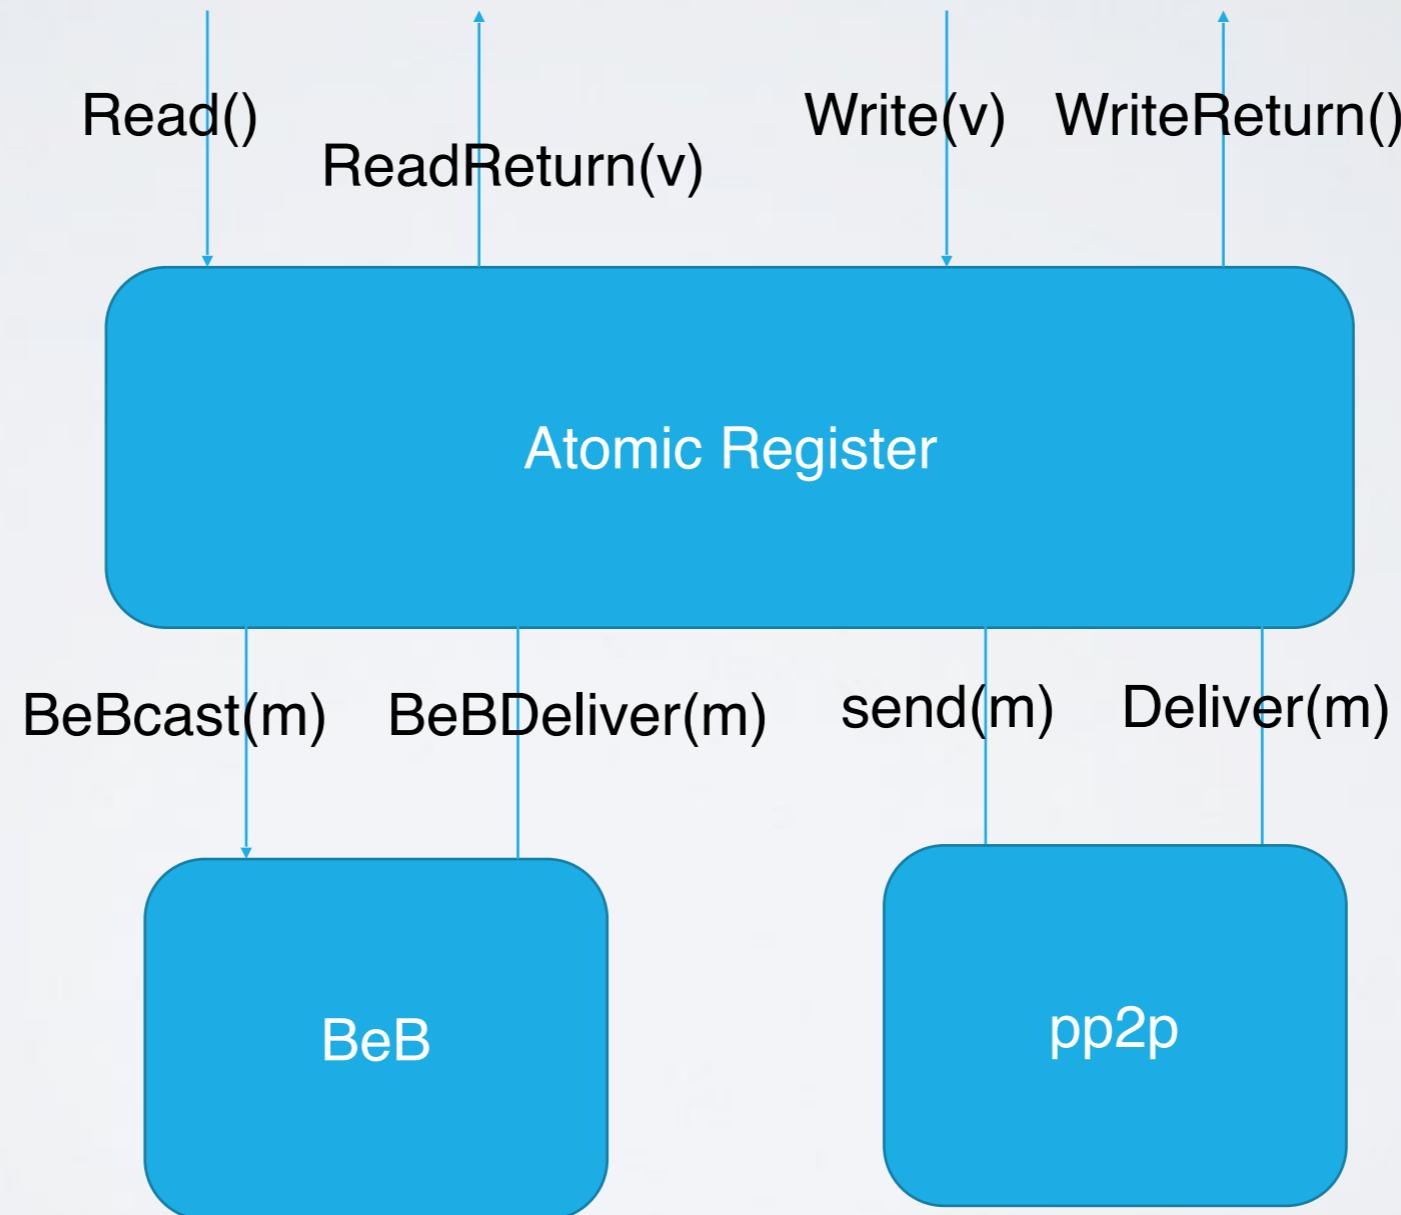


=> impose an additional phase using a broadcast to imp

READ-IMPOSE WRITE-MAJORITY

A majority of correct processes is assumed.

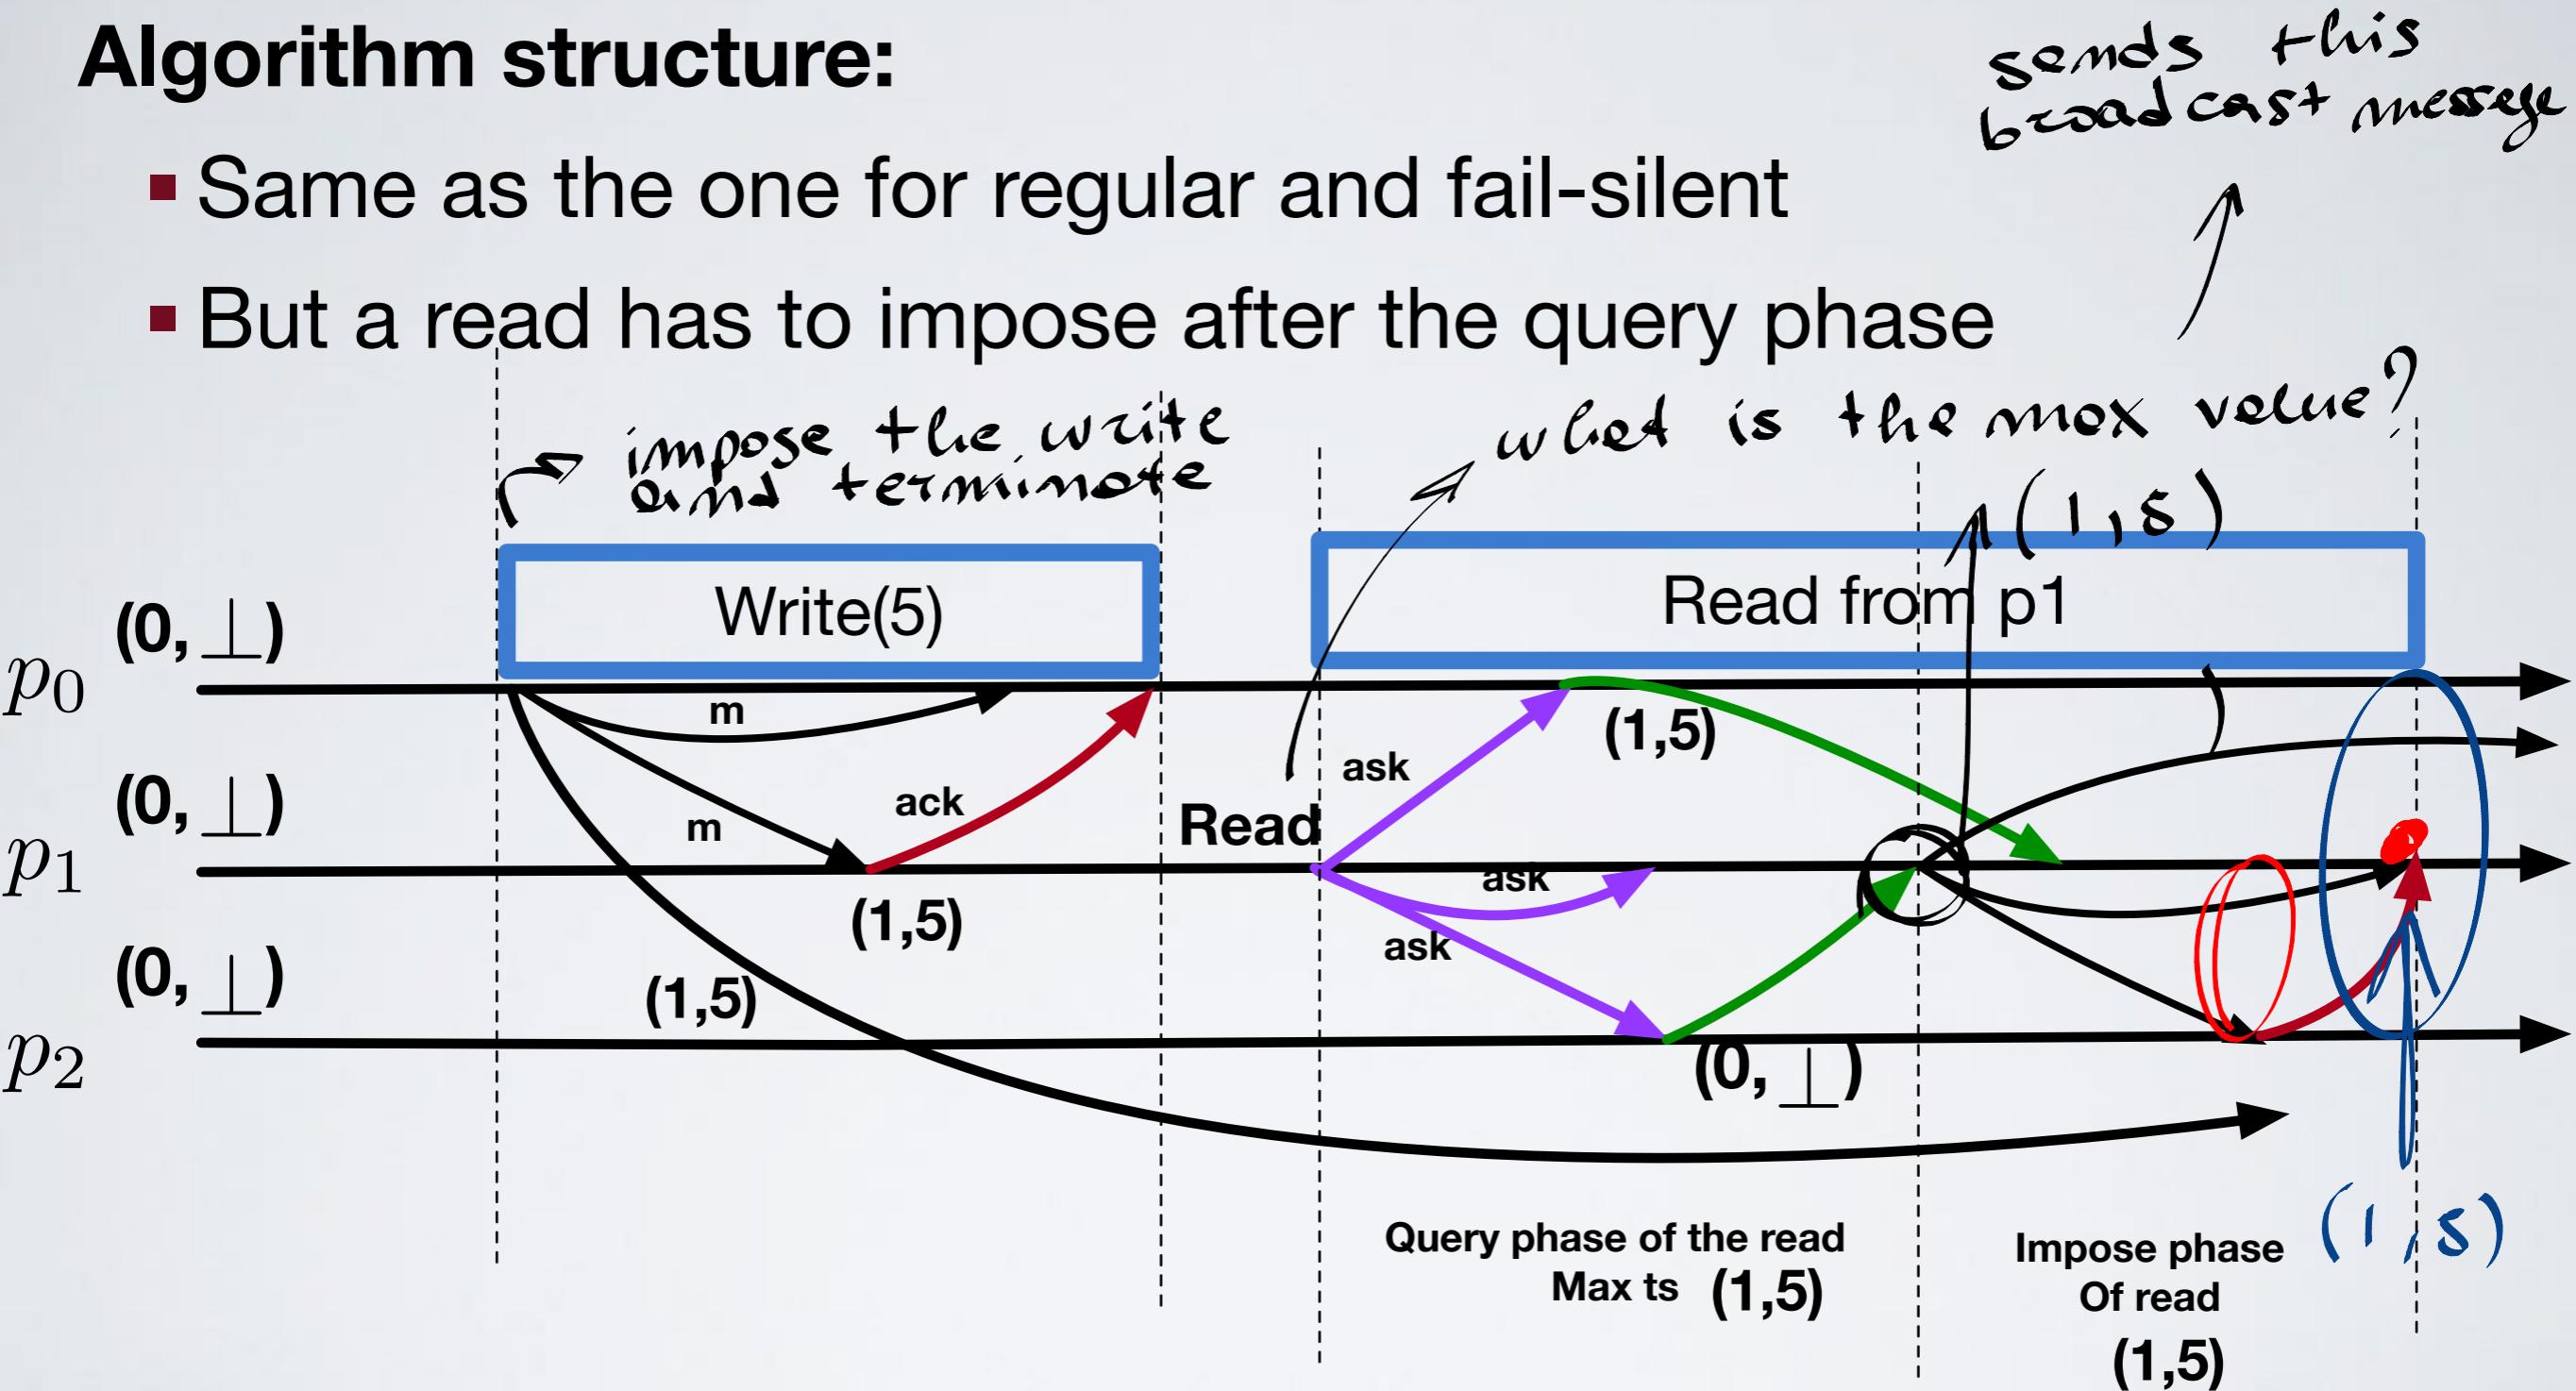
The algorithm is a variation of the Majority Voting (1,N) Regular Register



READ-IMPOSE WRITE-MAJORITY

Algorithm structure:

- Same as the one for regular and fail-silent
- But a read has to impose after the query phase



upon event $\langle onar, Init \rangle$ **do**

$(ts, val) := (0, \perp);$

$wts := 0;$

$acks := 0;$ \rightsquigarrow to end has to be the majority

$rid := 0;$

$readlist := [\perp]^N;$

$readval := \perp;$

$reading := \text{FALSE};$

upon event $\langle onar, Read \rangle$ **do**

$rid := rid + 1;$ \rightsquigarrow I have to distinguish the acks

$acks := 0;$

$readlist := [\perp]^N;$

$reading := \text{TRUE};$

trigger $\langle beb, Broadcast \mid [\text{READ}, rid] \rangle;$

upon event $\langle beb, Deliver \mid p, [\text{READ}, r] \rangle$ **do**

trigger $\langle pl, Send \mid p, [\text{VALUE}, r, ts, val] \rangle;$

upon event $\langle pl, Deliver \mid q, [\text{VALUE}, r, ts', v'] \rangle$ **such that** $r = rid$ **do**

$readlist[q] := (ts', v');$

if $\#(readlist) > N/2$ **then**

* $(maxts, readval) := \text{highest}(readlist);$ \rightarrow take the value with the highest timestamp

$readlist := [\perp]^N;$

trigger $\langle beb, Broadcast \mid [\text{WRITE}, rid, maxts, readval] \rangle;$

collecting
readings

upon event $\langle onar, \text{Write} | v \rangle$ **do**

$rid := rid + 1;$

$wts := wts + 1;$

$acks := 0;$

trigger $\langle beb, \text{Broadcast} | [\text{WRITE}, rid, wts, v] \rangle;$

use > ACVs
for the

make the other
processes to know
an older

that value is not
value

check
value is
older than
the

upon event $\langle beb, \text{Deliver} | p, [\text{WRITE}, r, ts', v'] \rangle$ **do**

if $ts' > ts$ **then**

$(ts, val) := (ts', v');$

trigger $\langle pl, \text{Send} | p, [\text{ACK}, r] \rangle;$

active?
1

upon event $\langle pl, \text{Deliver} | q, [\text{ACK}, r] \rangle$ **such that** $r = rid$ **do**

$acks := acks + 1;$

if $acks > N/2$ **then**

$acks := 0;$

if $reading = \text{TRUE}$ **then**

$reading := \text{FALSE};$

trigger $\langle onar, \text{ReadReturn} | readval \rangle;$



else

trigger $\langle onar, \text{WriteReturn} \rangle;$

collecting
ACKs

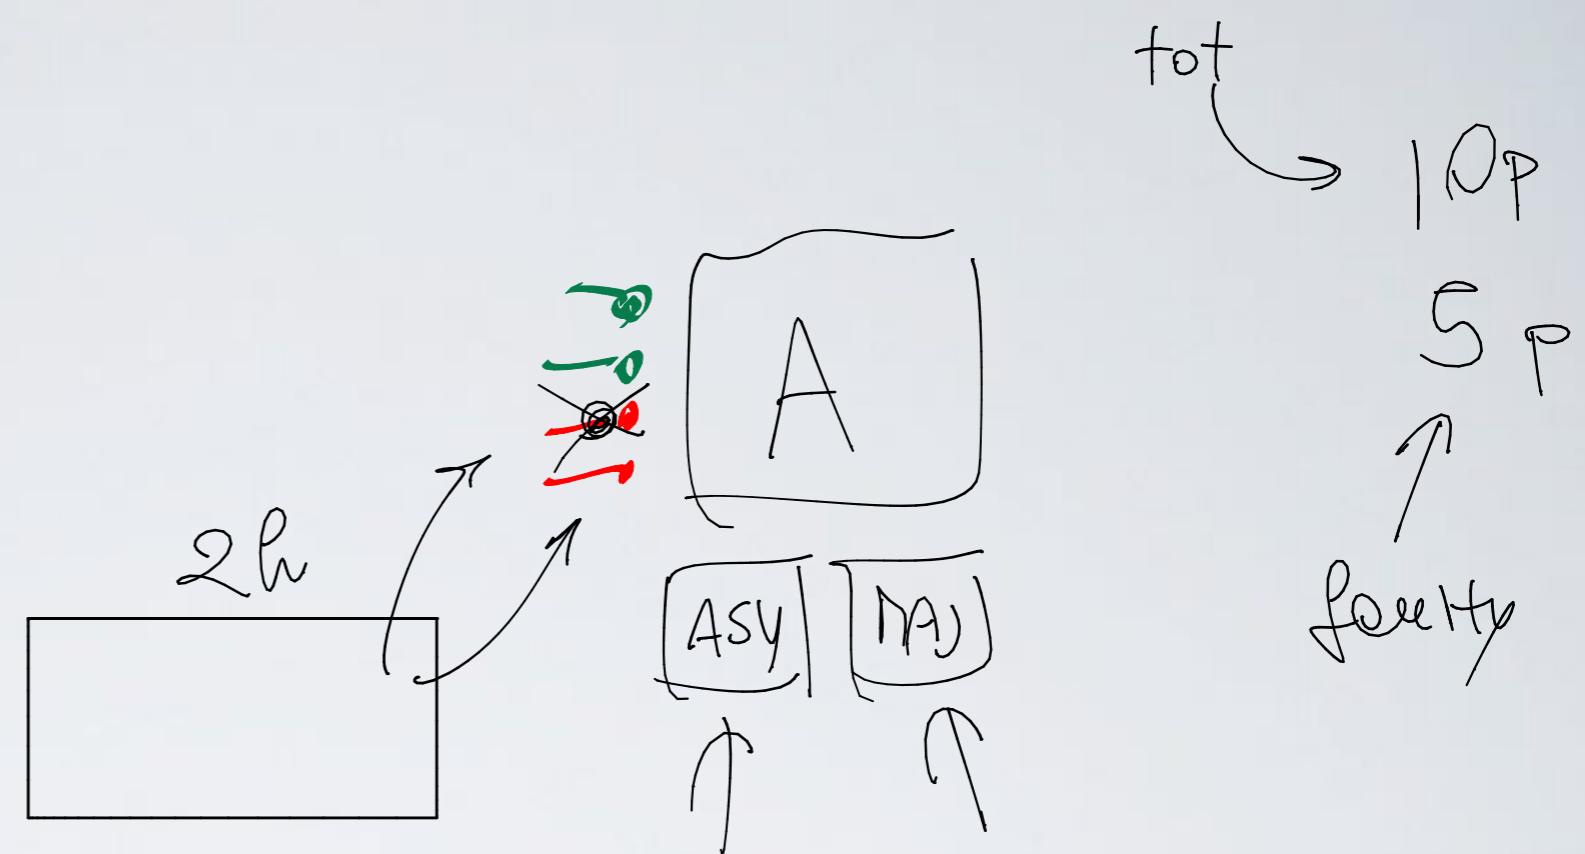
READ-IMPOSE WRITE-MAJORITY

Correctness:

- Termination – as Majority Voting (1,N) Regular Register
- Validity – as Majority Voting (1,N) Regular Register.
- Ordering – due to the fact that the read imposes the write of the value read to a majority of processes and to the property of intersection of quorums.

Performance:

- Write – at most $2N$ messages
- Read – at most $4N$ messages
- Step complexity: Write 2 steps, Read 4 steps.



COMPARISONS BETWEEN ALGORITHMS

REGULAR VS ATOMIC: PERFORMANCES

MESSAGE PASSING IMPLEMENTATIONS!

*number of
failures of
the system that
take zero com*

		(1,n)-Regular Register		(1,n)-Atomic Register	
		READ	WRITE	READ	WRITE
Fail-Stop (P)	Messages	0 (local)	2N	2N	2N
	Steps	0 (local)	2	2	2
	Resiliency	$f < N$		$f < N$	
Fail-Silent	Messages	2N	2N	4N	2N
	Steps	2	2	4	2
	Resiliency	$f < N/2$		$f < N/2$	

(N,N)- ATOMIC REGISTER

(N, N) -ATOMIC REGISTER SPEC

Module 4.3: Interface and properties of an (N, N) atomic register

Module:

Name: (N, N) -AtomicRegister, **instance** $nmar$.

Events:

Request: $\langle nmar, \text{Read} \rangle$: Invokes a read operation on the register.

Request: $\langle nmar, \text{Write} \mid v \rangle$: Invokes a write operation with value v on the register.

Indication: $\langle nmar, \text{ReadReturn} \mid v \rangle$: Completes a read operation on the register with return value v .

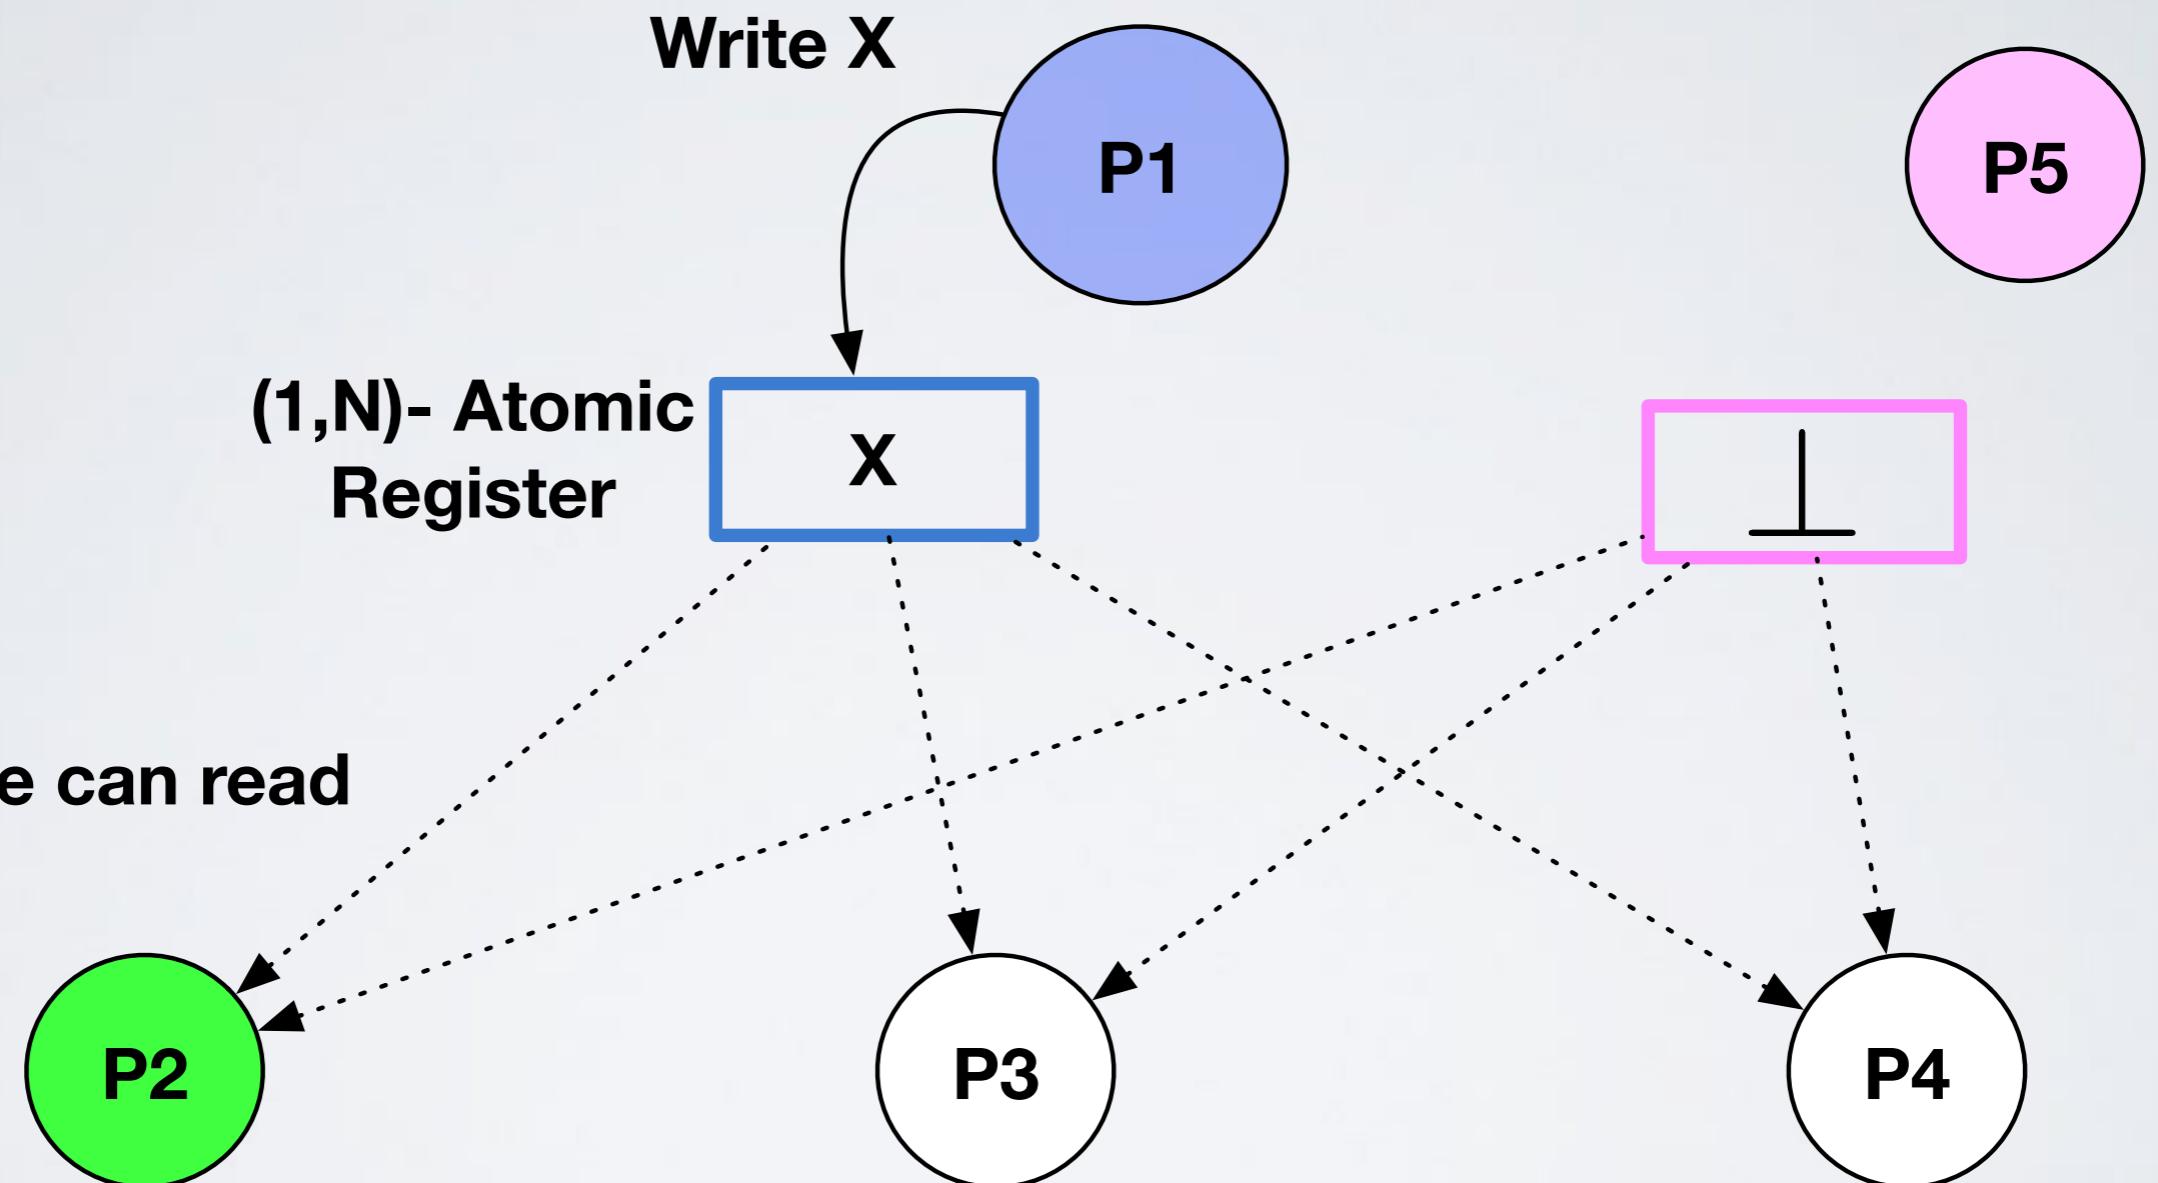
Indication: $\langle nmar, \text{WriteReturn} \rangle$: Completes a write operation on the register.

Properties:

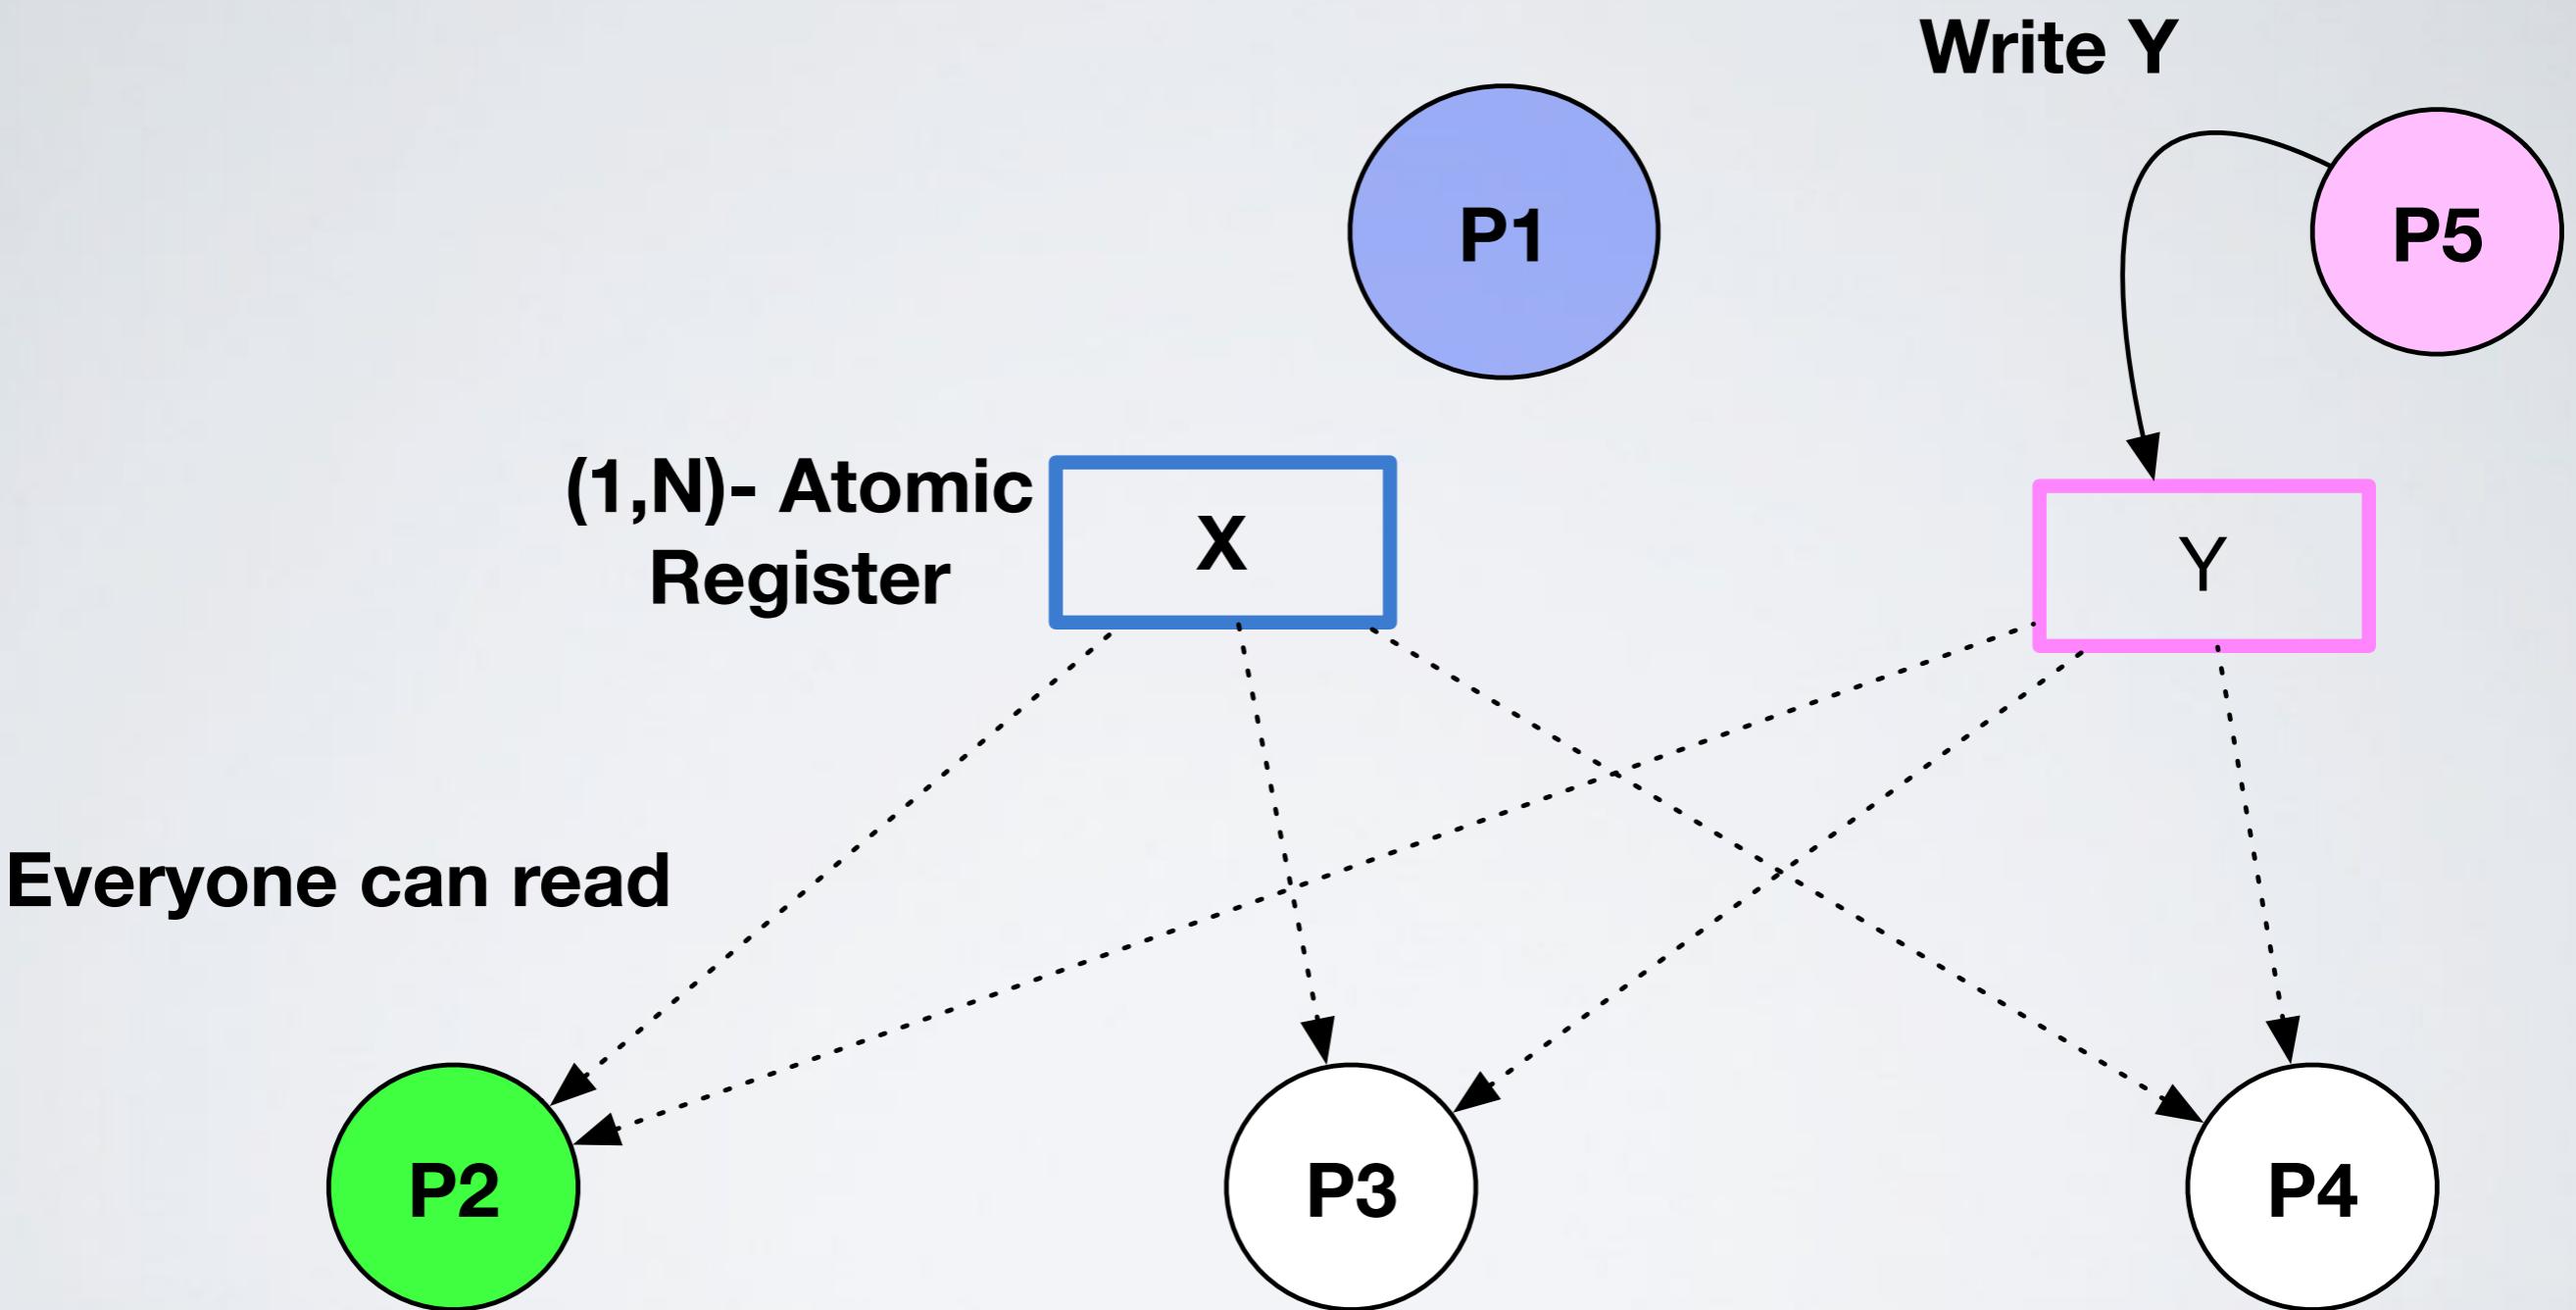
NNAR1: *Termination*: Same as property ONAR1 of a $(1, N)$ atomic register (Module 4.2).

NNAR2: *Atomicity*: Every read operation returns the value that was written most recently in a hypothetical execution, where every failed operation appears to be complete or does not appear to have been invoked at all, and every complete operation appears to have been executed at some instant between its invocation and its completion.

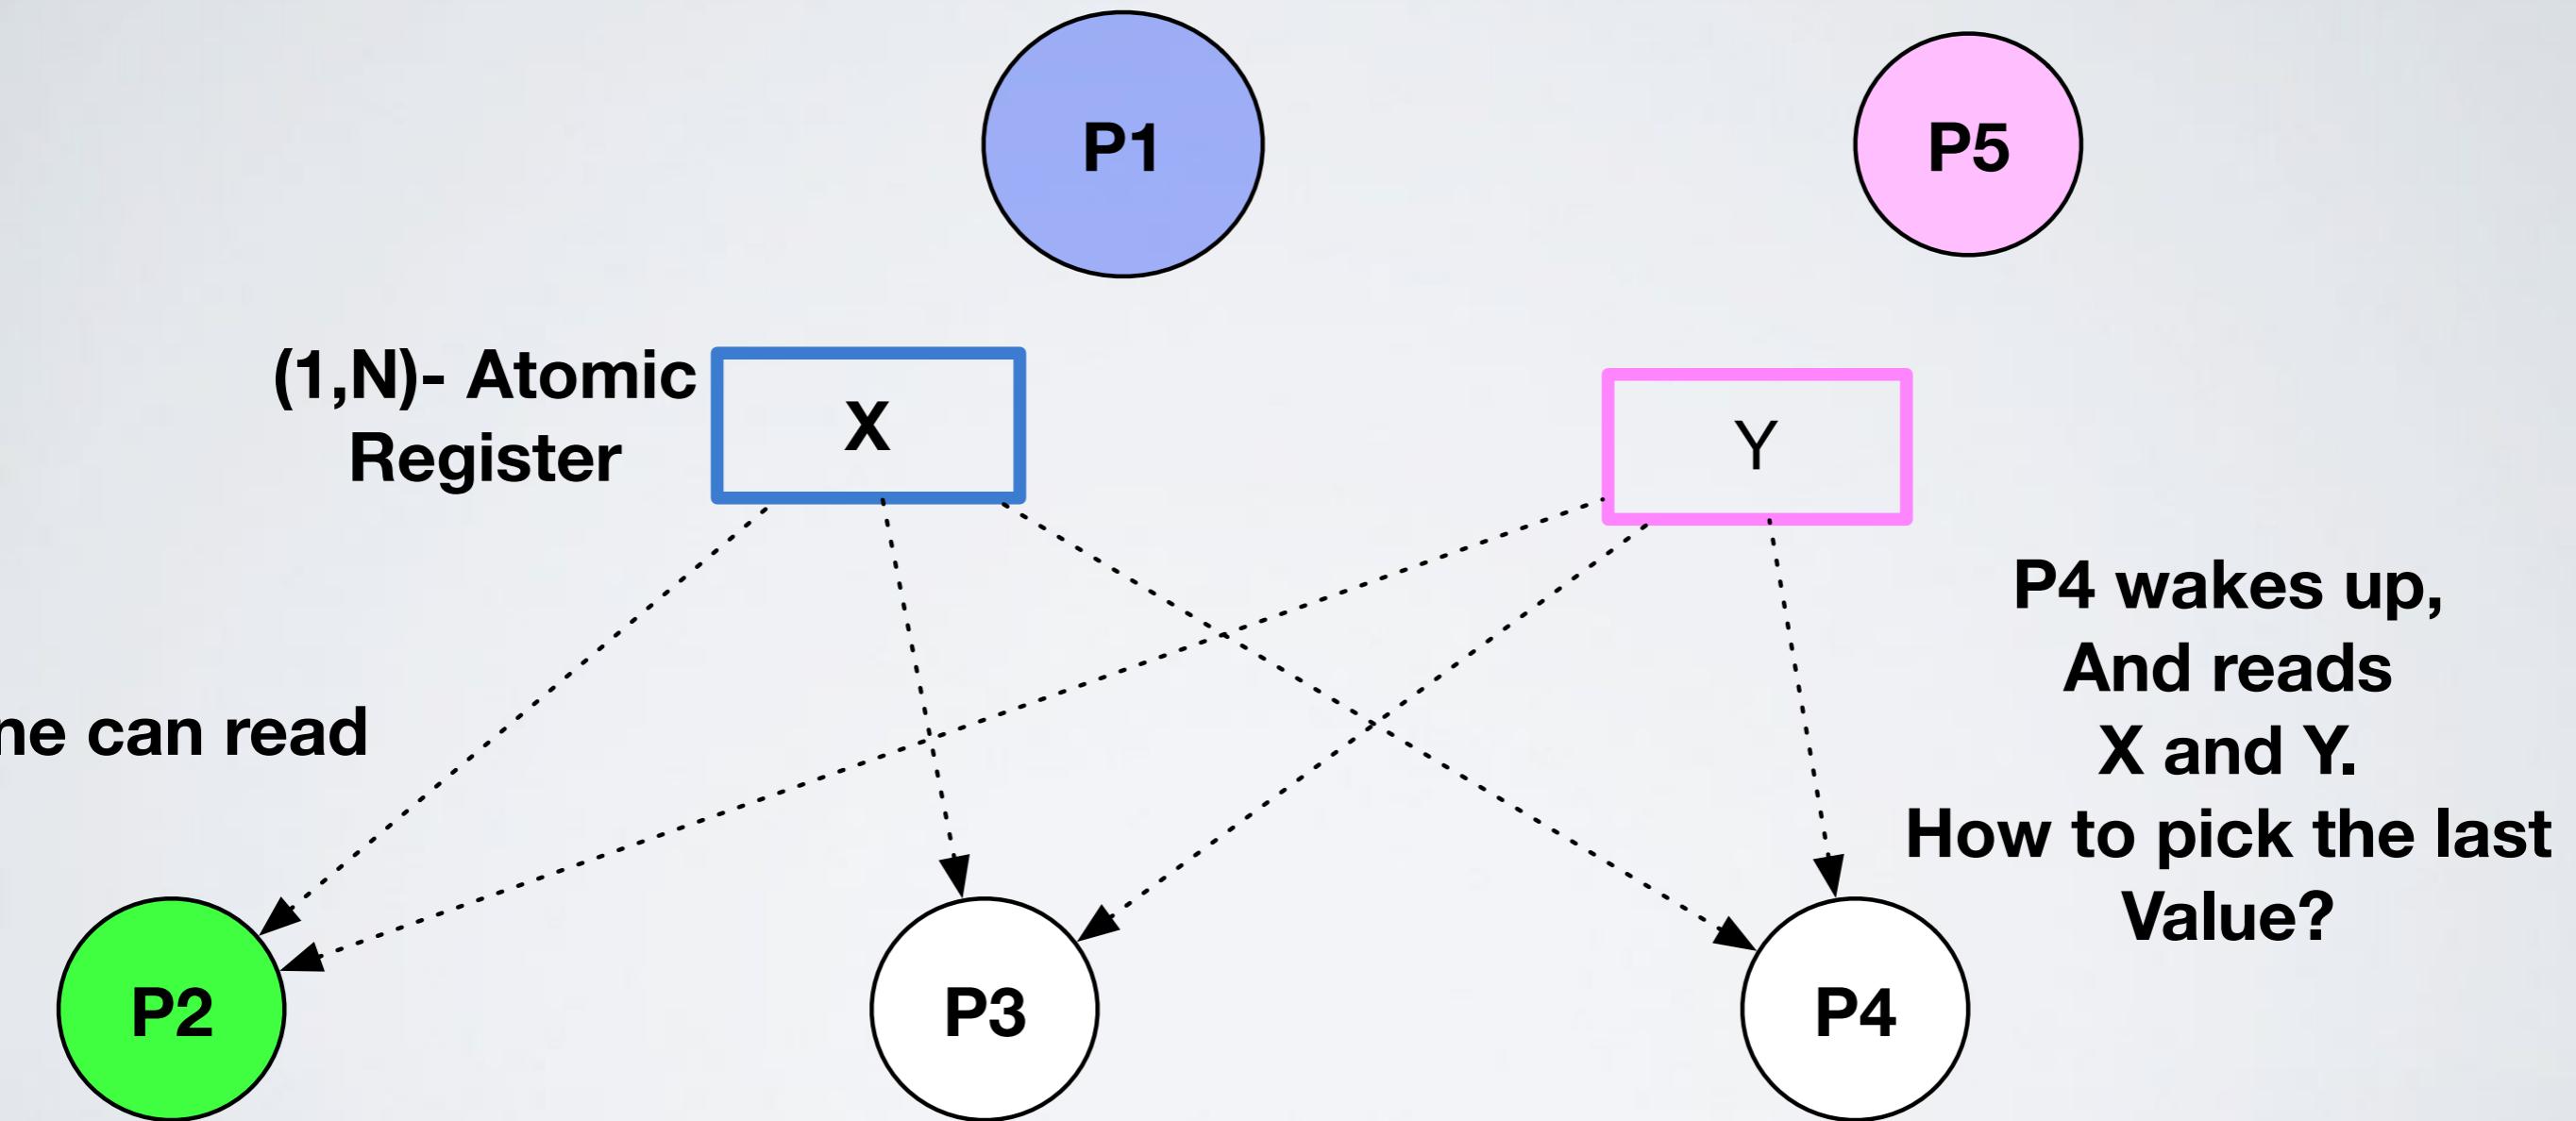
(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC

$ts = \langle wtd, id_writer \rangle$

Usual trick: we use IDs to break ties, if two TSs have equal wts, pick the smallest ID

Write ($(ts=(1,P1),X)$)

(1,N)- Atomic Register

(ts, X)

P1

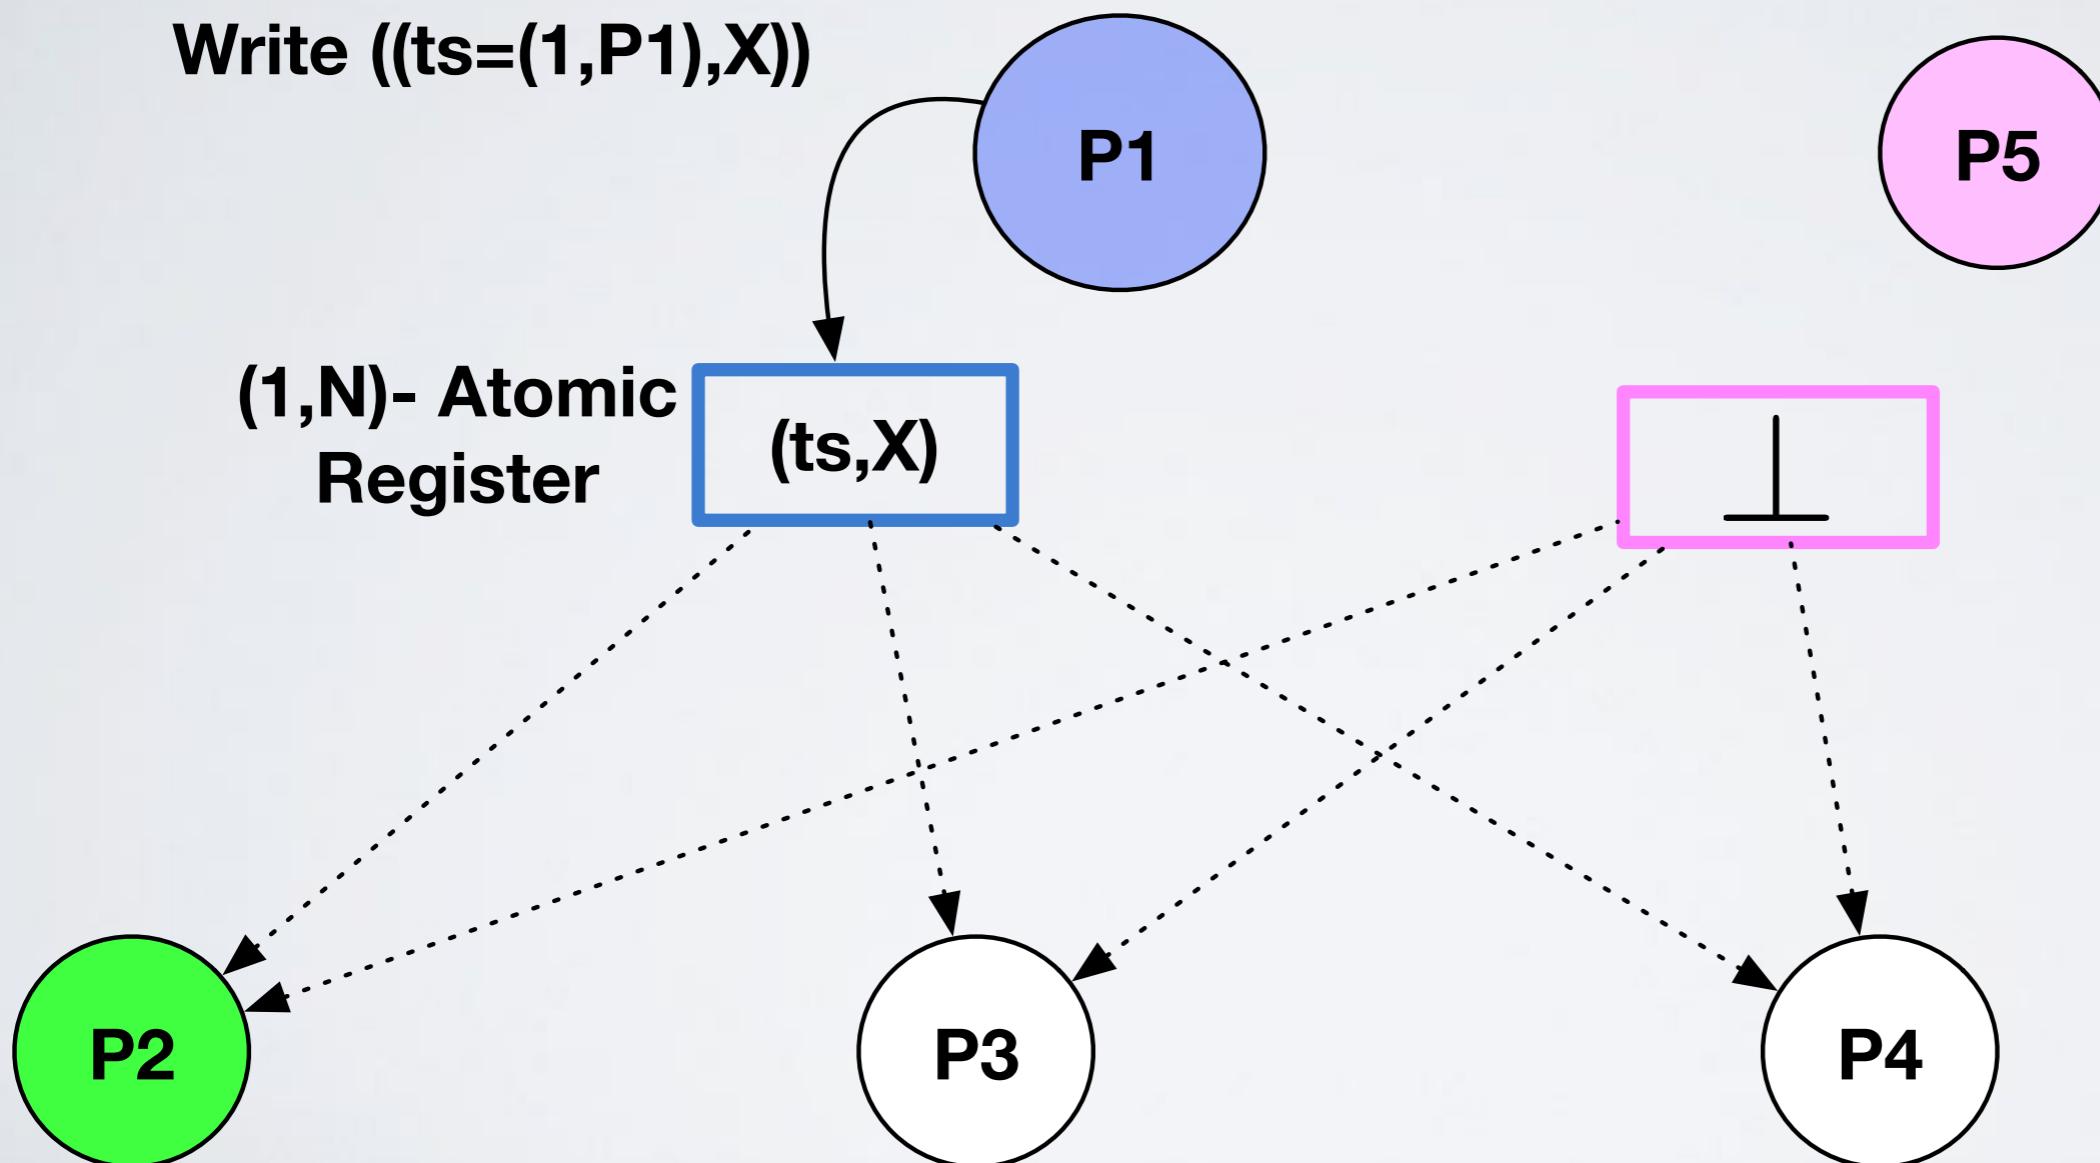
P5

P2

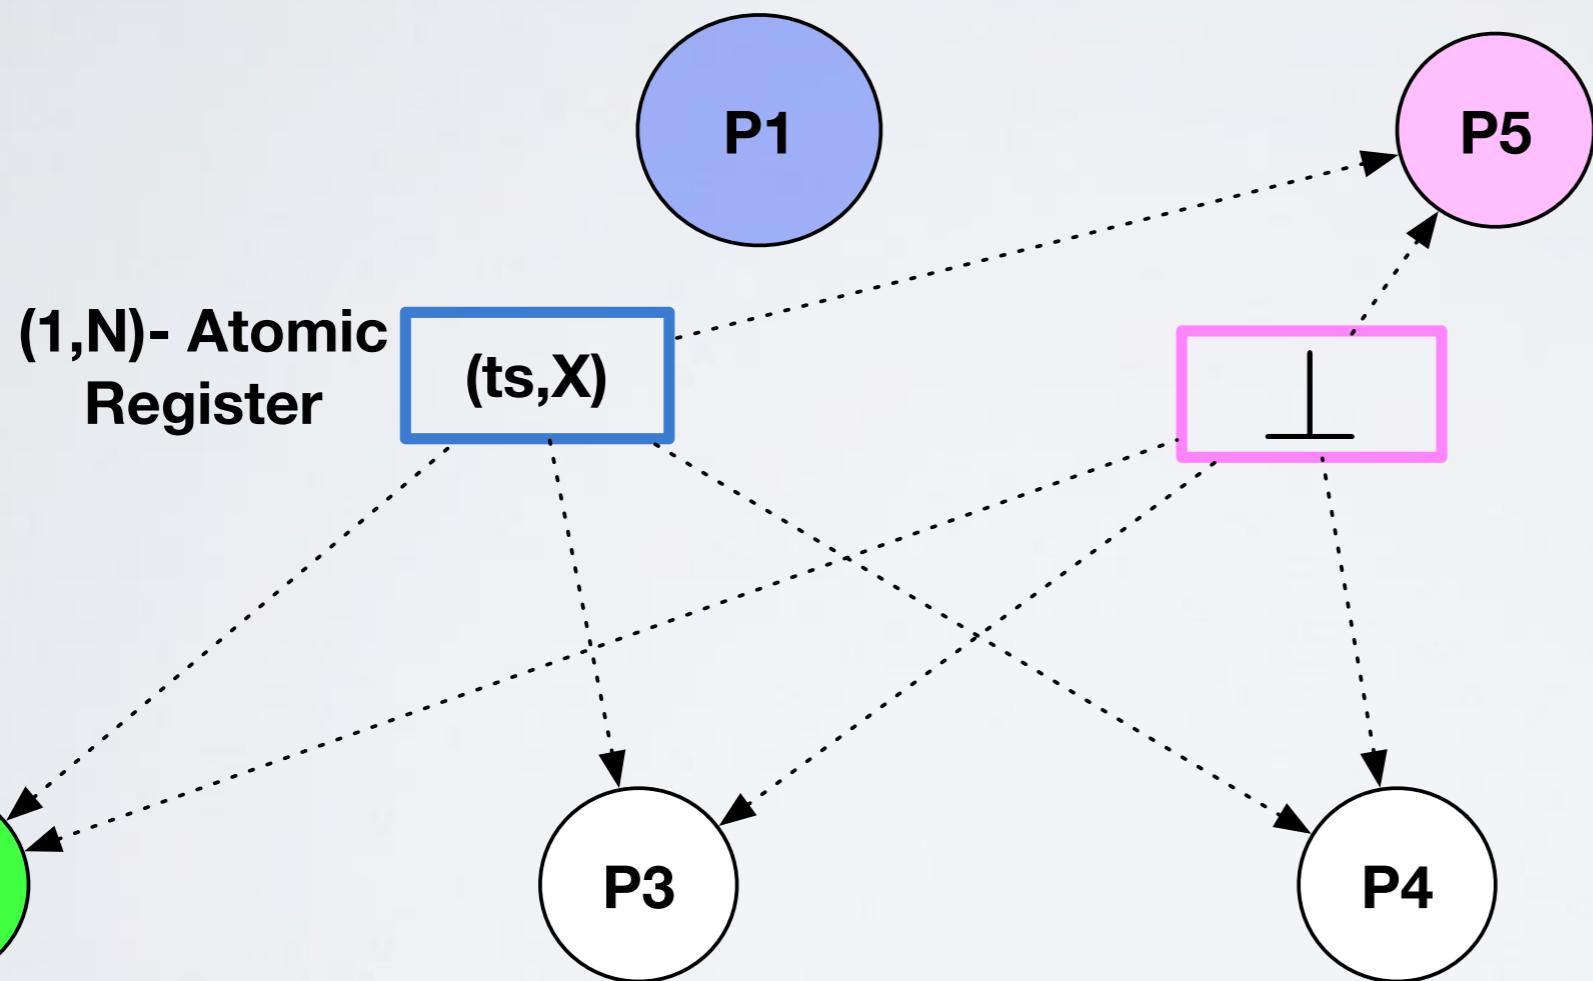
P3

P4

\perp



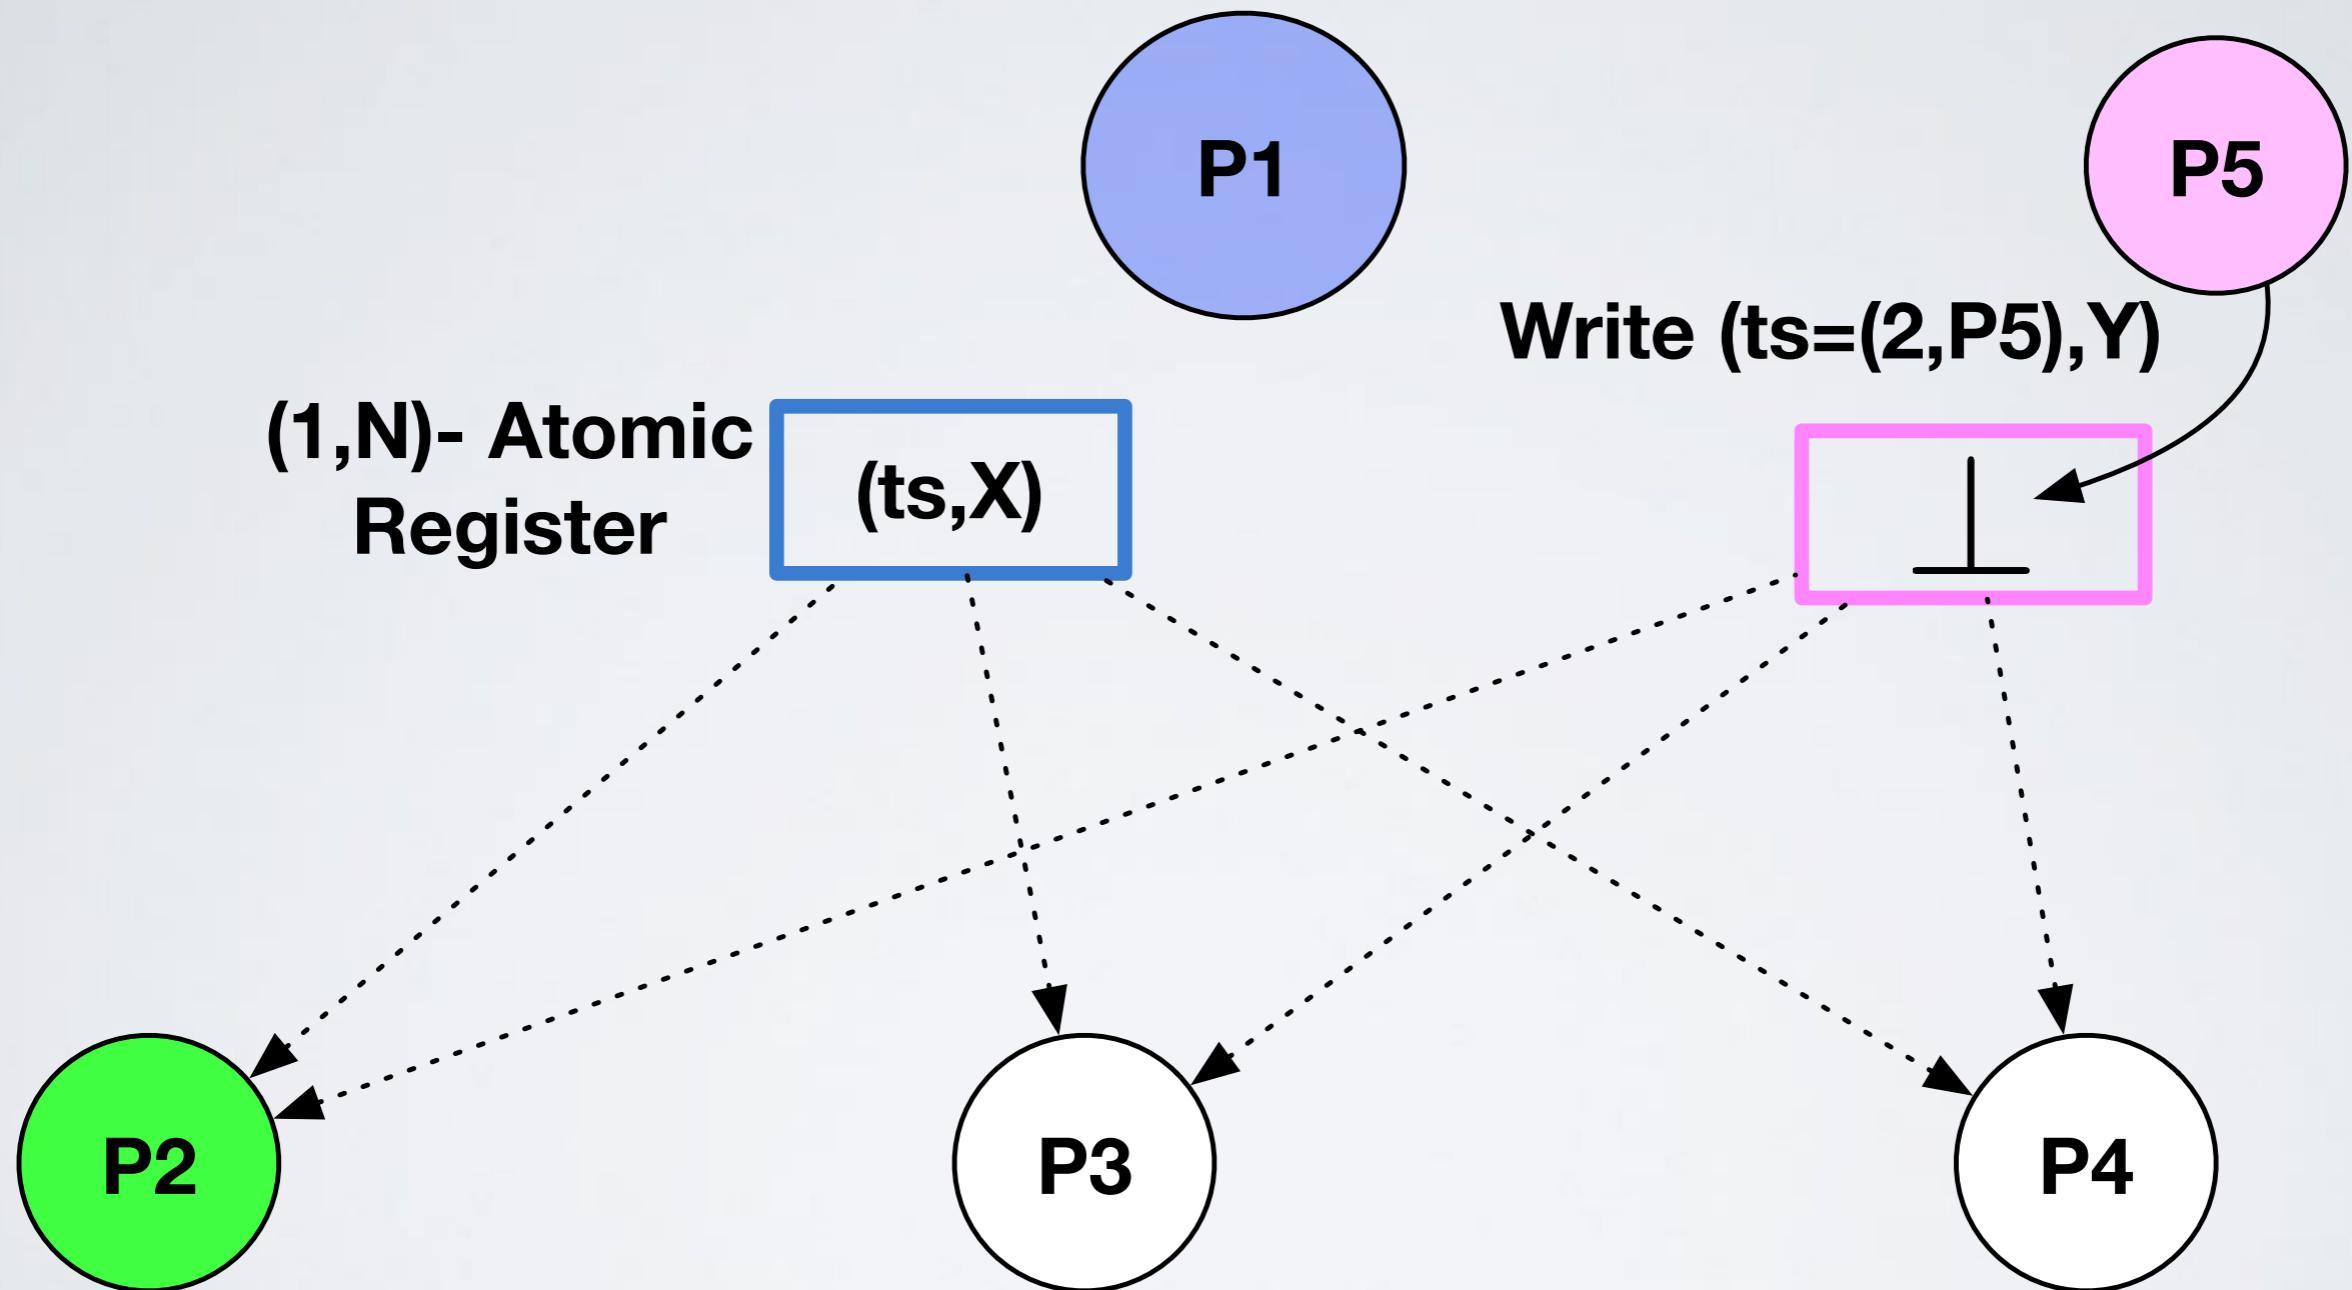
(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



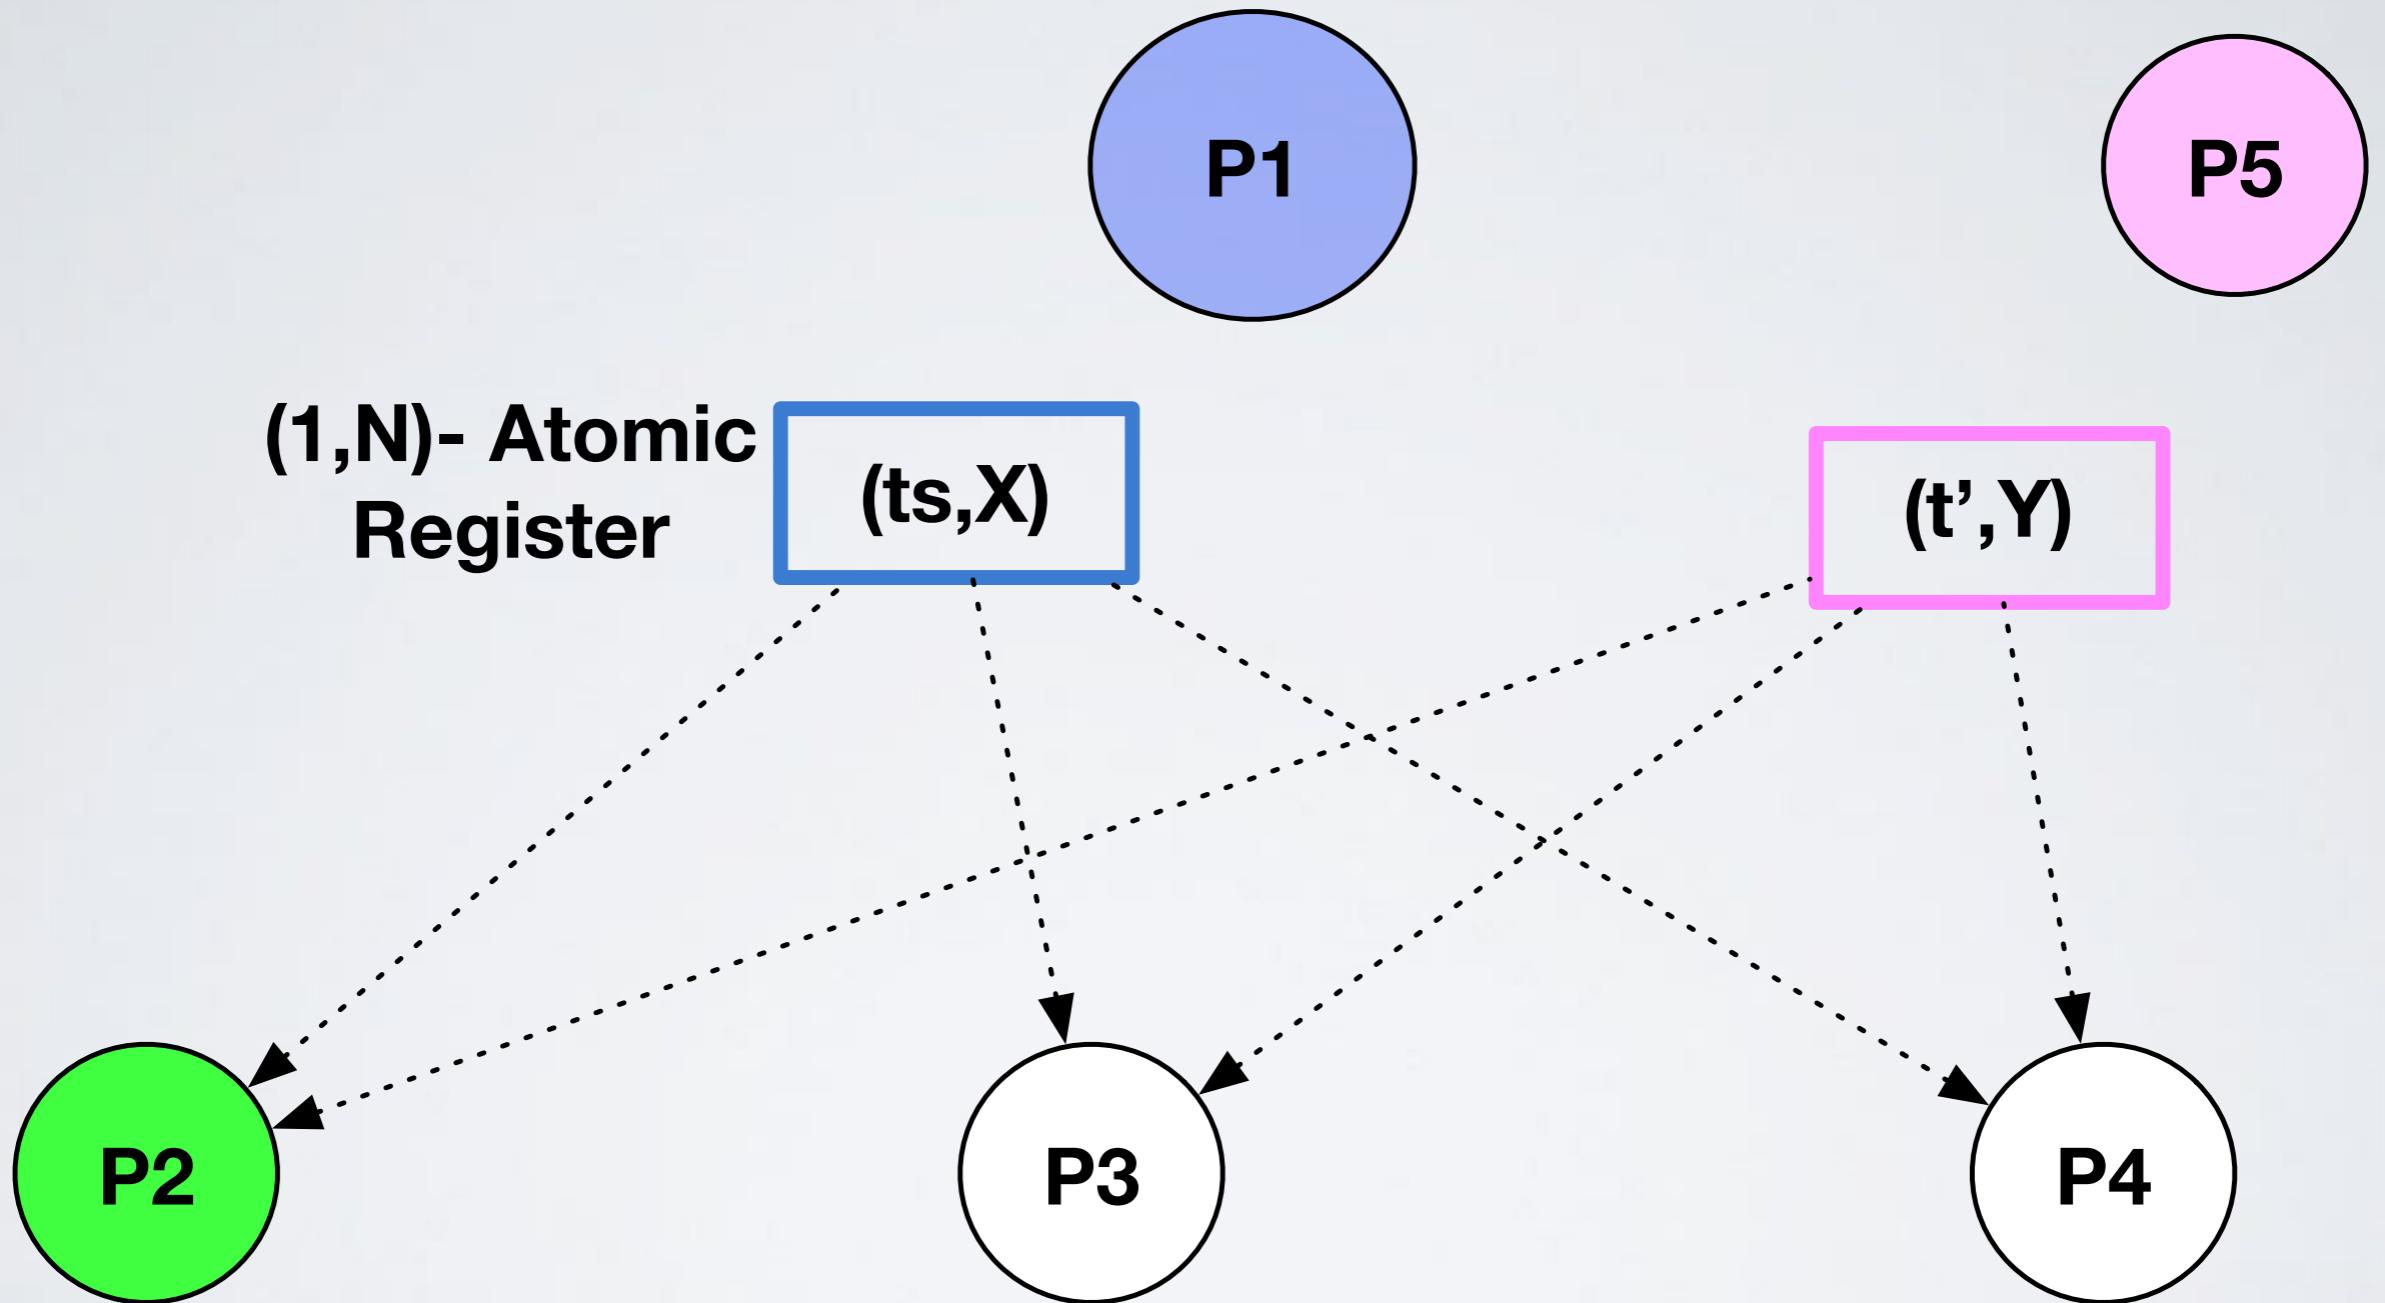
A writer has a consult phase
In which it reads all registers

wts_max=1

(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC



(N,N)-ATOMIC SIMULATION USING (1,N)-ATOMIC

upon event $\langle nnar, Init \rangle$ **do**

$val := \perp;$

$writing := \text{FALSE};$

$readlist := [\perp]^N;$

forall $q \in \Pi$ **do**

 Initialize a new instance $onar.q$ of (1, N)-AtomicRegister
 with writer q ;

upon event $\langle onar.q, ReadReturn \mid (ts', v') \rangle$ **do**

$readlist[q] := (ts', rank(q), v');$

if $\#(readlist) = N$ **then**

$(ts, v) := \text{highest}(readlist);$

$readlist := [\perp]^N;$

if $writing = \text{TRUE}$ **then**

$writing := \text{FALSE};$

trigger $\langle onar.self, Write \mid (ts + 1, val) \rangle$

else

trigger $\langle nnar, ReadReturn \mid v \rangle;$

upon event $\langle onar.self, WriteReturn \rangle$ **do**

trigger $\langle nnar, WriteReturn \rangle;$

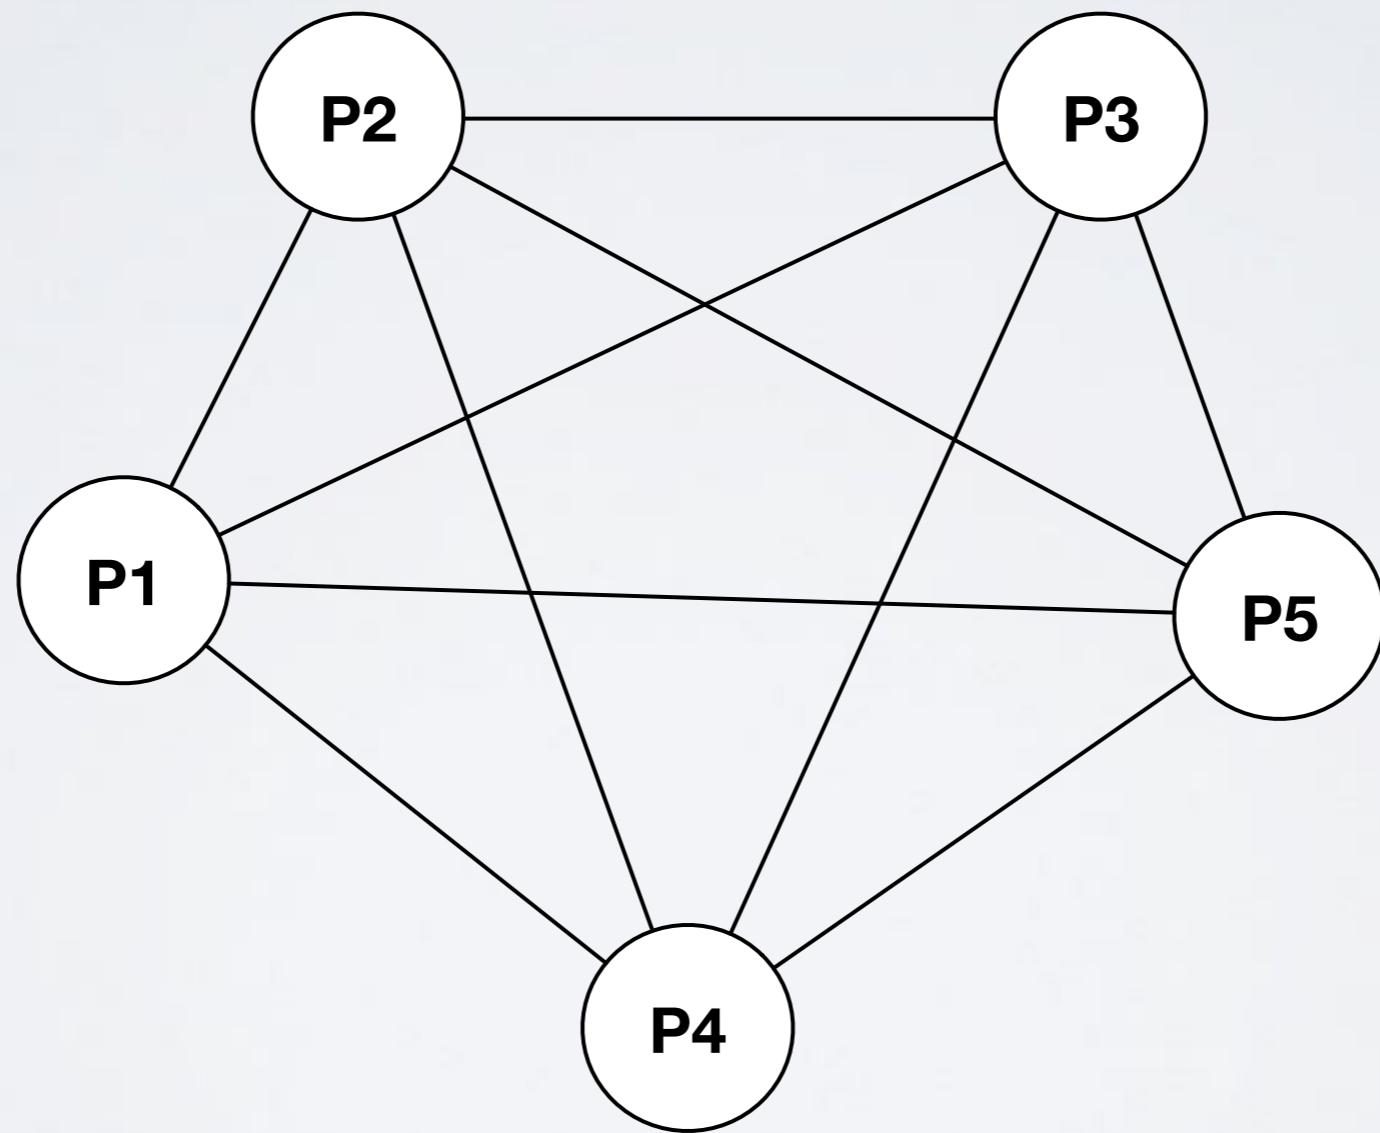
QUESTIONS

Is it possible to implement (N,N)-AR on message passing using P?

Is it possible to implement (N,N)-AR on fail-silent assuming majority of correct processes?

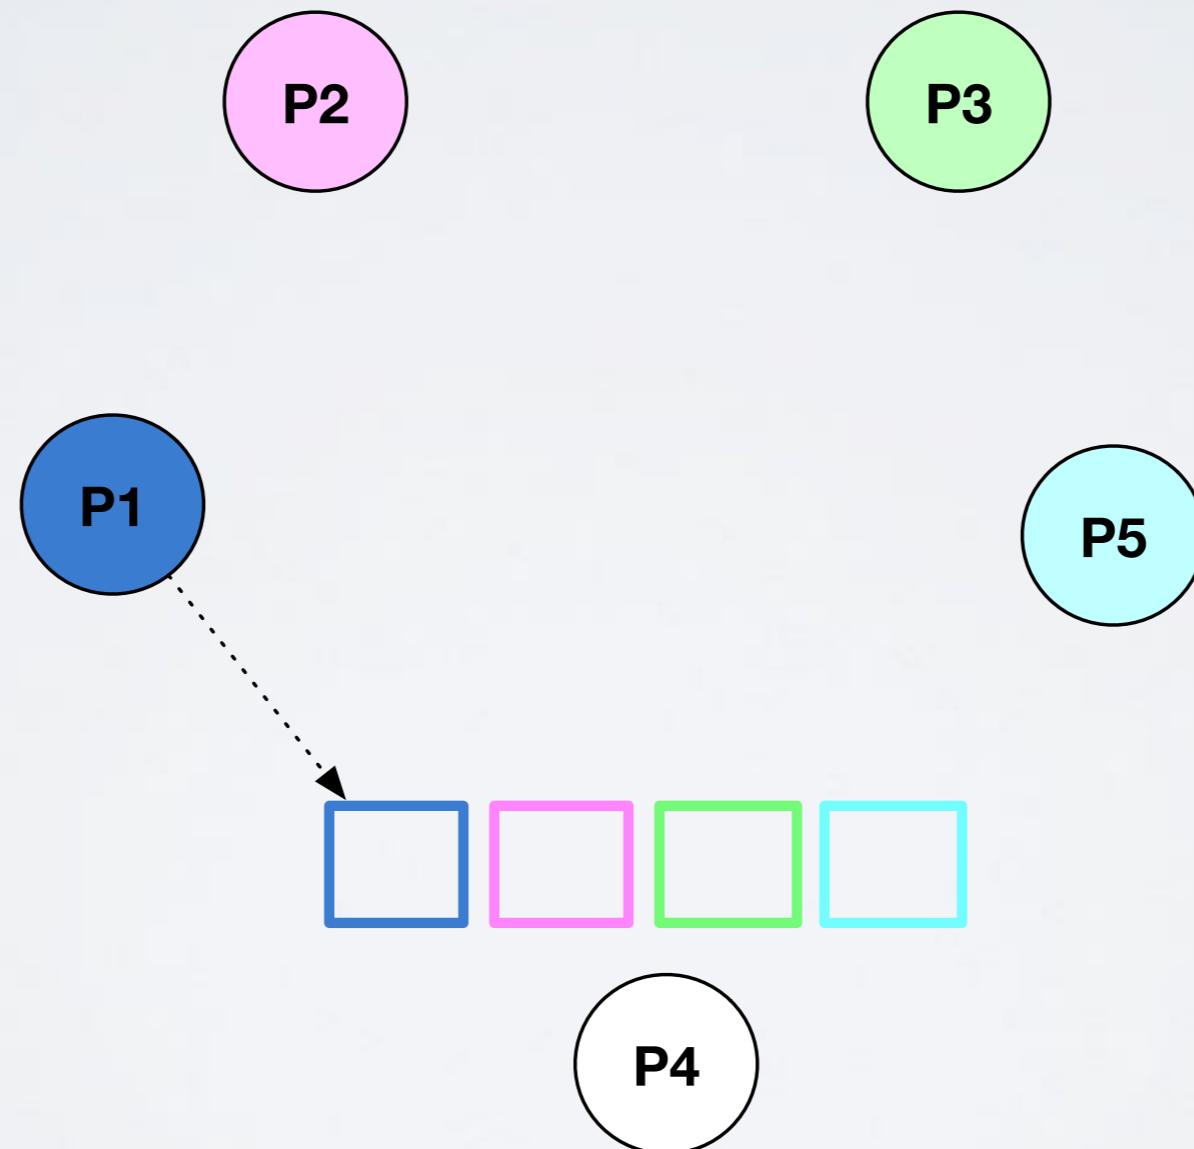
QUESTIONS

Could you show how to simulate message passing using a (1,1)-atomic registers?



QUESTIONS

Could you show how to simulate message passing using a (1,1)-atomic registers?



EXERCISE

Exercise 4.7: *Give an algorithm that implements a $(1, 1)$ atomic register in the fail-silent model and that is more efficient than the “Read-Impose Write-Majority” algorithm (Algorithm 4.6–4.7, which implements a $(1, N)$ atomic register in the fail-silent model).*

EXERCISE

Exercise 4.7: Give an algorithm that implements a $(1, 1)$ atomic register in the fail-silent model and that is more efficient than the “Read-Impose Write-Majority” algorithm (Algorithm 4.6–4.7, which implements a $(1, N)$ atomic register in the fail-silent model).

Algorithm 4.20: Modification of Majority Voting to Implement a $(1, 1)$ Atomic Register

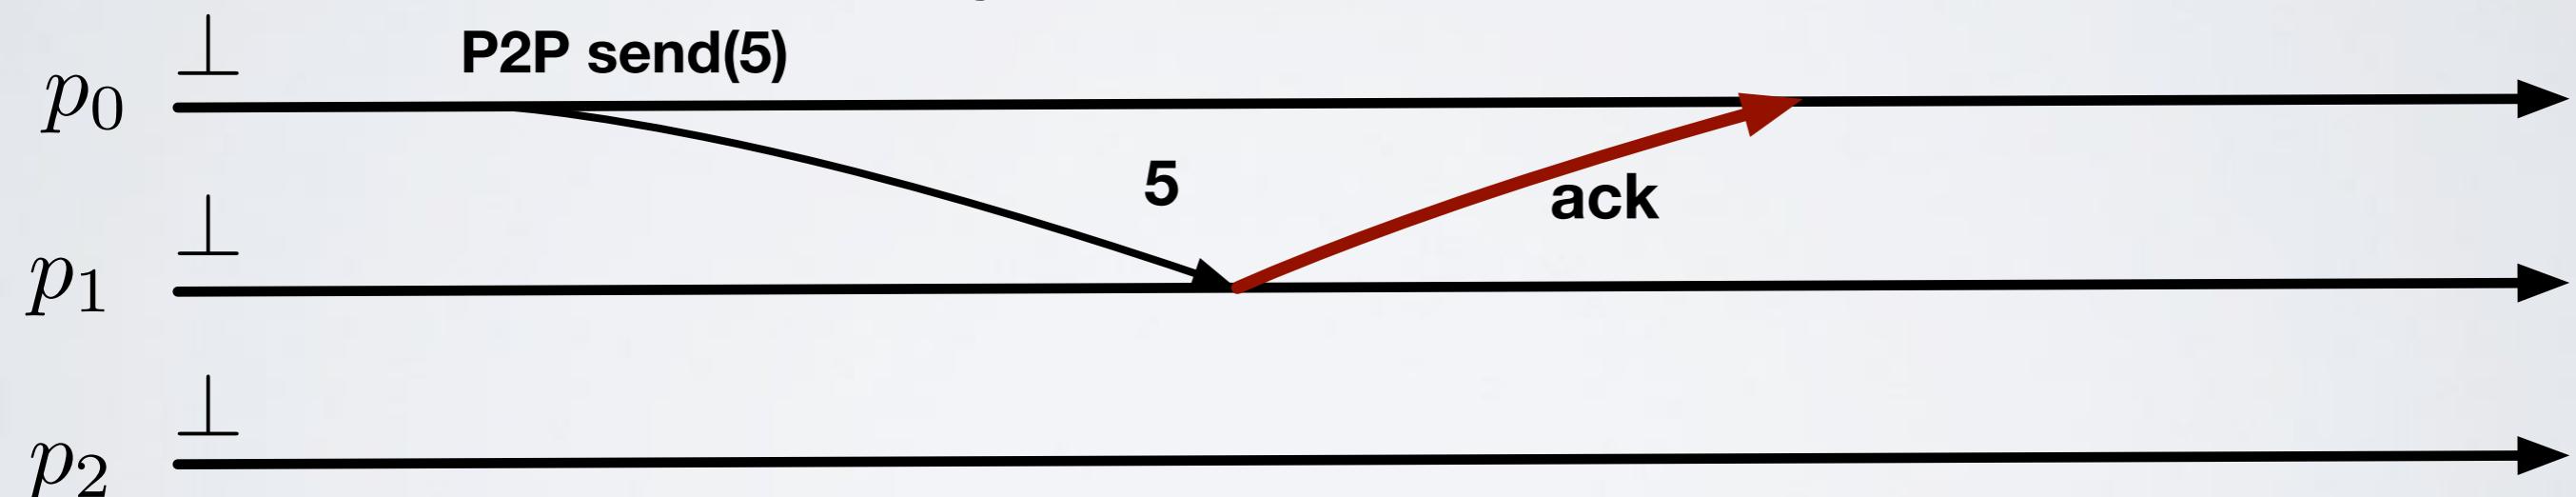
```
upon event < onrr, Read > do
    rid := rid + 1;
    readlist := [⊥]N;
    readlist[self] := (ts, val);
    trigger < beb, Broadcast | [READ, rid] >;
    
upon event < beb, Deliver | p, [READ, r] > do
    if p ≠ self then
        trigger < pl, Send | p, [VALUE, r, ts, val] >;
    
upon event < pl, Deliver | q, [VALUE, r, ts', v'] > such that r = rid do
    readlist[q] := (ts', v');
    if #(readlist) > N/2 then
        (ts, val) := highest(readlist);
        readlist := ∅;
        trigger < onrr, ReadReturn | val >;
```

EXERCISE

Exercise 4.7: Give an algorithm that implements a $(1, 1)$ atomic register in the fail-silent model and that is more efficient than the “Read-Impose Write-Majority” algorithm (Algorithm 4.6–4.7, which implements a $(1, N)$ atomic register in the fail-silent model).

Why not using just this pattern? It Is very efficient
Each write is 2 messages

Each write send to the only reader P1, and wait for ack.



EXERCISE

Exercise 4.6: Explain why the request identifier included in WRITE, ACK, READ, and VALUE messages of the “Majority Voting” algorithm (Algorithm 4.2) can be omitted in the “Read-One Write-All” algorithm (Algorithm 4.1).

Read-one Write-all Algorithm

upon event $\langle onrr, \text{Init} \rangle$ **do**

$val := \perp;$

$correct := \Pi;$

$writeset := \emptyset;$

upon event $\langle \mathcal{P}, \text{Crash} \mid p \rangle$ **do**

$correct := correct \setminus \{p\};$

upon event $\langle onrr, \text{Read} \rangle$ **do**

trigger $\langle onrr, \text{ReadReturn} \mid val \rangle;$

upon event $\langle onrr, \text{Write} \mid v \rangle$ **do**

trigger $\langle beb, \text{Broadcast} \mid [\text{WRITE}, v] \rangle;$

upon event $\langle beb, \text{Deliver} \mid q, [\text{WRITE}, v] \rangle$ **do**

$val := v;$

trigger $\langle pl, \text{Send} \mid q, \text{ACK} \rangle;$

upon event $\langle pl, \text{Deliver} \mid p, \text{ACK} \rangle$ **do**

$writeset := writeset \cup \{p\};$

upon $correct \subseteq writeset$ **do**

$writeset := \emptyset;$

trigger $\langle onrr, \text{WriteReturn} \rangle;$

EXERCISE

Exercise 4.6: Explain why the request identifier included in WRITE, ACK, READ, and VALUE messages of the “Majority Voting” algorithm (Algorithm 4.2) can be omitted in the “Read-One Write-All” algorithm (Algorithm 4.1).

Read-one Write-all Algorithm

upon event $\langle onrr, \text{Init} \rangle$ **do**

$(ts, val) := (0, \perp);$
 $wts := 0;$
 $acks := 0;$
 $rid := 0;$
 $readlist := [\perp]^N;$

upon event $\langle onrr, \text{Write} \mid v \rangle$ **do**

$wts := wts + 1;$
 $acks := 0;$
trigger $\langle beb, \text{Broadcast} \mid [\text{WRITE}, wts, v] \rangle;$

upon event $\langle beb, \text{Deliver} \mid p, [\text{WRITE}, ts', v'] \rangle$ **do**

if $ts' > ts$ **then**
 $(ts, val) := (ts', v');$
trigger $\langle pl, \text{Send} \mid p, [\text{ACK}, ts'] \rangle;$

upon event $\langle pl, \text{Deliver} \mid q, [\text{ACK}, ts'] \rangle$ **such that** $ts' = wts$ **do**

$acks := acks + 1;$
if $acks > N/2$ **then**
 $acks := 0;$
trigger $\langle onrr, \text{WriteReturn} \rangle;$

upon event $\langle onrr, \text{Read} \rangle$ **do**

$rid := rid + 1;$
 $readlist := [\perp]^N;$
trigger $\langle beb, \text{Broadcast} \mid [\text{READ}, rid] \rangle;$

upon event $\langle beb, \text{Deliver} \mid p, [\text{READ}, r] \rangle$ **do**

trigger $\langle pl, \text{Send} \mid p, [\text{VALUE}, r, ts, val] \rangle;$

upon event $\langle pl, \text{Deliver} \mid q, [\text{VALUE}, r, ts', v'] \rangle$ **such that** $r = rid$ **do**

$readlist[q] := (ts', v');$
if $\#(readlist) > N/2$ **then**
 $v := \text{highestval}(readlist);$
 $readlist := [\perp]^N;$
trigger $\langle onrr, \text{ReadReturn} \mid v \rangle;$