

# ORDERED COMMUNICATIONS 2 TOTAL ORDER

DISTRIBUTED SYSTEMS  
Master of Science in Cyber Security

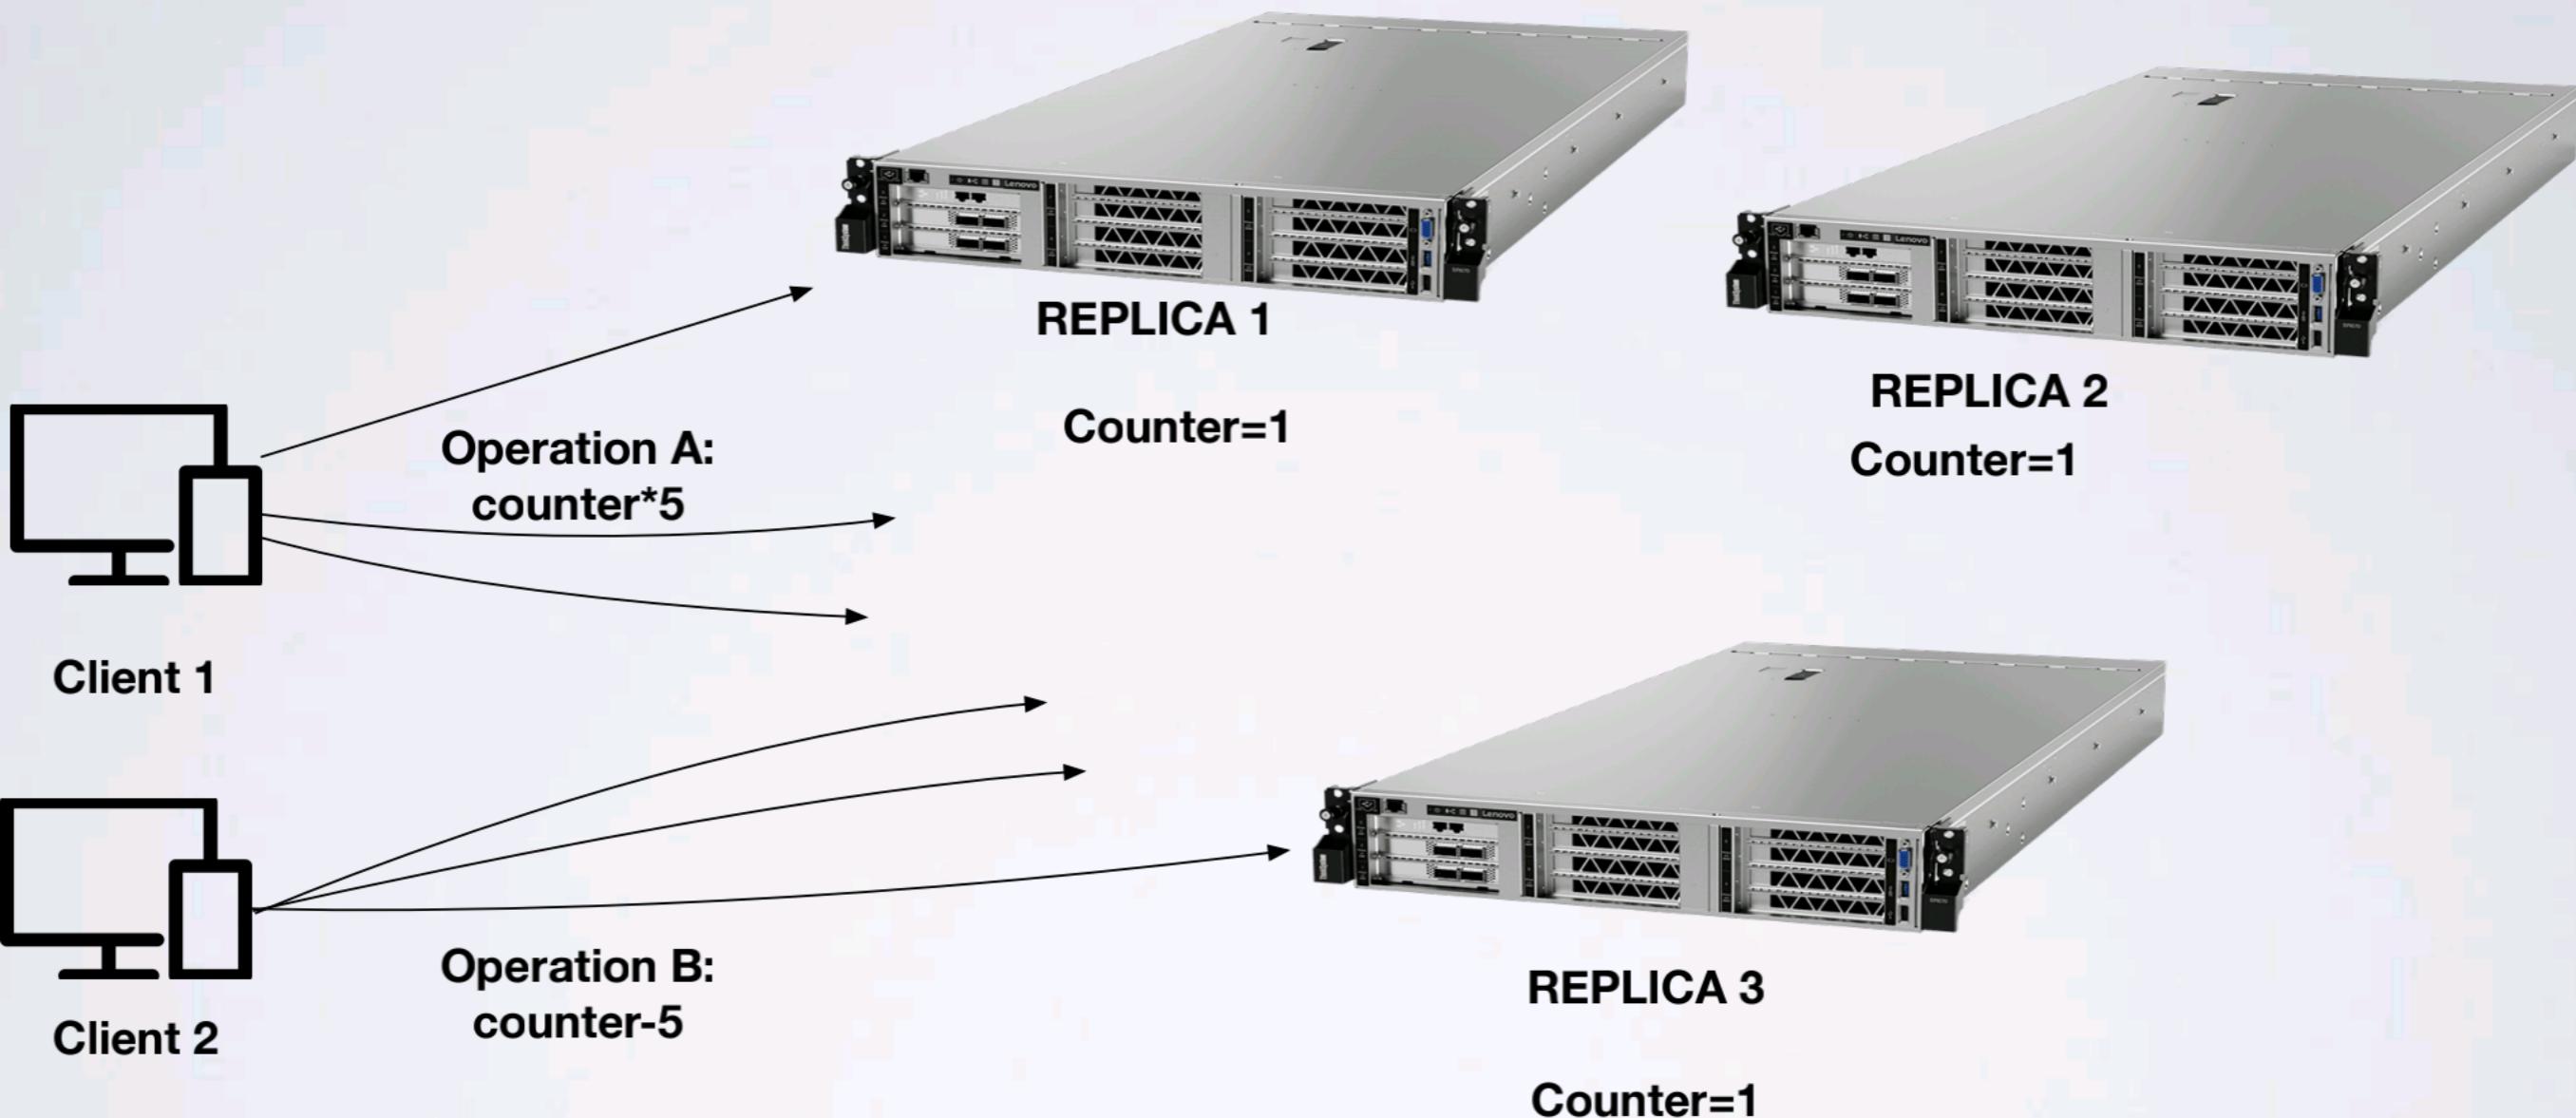


SAPIENZA  
UNIVERSITÀ DI ROMA

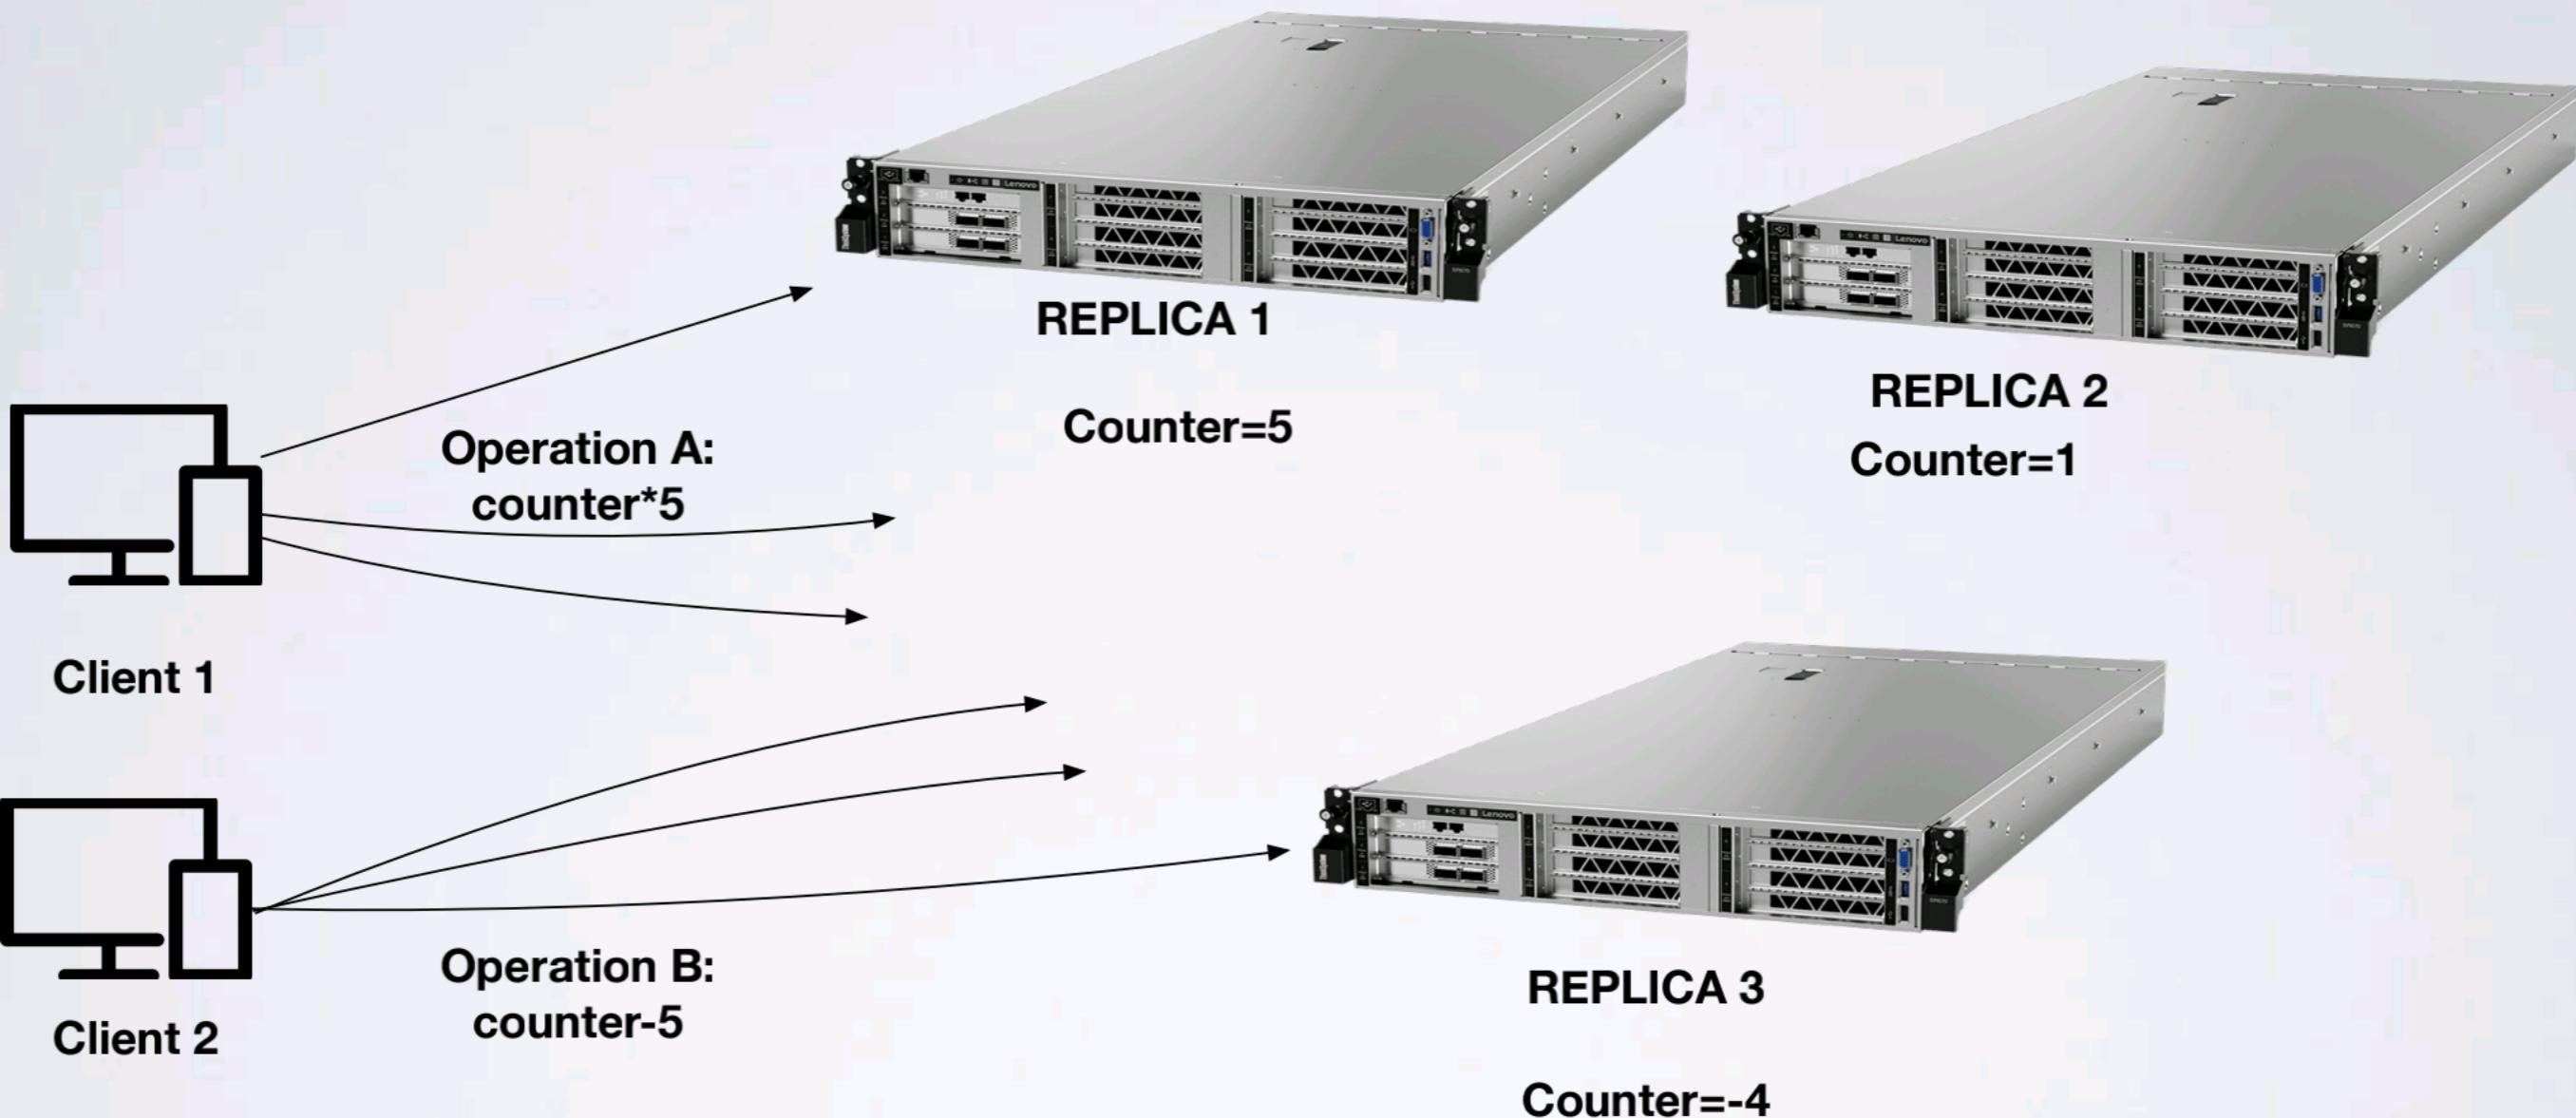


CIS SAPIENZA  
CYBER INTELLIGENCE AND INFORMATION SECURITY

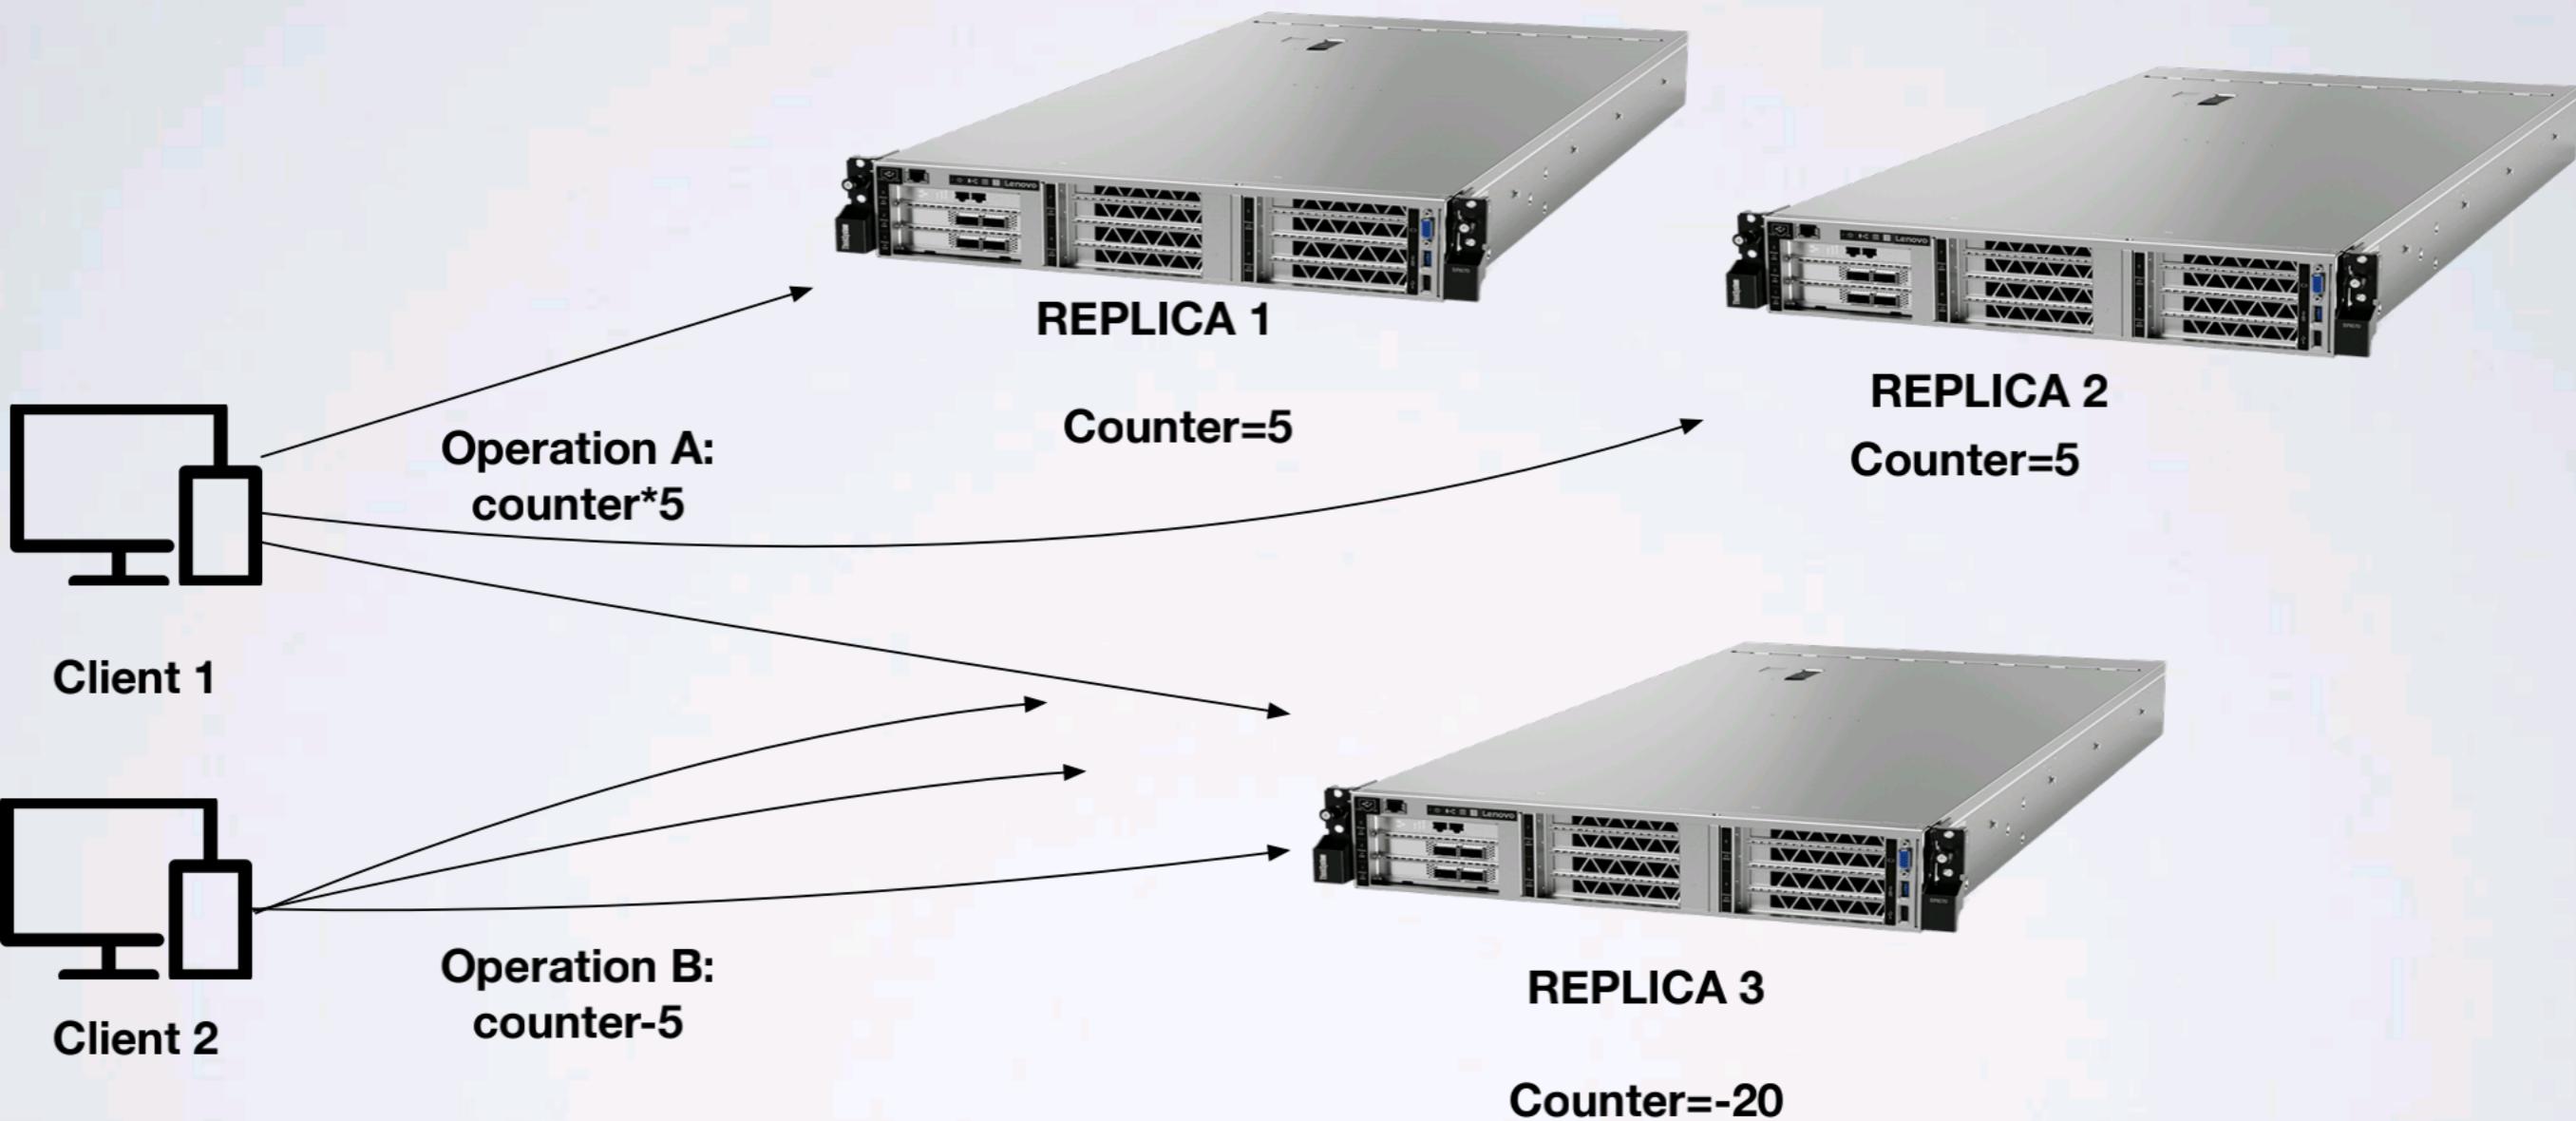
# USES OF TO-BROADCAST/(ATOMIC BROADCAST)



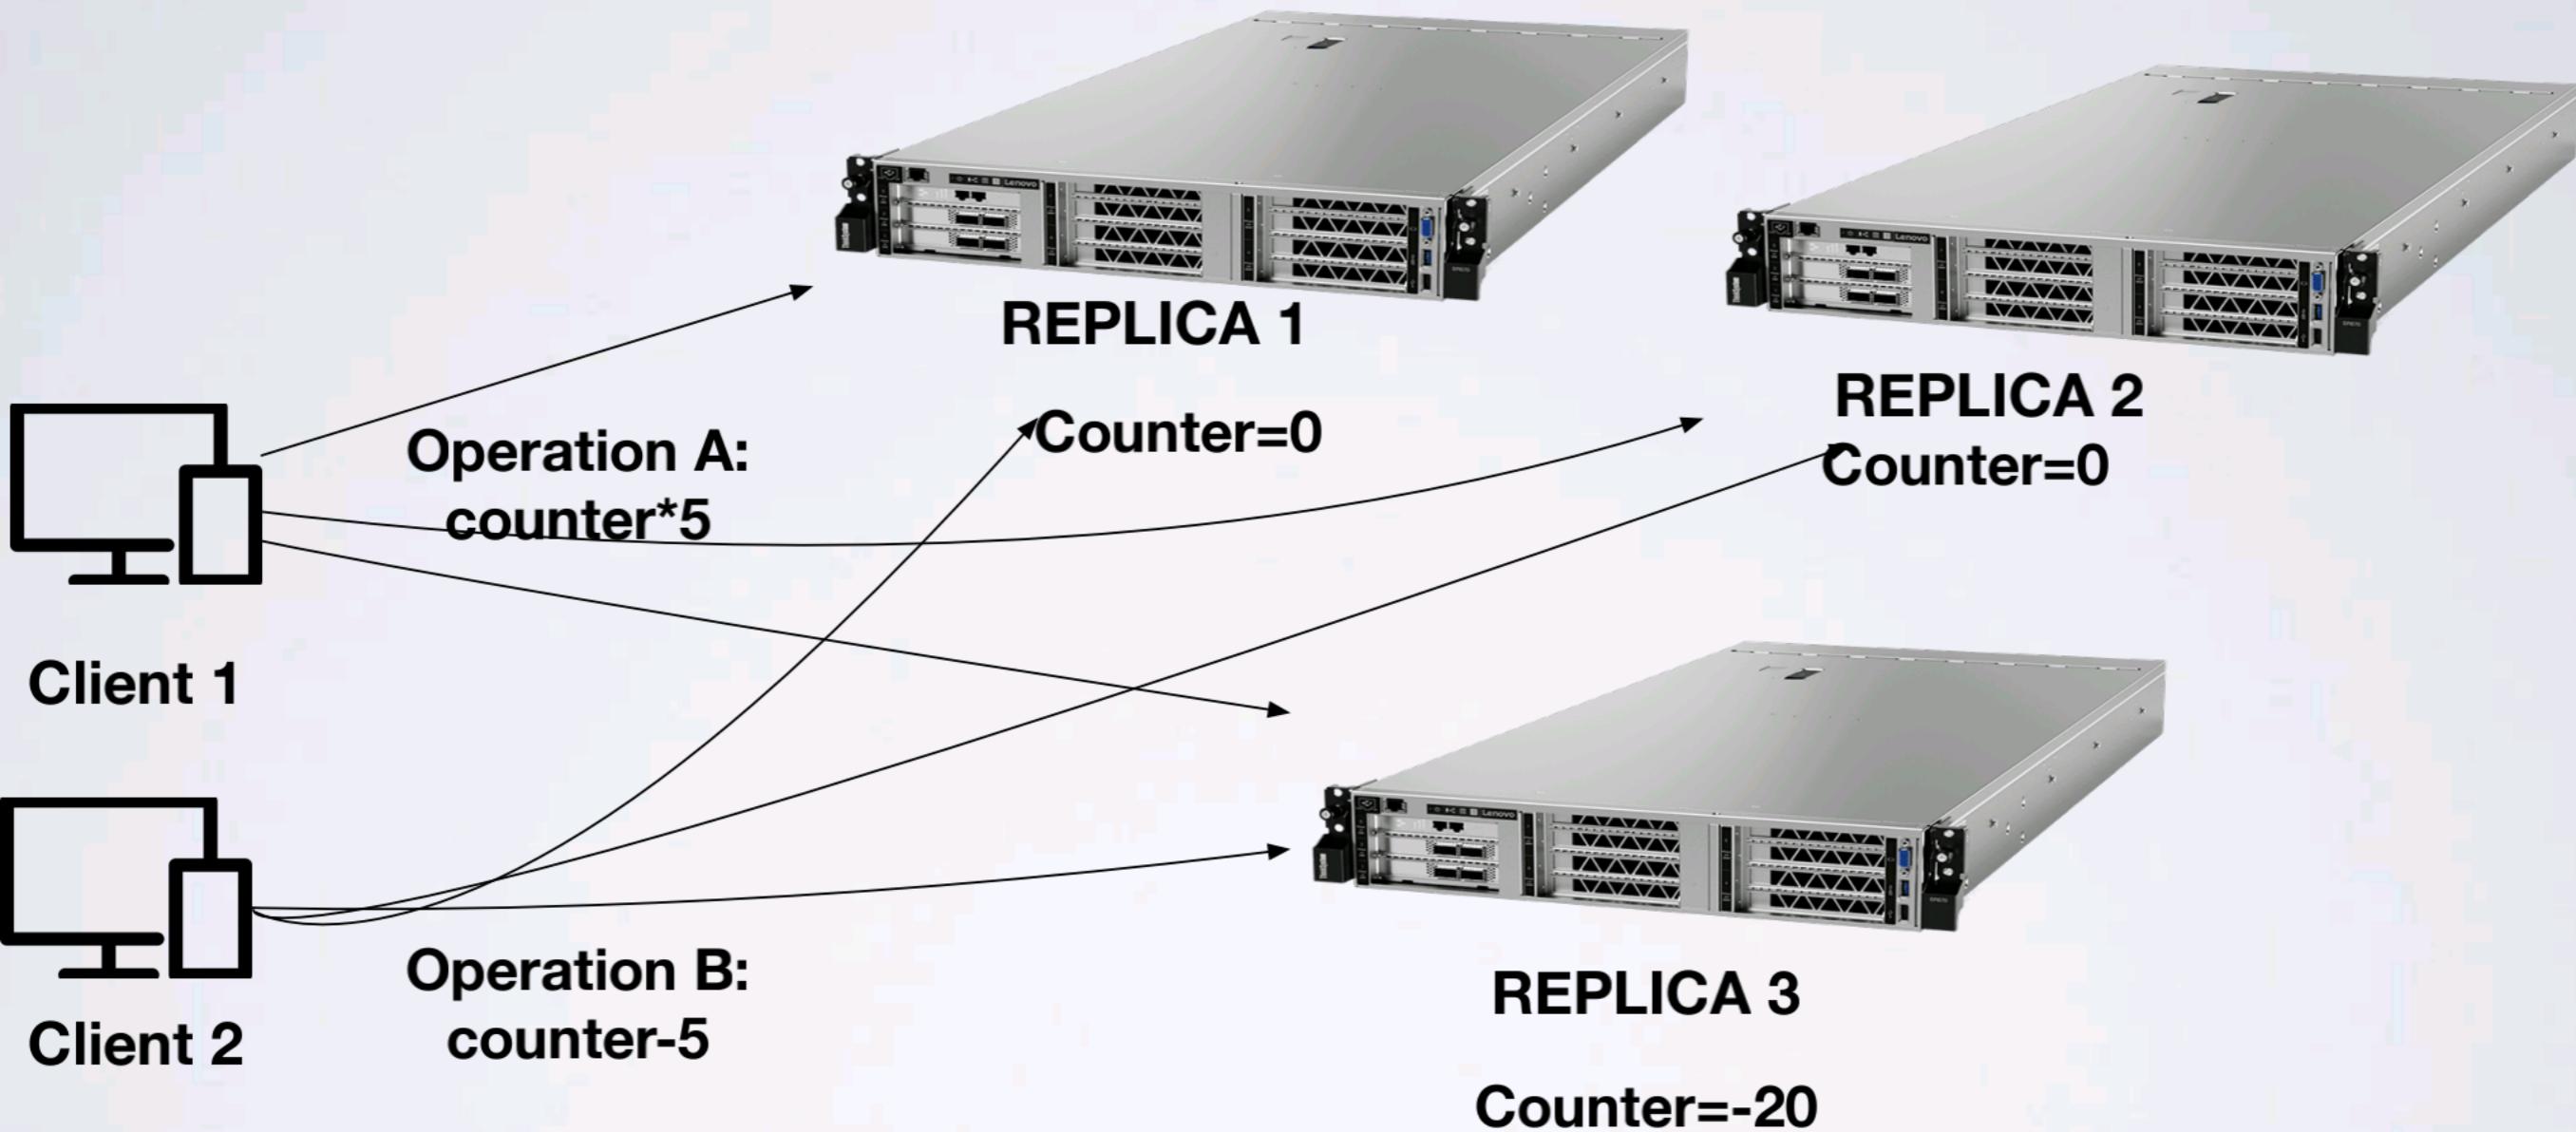
# USES OF TO-BROADCAST/(ATOMIC BROADCAST)



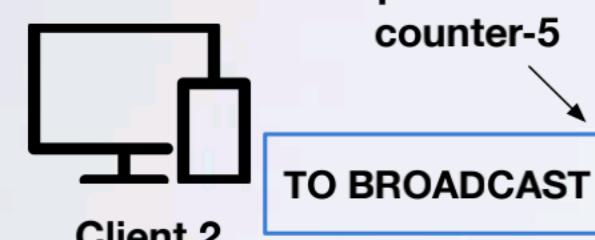
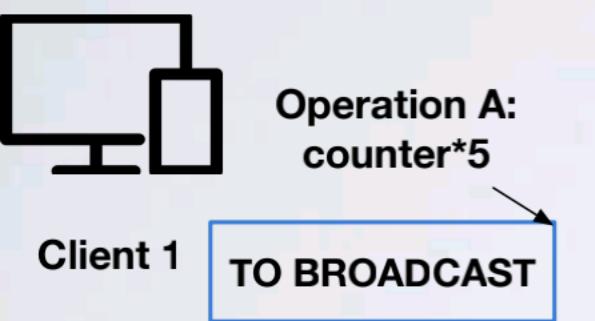
# USES OF TO-BROADCAST/(ATOMIC BROADCAST)



# USES OF TO-BROADCAST/(ATOMIC BROADCAST)



# USES OF TO-BROADCAST/(ATOMIC BROADCAST)



# TO SPECIFICATIONS

Total order specifications are usually composed by four properties:

- a **Validity** property guarantees that messages sent by correct processes will eventually be delivered at least by correct processes;
- an **Integrity** property guarantees that no spurious or duplicate messages are delivered;
- an **Agreement** property ensures that (at least correct) processes deliver the same set of messages;
- an **Order** property constrains (at least correct) processes delivering the same messages to deliver them in the same order.

# TO SPECIFICATIONS

## Total Order Broadcast

- V = Validity
- I = Integrity
- A = Agreement
- O = Order

Distinct specifications arise from distinct formulations of each property

- uniform vs non-uniform
- A uniform property imposes restrictions on the behavior of (at least) correct processes on the basis of events occurred in some process

# THE AGREEMENT PROPERTY

## Uniform Agreement (UA)

If a process (correct or not) TOTDelivers a message m, then all correct processes will eventually TOTDeliver m

## Non Uniform Agreement (NUA)

If a correct process TOTDelivers a message m, then all correct processes will eventually TOTDeliver m

### Constrains the set of delivered messages

Correct processes always deliver the same set of messages M

Each faulty process p delivers a set  $M_p$

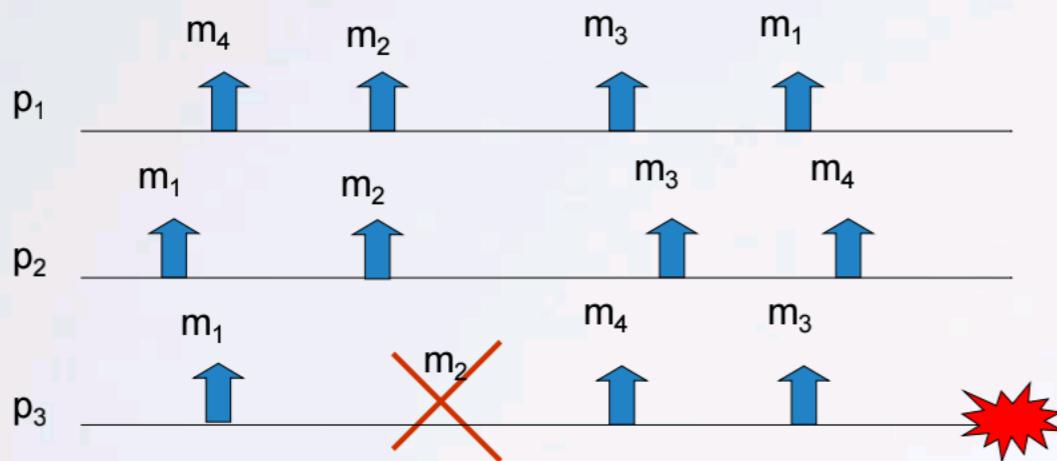
$$\text{UA: } M_p \subseteq M$$

$$\text{NUA: } M_p \text{ can be s.t. } M_p \cap M \neq \emptyset$$

# THE AGREEMENT PROPERTY

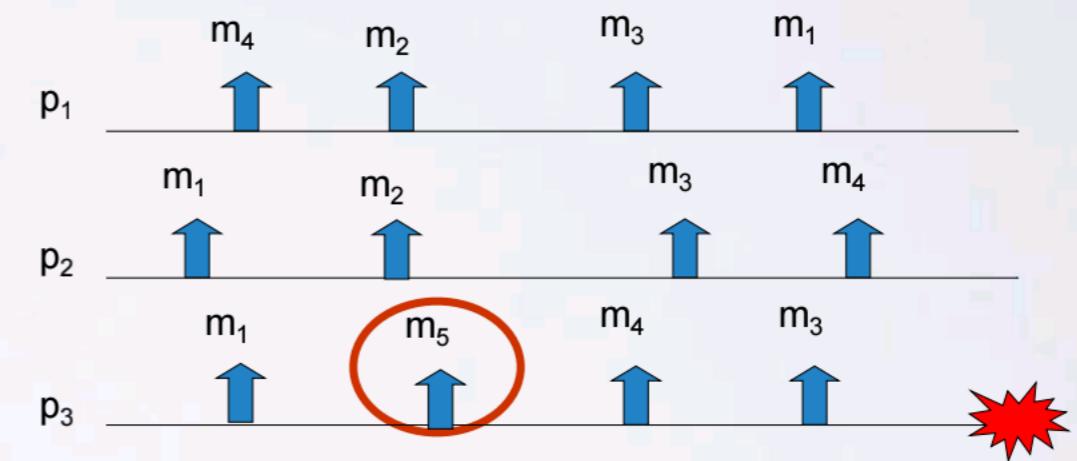
## Uniform Agreement (UA)

If a process (correct or not) `TODelivers` a message  $m$ , then all correct processes will eventually `TODeliver`  $m$



## Non Uniform Agreement (NUA)

If a correct process `TODelivers` a message  $m$ , then all correct processes will eventually `TODeliver`  $m$



## Module 6.1: Interface and properties of regular total-order broadcast

### Module:

**Name:** TotalOrderBroadcast, **instance**  $tob$ .

### Events:

$m_1, m_2$   
 $p \quad q$   
 $\pi_2 \quad \pi_1$

$m_1, m_2$   
 $q \quad p$   
 $\pi_2 \quad \pi_1$

**Request:**  $\langle tob, Broadcast \mid m \rangle$ : Broadcasts a message  $m$  to all processes.

**Indication:**  $\langle tob, Deliver \mid p, m \rangle$ : Delivers a message  $m$  broadcast by process  $p$ .

### Properties:

- **TOB1: Validity:** If a correct process  $p$  broadcasts a message  $m$ , then  $p$  eventually delivers  $m$ .

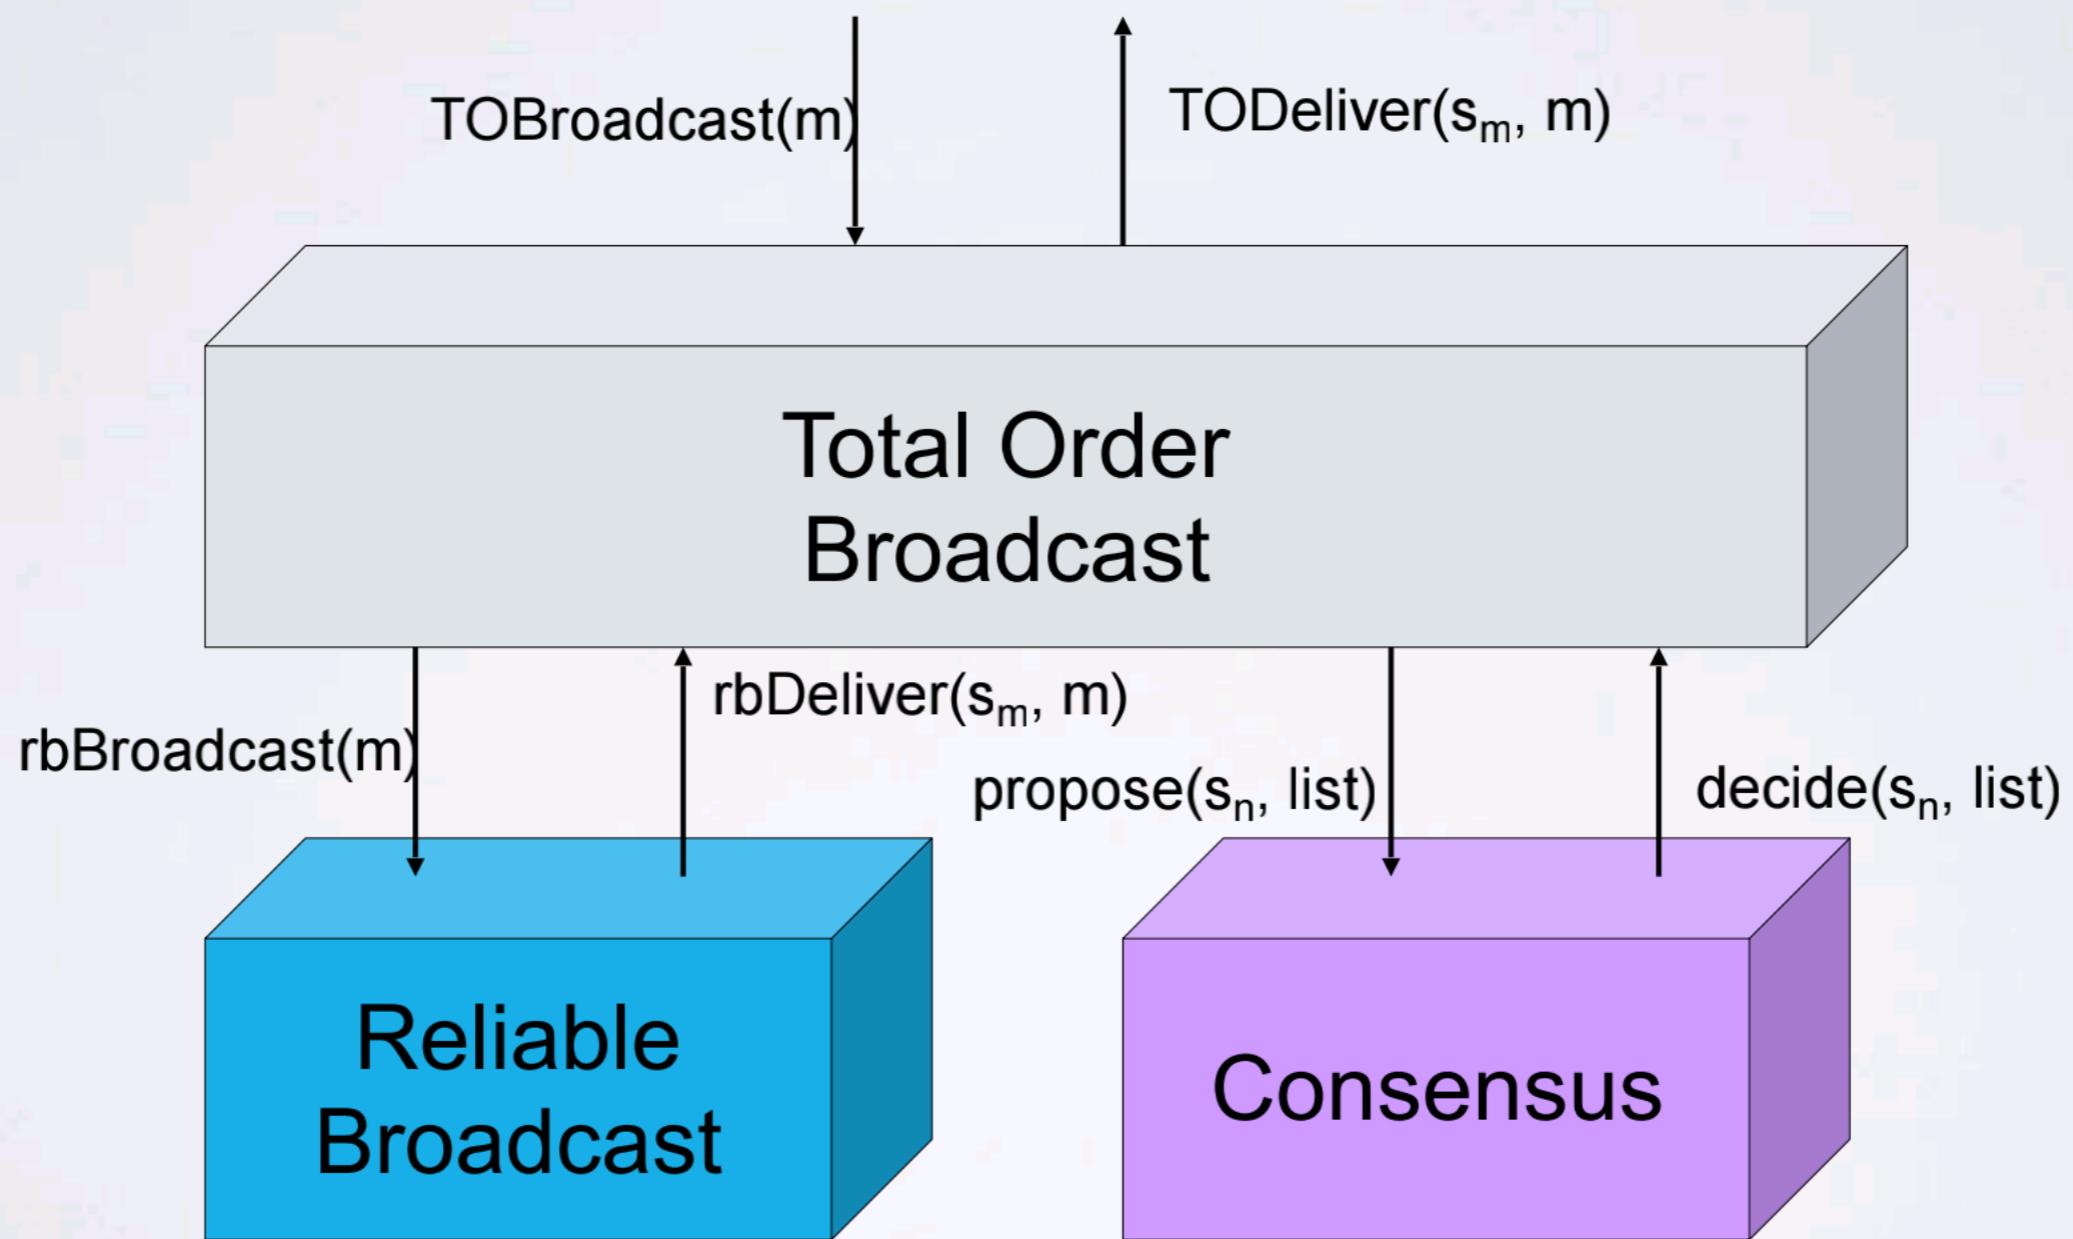
**TOB2: No duplication:** No message is delivered more than once.

**TOB3: No creation:** If a process delivers a message  $m$  with sender  $s$ , then  $m$  was previously broadcast by process  $s$ .

- **TOB4: Agreement:** If a message  $m$  is delivered by some correct process, then  $m$  is eventually delivered by every correct process.

**TOB5: Total order:** Let  $m_1$  and  $m_2$  be any two messages and suppose  $p$  and  $q$  are any two correct processes that deliver  $m_1$  and  $m_2$ . If  $p$  delivers  $m_1$  before  $m_2$ , then  $q$  delivers  $m_1$  before  $m_2$ .

# TOTAL ORDER IMPLEMENTATION



# TOTAL ORDER ALGORITHM

## Algorithm 6.1: Consensus-Based Total-Order Broadcast

Implements:

TotalOrderBroadcast, instance  $tob$ .

Uses:

ReliableBroadcast, instance  $rb$ ;  
Consensus (multiple instances).

**upon event**  $\langle tob, \text{Init} \rangle$  **do**

- $unordered := \emptyset;$
- $delivered := \emptyset;$
- $round := 1;$
- $wait := \text{FALSE};$

**upon event**  $\langle tob, \text{Broadcast} \mid m \rangle$  **do**

trigger  $\langle rb, \text{Broadcast} \mid m \rangle;$

**upon event**  $\langle rb, \text{Deliver} \mid p, m \rangle$  **do**

if  $m \notin delivered$  then → no duplication  
 $unordered := unordered \cup \{(p, m)\}$ ; || ← total order

**upon**  $unordered \neq \emptyset \wedge wait = \text{FALSE}$  **do**  
 $wait := \text{TRUE};$

Initialize a new instance  $c.\text{round}$  of consensus;  
trigger  $\langle c.\text{round}, \text{Propose} \mid \underline{unordered} \rangle;$

**upon event**  $\langle c.r, \text{Decide} \mid \underline{decided} \rangle$  such that  $r = \text{round}$  **do**

**forall**  $(s, m) \in \text{sort}(decided)$  **do** // by the order in  $decided$

trigger  $\langle tob, \text{Deliver} \mid s, m \rangle;$  // in the resulting sorted list

$delivered := delivered \cup decided;$

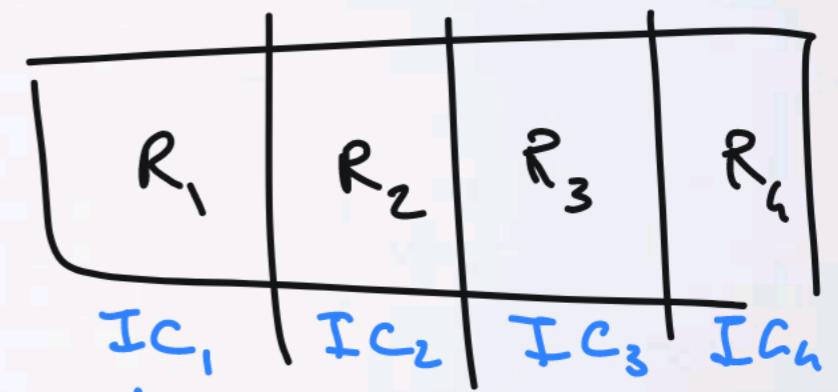
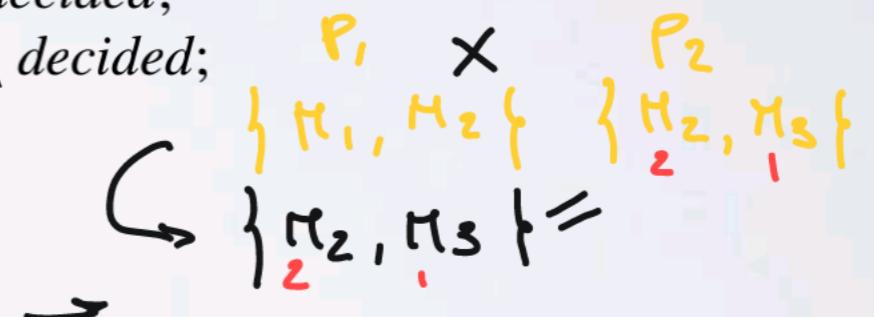
$unordered := unordered \setminus decided;$

$round := round + 1;$

$wait := \text{FALSE};$

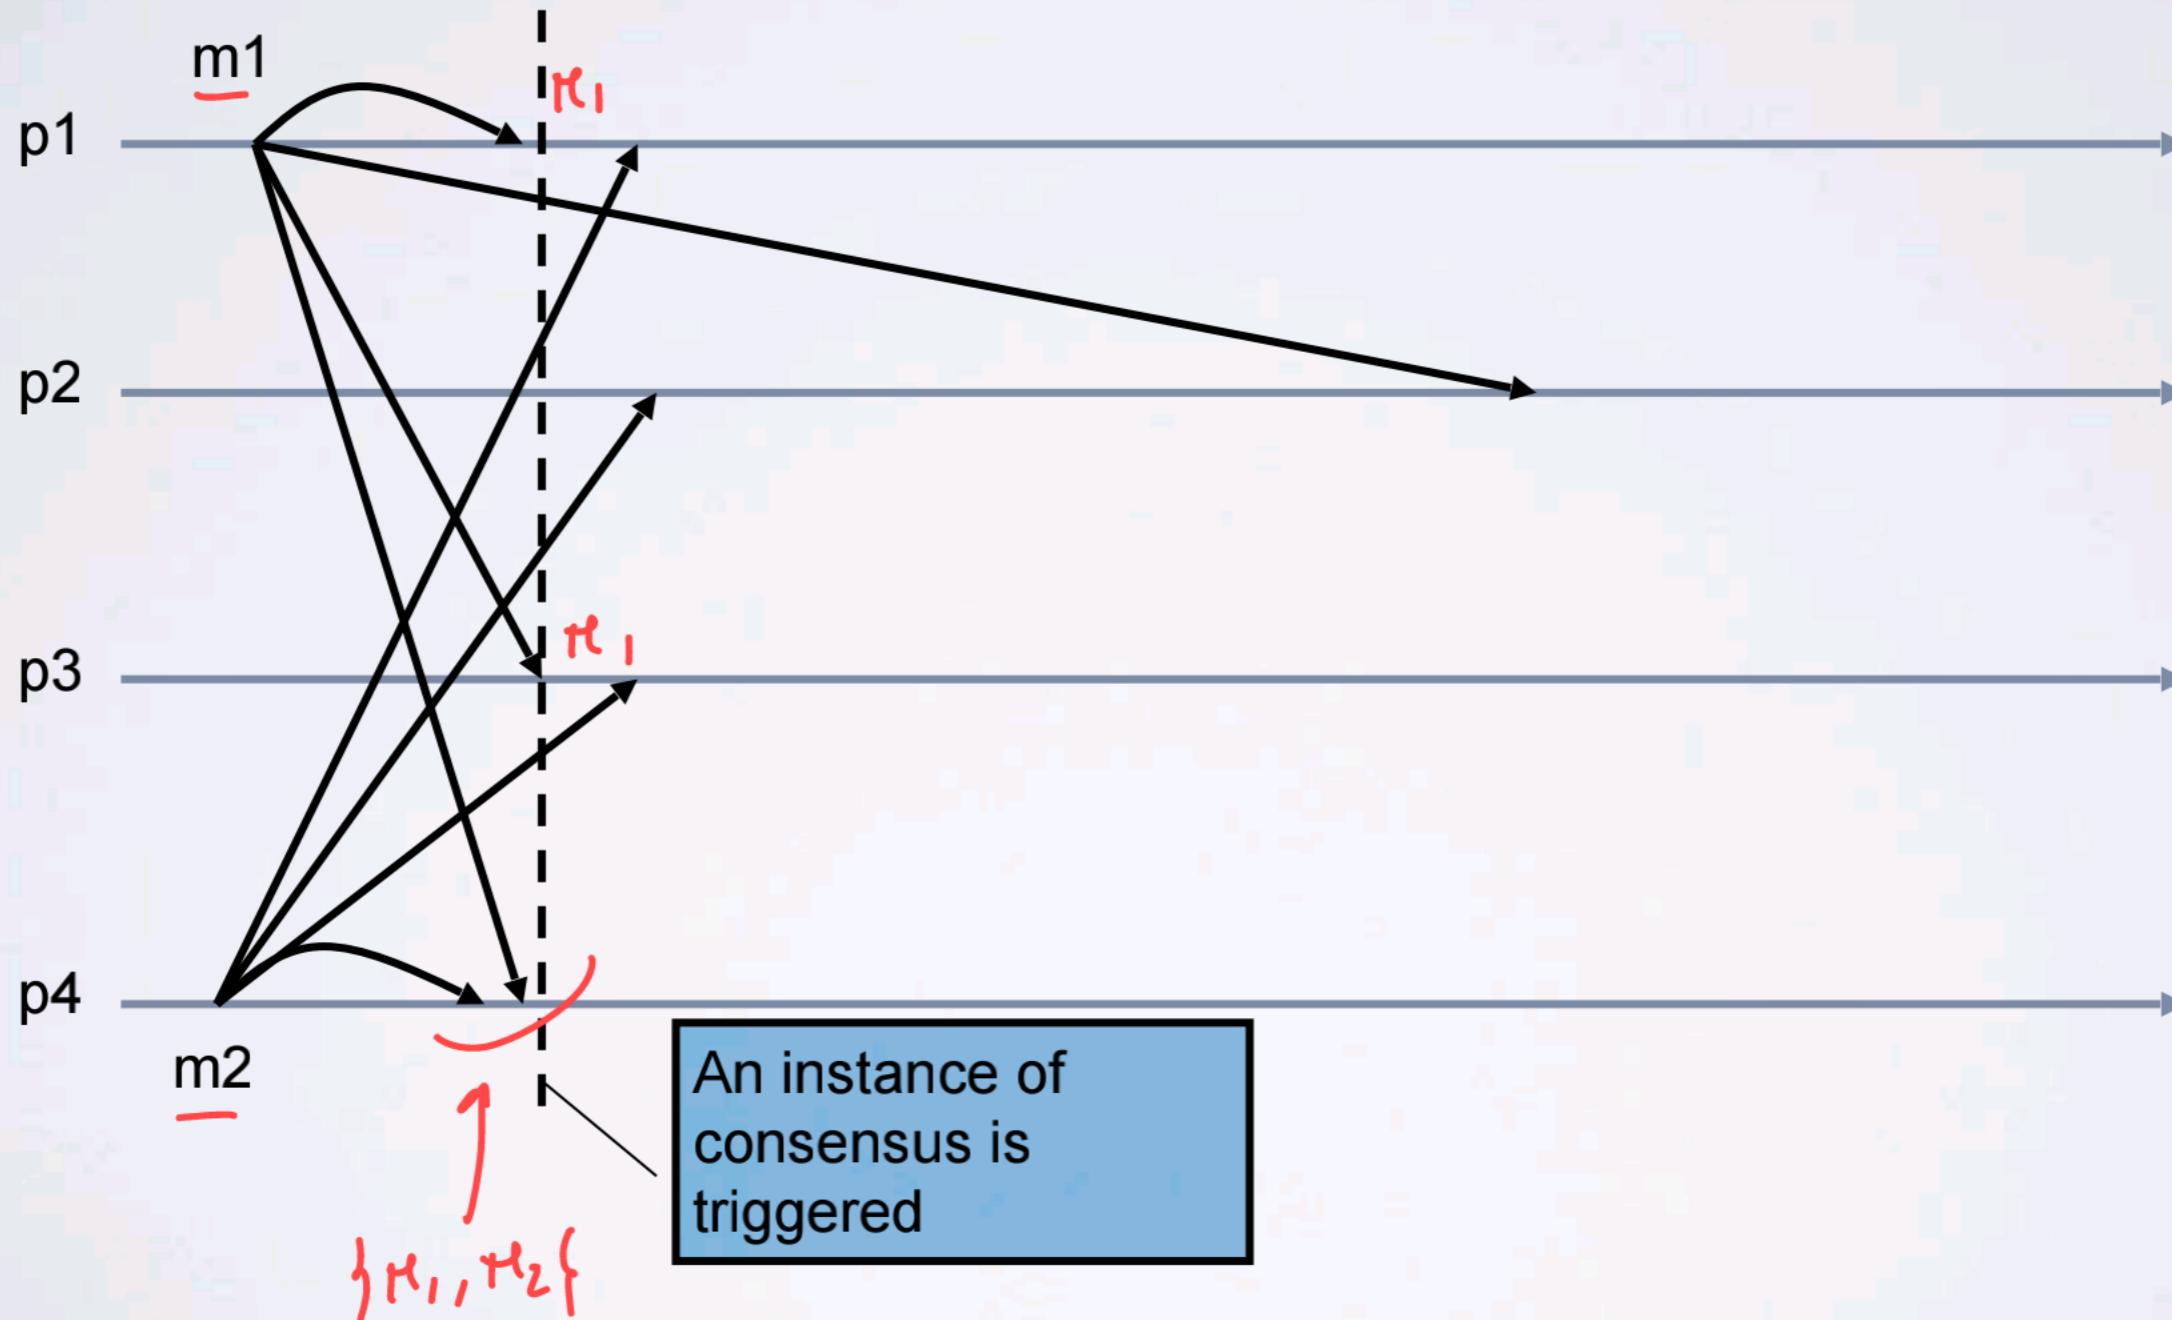
remove  $H_2$   
in the example

messages to send

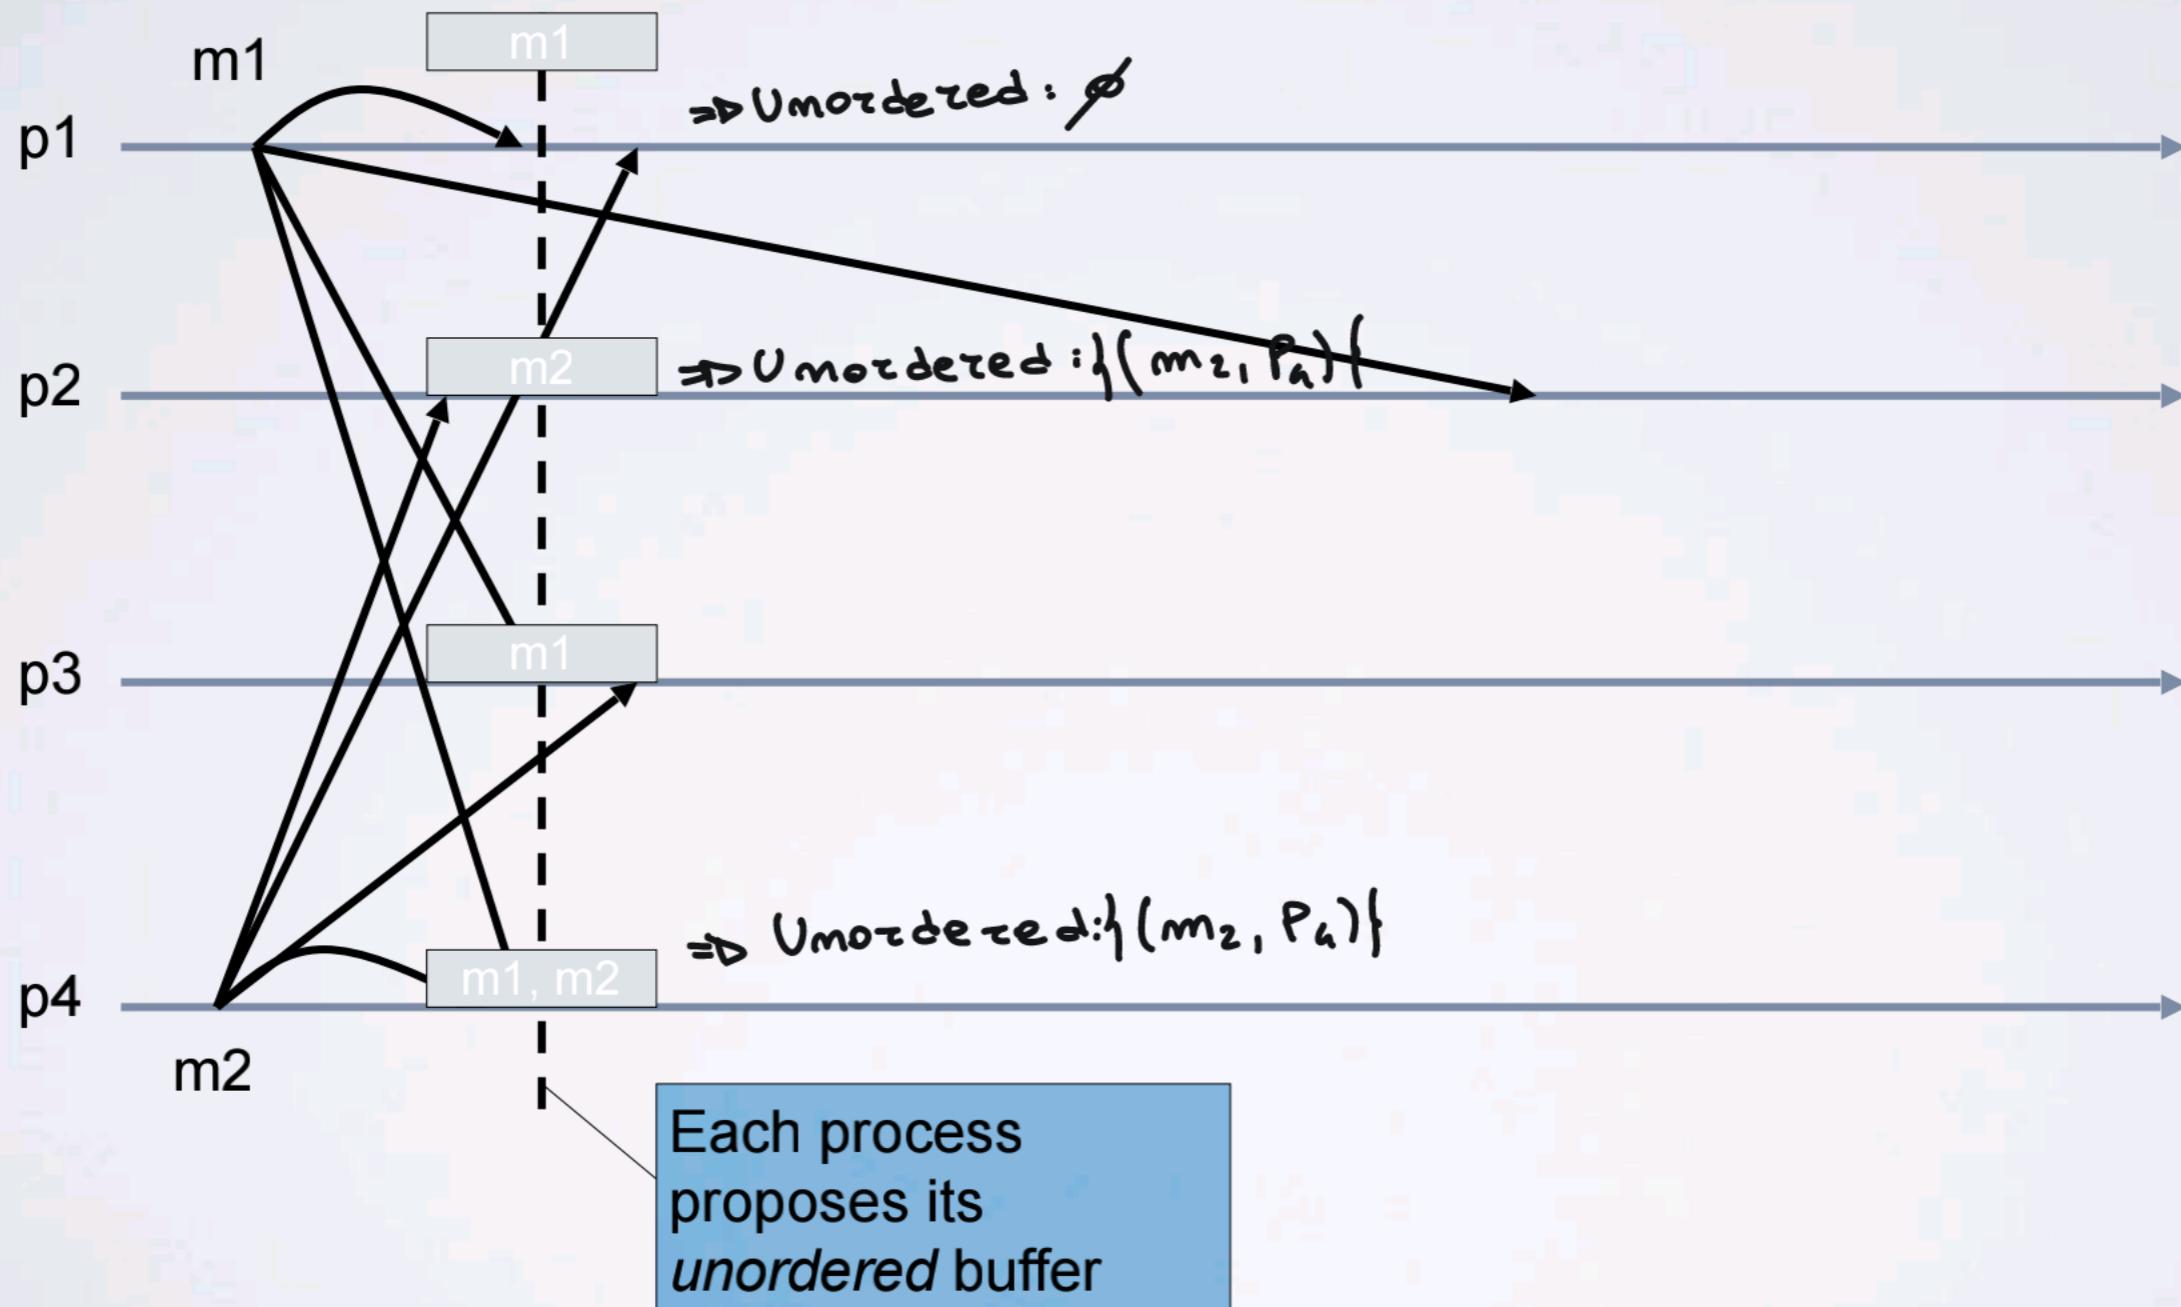


instance of  
consensus

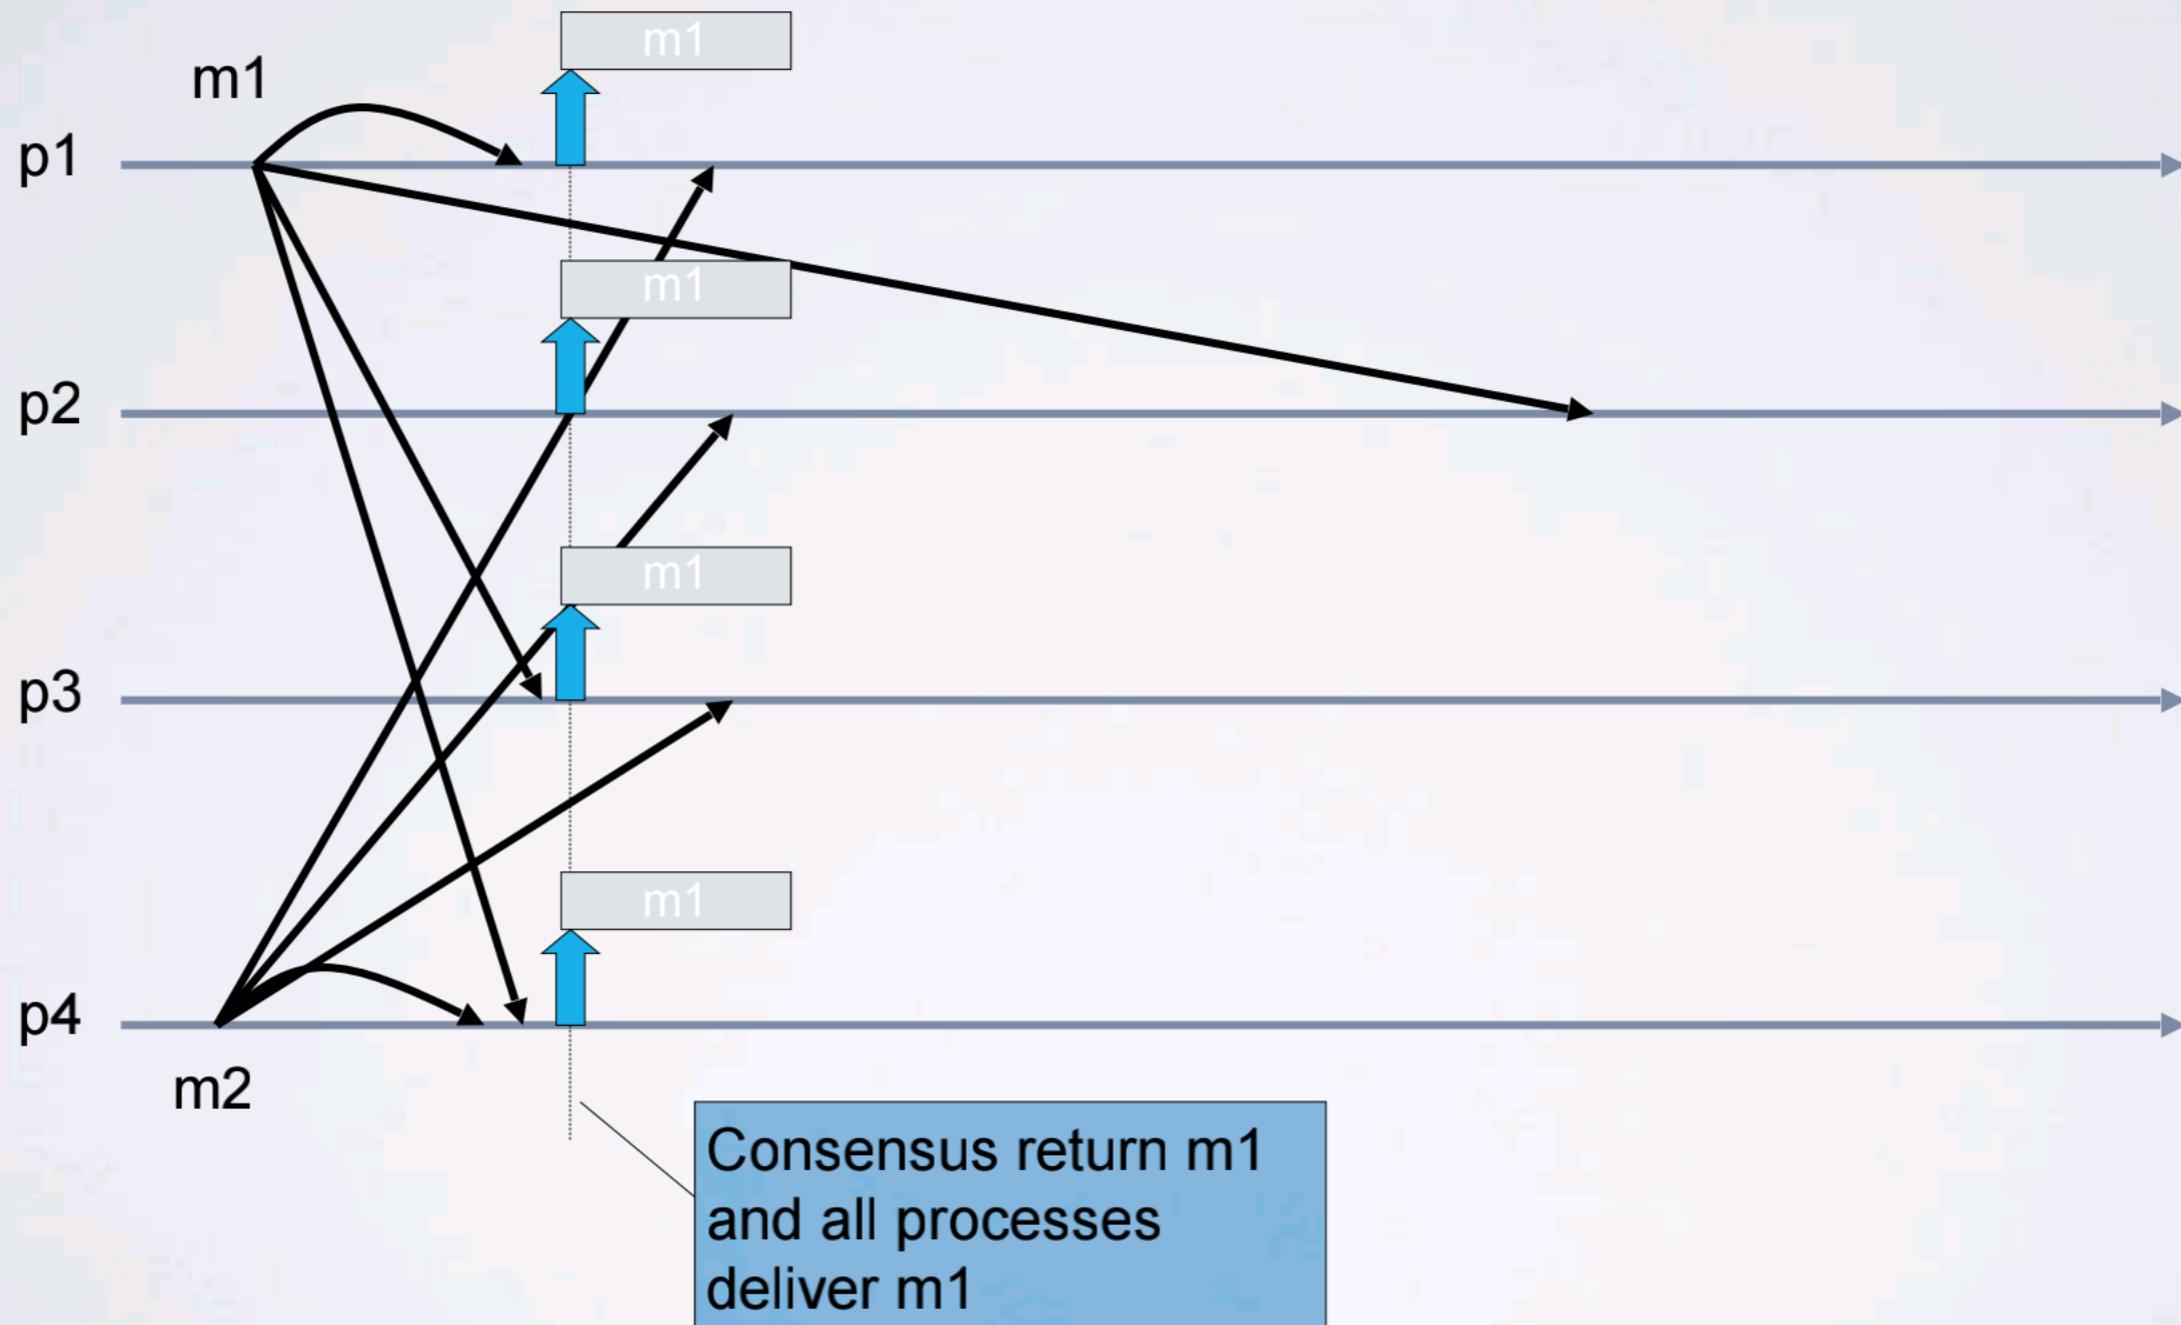
# EXAMPLE



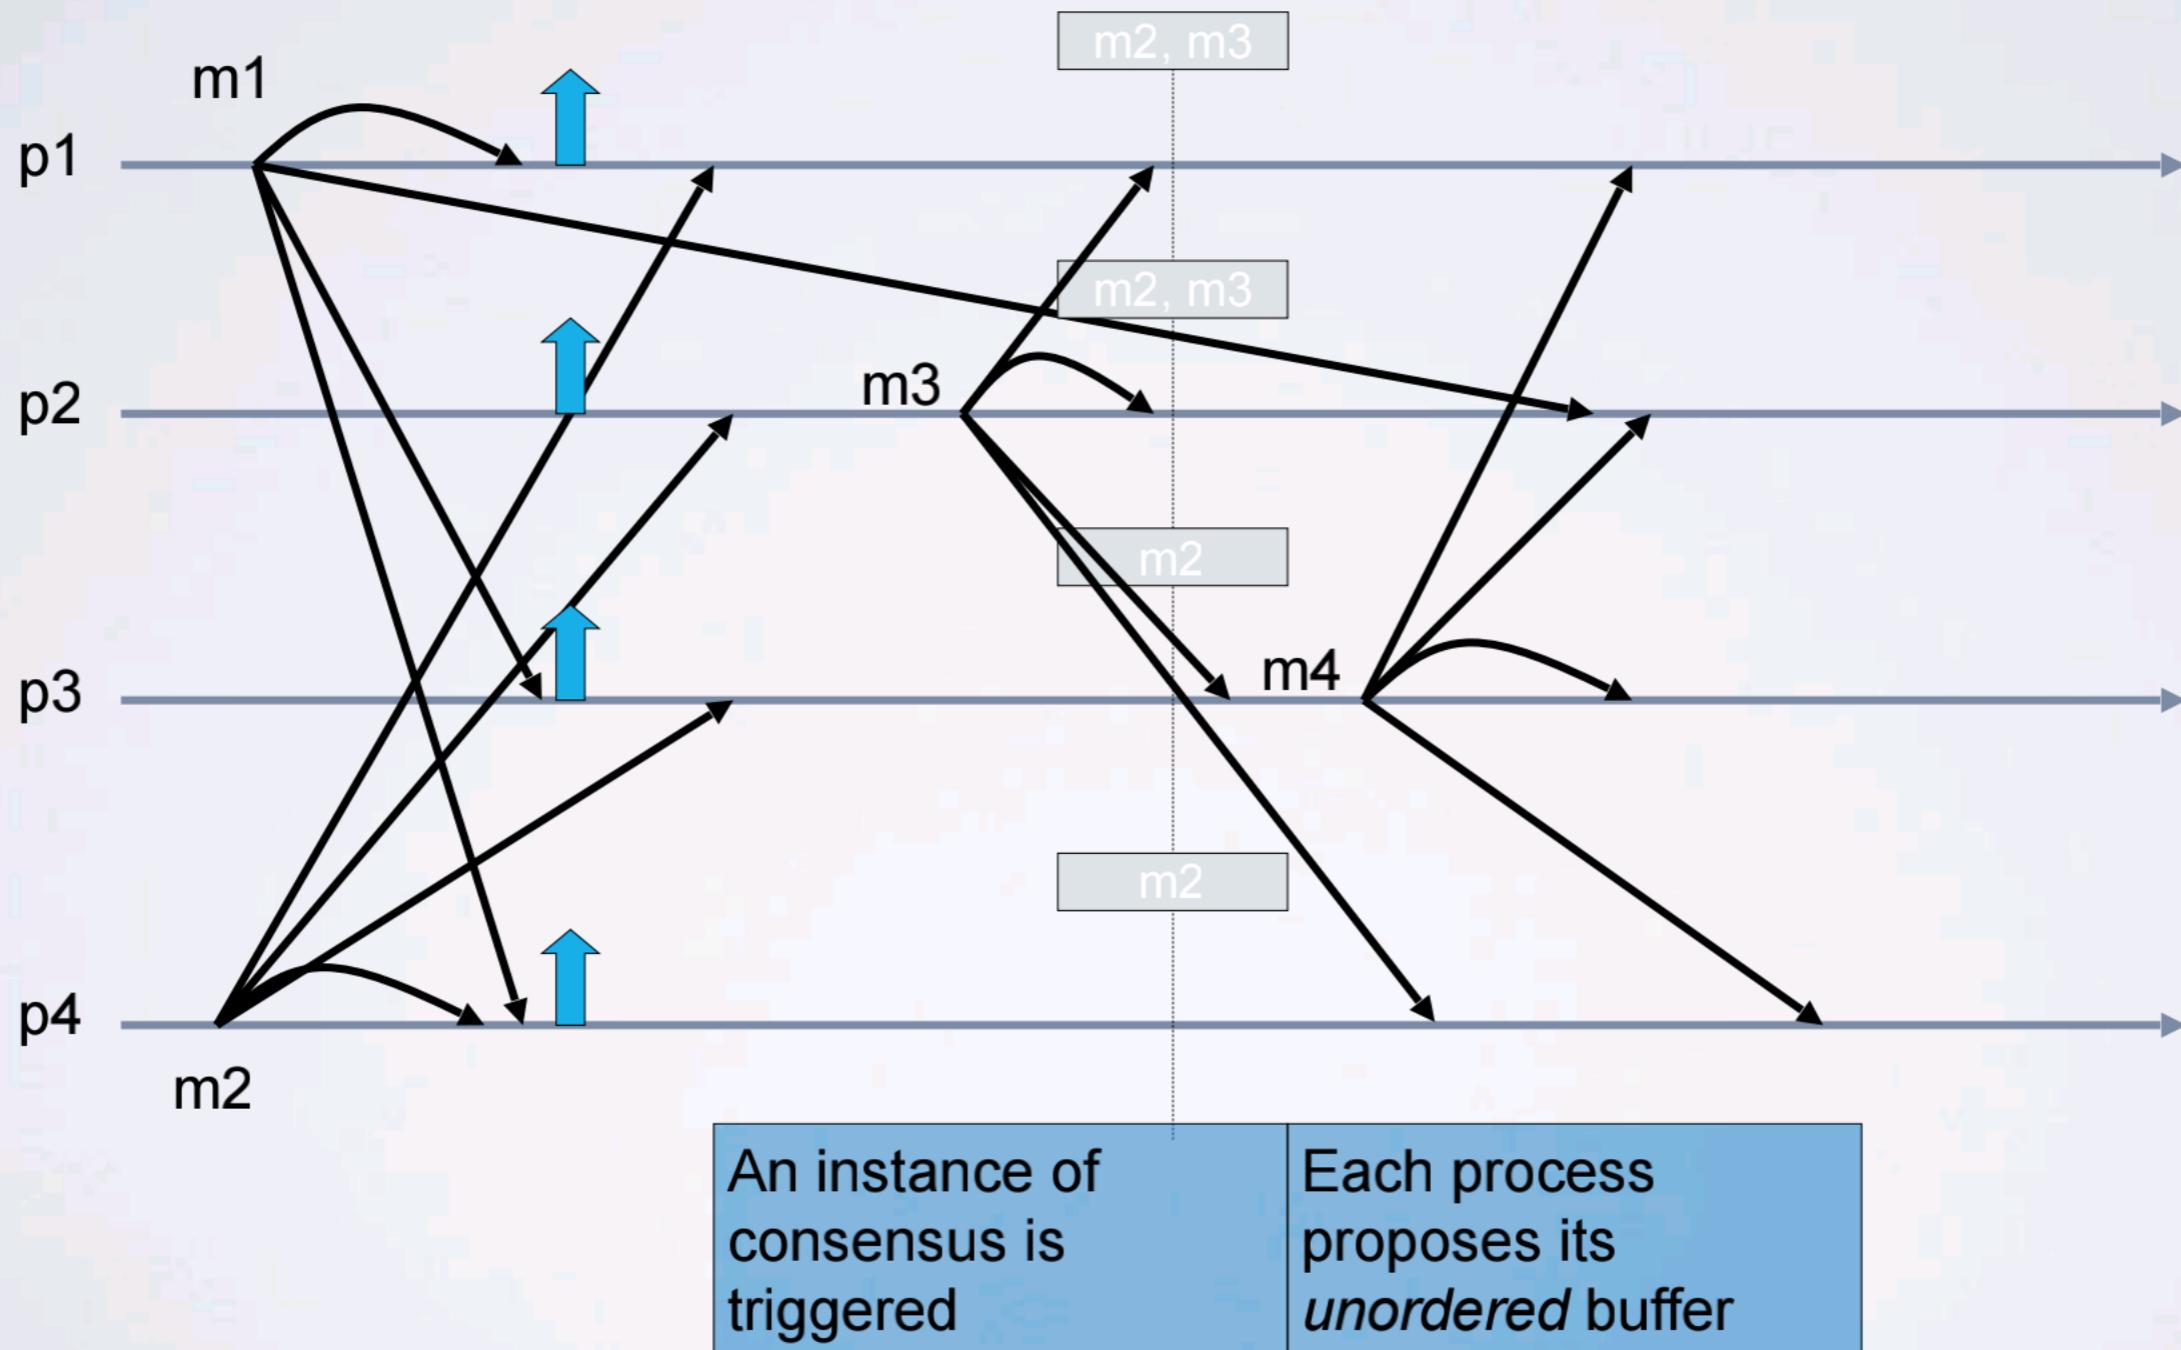
# EXAMPLE



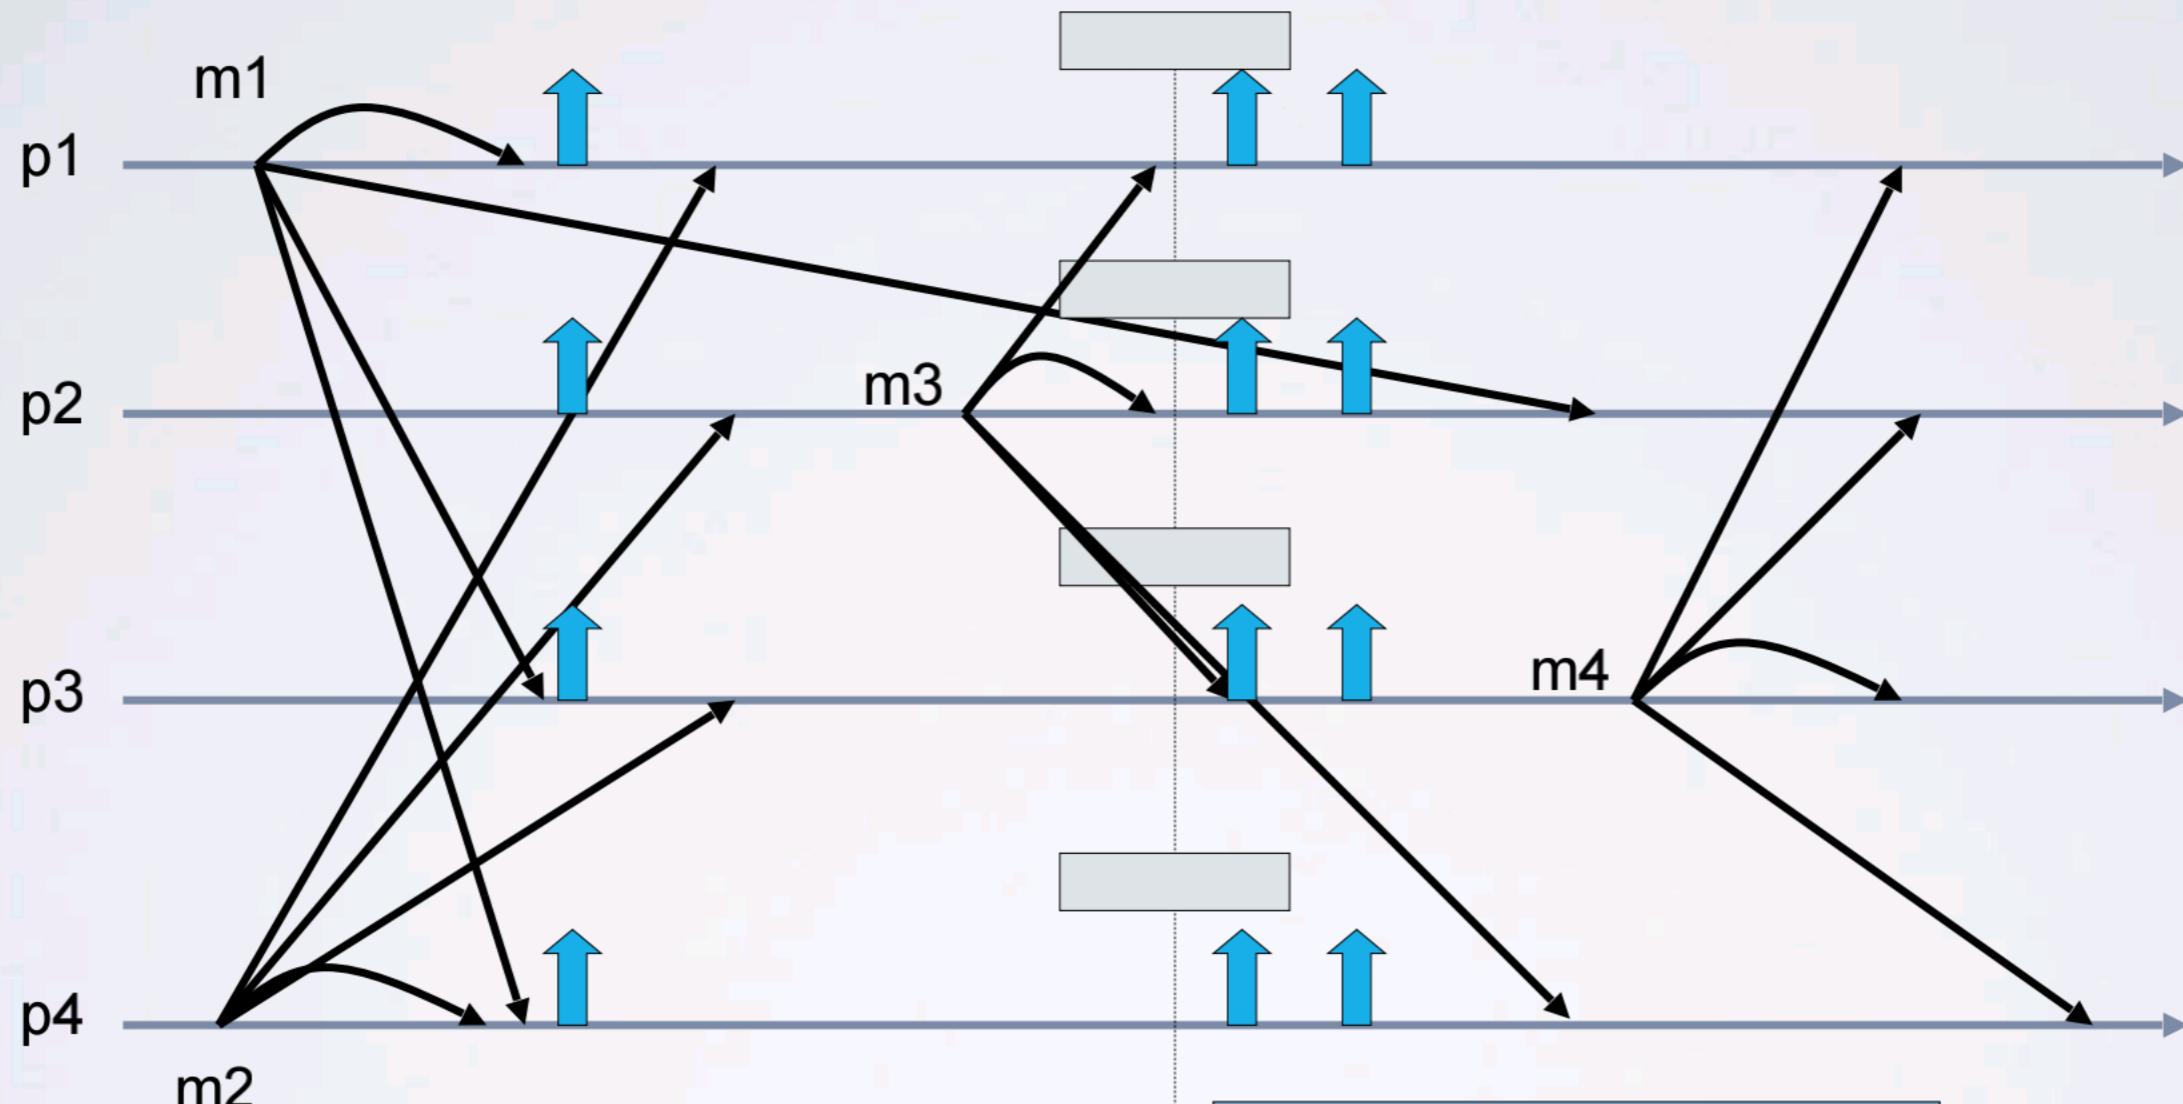
# EXAMPLE



# EXAMPLE



# EXAMPLE

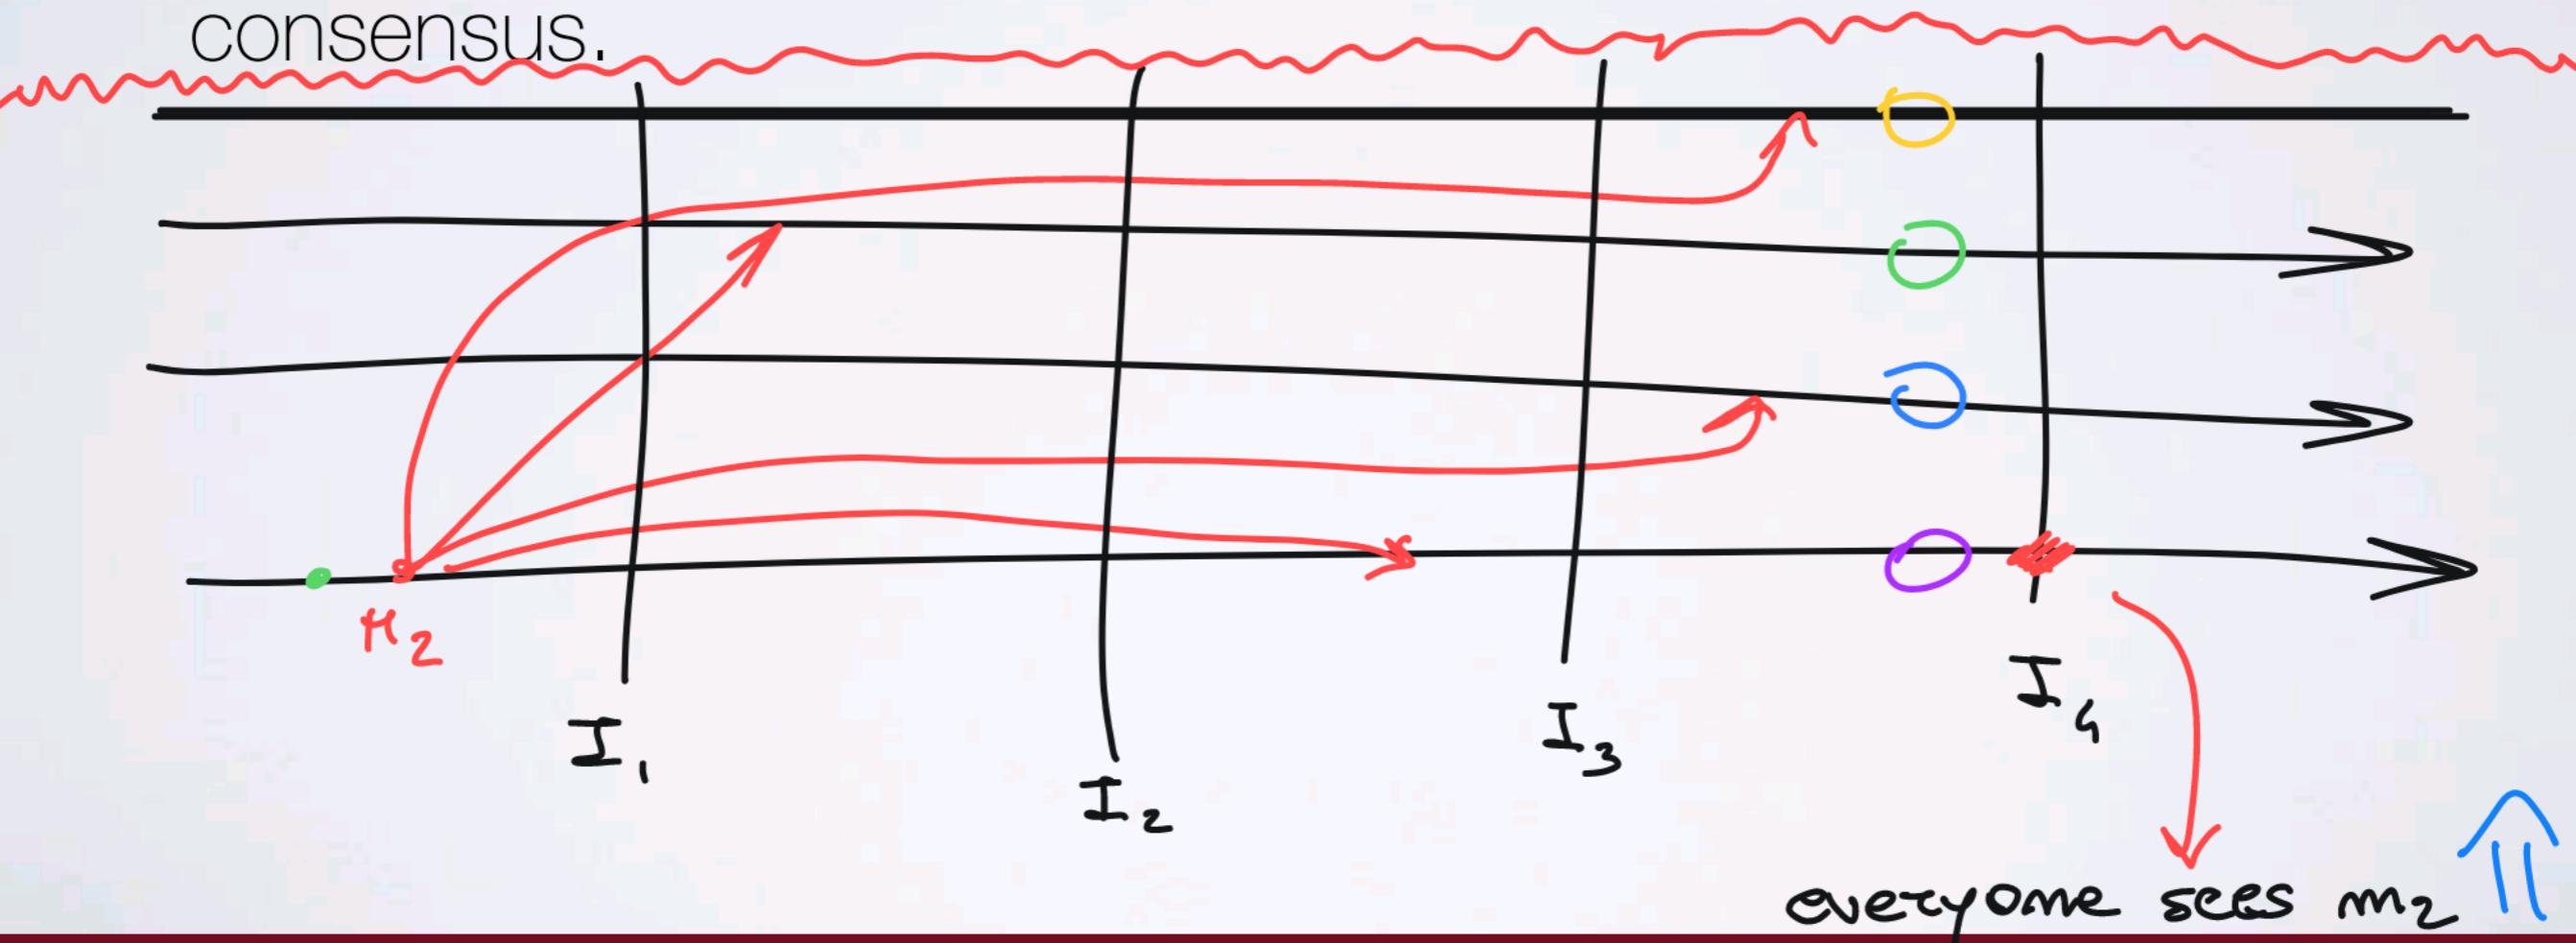


Consensus return  $m_2, m_3$   
and all processes deliver  
 $m_2$  and then  $m_3$

# FORMAL PROOF

**No creation:** There are two points at which a message may be created:

- $\langle \text{rb}, \text{deliver}, p, m \rangle$ . -> impossible by no creation of rb
- $\langle c, r, \text{decided}, \text{messages} \rangle$  -> impossible by validity of consensus.



# FORMAL PROOF

**No duplication:** Suppose by contradiction you deliver message m twice. Then you deliver it because it is in the decided of two different instances of consensus ( $r$ , and  $r'$  with  $r' > r$ ).

## Why?

But at round  $r$  you put m in delivered (and also the others), and you remove it from unordered. Now if at  $r'$  you see m in decided then someone, let it p, triggered  $\langle r', \text{Propose}, \text{unordered} \rangle$  with unordered containing m. **WHY?**

But p at the end of round  $r$ , p puts m in delivered and removes it from unordered. So it means it had to add it between the end of round  $r$  and the beginning of round  $r'$ . **But unordered is guarded by delivered!**

# FORMAL PROOF

**Agreement:** Suppose a correct process  $p_0$  delivers  $m$  at round  $r$ .

Then  $m$  is in decided of the consensus of round  $r$ .

If  $m$  is not in the decided of the consensus of round  $r$  of a correct process  $p_1$ , which **property of consensus are you violating?**

So  $m$  has to be in the decided of any correct process at round  $r$ .  $\rightarrow$  agreement.

# FORMAL PROOF

**Total order:** We have to divide the proof in two cases.

1) m and m' are delivered in the same round -> they must be delivered in the same round on each pair of correct process (WHY?)-> the order is fixed by the deterministic function order. Thus m and m' are delivered in the same order on each correct process.

1) m and m' are delivered in two different rounds r and r' on correct process p -> The only thing that we have to show is that each other correct delivers them at the exact same round r and r' -> but this is direct from the agreement of the consensus.

# QUESTIONS

- What happens if I have Beb instead of RB? Is the algorithm still correct or not?
- Why I need a broadcast?

## Module 6.2: Interface and properties of uniform total-order broadcast

**Module:**

**Name:** UniformTotalOrderBroadcast, **instance** *utob*.

**Events:**

**Request:**  $\langle \text{utob}, \text{Broadcast} \mid m \rangle$ : Broadcasts a message  $m$  to all processes.

**Indication:**  $\langle \text{utob}, \text{Deliver} \mid p, m \rangle$ : Delivers a message  $m$  broadcast by process  $p$ .

**Properties:**

**UTOB1–UTOB3:** Same as properties TOB1–TOB3 in regular total-order broadcast (Module 6.1).

**UTOB4:** *Uniform agreement*: If a message  $m$  is delivered by some process (whether correct or faulty), then  $m$  is eventually delivered by every correct process.

**UTOB5:** *Uniform total order*: Let  $m_1$  and  $m_2$  be any two messages and suppose  $p$  and  $q$  are any two processes that deliver  $m_1$  and  $m_2$  (whether correct or faulty). If  $p$  delivers  $m_1$  before  $m_2$ , then  $q$  delivers  $m_1$  before  $m_2$ .

# UNIFORM

## Algorithm 6.1: Consensus-Based Total-Order Broadcast

Implements:

TotalOrderBroadcast, **instance** *tob*.

Uses:

ReliableBroadcast, **instance** *rb*;

Consensus (multiple instances).

**upon event**  $\langle \text{tob}, \text{Init} \rangle$  **do**

*unordered* :=  $\emptyset$ ;

*delivered* :=  $\emptyset$ ;

*round* := 1;

*wait* := FALSE;

**upon** *unordered*  $\neq \emptyset \wedge \text{wait} = \text{FALSE}$  **do**

*wait* := TRUE;

        Initialize a new instance *c.round* of consensus;

**trigger**  $\langle \text{c.round}, \text{Propose} \mid \text{unordered} \rangle$ ;

**upon event**  $\langle \text{c.r}, \text{Decide} \mid \text{decided} \rangle$  **such that** *r* = *round* **do**

**forall**  $(s, m) \in \text{sort}(\text{decided})$  **do**

            // by the order in the list

**trigger**  $\langle \text{tob}, \text{Deliver} \mid s, m \rangle$ ;

            // in the resulting sorted list

*delivered* := *delivered*  $\cup \text{decided}$ ;

*unordered* := *unordered*  $\setminus \text{decided}$ ;

*round* := *round* + 1;

*wait* := FALSE;

**upon event**  $\langle \text{tob}, \text{Broadcast} \mid m \rangle$  **do**

**trigger**  $\langle \text{rb}, \text{Broadcast} \mid m \rangle$ ;

**upon event**  $\langle \text{rb}, \text{Deliver} \mid p, m \rangle$  **do**

**if** *m*  $\notin \text{delivered}$  **then**

*unordered* := *unordered*  $\cup \{(p, m)\}$ ;

# EXERCISE

Consider the previous algorithm. Which TO specification does it satisfy?

It depends from the assumptions about Reliable Broadcast and Consensus...

Consensus		
Reliable Broadcast	Uniform	Non Uniform
Uniform	?	?
Non Uniform	?	?

# EXERCISE

Consider the previous algorithm. Which TO specification does it satisfy?

It depends from the assumptions about Reliable Broadcast and Consensus...

Consensus		
Reliable Broadcast	Uniform	Non Uniform
Uniform	?	?
Non Uniform	?	Non-uniform TO

# EXERCISE

Consider the previous algorithm. Which TO specification does it satisfy?

Consensus		
Reliable Broadcast	Uniform	Non Uniform
Uniform	?	<b>Non-uniform TO</b>
Non Uniform	?	Non-uniform TO

# EXERCISE

Consider the previous algorithm. Which TO specification does it satisfy?

You only deliver the  
Decided things, however  
The decided are the same  
For each one (correct or  
not)!

Consensus		
Reliable Broadcast	Uniform	Non Uniform
Uniform	Uniform TO	Non-uniform TO
Non Uniform	?	Non-uniform TO

# EXERCISE

Consider the previous algorithm. Which TO specification does it satisfy?

**You only deliver the  
Decided things, however  
The decided are the same  
For each one (correct or  
not)!**

Consensus		
Reliable Broadcast	Uniform	Non Uniform
Uniform	Uniform TO	Non-uniform TO
Non Uniform	<b>Uniform TO</b>	Non-uniform TO

# EXERCISE

You have a total order broadcast algorithm A, is it also a causal broadcast? Or a FIFO?

# EXERCISE

You have a total order broadcast algorithm A, is it also a causal broadcast? Or a FIFO?

**NO!** Take one algorithm that uses as order function, a function that on purpose uses a reverse of causal and fifo order. It is still total!

# EXERCISE

\*\*Take the total order algorithm that we have seen in the lecture and modify it to implement a total order that is also causal.